

SPEC-DOC-002: Core Architecture Documentation

- Objectives
- Scope
  - In Scope
  - Out of Scope
- Deliverables
  - Primary Documentation
  - Supporting Materials
- Success Criteria
- Related SPECs
- Cargo Workspace Structure
  - Workspace Overview
  - Crate Categories
    - Core Application Crates
    - Backend Service Crates
    - MCP (Model Context Protocol) Crates
    - Protocol & Model Provider Crates
    - Utility & Tooling Crates
    - Execution & Sandbox Crates
    - Supporting Crates
- Dependency Graph
- Build Profiles
  - dev-fast (Default Development)
  - release (Production)
  - perf (Performance Profiling)
- Workspace-Level Configuration
  - Shared Dependencies
  - Workspace Lints
  - Patched Dependencies
- Build System
  - Fast Builds (Development)
  - Release Builds (Production)
  - Testing
  - Code Quality
- Crate Size Breakdown
- Adding New Crates
  - When to Create a New Crate
  - Creating a New Crate
- Next Steps
- Configuration System
  - Overview
  - 5-Tier Precedence
    - Precedence Order
    - Precedence Examples
  - Configuration Loading
    - Loader Architecture
    - Environment Variable Mapping
- Hot-Reload System
  - Architecture
  - HotReloadWatcher
  - Debouncing
  - Lock Contention
- Performance Metrics
  - Hot-Reload Latency
  - Lock Timing
- Validation
  - Schema Validation
  - Type Safety
- Profile System

- Profile Definition
- Profile Selection
- Profile Merging
- Registry System
  - Command Registry (Spec-Kit)
- Summary
- Core Execution System
  - Overview
  - Agent Orchestration
    - ConversationManager
    - Agent Spawning Pattern
    - Multi-Agent Orchestration (Spec-Kit)
  - Model Providers
    - Provider Architecture
    - OpenAI Provider
    - Anthropic Provider (CLI)
    - Google Provider (CLI)
  - Protocol Implementation
    - Request/Response Types
    - Streaming Events
  - Retry Logic
    - Exponential Backoff
    - Error Classification
    - Retry Execution
  - Timeout Management
    - Per-Operation Timeouts
    - Configurable Timeouts
  - Tool Execution
    - Tool Call Flow
    - Sandbox Enforcement
  - Performance Optimizations
    - Connection Pooling
    - Concurrent Execution
  - Error Handling
    - Error Hierarchy
    - Graceful Degradation
  - Summary
- Database Layer
  - Overview
  - Architecture
  - Connection Pooling
    - R2D2 Configuration
    - Connection Customizer (Pragma Optimization)
    - Performance Impact
  - Schema
    - consensus\_runs
    - agent\_outputs
    - consensus\_artifacts (legacy)
  - Transaction Handling
    - Transaction Behavior
    - Transaction Helpers
    - Batch Operations
  - Async Wrapper
  - Consensus Database (Spec-Kit)
    - ConsensusDb
    - Dual-Schema Migration (SPEC-945B)
  - Auto-Vacuum Strategy
    - Incremental Auto-Vacuum
  - Retry Logic (Database Operations)
    - Sync Retry

- Error Handling
  - Database Errors
  - Error Classification
- Schema Migrations
  - Migration System
- Summary
- MCP Integration
  - Overview
  - Architecture
    - McpClient Implementation
      - Core Structure
      - JSON-RPC Protocol
      - Tool Invocation
    - McpConnectionManager
      - Tool Name Qualification
      - App-Level Shared Connection Manager
  - Performance: Native vs Subprocess
    - Benchmark Results
  - Server Lifecycle
    - Initialization Sequence
    - Shutdown Sequence
    - Health Monitoring
  - Error Handling
    - Error Types
    - Retry Logic
    - Graceful Degradation
  - Tool Schema Validation
    - Input Schema
  - Configuration
    - Per-Server Config
    - Excluded Tools
  - Summary
- System Overview & Architecture
  - Table of Contents
  - High-Level Overview
  - Design Philosophy
    - 1. Separation of Concerns
    - 2. Async/Sync Hybrid Architecture
    - 3. Modularity via Cargo Workspace
    - 4. Fork Isolation Strategy
    - 5. Performance-First Design
    - 6. Fault Tolerance
  - Component Architecture
    - Core Components
  - Data Flow
    - Conversation Flow
    - Spec-Kit Automation Flow
    - MCP Tool Invocation Flow
  - Technology Stack
    - Core Technologies
    - Supporting Libraries
  - Fork-Specific Additions
    - Isolation Strategy
    - New Crates
  - Integration Points
    - 1. TUI ↔ Core Services
    - 2. Core ↔ MCP Servers
    - 3. Core ↔ Model Providers
    - 4. Spec-Kit ↔ Database
    - 5. Configuration ↔ Filesystem

- Summary
- TUI Architecture
  - Overview
  - Component Hierarchy
  - ChatWidget Structure
  - Async/Sync Boundary Pattern
    - The Problem
    - The Solution: Channel Bridge
    - Operation Types (Op enum)
  - Event Loop
- Rendering System
  - Immediate Mode Rendering
  - Widget Composition
  - HistoryCell Trait
- Spec-Kit Integration (Friend Module)
  - Friend Module Pattern
  - Context Trait Abstraction
- Streaming & Interrupts
  - StreamController
  - InterruptManager
- Input Handling
  - BottomPane (Composer)
- File Search (@trigger)
- Performance Considerations
  - Rendering Optimizations
  - Event Processing
- Summary

## SPEC-DOC-002: Core Architecture Documentation

**Status:** Pending **Priority:** P0 (High) **Estimated Effort:** 16-20 hours  
**Target Audience:** Contributors, system architects **Created:** 2025-11-17

---

## Objectives

Document the complete internal architecture of the theturtlecsz/code project:  
1. System overview and design philosophy  
2. Cargo workspace structure (24 crates, dependencies)  
3. TUI architecture (Ratatui, async/sync boundaries)  
4. Core execution system (agent orchestration, tmux management)  
5. MCP integration (native client, connection management)  
6. Database layer (SQLite, schema, transactions)  
7. Configuration system (5-tier precedence, hot-reload)

---

## Scope

### In Scope

**System Architecture:** - High-level component diagram - Data flow diagrams - Integration points between subsystems - Fork-specific additions (98.8% isolation from upstream)

**Cargo Workspace** (24 crates): - Crate dependency graph - Purpose of each crate - Inter-crate dependencies - Build profiles (dev-fast, release, perf)

**TUI System:** - Ratatui architecture - ChatWidget structure (912K LOC mod.rs) - Async/sync boundary handling (Handle::block\_on pattern) - Friend module pattern for spec-kit isolation - Widget lifecycle and rendering

**Core Execution:** - Agent spawning and orchestration - Tmux session management - Model provider clients (OpenAI, Anthropic, Google) - Protocol implementation - Timeout and retry logic (AR-1 through AR-4)

**MCP Integration:** - Native client implementation (5.3x speedup) - App-level shared connection manager (ARCH-005) - Server lifecycle management - Tool invocation patterns

**Database Layer:** - SQLite schema (consensus\_artifacts.db) - Transaction handling (IMMEDIATE mode) - File locking (fs2 crate, ARCH-007) - Retry logic and error handling - Auto-vacuum strategy (99.95% reduction)

**Configuration System:** - 5-tier precedence (CLI > shell > profile > TOML > defaults) - Hot-reload mechanism (config\_reload.rs, 300ms debounce) - Profile system - Environment variable overrides

## Out of Scope

- User-facing configuration guide (see SPEC-DOC-006)
  - Spec-kit framework details (see SPEC-DOC-003)
  - Testing infrastructure (see SPEC-DOC-004)
  - Security implementation (see SPEC-DOC-007)
- 

## Deliverables

### Primary Documentation

1. **content/system-overview.md** - Architecture overview, component diagram
2. **content/cargo-workspace.md** - Workspace structure, crate guide
3. **content/tui-architecture.md** - Ratatui, async/sync, ChatWidget
4. **content/core-execution.md** - Agent orchestration, tmux, providers
5. **content/mcp-integration.md** - Native client, connection manager
6. **content/database-layer.md** - SQLite, schema, transactions
7. **content/configuration-system.md** - 5-tier precedence, hot-reload

### Supporting Materials

- **evidence/diagrams/** - Architecture diagrams, data flow charts
  - **adr/** - Key architectural decisions (if new decisions documented)
-

## Success Criteria

- Complete component diagram created
  - All 24 crates documented with purposes
  - Async/sync boundary patterns explained with examples
  - MCP integration fully documented (connection lifecycle, retry logic)
  - SQLite schema documented with ER diagram
  - Configuration precedence clearly illustrated
  - All file paths reference actual source code locations
- 

## Related SPECs

- SPEC-DOC-000 (Master)
  - SPEC-DOC-001 (User Onboarding - references architecture concepts)
  - SPEC-DOC-003 (Spec-Kit - detailed spec-kit architecture)
  - SPEC-DOC-004 (Testing - test infrastructure architecture)
  - SPEC-DOC-005 (Development - build system details)
- 

**Status:** Structure defined, content pending

---

## Cargo Workspace Structure

Complete guide to the 24-crate workspace architecture.

---

## Workspace Overview

**Location:** /home/user/code/codex-rs/Cargo.toml

**Statistics:** - 24 member crates - 226,607 lines of Rust code - 538 source files - Edition 2024 (Rust 2024 edition features)

---

## Crate Categories

### Core Application Crates

#### tui - Terminal User Interface

**Purpose:** Main application binary, Ratatui-based TUI

**Key Modules:** - app.rs: Application state and event loop - chatwidget/: Conversation interface (912K LOC mod.rs) - chatwidget/spec\_kit/: Fork feature integration (55 modules)

**Dependencies:** ratatui, codex-core, spec-kit, mcp-client, tokio

**Binary:** code-tui

**File:** codex-rs/tui/Cargo.toml

---

### cli - CLI Wrapper

**Purpose:** Command-line interface entry point

**Responsibilities:** - Argument parsing (clap) - Shell completion generation - Delegates to code-tui binary

**Dependencies:** codex-tui, codex-core, clap, clap\_complete

**Binary:** code

**File:** codex-rs/cli/Cargo.toml

---

### spec-kit - Multi-Agent Automation Framework (Fork)

**Purpose:** Reusable library for spec-kit automation

**Modules:** - config/: Configuration system (hot-reload, 5-tier precedence) - retry/: Exponential backoff retry logic - types.rs: Core types (SpecStage, QualityCheckpoint) - error.rs: Error handling

**Dependencies:** codex-core, mcp-types, rusqlite, notify, tokio

**Note:** Can be extracted as standalone crate (MAINT-10 future work)

**File:** codex-rs/spec-kit/Cargo.toml

---

## Backend Service Crates

### core - Backend Services

**Purpose:** Backend orchestration (conversation, MCP, DB, config)

**Key Modules:** - conversation\_manager.rs: Agent lifecycle - mcp\_connection\_manager.rs: MCP server aggregation - db/: SQLite connection pooling, transactions - config\_types.rs: Configuration data structures - protocol.rs: Model provider abstraction

**Dependencies:** mcp-client, codex-protocol, rusqlite, r2d2, tokio

**File:** codex-rs/core/Cargo.toml

---

## MCP (Model Context Protocol) Crates

### mcp-client - Async MCP Client

**Purpose:** Subprocess communication with MCP servers

**Key Features:** - Line-delimited JSON-RPC over stdio - Concurrent reader/writer tasks (prevent deadlock) - Request/response pairing via HashMap - 1MB buffer for large responses

**Dependencies:** mcp-types, tokio, tokio-util, serde\_json

**File:** codex-rs/mcp-client/src/mcp\_client.rs:63-150

---

### **mcp-server - MCP Server Implementation**

**Purpose:** Implements MCP server for tool exposure

**Dependencies:** codex-core, codex-protocol, mcp-types, tokio

**File:** codex-rs/mcp-server/Cargo.toml

---

### **mcp-types - Protocol Types**

**Purpose:** JSON-RPC and MCP protocol definitions

**Key Types:** - JSONRPCMessage: Request/Response/Notification - ToolInfo: Tool metadata - MpcServerConfig: Server configuration

**Dependencies:** serde, serde\_json, ts-rs (TypeScript bindings)

**File:** codex-rs/mcp-types/Cargo.toml

---

## **Protocol & Model Provider Crates**

### **protocol - OpenAI Protocol**

**Purpose:** OpenAI API client and types

**Features:** - Responses API support - Chat Completions API support - Streaming SSE support - Request/response types

**Dependencies:** serde, serde\_json, reqwest

**File:** codex-rs/protocol/Cargo.toml

---

### **chatgpt - ChatGPT Auth**

**Purpose:** ChatGPT authentication flow

**Dependencies:** reqwest, serde

**File:** codex-rs/chatgpt/Cargo.toml

---

## **Utility & Tooling Crates**

### **common - Shared Utilities**

**Purpose:** Common utilities across crates

**Modules:** - Model presets (GPT-5, Claude, Gemini) - Elapsed time formatting - CLI helpers

**Dependencies:** serde, reqwest, clap

**File:** codex-rs/common/Cargo.toml

---

## **git-tooling - Git Operations**

**Purpose:** Git integration helpers

**File:** codex-rs/git-tooling/Cargo.toml

---

## **file-search - File Search**

**Purpose:** Fuzzy file search (@ trigger in composer)

**Dependencies:** nucleo-matcher (fuzzy matching)

**File:** codex-rs/file-search/Cargo.toml

---

## **Execution & Sandbox Crates**

### **exec - Command Execution**

**Purpose:** Sandboxed command execution

**File:** codex-rs/exec/Cargo.toml

---

### **execpolicy - Execution Policy**

**Purpose:** Approval policy enforcement

**File:** codex-rs/execpolicy/Cargo.toml

---

### **linux-sandbox - Linux Sandboxing**

**Purpose:** Landlock, seccomp sandboxing

**Dependencies:** landlock, seccompiler

**File:** codex-rs/linux-sandbox/Cargo.toml

---

## **Supporting Crates**

**login:** Authentication helpers **browser:** Browser integration (CDP)

**ollama:** Ollama model support **arg0:** Binary name detection **apply-**

**patch:** Diff application **ansi-escape:** ANSI escape code handling

**codex-version:** Version utilities

---

## **Dependency Graph**

```
tui (main binary)
  └── spec-kit (fork feature)
    ├── mcp-types
    ├── codex-core
    │   └── mcp-client
```

```

    |   |   └── mcp-types
    |   ├── codex-protocol
    |   ├── rusqlite
    |   └── r2d2
    └── rusqlite (direct)
        ├── codex-core
        ├── mcp-client
        ├── codex-protocol
        ├── codex-common
        ├── ratatui (v0.29.0 patched)
        └── tokio

cli (entry point)
├── codex-tui
└── clap

core (backend)
├── mcp-client
├── codex-protocol
├── rusqlite
├── r2d2
└── tokio

```

**Key Observations:** - **spec-kit** depends on **core** (reuses DB, MCP) -  
**tui** depends on **spec-kit** (friend module pattern) - **mcp-client** is  
lightweight (only mcp-types, tokio) - **core** aggregates backend services

---

## Build Profiles

### dev-fast (Default Development)

```
[profile.dev-fast]
inherits = "dev"
opt-level = 1                      # Basic optimization
debug = 1                           # Line tables only
incremental = true                  # Incremental compilation
codegen-units = 256                 # Parallel codegen
lto = "off"                         # No LTO (faster builds)
```

**Use:** ./build-fast.sh or cargo build --profile dev-fast

**Build Time:** ~30-60 seconds (incremental)

**Binary:** codex-rs/target/dev-fast/code

---

### release (Production)

```
[profile.release]
lto = "fat"                         # Full LTO (link-time optimization)
strip = "symbols"                    # Remove debug symbols
codegen-units = 1                   # Single codegen unit (max optimization)
```

**Use:** cargo build --release

**Build Time:** ~5-10 minutes (full rebuild)

**Binary Size:** ~15-20 MB (stripped)

**Binary:** codex-rs/target/release/code

---

## perf (Performance Profiling)

```
[profile.perf]
inherits = "release"
incremental = true
codegen-units = 256
lto = "off"
debug = 2                      # Full debug info
strip = "none"                  # Keep symbols
split-debuginfo = "packed" # Separate debug file
```

**Use:** cargo build --profile perf

**Purpose:** Performance profiling with perf, flamegraph

---

# Workspace-Level Configuration

## Shared Dependencies

```
[workspace.dependencies]
# Internal crates (path dependencies)
codex-core = { path = "core" }
codex-tui = { path = "tui" }
spec-kit = { path = "spec-kit" }
mcp-client = { path = "mcp-client" }
mcp-types = { path = "mcp-types" }
# ... 20+ internal crates

# External crates (version pinning)
tokio = "1"
ratatui = "0.29.0"
rusqlite = { version = "0.37", features = ["bundled"] }
reqwest = "0.12"
serde = "1"
serde_json = "1"
anyhow = "1"
thiserror = "2.0.16"
# ... 100+ external dependencies
```

**Benefits:** - Single version source of truth - Automatic version consistency - Easier dependency updates

---

## Workspace Lints

```
[workspace.lints.clippy]
expect_used = "deny"          # Forbid .expect()
unwrap_used = "deny"          # Forbid .unwrap()
manual_ok_or = "deny"          # Enforce .ok_or()
needless_borrow = "deny"       # Remove unnecessary borrows
redundant_clone = "deny"       # Remove unnecessary clones
uninlined_format_args = "deny" # Use format!("{}") not
format!("{}", var)
# ... 30+ lints
```

**Enforcement:** All crates inherit workspace lints unless overridden

---

## Patched Dependencies

**Ratatui Fork:**

```
[patch.crates-io]
ratatui = { git = "https://github.com/nornagon/ratatui", branch =
"nornagon-v0.29.0-patch" }
```

**Reason:** Custom patches for TUI improvements

---

## Build System

### Fast Builds (Development)

```
# Use build-fast.sh
./build-fast.sh

# Or directly
cd codex-rs
cargo build --profile dev-fast --bin code --bin code-tui
```

**Output:** codex-rs/target/dev-fast/code

---

### Release Builds (Production)

```
# Full release build
cd codex-rs
cargo build --release --bin code --bin code-tui --bin code-exec

# Quick build (code only)
npm run build:quick
```

**Output:** codex-rs/target/release/code

---

## Testing

```
cd codex-rs

# All tests
cargo test

# Specific crate
cargo test -p spec-kit

# Specific test
cargo test -p codex-tui spec_auto
```

---

## Code Quality

```
cd codex-rs
```

---

```

# Format
cargo fmt --all

# Lint
cargo clippy --workspace --all-targets --all-features -- -D warnings

# Check (no codegen)
cargo check --workspace

```

---

## Crate Size Breakdown

Crate	LOC (Rust)	Files	Key Responsibility
<b>tui</b>	~45,000	120	UI, ChatWidget, Spec-Kit integration
<b>spec-kit</b> (lib)	~8,000	25	Config, retry, types (reusable)
<b>spec-kit</b> (tui integration)	~35,000	55	Command handlers, consensus, pipeline
<b>core</b>	~30,000	85	Backend services, MCP, DB, config
<b>mcp-client</b>	~3,500	12	Async MCP client
<b>mcp-server</b>	~5,000	18	MCP server implementation
<b>protocol</b>	~12,000	35	OpenAI protocol client
<b>Others</b>	~88,000	~230	Utilities, tooling, sandbox

**Total:** ~226,607 LOC across 538 files

---

## Adding New Crates

### When to Create a New Crate

**Good reasons:** - ✓ Reusable component (can extract to standalone library) - ✓ Clear responsibility boundary - ✓ Independent compilation (speeds up builds) - ✓ Different dependency requirements - ✓ Potential for external use (spec-kit library)

**Bad reasons:** - ✗ Small utility module (use common instead) - ✗ Tightly coupled to single crate - ✗ No clear abstraction boundary

---

### Creating a New Crate

```

cd codex-rs

# Create new crate in workspace
cargo new --lib my-new-crate

# Or with specific edition
cargo new --lib --edition 2024 my-new-crate

```

**Add to workspace** (Cargo.toml):

```
[workspace]
members = [
    "ansi-escape",
    # ... existing crates
    "my-new-crate", # Add here
]
```

**Add workspace dependency:**

```
[workspace.dependencies]
my-new-crate = { path = "my-new-crate" }
```

**Use in other crates:**

```
# In another crate's Cargo.toml
[dependencies]
my-new-crate = { workspace = true }
```

---

## Next Steps

- [TUI Architecture](#) - Detailed TUI and ChatWidget design
  - [Core Execution](#) - Agent orchestration and providers
  - [MCP Integration](#) - Native MCP client details
  - [Database Layer](#) - SQLite optimization
- 

**File References:** - Workspace: codex-rs/Cargo.toml:1-234 - Build profiles: codex-rs/Cargo.toml:201-234 - TUI: codex-rs/tui/Cargo.toml - Spec-Kit: codex-rs/spec-kit/Cargo.toml - Core: codex-rs/core/Cargo.toml - MCP Client: codex-rs/mcp-client/Cargo.toml

---

## Configuration System

Layered configuration with hot-reload and 5-tier precedence.

---

### Overview

The configuration system implements the **12-factor app pattern**: 1. **Defaults**: Built-in fallback values (in code) 2. **Config file**: `~/.code/config.toml` 3. **Environment variables**: `OPENAI_API_KEY`, etc. 4. **Profiles**: Named configurations (`[profiles.premium]`) 5. **CLI flags**: `-model gpt-5, --config key=value`

**Hot-Reload**: File changes detected and applied without restart (<100ms latency)

**Location**: codex-rs/spec-kit/src/config/

---

### 5-Tier Precedence

## Precedence Order

**Highest to Lowest** (higher overrides lower):

1. **CLI Flags** (highest priority)

```
code --model gpt-5 --config approval_policy=never
```

2. **Shell Environment**

```
export OPENAI_API_KEY="sk-proj-..."  
export CODEX_HOME="/custom/path"
```

3. **Profile** (selected via --profile or profile = "name")

```
[profiles.premium]  
model = "o3"  
model_reasoning_effort = "high"
```

4. **Config File** (~/.code/config.toml)

```
model = "gpt-5"  
approval_policy = "on_request"
```

5. **Defaults** (lowest priority, built-in)

```
impl Default for Config {  
    fn default() -> Self {  
        Self {  
            model: "gpt-5-codex".to_string(),  
            approval_policy: ApprovalPolicy::OnRequest,  
            ...  
        }  
    }  
}
```

---

## Precedence Examples

### Example 1: Model selection

```
# Default: gpt-5-codex  
# config.toml: model = "gpt-5"  
# --model o3  
  
Effective value: "o3" (CLI flag wins)
```

### Example 2: API key

```
# Default: none  
# config.toml: (none)  
# export OPENAI_API_KEY="sk-proj-123"  
# --config model_provider.openai.api_key="sk-proj-456"  
  
Effective value: "sk-proj-456" (CLI flag wins)
```

### Example 3: Profile

```
# Default: approval_policy = "on_request"  
# config.toml: approval_policy = "untrusted"  
# [profiles.premium]: approval_policy = "never"  
# --profile premium
```

Effective value: "never" (profile wins over config.toml)

---

## Configuration Loading

### Loader Architecture

**Location:** codex-rs/spec-kit/src/config/loader.rs

```
pub struct ConfigLoader {
    defaults: AppConfig,
    file_path: Option<PathBuf>,
    env_overrides: HashMap<String, String>,
    cli_overrides: HashMap<String, String>,
}

impl ConfigLoader {
    pub fn new() -> Self {
        Self {
            defaults: AppConfig::default(),
            file_path: None,
            env_overrides: HashMap::new(),
            cli_overrides: HashMap::new(),
        }
    }

    pub fn with_file(mut self, path: impl Into<PathBuf>) -> Self {
        self.file_path = Some(path.into());
        self
    }

    pub fn with_env_overrides(mut self, overrides: HashMap<String, String>) -> Self {
        self.env_overrides = overrides;
        self
    }

    pub fn with_cli_overrides(mut self, overrides: HashMap<String, String>) -> Self {
        self.cli_overrides = overrides;
        self
    }

    pub fn load(self) -> Result<AppConfig> {
        // Layer 1: Defaults
        let mut config = self.defaults;

        // Layer 2: Config file
        if let Some(path) = self.file_path {
            if path.exists() {
                let file_config: AppConfig =
                    toml::from_str(&std::fs::read_to_string(path)?);
                config.merge(file_config);
            }
        }

        // Layer 3: Environment variables
        for (key, value) in self.env_overrides {
            config.set(key, value);
        }
    }
}
```

```

        config.set_from_string(&key, &value)?;
    }

    // Layer 4: CLI overrides
    for (key, value) in self.cli_overrides {
        config.set_from_string(&key, &value)?;
    }

    // Layer 5: Profile (if selected)
    if let Some(profile_name) = &config.profile {
        if let Some(profile) = config.profiles.get(profile_name) {
            config.merge(profile.clone());
        }
    }

    config.validate()?;
    Ok(config)
}

```

---

## Environment Variable Mapping

**Pattern:** SPECKIT\_<SECTION>\_<KEY>

**Examples:**

```

# Model configuration
SPECKIT_MODEL_NAME="gpt-5"
SPECKIT_MODEL_REASONING EFFORT="high"

# Quality gates
SPECKIT_QUALITY_GATES_PLAN="gemini,claude,code"
SPECKIT_QUALITY_GATES_TASKS="gemini"

# Evidence configuration
SPECKIT_EVIDENCE_BASE_DIR="/custom/evidence"
SPECKIT_EVIDENCE_MAX_SIZE_MB="50"

```

**Parsing:**

```

fn parse_env_key(key: &str) -> Option<(String, String)> {
    if let Some(stripped) = key.strip_prefix("SPECKIT_") {
        let parts: Vec<&str> = stripped.split('_').collect();
        if parts.len() >= 2 {
            let section =
parts[0..parts.len()-1].join("_").to_lowercase();
            let key = parts[parts.len()-1].to_lowercase();
            return Some((section, key));
        }
    }
    None
}

```

---

## Hot-Reload System

### Architecture

```

File Change → notify crate → Debouncer (300ms) → Validate → Lock →
Replace → Event
                                ↓ Fail
                                Preserve Old Config

```

**Location:** codex-rs/spec-kit/src/config/hot\_reload.rs:1-100

---

## HotReloadWatcher

```

use notify::{Watcher, RecursiveMode, Event};
use notify_debouncer_full::{new_debouncer, Debouncer, FileIdMap};
use std::sync::{Arc, RwLock, Mutex};
use tokio::sync::mpsc;

pub enum ConfigReloadEvent {
    FileChanged(PathBuf),           // File change detected (pre-validation)
    ReloadSuccess,                 // Config reloaded successfully
    ReloadFailed(String),          // Validation error (old config preserved)
}

pub struct HotReloadWatcher {
    debouncer: Arc<Mutex<Debouncer<...>>,
    config: Arc<RwLock<AppConfig>>,
    event_tx: mpsc::Sender<ConfigReloadEvent>,
}

impl HotReloadWatcher {
    pub async fn new(
        config_path: impl Into<PathBuf>,
        debounce_duration: Duration,
    ) -> Result<Self> {
        let config_path = config_path.into();
        let config = Arc::new(RwLock::new(ConfigLoader::new()
            .with_file(&config_path)
            .load()?));

        let (event_tx, _event_rx) = mpsc::channel(32);

        // Create debouncer (buffers events for 300ms)
        let file_id_map = FileIdMap::new();
        let config_clone = Arc::clone(&config);
        let event_tx_clone = event_tx.clone();
        let path_clone = config_path.clone();

        let debouncer = new_debouncer(
            debounce_duration,
            Some(Duration::from_secs(1)), // Tick interval
            move |events: DebounceEventResult| {
                match events {
                    Ok(events) => {
                        for event in events {
                            if event.paths.contains(&path_clone) {
                                // Config file changed, attempt
                                reload
                                match reload_config(&config_clone,
                                    &path_clone) {
                                    Ok(_) => {

```

```

        let _ =
event_tx_clone.try_send(
    ConfigReloadEvent::ReloadSuccess
        );
    },
    Err(e) => {
        let _ =
event_tx_clone.try_send(
    ConfigReloadEvent::ReloadFailed(e.to_string())
        );
    },
}
},
Err(e) => {
    eprintln!("Watch error: {:?}", e);
},
},
),
)?;

// Watch config file
debouncer.watcher().watch(&config_path,
RecursiveMode::NonRecursive)?;

Ok(Self {
    debouncer: Arc::new(Mutex::new(debouncer)),
    config,
    event_tx,
})
}

pub fn get_config(&self) -> Arc<AppConfig> {
    // Read lock held briefly (<1μs) to clone Arc
    Arc::clone(&*self.config.read().unwrap())
}

fn reload_config(
    config: &Arc<RwLock<AppConfig>>,
    path: &Path,
) -> Result<()> {
    // Parse and validate new config
    let new_config = ConfigLoader::new()
        .with_file(path)
        .load()?;
    // Atomic write lock (<1ms)
    *config.write().unwrap() = new_config;
    Ok(())
}

```

---

## Debouncing

**Purpose:** Prevent reload storms (e.g., text editor save creates multiple events)

**Configuration:** 300ms debounce window

**Behavior:**

```
t=0ms:    File change event 1
t=50ms:   File change event 2 (reset timer)
t=100ms:  File change event 3 (reset timer)
t=400ms:  No more events for 300ms → Trigger reload
```

**Result:** Only one reload despite multiple filesystem events

---

## Lock Contention

**Read Locks** (frequent, fast):

```
let config = watcher.get_config(); // Arc::clone, <1μs
```

**Write Locks** (rare, fast):

```
*config.write().unwrap() = new_config; // <1ms
```

**Performance:** - Read locks: Non-blocking (RwLock allows concurrent readers) - Write lock: Blocks readers briefly (<1ms) - Reload frequency: Low (manual file edits only)

**CPU overhead:** <0.5% (idle, file watching)

---

## Performance Metrics

### Hot-Reload Latency

**Measured end-to-end:**

```
File save → Filesystem event → Debounce wait → Parse TOML → Validate
→ Write lock → Event
  0ms           ~10ms          300ms        ~20ms       ~5ms
<1ms         ~1ms
```

Total: ~336ms (p50)
~420ms (p95)

**Acceptable:** Sub-second reload for manual config edits

---

### Lock Timing

**Read lock** (get\_config):

```
Acquire read lock: <1μs
Clone Arc: <100ns
Release read lock: <100ns
```

Total: <1μs

### Write lock (reload\_config):

```
Acquire write lock: <500μs (wait for readers to finish)
Replace config: <100ns
Release write lock: <100ns

Total: <1ms
```

---

## Validation

### Schema Validation

```
impl AppConfig {
    pub fn validate(&self) -> Result<()> {
        // Model provider must exist
        if self.model_providers.is_empty() {
            return Err(Error::NoModelProviders);
        }

        // Selected provider must be configured
        if !self.model_providers.contains_key(&self.model_provider) {
            return Err(Error::ProviderNotFound(self.model_provider.clone()));
        }

        // Quality gate agents must be valid
        for agents in self.quality_gates.values() {
            for agent in agents {
                if !self.agents.iter().any(|a| a.canonical_name == *agent) {
                    return Err(Error::AgentNotFound(agent.clone()));
                }
            }
        }

        // Evidence max size must be reasonable
        if self.evidence.max_size_mb > 1000 {
            return Err(Error::EvidenceSizeTooLarge(self.evidence.max_size_mb));
        }

        Ok(())
    }
}
```

**On validation failure:** Preserve old config, emit ReloadFailed event

---

### Type Safety

TOML parsing uses serde:

```
#[derive(Deserialize, Serialize, Clone)]
pub struct AppConfig {
    pub model: String,
    pub model_provider: String,
```

```

    pub approval_policy: ApprovalPolicy, // Enum (type-safe)
    pub quality_gates: QualityGateConfig,
    pub evidence: EvidenceConfig,
    // ... 20+ fields
}

#[derive(Deserialize, Serialize, Clone)]
#[serde(rename_all = "snake_case")]
pub enum ApprovalPolicy {
    Untrusted,
    OnFailure,
    OnRequest,
    Never,
}

```

**Benefits:** - Compile-time type checking - Automatic deserialization - Invalid values rejected at parse time

---

## Profile System

### Profile Definition

```

# ~/.code/config.toml

# Default configuration
model = "gpt-5"
approval_policy = "on_request"

# Profile for premium reasoning
[profiles.premium]
model = "o3"
model_reasoning_effort = "high"
model_reasoning_summary = "detailed"
approval_policy = "never"

# Profile for fast iteration
[profiles.fast]
model = "gpt-4o-mini"
model_reasoning_effort = "low"
approval_policy = "never"

# Profile for automation/CI
[profiles.ci]
model = "gpt-4o"
approval_policy = "never"
sandbox_mode = "read-only"

```

---

### Profile Selection

#### Via config:

```
profile = "premium" # Active profile
```

#### Via CLI:

```
code --profile premium "complex task"
code --profile fast "simple formatting"
```

```
code --profile ci "generate report"
```

**Precedence:** CLI --profile > config profile field > no profile

---

## Profile Merging

```
impl AppConfig {
    pub fn merge(&mut self, other: AppConfig) {
        // Non-Option fields: other wins
        self.model = other.model;
        self.approval_policy = other.approval_policy;

        // Option fields: other overwrites if Some
        if other.model_reasoning_effort.is_some() {
            self.model_reasoning_effort =
                other.model_reasoning_effort;
        }

        // Collections: extend (not replace)
        self.quality_gates.extend(other.quality_gates);
        self.agents.extend(other.agents);
    }
}
```

---

## Registry System

### Command Registry (Spec-Kit)

**Location:** codex-rs/tui/src/chatwidget/spec\_kit/command\_registry.rs

```
pub trait SpecKitCommand: Send + Sync {
    fn name(&self) -> &'static str;
    fn aliases(&self) -> &[&'static str];
    fn description(&self) -> &'static str;
    fn execute(&self, widget: &mut ChatWidget, args: String);
}

pub struct CommandRegistry {
    commands: HashMap<String, Box<dyn SpecKitCommand>>,
    by_alias: HashMap<String, String>, // Alias → primary name
}

impl CommandRegistry {
    pub fn register(&mut self, command: Box<dyn SpecKitCommand>) {
        let name = command.name().to_string();

        // Register primary name
        self.commands.insert(name.clone(), command);

        // Register aliases
        for alias in command.aliases() {
            self.by_alias.insert(alias.to_string(), name.clone());
        }
    }

    pub fn find(&self, name: &str) -> Option<&dyn SpecKitCommand> {
```

```

    // Resolve alias → primary name
    let resolved_name = self.by_alias.get(name).unwrap_or(name);

    // Find command
    self.commands.get(resolved_name).map(|b| &*b)
}
}

```

**Benefits:** - Dynamic dispatch (no enum growth) - Alias support (backward compatibility) - Decoupled from upstream SlashCommand enum

---

## Summary

### Configuration System Highlights:

1. **5-Tier Precedence**: CLI > Shell > Profile > TOML > Defaults
2. **Hot-Reload**: <100ms latency (p50), <0.5% CPU overhead
3. **Debouncing**: 300ms window prevents reload storms
4. **Lock Performance**: <1µs read locks, <1ms write locks
5. **Validation**: Schema validation preserves old config on error
6. **Type Safety**: Serde deserialization with enum validation
7. **Profile System**: Named configurations for different workflows
8. **Registry Pattern**: Dynamic command dispatch (Spec-Kit)

**Architecture:** - **notify crate**: Filesystem watching - **Arc**: Atomic config updates - **Debouncer**: Event buffering - **ConfigLoader**: Layered loading with validation

---

**File References:** - Loader: codex-rs/spec-kit/src/config/loader.rs - Hot-reload: codex-rs/spec-kit/src/config/hot\_reload.rs:1-100 - Validation: codex-rs/spec-kit/src/config/validator.rs - Registry: codex-rs/tui/src/chatwidget/spec\_kit/command\_registry.rs

---

## Core Execution System

Agent orchestration, model providers, and execution management.

---

## Overview

The Core Execution system manages:

- **Agent Lifecycle**: Spawning, tracking, cleanup
- **Model Providers**: OpenAI, Anthropic, Google integration
- **Conversation Management**: Request/response handling
- **Retry Logic**: Exponential backoff for failures
- **Timeout Management**: Per-operation deadlines

**Location:** codex-rs/core/src/

---

## Agent Orchestration

## ConversationManager

**Purpose:** Central hub for agent conversation lifecycle

**Location:** codex-rs/core/src/conversation\_manager.rs

```
pub struct ConversationManager {
    conversations: Arc<Mutex<HashMap<ConversationId,
Conversation>>>,
    provider_clients: Arc<ModelProviderClients>,
    config: Arc<RwLock<Config>>,
}

impl ConversationManager {
    pub async fn new_conversation(&self, config: Config) ->
Result<NewConversation> {
        let conversation = Conversation::new(
            config,
            self.provider_clients.clone(),
        ).await?;

        Ok(NewConversation { conversation, id })
    }
}
```

**Responsibilities:** - Create new conversations - Manage conversation lifecycle - Coordinate with model providers - Handle configuration updates

---

## Agent Spawning Pattern

**From TUI:** codex-rs/tui/src/chatwidget/agent.rs:16-62

```
pub(crate) fn spawn_agent(
    config: Config,
    app_event_tx: AppEventSender,
    server: Arc<ConversationManager>,
) -> UnboundedSender<Op> {
    let (codex_op_tx, mut codex_op_rx) = unbounded_channel::<Op>();

    tokio::spawn(async move {
        // Create conversation
        let conversation = server.new_conversation(config).await?;

        // Operation processor
        tokio::spawn(async move {
            while let Some(op) = codex_op_rx.recv().await {
                conversation.submit(op).await;
            }
        });
    });

    // Event forwarder
    while let Ok(event) = conversation.next_event().await {
        app_event_tx.send(AppEvent::CodexEvent(event))?;
    }
);

    codex_op_tx
}
```

**Key Points:** - **Async task:** Runs on Tokio runtime - **Channel-based:** UnboundedSender for sync → async bridge - **Concurrent processing:** Separate op handler and event forwarder - **Automatic cleanup:** Tasks exit when conversation ends

---

## Multi-Agent Orchestration (Spec-Kit)

**Location:** codex-  
rs/tui/src/chatwidget/spec\_kit/agent\_orchestrator.rs

```
pub struct AgentOrchestrator {
    active_agents: HashMap<String, AgentHandle>,
    agent_configs: Vec<AgentConfig>,
}

pub struct AgentHandle {
    agent_id: String,
    agent_name: String,
    process: Child, // Subprocess handle
    stdout_reader: tokio::task::JoinHandle<()>,
    stderr_reader: tokio::task::JoinHandle<()>,
}

impl AgentOrchestrator {
    pub async fn spawn_agents(
        &mut self,
        spec_id: &str,
        stage: SpecStage,
        prompt: &str,
    ) -> Result<Vec<String>> {
        let mut agent_ids = Vec::new();

        for agent_config in &self.agent_configs {
            let agent_id = format!("{}-{}-{}", spec_id, stage, agent_config.name);

            // Spawn subprocess
            let mut child = Command::new(&agent_config.command)
                .args(&agent_config.args)
                .stdin(Stdio::piped())
                .stdout(Stdio::piped())
                .stderr(Stdio::piped())
                .spawn()?;

            // Read stdout/stderr asynchronously
            let stdout_reader =
                tokio::spawn(read_stream(child.stdout.take()));
            let stderr_reader =
                tokio::spawn(read_stream(child.stderr.take()));

            let handle = AgentHandle {
                agent_id: agent_id.clone(),
                agent_name: agent_config.name.clone(),
                process: child,
                stdout_reader,
                stderr_reader,
            };

            self.active_agents.insert(agent_id.clone(), handle);
        }
    }
}
```

```

        agent_ids.push(agent_id);
    }

    Ok(agent_ids)
}

pub async fn wait_for_completion(&mut self, timeout: Duration) -> Result<Vec<AgentOutput>> {
    let deadline = Instant::now() + timeout;
    let mut outputs = Vec::new();

    for (agent_id, handle) in self.active_agents.drain() {
        let remaining =
            deadline.saturating_duration_since(Instant::now());

        match tokio::time::timeout(remaining,
handle.process.wait()).await {
            Ok(Ok(status)) => {
                let stdout = handle.stdout_reader.await?;
                outputs.push(AgentOutput {
                    agent_id,
                    agent_name: handle.agent_name,
                    stdout,
                    exit_code: status.code(),
                });
            },
            Ok(Err(e)) => { /* process error */ },
            Err(_) => { /* timeout */ },
        }
    }
}

Ok(outputs)
}
}

```

**Features:** - **Concurrent spawning:** Launch multiple agents in parallel - **Timeout enforcement:** Per-agent deadlines - **Stream capture:** Async stdout/stderr reading - **Graceful cleanup:** Kill on timeout or error

---

## Model Providers

### Provider Architecture

```

Application
  ↓
ModelProviderClients (registry)
    ↳ OpenAIProvider (Responses API)
    ↳ AnthropicProvider (CLI subprocess)
    ↳ GoogleProvider (CLI subprocess)

```

---

### OpenAI Provider

**Location:** codex-*rs*/protocol/src/openai\_client.rs

```
pub struct OpenAIClient {
```

```

        base_url: String,
        api_key: String,
        http_client: reqwest::Client,
        retry_config: RetryConfig,
    }

impl OpenAIClient {
    pub async fn chat_completion(
        &self,
        request: ChatCompletionRequest,
    ) -> Result<impl Stream<Item = Result<ChatCompletionChunk>>> {
        let url = format!("{}chat/completions", self.base_url);

        let response = self.http_client
            .post(&url)
            .header("Authorization", format!("Bearer {}", self.api_key))
            .json(&request)
            .send()
            .await?;

        // Server-Sent Events stream
        let stream = response
            .bytes_stream()
            .events_source()
            .map(|event| parse_sse_event(event));

        Ok(stream)
    }

    pub async fn responses_api(
        &self,
        request: ResponsesRequest,
    ) -> Result<impl Stream<Item = Result<ResponseEvent>>> {
        // Similar but for Responses API
    }
}

```

**Features:** - **Streaming:** Server-Sent Events (SSE) - **Retry logic:** Exponential backoff on failures - **Rate limiting:** 429 response handling - **Timeout:** Per-request deadlines

---

## Anthropic Provider (CLI)

**Location:** codex-  
rs/tui/src/chatwidget/spec\_kit/agent\_orchestrator.rs

```

// Anthropic via CLI subprocess
let mut child = Command::new("claude")
    .arg("--model").arg("claude-sonnet-4-5")
    .arg(prompt)
    .env("ANTHROPIC_API_KEY", api_key)
    .stdout(Stdio::piped())
    .spawn()?;

```

```

let output = child.wait_with_output().await?;
let response = String::from_utf8(output.stdout)?;

```

**Note:** Uses CLI subprocess, not direct API (simpler integration for multi-agent)

---

## Google Provider (CLI)

```
// Google via CLI subprocess
let mut child = Command::new("gemini")
    .arg("-i").arg(prompt)
    .env("GOOGLE_API_KEY", api_key)
    .stdout(Stdio::piped())
    .spawn()?;

let output = child.wait_with_output().await?;
let response = String::from_utf8(output.stdout)?;
```

---

## Protocol Implementation

### Request/Response Types

**Location:** codex-rs/protocol/src/types.rs

```
// Chat Completions API
pub struct ChatCompletionRequest {
    pub model: String,
    pub messages: Vec<Message>,
    pub temperature: Option<f32>,
    pub max_tokens: Option<u32>,
    pub stream: bool,
}

pub struct Message {
    pub role: Role,      // system, user, assistant, tool
    pub content: String,
    pub name: Option<String>,
    pub tool_calls: Option<Vec<ToolCall>>,
}

// Responses API (newer)
pub struct ResponsesRequest {
    pub model: String,
    pub messages: Vec<Message>,
    pub reasoning: Option<ReasoningConfig>,
    pub response_format: Option<ResponseFormat>,
}

pub struct ReasoningConfig {
    pub effort: ReasoningEffort, // minimal, low, medium, high
    pub summary: SummaryLevel,   // none, auto, concise, detailed
}
```

---

### Streaming Events

```
pub enum ResponseEvent {
    Start { id: String },
    Token { content: String },
```

```

ReasoningToken { content: String },
ToolCall { call: ToolCall },
Complete { finish_reason: FinishReason },
Error { error: String },
}

```

### Processing:

```

while let Some(event) = stream.next().await {
    match event? {
        ResponseEvent::Token { content } => {
            // Forward to UI for rendering
            app_event_tx.send(AppEvent::Token(content))?;
        },
        ResponseEvent::ToolCall { call } => {
            // Execute tool and inject result
            let result = execute_tool(call).await?;
            conversation.submit(Op::ToolResponse(call.id,
result)).await?;
        },
        ResponseEvent::Complete { finish_reason } => {
            break;
        },
        _ => {}
    }
}

```

---

## Retry Logic

### Exponential Backoff

**Location:** codex-rs/spec-kit/src/retry/strategy.rs

```

pub struct RetryConfig {
    pub max_attempts: usize,
    pub initial_backoff_ms: u64,
    pub max_backoff_ms: u64,
    pub backoff_multiplier: f64,
    pub jitter_factor: f64,
}

impl Default for RetryConfig {
    fn default() -> Self {
        Self {
            max_attempts: 3,
            initial_backoff_ms: 100,
            max_backoff_ms: 10_000,
            backoff_multiplier: 2.0,
            jitter_factor: 0.5,
        }
    }
}

```

### Backoff Calculation:

Attempt 1: 100ms + jitter( $\pm 50\text{ms}$ )  
 Attempt 2: 200ms + jitter( $\pm 100\text{ms}$ )  
 Attempt 3: 400ms + jitter( $\pm 200\text{ms}$ )

```
Attempt 4: 800ms + jitter(±400ms)
Attempt 5: 1600ms + jitter(±800ms)
...
Max: 10,000ms (10s)
```

---

## Error Classification

```
pub trait RetryClassifiable {
    fn is_retryable(&self) -> bool;
}

impl RetryClassifiable for ApiError {
    fn is_retryable(&self) -> bool {
        match self {
            // Transient errors (retry)
            ApiError::RateLimitExceeded => true,
            ApiError::ServiceUnavailable => true,
            ApiError::Timeout => true,
            ApiError::NetworkError(_) => true,

            // Permanent errors (don't retry)
            ApiError::AuthenticationFailed => false,
            ApiError::InvalidRequest(_) => false,
            ApiError::InsufficientQuota => false,
            _ => false,
        }
    }
}
```

---

## Retry Execution

```
pub async fn execute_with_backoff<F, Fut, T, E>(
    mut operation: F,
    config: &RetryConfig,
) -> Result<T>
where
    F: FnMut() -> Fut,
    Fut: Future<Output = Result<T, E>>,
    E: Error + RetryClassifiable,
{
    let mut attempts = 0;
    let mut backoff_ms = config.initial_backoff_ms;

    loop {
        attempts += 1;

        match operation().await {
            Ok(value) => return Ok(value),
            Err(err) if !err.is_retryable() => {
                return Err(RetryError::PermanentError(err.to_string()));
            },
            Err(_) if attempts >= config.max_attempts => {
                return Err(RetryError::MaxAttemptsExceeded(attempts));
            },
        }
    }
}
```

```
    Err(_) => {
        // Calculate jittered backoff
        let jitter = (rand::random::f64() - 0.5) * 2.0 *
config.jitter_factor;
                    let delay_ms = (backoff_ms as f64 * (1.0 + jitter))
as u64;

tokio::time::sleep(Duration::from_millis(delay_ms)).await;

        // Exponential increase
        backoff_ms = (backoff_ms as f64 *
config.backoff_multiplier) as u64;
                    backoff_ms = backoff_ms.min(config.max_backoff_ms);
    }
}
}
```

## **Timeout Management**

## Per-Operation Timeouts

```
pub async fn execute_with_timeout<F, T>(
    operation: F,
    timeout: Duration,
) -> Result<T>
where
    F: Future<Output = Result<T>>,
{
    match tokio::time::timeout(timeout, operation).await {
        Ok(Ok(value)) => Ok(value),
        Ok(Err(e)) => Err(e),
        Err(_) => Err(Error::Timeout),
    }
}
```

## Usage:

```
// API request with 30s timeout
let response = execute_with_timeout(
  openai_client.chat_completion(request),
  Duration::from_secs(30),
).await?;

// Agent execution with 5min timeout
let outputs = execute_with_timeout(
  agent_orchestrator.wait_for_completion(),
  Duration::from_secs(300),
).await?;
```

## Configurable Timeouts

## From config.toml:

## [mcp servers.filesystem]

```

startup_timeout_sec = 10    # Server initialization
tool_timeout_sec = 60       # Per-tool execution

[model_providers.openai]
stream_idle_timeout_ms = 300000 # 5min idle timeout

```

---

## Tool Execution

### Tool Call Flow

1. Model returns tool\_use event  
↳ { "type": "tool\_use", "name": "filesystem\_\_read\_file", "arguments": {...} }
  2. Conversation extracts tool call  
↳ ToolCall { id, name, arguments }
  3. Submit to MCP manager  
↳ mcp\_manager.invoke\_tool(name, arguments).await
  4. MCP client executes  
↳ Server subprocess processes request
  5. Result returned  
↳ ToolResult { id, content: "file contents..." }
  6. Inject into conversation  
↳ conversation.submit(Op::ToolResponse(id, result)).await
  7. Model continues with tool result
- 

### Sandbox Enforcement

```

pub enum SandboxMode {
    ReadOnly,           // No writes, no network
    WorkspaceWrite,     // Write to workspace, no network
    DangerFullAccess,   // Full access (use in Docker)
}

pub fn execute_sandboxed(
    command: &str,
    args: &[&str],
    sandbox: SandboxMode,
) -> Result<Output> {
    match sandbox {
        SandboxMode::ReadOnly => {
            // Landlock: deny all writes
            apply_landlock_READONLY()?;
        },
        SandboxMode::WorkspaceWrite => {
            // Landlock: allow workspace writes only
            apply_landlock_workspace(workspace_path)?;
        },
        SandboxMode::DangerFullAccess => {
            // No sandboxing
        },
    }
}

```

```

    }

    // Execute command
    Command::new(command)
        .args(args)
        .output()
    }
}

```

**File:** codex-rs/linux-sandbox/src/lib.rs

---

## Performance Optimizations

### Connection Pooling

**HTTP Client:**

```

let http_client = reqwest::Client::builder()
    .pool_max_idle_per_host(10)      // Reuse connections
    .timeout(Duration::from_secs(30))
    .build()?;

```

**Benefits:** - Reuse TCP connections (avoid handshake overhead) -  
Reduce latency for subsequent requests

---

### Concurrent Execution

**Multi-Agent:**

```

// Spawn all agents concurrently
let mut join_set = JoinSet::new();
for agent_config in agent_configs {
    join_set.spawn(spawn_agent(agent_config, prompt.clone()));
}

// Wait for all to complete
let mut outputs = Vec::new();
while let Some(result) = join_set.join_next().await {
    outputs.push(result??);
}

```

**Performance:** 3 agents finish in ~10s instead of ~30s (3× speedup)

---

## Error Handling

### Error Hierarchy

```

pub enum Error {
    // Network errors (retryable)
    NetworkError(reqwest::Error),
    Timeout,

    // API errors (some retryable)
    RateLimitExceeded,
    ServiceUnavailable,
}

```

```

        AuthenticationFailed,           // Not retryable
        InvalidRequest(String),       // Not retryable

        // Agent errors
        AgentSpawnFailed(io::Error),
        AgentTimeout,
        AgentCrashed(i32),           // Exit code

        // Tool errors
        ToolNotFound(String),
        ToolExecutionFailed(String),
    }

```

---

## Graceful Degradation

### Multi-Agent Consensus:

```

// If 1/3 agents fail, continue with 2/3
let outputs = agent_orchestrator.wait_for_completion().await?;

if outputs.len() >= 2 {
    // Consensus still valid with 2/3 agents
    let synthesis = synthesize_consensus(&outputs);
    return Ok((synthesis, degraded: true));
} else {
    return Err(Error::InsufficientAgents);
}

```

---

## Summary

### Core Execution Highlights:

- Agent Orchestration:** Async task spawning with channel-based communication
- Multi-Provider:** OpenAI (HTTP), Anthropic/Google (CLI subprocess)
- Streaming:** Real-time token delivery via SSE
- Retry Logic:** Exponential backoff with jitter (100ms → 10s)
- Timeout Management:** Per-operation deadlines
- Tool Execution:** MCP integration with sandbox enforcement
- Error Handling:** Permanent vs transient classification
- Performance:** Connection pooling, concurrent agent execution

**Next Steps:** - [MCP Integration](#) - Native client details - [Database Layer](#) - SQLite optimization - [Configuration System](#) - Hot-reload

---

**File References:** - Conversation manager: codex-rs/core/src/conversation\_manager.rs - Agent spawner: codex-rs/tui/src/chatwidget/agent.rs:16-62 - Agent orchestrator: codex-rs/tui/src/chatwidget/spec\_kit/agent\_orchestrator.rs - OpenAI client: codex-rs/protocol/src/openai\_client.rs - Retry logic: codex-rs/spec-kit/src/retry/strategy.rs - Sandbox: codex-rs/linux-sandbox/src/lib.rs

---

# Database Layer

SQLite storage with optimized performance for consensus artifacts.


---

## Overview

**Database:** SQLite (embedded, ACID transactions) **Primary Use:** Consensus artifact storage (agent outputs, synthesis) **Performance:** **6.6x read speedup, 2.3x write speedup** (via optimizations)  
**Location:** codex-rs/core/src/db/, codex-rs/tui/src/chatwidget/spec\_kit/consensus\_db.rs

---

## Architecture



## Connection Pooling

### R2D2 Configuration

**Location:** codex-rs/core/src/db/connection.rs:39-105

```
use r2d2::{Pool, PooledConnection};
use r2d2_sqlite::SqliteConnectionManager;
use rusqlite::Connection;

pub fn initialize_pool(
    db_path: &Path,
    pool_size: u32,
) -> Result<Pool<SqliteConnectionManager>> {
    let manager = SqliteConnectionManager::file(db_path);

    let pool = Pool::builder()
        .max_size(pool_size) // Max connections (default:
8)
        .min_idle(Some(2)) // Keep 2 connections warm
        .connection_customizer(Box::new(ConnectionCustomizer))
        .test_on_check_out(true) // Health check before use
        .build(manager)?;

    // Verify pragmas on first connection
    let conn = pool.get()?;
}
```

```

    verify_pragmas(&conn);  

    Ok(pool)  

}

```

**Benefits:** - **Connection reuse:** Eliminate open/close overhead - **Concurrency:** Multiple connections for parallel access - **Health checks:** Detect broken connections before use - **Warm pool:** Keep minimum connections ready

---

## Connection Customizer (Pragma Optimization)

```

struct ConnectionCustomizer;  

impl r2d2::CustomizeConnection<Connection, rusqlite::Error> for  

ConnectionCustomizer {
    fn on_acquire(&self, conn: &mut Connection) -> Result<(),  

rusqlite::Error> {
        conn.execute_batch(
            "PRAGMA journal_mode = WAL"; -- Write-Ahead  

Logging (6.6x read speedup)
            PRAGMA synchronous = NORMAL; -- 2-3x write  

speedup (safe with WAL)
            PRAGMA foreign_keys = ON; -- Referential  

integrity
            PRAGMA cache_size = -32000; -- 32MB page  

cache
            PRAGMA temp_store = MEMORY; -- In-memory  

temporary tables
            PRAGMA auto_vacuum = INCREMENTAL; -- Prevent  

unbounded growth
            PRAGMA mmap_size = 1073741824; -- 1GB memory-
mapped I/O
            PRAGMA busy_timeout = 5000; -- 5s deadlock  

wait
            "
        )
    }
}

```

### Pragma Explanations:

Pragma	Value	Effect
journal_mode	WAL	Write-Ahead Logging: Concurrent reads during writes
synchronous	NORMAL	Fewer fsync calls (safe with WAL)
foreign_keys	ON	Enforce foreign key constraints
cache_size	-32000	32MB in-memory page cache
temp_store	MEMORY	Temporary tables in RAM
auto_vacuum	INCREMENTAL	Gradual space reclamation
		Memory-mapped I/O

mmap_size	1GB	for reads
busy_timeout	5s	Retry on lock for 5 seconds

---

## Performance Impact

### Before optimizations:

Single read: 850µs  
 Single write: 2.1ms  
 100-read batch: 78ms

### After optimizations:

Single read: 129µs (6.6x faster)  
 Single write: 0.9ms (2.3x faster)  
 100-read batch: 12ms (6.5x faster)

**Total improvement:** 6.6x read, 2.3x write

**Benchmark:** codex-rs/core/tests/db\_benchmark.rs

---

## Schema

### consensus\_runs

**Purpose:** Track workflow execution metadata

```
CREATE TABLE IF NOT EXISTS consensus_runs (
    run_id INTEGER PRIMARY KEY AUTOINCREMENT,
    spec_id TEXT NOT NULL,
    stage TEXT NOT NULL,
    consensus_ok BOOLEAN NOT NULL,
    degraded BOOLEAN NOT NULL,
    synthesis_json TEXT,
    created_at TEXT NOT NULL DEFAULT CURRENT_TIMESTAMP,
    UNIQUE(spec_id, stage)
);

CREATE INDEX IF NOT EXISTS idx_consensus_runs_spec
ON consensus_runs(spec_id, stage);
```

**Columns:** - run\_id: Auto-increment primary key - spec\_id: SPEC identifier (e.g., "SPEC-KIT-065") - stage: Workflow stage ("plan", "tasks", "implement", etc.) - consensus\_ok: Consensus achieved (true/false) - degraded: Some agents failed (true/false) - synthesis\_json: Synthesized consensus result (JSON) - created\_at: Timestamp

**Constraint:** One run per (spec\_id, stage) pair (UPSERT semantics)

---

### agent\_outputs

**Purpose:** Store individual agent responses

```

CREATE TABLE IF NOT EXISTS agent_outputs (
    output_id INTEGER PRIMARY KEY AUTOINCREMENT,
    run_id INTEGER NOT NULL,
    agent_name TEXT NOT NULL,
    agent_version TEXT,
    content_json TEXT NOT NULL,
    response_text TEXT,
    created_at TEXT NOT NULL DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (run_id) REFERENCES consensus_runs(run_id) ON DELETE
CASCADE
);
CREATE INDEX IF NOT EXISTS idx_agent_outputs_run
ON agent_outputs(run_id);

```

**Columns:** - output\_id: Auto-increment primary key - run\_id: Foreign key to consensus\_runs - agent\_name: Agent identifier ("gemini", "claude", "code") - agent\_version: Model version ("gemini-flash-1.5", "claude-sonnet-4-5") - content\_json: Structured agent output (JSON) - response\_text: Raw response text - created\_at: Timestamp

**Relationship:** Many agent\_outputs per consensus\_run (1:N)

---

## consensus\_artifacts (legacy)

**Purpose:** Old schema (being phased out)

```

CREATE TABLE IF NOT EXISTS consensus_artifacts (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    spec_id TEXT NOT NULL,
    stage TEXT NOT NULL,
    agent_name TEXT NOT NULL,
    content_json TEXT NOT NULL,
    response_text TEXT,
    run_id TEXT,
    created_at TEXT NOT NULL DEFAULT CURRENT_TIMESTAMP
);

```

**Status:** Deprecated, replaced by consensus\_runs + agent\_outputs

---

## Transaction Handling

### Transaction Behavior

```

use rusqlite::TransactionBehavior;

pub enum TransactionBehavior {
    Deferred,      // Lock on first read/write
    Immediate,     // Write lock immediately
    Exclusive,     // Exclusive lock (blocks all)
}

```

**Recommendation:** Use IMMEDIATE for writes (avoid write-write deadlock)

---

## Transaction Helpers

**Location:** codex-rs/core/src/db/transactions.rs:40-119

```
pub fn execute_in_transaction<F, T>(
    conn: &mut Connection,
    behavior: TransactionBehavior,
    operation: F,
) -> Result<T>
where
    F: FnOnce(&Transaction) -> Result<T>,
{
    let tx = conn.transaction_with_behavior(behavior)?;

    match operation(&tx) {
        Ok(result) => {
            tx.commit()?;
            Ok(result)
        },
        Err(e) => {
            // Automatic rollback via Drop trait
            Err(e)
        },
    }
}
```

**Usage:**

```
execute_in_transaction(conn, TransactionBehavior::Immediate, |tx| {
    tx.execute("INSERT INTO consensus_runs (...) VALUES (?)",
    params)?;
    tx.execute("INSERT INTO agent_outputs (...) VALUES (?)",
    params)?;
    Ok(())
})?;
```

**ACID guarantees:** - **Atomicity:** All or nothing (rollback on error) -  
**Consistency:** Foreign keys enforced - **Isolation:** IMMEDIATE locks  
prevent conflicts - **Durability:** WAL ensures persistence

---

## Batch Operations

```
pub fn batch_insert<T>(
    conn: &mut Connection,
    _table: &str,
    _columns: &[&str],
    rows: &[T],
    bind_fn: impl Fn(&Transaction, &T) -> Result<()>,
) -> Result<usize> {
    execute_in_transaction(conn, TransactionBehavior::Immediate,
    |tx| {
        for row in rows {
            bind_fn(tx, row)?;
        }
        Ok(rows.len())
    })
}
```

**Performance:** 100 inserts in single transaction ~12ms (vs ~2s for 100 individual commits)

---

## Async Wrapper

**Problem:** SQLite is synchronous, Tokio is async

**Solution:** tokio::task::spawn\_blocking wrapper

**Location:** codex-rs/core/src/db/async\_wrapper.rs:69-150

```
pub async fn with_connection<F, T>(
    pool: &Pool<SqliteConnectionManager>,
    f: F,
) -> Result<T>
where
    F: FnOnce(&mut Connection) -> Result<T> + Send + 'static,
    T: Send + 'static,
{
    let pool = pool.clone();

    // Run synchronous SQLite code in blocking thread pool
    tokio::task::spawn_blocking(move || {
        let mut conn = pool.get()?;
        f(&mut conn)
    })
    .await?
}
```

**Usage:**

```
// Async function calling sync SQLite
pub async fn store_consensus(
    pool: &Pool<SqliteConnectionManager>,
    spec_id: &str,
    stage: &str,
    synthesis: &str,
) -> Result<i64> {
    with_connection(pool, move |conn| {
        execute_in_transaction(conn, TransactionBehavior::Immediate,
            |tx| {
                upsert_consensus_run(tx, spec_id, stage, synthesis)
            }
        }).await
}
```

**Key Points:** - Doesn't block Tokio runtime: Runs on separate thread pool - Connection pooling: Still benefits from R2D2 pool - Error propagation: Propagates SQLite errors to async context

---

## Consensus Database (Spec-Kit)

### ConsensusDb

**Location:** codex-rs/tui/src/chatwidget/spec\_kit/consensus\_db.rs:34-150

```

pub struct ConsensusDb {
    conn: Arc<Mutex<Connection>>, // Legacy
single connection
    pool: Option<Pool<SqliteConnectionManager>>, // New
pooled connections
}

impl ConsensusDb {
    pub fn init(db_path: &Path) -> Result<Self> {
        // Open single connection (legacy)
        let conn = Connection::open(db_path)?;

        // Create schema
        conn.execute(
            "CREATE TABLE IF NOT EXISTS agent_executions (
                agent_id TEXT PRIMARY KEY,
                spec_id TEXT NOT NULL,
                stage TEXT NOT NULL,
                phase_type TEXT NOT NULL,
                agent_name TEXT NOT NULL,
                run_id TEXT,
                spawned_at TEXT NOT NULL,
                completed_at TEXT,
                response_text TEXT,
                extraction_error TEXT
            )",
            [],
        )?;

        // Create indices
        conn.execute(
            "CREATE INDEX IF NOT EXISTS idx_agent_executions_spec
            ON agent_executions(spec_id, stage)",
            [],
        )?;

        // Initialize new schema pool (SPEC-945B)
        let pool = initialize_pool(db_path, 8)?;

        Ok(Self {
            conn: Arc::new(Mutex::new(conn)),
            pool: Some(pool),
        })
    }

    pub fn upsert_consensus_run(
        &self,
        spec_id: &str,
        stage: &str,
        consensus_ok: bool,
        degraded: bool,
        synthesis_json: Option<&str>,
    ) -> Result<i64> {
        let pool =
self.pool.as_ref().ok_or(Error::PoolNotInitialized)?;
        let conn = pool.get()?;

        execute_in_transaction(&mut conn,
TransactionBehavior::Immediate, |tx| {
            tx.execute(

```

```

        "INSERT INTO consensus_runs (spec_id, stage,
consensus_ok, degraded, synthesis_json)
VALUES (?1, ?2, ?3, ?4, ?5)
ON CONFLICT(spec_id, stage) DO UPDATE SET
    consensus_ok = excluded.consensus_ok,
    degraded = excluded.degraded,
    synthesis_json = excluded.synthesis_json,
    created_at = CURRENT_TIMESTAMP",
    params![spec_id, stage, consensus_ok, degraded,
synthesis_json],
)?:;

    let run_id = tx.last_insert_rowid();
    Ok(run_id)
}
}

pub fn insert_agent_output(
    &self,
    run_id: i64,
    agent_name: &str,
    agent_version: Option<&str>,
    content_json: &str,
    response_text: Option<&str>,
) -> Result<i64> {
    let pool =
self.pool.as_ref().ok_or(Error::PoolNotInitialized)?;
    let conn = pool.get()?;

    conn.execute(
        "INSERT INTO agent_outputs (run_id, agent_name,
agent_version, content_json, response_text)
VALUES (?1, ?2, ?3, ?4, ?5)",
        params![run_id, agent_name, agent_version, content_json,
response_text],
    )?;

    Ok(conn.last_insert_rowid())
}
}

```

---

## Dual-Schema Migration (SPEC-945B)

**Phase 1:** Old schema only (agent\_executions) **Phase 2:** Dual-write to both schemas (current) **Phase 3:** New schema only (consensus\_runs + agent\_outputs)

**Current state:** Dual-write active

```

// Write to both old and new schema
pub fn store_consensus(&self, spec_id: &str, stage: &str) ->
Result<()> {
    // Old schema (legacy)
    self.conn.lock().unwrap().execute(
        "INSERT INTO agent_executions (...) VALUES (?)",
        params![...],
    )?;

    // New schema (pooled)

```

```

    if let Some(pool) = &self.pool {
        self.upsert_consensus_run(spec_id, stage, true, false,
None)?;
    }

    Ok(())
}

```

---

## Auto-Vacuum Strategy

### Incremental Auto-Vacuum

**Purpose:** Prevent unbounded database growth

**Configuration:** PRAGMA auto\_vacuum = INCREMENTAL;

**How it works:** - Database tracks free pages internally - On PRAGMA incremental\_vacuum(N), reclaim N pages - No blocking full-vacuum required

**Usage:**

```

pub fn compact_database(conn: &Connection, max_pages: u32) ->
Result<()> {
    conn.execute(&format!("PRAGMA incremental_vacuum({})", max_pages),
[])?;
    Ok(())
}

```

**Result:** 99.95% size reduction after cleanup (multi-GB → few MB)

**Evidence:** docs/SPEC-OPS-004-integrated-coder-hooks/evidence/database-cleanup.log

---

## Retry Logic (Database Operations)

### Sync Retry

**Location:** codex-rs/spec-kit/src/retry/strategy.rs

```

pub fn execute_with_backoff_sync<F, T, E>(
    mut operation: F,
    config: &RetryConfig,
) -> Result<T>
where
    F: FnMut() -> Result<T, E>,
    E: Error + RetryClassifiable,
{
    let mut attempts = 0;
    let mut backoff_ms = config.initial_backoff_ms;

    loop {
        attempts += 1;

        match operation() {
            Ok(value) => return Ok(value),

```

```
    Err(err) if err.error_code() ==  
Some(rusqlite::ErrorCode::DatabaseBusy) => {  
    // Database locked, retry with backoff  
    if attempts >= config.max_attempts {  
        return  
    Err(RetryError::MaxAttemptsExceeded(attempts));  
    }  
  
    std::thread::sleep(Duration::from_millis(backoff_ms));  
    backoff_ms = (backoff_ms as f64 *  
config.backoff_multiplier) as u64;  
    backoff_ms = backoff_ms.min(config.max_backoff_ms);  
},  
    Err(err) => {  
        // Permanent error, don't retry  
        return  
    Err(RetryError::PermanentError(err.to_string()));  
},  
}  
}  
}
```

## Usage:

```
let result = execute_with_backoff_sync(
    || conn.execute("INSERT INTO ...", params),
    &RetryConfig::default(),
)?;
```

# Error Handling

## Database Errors

```
pub enum DbError {
    // Connection errors
    PoolExhausted,
    ConnectionFailed(r2d2::Error),

    // SQLite errors
    DatabaseBusy,           // Retryable
    DatabaseLocked,         // Retryable
    ConstraintViolation(String), // Not retryable
    SchemaError(String),    // Not retryable

    // Application errors
    InvalidSchema(String),
    MigrationFailed(String),
}
```

## Error Classification

```
impl RetryClassifiable for DbError {
    fn is_retryable(&self) -> bool {
        match self {
            DbError::DatabaseBusy => true,
```

```

        DbError::DatabaseLocked => true,
        DbError::PoolExhausted => true,

        DbError::ConstraintViolation(_) => false,
        DbError::SchemaError(_) => false,
        DbError::InvalidSchema(_) => false,
        _ => false,
    }
}
}

```

---

## Schema Migrations

### Migration System

**Location:** codex-rs/core/src/db/migrations.rs:9-87

```

pub const SCHEMA_VERSION: i32 = 2;

pub fn migrate_to_latest(conn: &mut Connection) -> Result<()> {
    let current_version = get_schema_version(conn)?;

    if current_version == SCHEMA_VERSION {
        return Ok(()); // Already up-to-date
    }

    let tx =
        conn.transaction_with_behavior(TransactionBehavior::Exclusive)?;

    for version in (current_version + 1)..=SCHEMA_VERSION {
        apply_migration(&tx, version)?;
    }

    // Update schema version
    tx.execute(&format!("PRAGMA user_version = {}", SCHEMA_VERSION),
    [])?;
    tx.commit()?;
}

Ok(())
}

fn get_schema_version(conn: &Connection) -> Result<i32> {
    let version: i32 = conn.query_row("PRAGMA user_version", [],
    |row| row.get(0))?;
    Ok(version)
}

fn apply_migration(conn: &Connection, version: i32) -> Result<()> {
    match version {
        1 => migration_v1(conn), // Create consensus_runs,
        agent_outputs
        2 => migration_v2(conn), // Add indices
        _ => Err(Error::UnknownMigrationVersion(version)),
    }
}

```

**Migration v1:**

```

fn migration_v1(conn: &Connection) -> Result<()> {
    conn.execute_batch(
        "CREATE TABLE IF NOT EXISTS consensus_runs (...);  

        CREATE TABLE IF NOT EXISTS agent_outputs (...);  

        CREATE INDEX IF NOT EXISTS idx_consensus_runs_spec ON  

        consensus_runs(spec_id, stage);  

        CREATE INDEX IF NOT EXISTS idx_agent_outputs_run ON  

        agent_outputs(run_id);"  

    )?  

    Ok(())
}

```

**Forward-only:** Migrations never rollback (destructive changes prohibited)

---

## Summary

### Database Layer Highlights:

1. **Performance:** 6.6× read speedup, 2.3× write speedup (WAL + pragmas)
2. **Connection Pooling:** R2D2 with 2-8 connections
3. **Async Wrapper:** `tokio::task::spawn_blocking` for Tokio integration
4. **ACID Transactions:** IMMEDIATE mode for write consistency
5. **Schema:** `consensus_runs` + `agent_outputs` (normalized)
6. **Auto-Vacuum:** Incremental strategy (99.95% size reduction)
7. **Retry Logic:** Exponential backoff for database busy errors
8. **Migrations:** Forward-only schema evolution

**Next Steps:** - Configuration System - Hot-reload and 5-tier precedence

---

**File References:** - Connection pool: codex-  
`rs/core/src/db/connection.rs:39-105` - Transactions: codex-  
`rs/core/src/db/transactions.rs:40-119` - Async wrapper: codex-  
`rs/core/src/db/async_wrapper.rs:69-150` - Consensus DB: codex-  
`rs/tui/src/chatwidget/spec_kit/consensus_db.rs:34-150` - Migrations:  
`codex-  
 rs/core/src/db/migrations.rs:9-87` - Retry logic: codex-  
`rs/spec-kit/src/retry/strategy.rs`

---

## MCP Integration

Native Model Context Protocol client and server integration.

---

## Overview

**MCP (Model Context Protocol)** allows AI models to access external tools:  
 - File operations - Database queries - API calls - Custom tools

**Native Implementation:** 5.3× faster than subprocess approach  
 (8.7ms typical vs 46ms)

**Location:** codex-rs/mcp-client/, codex-rs/mcp-types/, codex-rs/core/src/mcp\_connection\_manager.rs

---

## Architecture

```
Application (TUI/Core)
  ↓
McpConnectionManager (singleton, shared across widgets)
  ↳ McpClient("filesystem")
    ↳ Writer Task (outgoing → stdin)
    ↳ Reader Task (stdout → pending HashMap)
      ↳ Subprocess (MCP server)
  ↳ McpClient("git-status")
  ↳ McpClient("local-memory")

MCP Servers (subprocesses)
  ↳ @modelcontextprotocol/server-filesystem
  ↳ @modelcontextprotocol/server-git-status
  ↳ @modelcontextprotocol/server-local-memory
```

---

## McpClient Implementation

### Core Structure

**Location:** codex-rs/mcp-client/src/mcp\_client.rs:63-150

```
pub struct McpClient {
    child: tokio::process::Child,                                     // Subprocess
handle
    outgoing_tx: mpsc::Sender<JSONRPCMessage>,                  // Send requests
(128 capacity)
    pending: Arc<Mutex<HashMap<i64, PendingSender>>>, // Request ID
→ response oneshot
    id_counter: AtomicI64,                                         // Monotonic ID
generator
}

type PendingSender = oneshot::Sender<JSONRPCMessage>;
```

```
impl McpClient {
    pub async fn new_stdio_client(
        program: OsString,
        args: Vec<OsString>,
        env: Option<HashMap<String, String>>,
    ) -> io::Result<Self> {
        // Spawn MCP server subprocess
        let mut child = Command::new(program)
            .args(args)
            .env_clear()
            .envs(env.unwrap_or_default())
            .stdin(Stdio::piped())
            .stdout(Stdio::piped())
            .stderr(Stdio::inherit())
            .kill_on_drop(true)
            .spawn()?;
    }
}
```

```

let stdin = child.stdin.take().unwrap();
let stdout = child.stdout.take().unwrap();

let (outgoing_tx, mut outgoing_rx) = mpsc::channel(128);
let pending = Arc::new(Mutex::new(HashMap::new()));

// Writer task: outgoing_rx → stdin (JSON-RPC requests)
tokio::spawn(async move {
    let mut stdin = stdin;
    while let Some(msg) = outgoing_rx.recv().await {
        let json = serde_json::to_string(&msg)?;
        stdin.write_all(json.as_bytes()).await?;
        stdin.write_all(b"\n").await?; // Newline-delimited
    }
    Ok::<(), io::Error>(())
});

// Reader task: stdout → pending HashMap (JSON-RPC
responses)
let pending_clone = Arc::clone(&pending);
tokio::spawn(async move {
    let mut lines = BufReader::with_capacity(1024 * 1024,
stdout).lines();

    while let Ok(Some(line)) = lines.next_line().await {
        let msg: JSONRPCMessage =
serde_json::from_str(&line)?;

        // Match response to request
        if let Some(id) = msg.id {
            if let Some(tx) =
pending_clone.lock().unwrap().remove(&id) {
                let _ = tx.send(msg);
            }
        }
        Ok::<(), anyhow::Error>(())
    });
    Ok(Self {
        child,
        outgoing_tx,
        pending,
        id_counter: AtomicI64::new(1),
    })
})
}

```

**Key Design Choices:** - **1MB buffer:** BufReader::with\_capacity(1024 \* 1024) handles large tool responses - **Line-delimited JSON:** Each JSON-RPC message on one line - **Concurrent I/O:** Separate reader/writer tasks prevent deadlock - **kill\_on\_drop:** Subprocess cleaned up automatically

---

## JSON-RPC Protocol

**Types:** codex-rs/mcp-types/src/jsonrpc.rs

```

#[derive(Serialize, Deserialize)]
#[serde(untagged)]
pub enum JSONRPCMessage {
    Request {
        jsonrpc: String, // "2.0"
        id: i64,
        method: String,
        params: Option<Value>,
    },
    Response {
        jsonrpc: String,
        id: i64,
        result: Option<Value>,
        error: Option<JSONRPCError>,
    },
    Notification {
        jsonrpc: String,
        method: String,
        params: Option<Value>,
    },
}

```

**Wire Format** (newline-delimited):

```

>{"jsonrpc":"2.0","id":1,"method":"tools/list","params":null}
 {"jsonrpc":"2.0","id":1,"result":{"tools":
 [{"name":"read_file","description":"..."}]}}

```

---

## Tool Invocation

```

impl McpClient {
    pub async fn call_tool(
        &self,
        tool_name: &str,
        arguments: Value,
    ) -> Result<Value> {
        // Generate unique request ID
        let id = self.id_counter.fetch_add(1, Ordering::SeqCst);

        // Create oneshot channel for response
        let (tx, rx) = oneshot::channel();
        self.pending.lock().unwrap().insert(id, tx);

        // Send request
        let request = JSONRPCMessage::Request {
            jsonrpc: "2.0".to_string(),
            id,
            method: "tools/call".to_string(),
            params: Some(json!({
                "name": tool_name,
                "arguments": arguments,
            })),
        };

        self.outgoing_tx.send(request).await?;

        // Wait for response (with timeout)
        let response = tokio::time::timeout(
            Duration::from_secs(60),

```

```
    rx
).await??;

// Extract result
match response {
    JSONRPCMessage::Response { result: Some(result), .. } =>
Ok(result),
    JSONRPCMessage::Response { error: Some(err), .. } =>
Err(err.into()),
    _ => Err(Error::InvalidResponse),
}
}
```

**Flow:** 1. Generate unique ID (atomic counter) 2. Create oneshot channel, store in pending HashMap 3. Send JSON-RPC request to server (via outgoing\_tx) 4. Writer task writes to stdin 5. Server processes request, writes to stdout 6. Reader task parses response 7. Match ID, send via oneshot channel 8. Remove from pending HashMap 9. Return result to caller

## McpConnectionManager

**Purpose:** Central hub for all MCP servers, tool aggregation

**Location:** codex-rs/core/src/mcp\_connection\_manager.rs:84-150

```
pub struct McpConnectionManager {
    clients: HashMap<String, Arc<McpClient>>, // Server name →
client
    tools: HashMap<String, ToolInfo>, // Qualified tool
name → metadata
}

pub struct ToolInfo {
    pub server_name: String,
    pub tool_name: String,
    pub qualified_name: String, // "server_tool_name"
    pub description: String,
    pub input_schema: Value,
}

impl McpConnectionManager {
    pub async fn new(
        mcp_servers: HashMap<String, McpServerConfig>,
        excluded_tools: HashSet<(String, String)>,
    ) -> Result<(Self, ClientStartErrors)> {
        let mut clients = HashMap::new();
        let mut tools = HashMap::new();
        let mut errors = ClientStartErrors::default();

        // Spawn all servers concurrently
        let mut join_set = JoinSet::new();

        for (server_name, config) in mcp_servers {
            let server_name_clone = server_name.clone();
            join_set.spawn(async move {
                let client = McpClient::new_stdio_client(

```

```

        config.command.into(),

config.args.into_iter().map(OsString::from).collect(),
    config.env,
).await?;

        // Initialize server
        client.call_method("initialize", json!
        {"protocolVersion": "1.0"}).await?;

        // List tools
        let response = client.call_method("tools/list",
None).await?;
        let server_tools: Vec<Tool> =
serde_json::from_value(response["tools"].clone())?;

        Ok::<(>(String, Arc<McClient>, Vec<Tool>), Error>((
            server_name_clone,
            Arc::new(client),
            server_tools,
        )))
    );
}

// Collect results
while let Some(result) = join_set.join_next().await {
    match result? {
        Ok((server_name, client, server_tools)) => {
            clients.insert(server_name.clone(), client);

            // Qualify and aggregate tools
            for tool in server_tools {
                let qualified_name =
qualify_tool_name(&server_name, &tool.name);

                if !excluded_tools.contains(&
                    (server_name.clone(), tool.name.clone())) {
                    tools.insert(qualified_name.clone(),
ToolInfo {
                        server_name: server_name.clone(),
                        tool_name: tool.name.clone(),
                        qualified_name,
                        description: tool.description,
                        input_schema: tool.input_schema,
                    });
                }
            }
        },
        Err(e) => {
            errors.add(server_name, e);
        },
    }
}

Ok((Self { clients, tools }, errors))
}

pub async fn invoke_tool(
    &self,
    qualified_name: &str,

```

```

        arguments: Value,
    ) -> Result<Value> {
    // Look up tool info
    let tool_info = self.tools.get(qualified_name)

    .ok_or(Error::ToolNotFound(qualified_name.to_string()))?;

    // Get client for server
    let client = self.clients.get(&tool_info.server_name)

    .ok_or(Error::ServerNotFound(tool_info.server_name.clone()))?;

    // Invoke tool on server
    client.call_tool(&tool_info.tool_name, arguments).await
}
}

```

**Key Features:** - **Concurrent initialization:** Spawn all servers in parallel (JoinSet) - **Tool qualification:** filesystem\_read\_file (prevents name collisions) - **Excluded tools:** Filter out unwanted tools - **Error collection:** Track which servers failed to start - **Shared clients:** Arc<McpClient> allows concurrent access

---

## Tool Name Qualification

**Purpose:** Prevent tool name collisions across servers

**Strategy:** Prefix with server name

```

fn qualify_tool_name(server_name: &str, tool_name: &str) -> String {
    let combined = format!("{}_{ }", server_name, tool_name);

    // Limit to 64 chars (OpenAI tool name limit)
    if combined.len() > 64 {
        // Hash collision avoidance for long names
        let hash = sha1::Sha1::digest(combined.as_bytes());
        let prefix = &combined[..40];
        format!("{}_{:x}", prefix, hash)
    } else {
        combined.replace('-', '_') // Normalize separators
    }
}

```

**Examples:** - read\_file (filesystem) → filesystem\_read\_file - query (database) → database\_query - very-long-server-name\_very-long-tool-name → very-long-server-name\_very-long-tool\_<hash>

---

## App-Level Shared Connection Manager

**Purpose:** Prevent MCP server process multiplication

**Problem:** Each ChatWidget spawning its own MCP connections would create  $N \times M$  processes ( $N$  widgets  $\times$   $M$  servers)

**Solution:** Singleton shared manager

**Location:** codex-rs/tui/src/app.rs:105-107

```

pub(crate) struct App<'a> {
    chat_widgets: Vec<ChatWidget<'a>>,
    mcp_manager: Arc<tokio::sync::Mutex<Option<Arc<McConnectionManager>>>,
        // ↑ Shared singleton across all ChatWidgets
}

impl App {
    pub fn new() -> Self {
        // Initialize MCP manager once
        let mcp_manager = Arc::new(tokio::sync::Mutex::new(None));

        tokio::spawn({
            let mcp_manager = Arc::clone(&mcp_manager);
            let config = load_config();

            async move {
                let manager = McConnectionManager::new(
                    config.mcp_servers,
                    HashSet::new(),
                ).await?;

                *mcp_manager.lock().await = Some(Arc::new(manager));
            }
        });
    }

    Self {
        chat_widgets: Vec::new(),
        mcp_manager,
    }
}

```

**Result:** Only M MCP server processes (one per configured server), regardless of widget count

---

## Performance: Native vs Subprocess

### Benchmark Results

#### Subprocess approach (pre-2025-10-18):

Tool invocation via local-memory MCP:  
- Spawn process: ~20ms  
- Execute tool: ~15ms  
- Parse output: ~5ms  
- Process cleanup: ~6ms  
Total: ~46ms per tool call

#### Native approach (current):

Tool invocation via native client:  
- Subprocess already running  
- JSON-RPC roundtrip: ~7ms  
- Parse response: ~1.7ms  
Total: ~8.7ms per tool call

**Speedup: 5.3x faster** (8.7ms vs 46ms)

**Additional benefits:** - No process spawn overhead - Persistent connection (reuse stdout/stdin) - Lower memory footprint (shared subprocess)

---

## Server Lifecycle

### Initialization Sequence

1. Spawn subprocess
    - ↳ `Command::new(program).spawn()`
  2. Send initialize request
    - ↳ `{"method": "initialize", "params": {"protocolVersion": "1.0"}}`
  3. Receive initialize response
    - ↳ `{"result": {"capabilities": {...}, "serverInfo": {...}}}`
  4. Send initialized notification
    - ↳ `{"method": "notifications/initialized"}`
  5. List tools
    - ↳ `{"method": "tools/list"}`
    - ↳ `{"result": {"tools": [...]}}`
  6. Ready for tool calls
- 

### Shutdown Sequence

1. Send shutdown request
    - ↳ `{"method": "shutdown"}`
  2. Receive shutdown response
    - ↳ `{"result": null}`
  3. Send exit notification
    - ↳ `{"method": "exit"}`
  4. Wait for process exit
    - ↳ `child.wait().await`
  5. Cleanup (automatic via `kill_on_drop`)
- 

## Health Monitoring

### Configurable timeouts:

```
# ~/.code/config.toml

[mcp_servers.filesystem]
startup_timeout_sec = 10    # Server initialization timeout
tool_timeout_sec = 60        # Per-tool call timeout
```

### Timeout enforcement:

```
// Startup timeout
```

---

```

let client = tokio::time::timeout(
    Duration::from_secs(config.startup_timeout_sec),
    McpClient::new_stdio_client(...))
.await??;

// Tool call timeout
let result = tokio::time::timeout(
    Duration::from_secs(config.tool_timeout_sec),
    client.call_tool(name, args)
).await?;;

```

---

## Error Handling

### Error Types

```

pub enum McpError {
    // Connection errors
    ServerSpawnFailed(io::Error),
    ServerNotResponding,
    ServerCrashed(i32),           // Exit code

    // Protocol errors
    InvalidMessage(serde_json::Error),
    UnexpectedResponse,
    RequestTimeout,

    // Tool errors
    ToolNotFound(String),
    ToolExecutionFailed(String),
    InvalidArguments(String),
}

```

---

### Retry Logic

**Transient errors** (retry): - Server not responding (may be overloaded) - Request timeout (network issue)

**Permanent errors** (don't retry): - Tool not found (invalid tool name) - Invalid arguments (schema mismatch) - Server crashed (exit code non-zero)

```

let result = execute_with_backoff(
    || client.call_tool(name, args),
    &RetryConfig::default(),
).await?;;

```

---

### Graceful Degradation

If MCP server fails:

```

match mcp_manager.invoke_tool(name, args).await {
    Ok(result) => {
        // Tool executed successfully
        use_tool_result(result);
    },
}

```

```

Err(McpError::ToolNotFound(_)) => {
    // Tool doesn't exist, inform model
    return "Tool not available. Please try another approach.";
},
Err(McpError::ServerCrashed(_)) => {
    // Server crashed, disable MCP for this conversation
    disable_mcp_for_conversation();
    return "MCP server unavailable. Continuing without tools.";
},
Err(e) => {
    // Other error, retry or fail
    return format!("Tool execution failed: {}", e);
},
}

```

---

## Tool Schema Validation

### Input Schema

**From MCP server:**

```
{
    "name": "read_file",
    "description": "Read contents of a file",
    "inputSchema": {
        "type": "object",
        "properties": {
            "path": {
                "type": "string",
                "description": "File path to read"
            }
        },
        "required": ["path"]
    }
}
```

**Validation before invocation:**

```

pub fn validate_tool_arguments(
    tool_info: &ToolInfo,
    arguments: &Value,
) -> Result<()> {
    let schema = &tool_info.input_schema;

    // Use jsonschema crate for validation
    let compiled_schema = JSONSchema::compile(schema)?;

    compiled_schema.validate(arguments)
        .map_err(|errors| {
            let messages: Vec<_> = errors.map(|e|
e.to_string()).collect();
            Error::InvalidArguments(messages.join(", "))
        })
}

```

---

## Configuration

## Per-Server Config

```
# ~/.code/config.toml

[mcp_servers.filesystem]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-filesystem",
"/home/user/project"]
env = { "NODE_ENV" = "production" }
startup_timeout_sec = 10
tool_timeout_sec = 60

[mcp_servers.local-memory]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-local-memory"]
startup_timeout_sec = 10
tool_timeout_sec = 30

[mcp_servers.git-status]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-git-status"]
```

---

## Excluded Tools

Filter out specific tools:

```
[mcp_servers.filesystem]
excluded_tools = ["write_file", "delete_file"] # Read-only mode
```

Programmatic exclusion:

```
let excluded = HashSet::from([
    ("filesystem".to_string(), "write_file".to_string()),
    ("filesystem".to_string(), "delete_file".to_string()),
]);
let (manager, errors) = McpConnectionManager::new(
    config.mcp_servers,
    excluded,
).await?;
```

---

## Summary

### MCP Integration Highlights:

1. **Native Client:** 5.3× faster than subprocess (8.7ms vs 46ms)
2. **Concurrent I/O:** Separate reader/writer tasks prevent deadlock
3. **Shared Manager:** App-level singleton prevents process multiplication
4. **Tool Qualification:** server\_tool\_name prevents collisions
5. **Timeout Enforcement:** Per-server startup and per-tool timeouts
6. **Error Handling:** Retry transient errors, fail fast on permanent
7. **Schema Validation:** Validate arguments before invocation
8. **Graceful Degradation:** Continue without MCP if servers fail

**Next Steps:** - Database Layer - SQLite optimization - Configuration System - Hot-reload

---

**File References:** - MCP client: codex-rs/mcp-client/src/mcp\_client.rs:63-150 - Connection manager: codex-rs/core/src/mcp\_connection\_manager.rs:84-150 - JSON-RPC types: codex-rs/mcp-types/src/jsonrpc.rs - App integration: codex-rs/tui/src/app.rs:105-107

---

# System Overview & Architecture

Comprehensive overview of the **theturtlecsz/code** architecture.


---

## Table of Contents

1. [High-Level Overview](#)
  2. [Design Philosophy](#)
  3. [Component Architecture](#)
  4. [Data Flow](#)
  5. [Technology Stack](#)
  6. [Fork-Specific Additions](#)
  7. [Integration Points](#)
- 

## High-Level Overview

**theturtlecsz/code** is a terminal-based AI coding assistant built on a multi-layered architecture:



**Key Characteristics:** - **226,607 lines of Rust** across 538 source files - **24-crate Cargo workspace** with modular design - **Async/sync hybrid** (Tokio async core, Ratatui sync UI) - **Native MCP integration** (5.3× faster than subprocess) - **98.8% upstream isolation** (fork features in isolated modules)

---

## Design Philosophy


### 1. Separation of Concerns

**Layer Boundaries:** - **UI Layer:** Ratatui TUI, user interaction, rendering - **Application Layer:** Business logic, workflow orchestration - **Service Layer:** MCP, database, configuration, agents - **Infrastructure Layer:** SQLite, stdio, HTTP clients

### 2. Async/Sync Hybrid Architecture

**Rationale:** Ratatui requires synchronous event loop, but backend services benefit from async I/O.

**Solution:** Clear async/sync boundary using channels:



**Pattern:** UnboundedSender<Op> bridges sync UI to async backend.

**File:** codex.rs/tui/src/chatwidget/agent.rs:16-62

### 3. Modularity via Cargo Workspace

**Benefits:** - Clear dependency boundaries - Independent compilation units - Parallel builds (faster CI) - Reusable components (spec-kit as separate crate)

**Workspace Size:** 24 crates, ~220K LOC Rust

### 4. Fork Isolation Strategy

**Goal:** Minimize rebase conflicts with upstream (just-every/code)

**Implementation:** - Fork features in isolated modules (tui/src/chatwidget/spec\_kit/) - “Friend module” pattern (access ChatWidget private fields without public API) - Dynamic command registry (avoids enum growth in SlashCommand) - Trait-based abstractions (SpecKitContext decouples from ChatWidget)

**Result:** 98.8% isolation (55 modules under spec\_kit/, minimal upstream changes)

**Evidence:** docs/spec-kit/REFACTORING\_COMPLETE\_SUMMARY.md

## 5. Performance-First Design

**Database:** - SQLite with WAL mode: **6.6× read speedup** (850µs → 129µs) - R2D2 connection pooling: Eliminates connection overhead - Incremental auto-vacuum: Prevents unbounded growth

**MCP Integration:** - Native client: **5.3× faster** than subprocess approach - Shared connection manager: Prevents process multiplication - 1MB buffer for large tool responses

**Configuration:** - Hot-reload with 300ms debounce - Arc for atomic updates (<1ms lock time)

## 6. Fault Tolerance

**Retry Logic** (SPEC-945C): - Exponential backoff: 100ms → 200ms → 400ms → 800ms - Jitter to prevent thundering herd - Permanent vs transient error classification

**Graceful Degradation:** - Multi-agent consensus works with 2/3 agents (if 1 fails) - MCP server failures don't crash app - Config reload failures preserve old config

---

# Component Architecture


## Core Components

### 1. TUI Layer (`codex-tui` crate)

**Purpose:** Terminal user interface using Ratatui framework

**Key Components:** - **App:** Top-level application state, event loop - **ChatWidget:** Main conversation interface (912K LOC) - **BottomPane:** Input composer, status bar - **HistoryCell:** Message rendering (user, assistant, tool, exec) - **StreamController:** Real-time token streaming

#### Event Flow:




**File:** `codex-rs/tui/src/app.rs`, `codex-rs/tui/src/chatwidget/mod.rs`

---

### 2. Spec-Kit Framework (`spec-kit` crate + TUI integration)

**Purpose:** Multi-agent automation pipeline with quality gates

#### Architecture:



```

↓
Pipeline Coordinator (state machine)
↓
Agent Orchestrator (spawn agents)
↓
Consensus Coordinator (aggregate results)
↓
Quality Gate Broker (checkpoints)
↓
Evidence Repository (artifacts)

```

**Key Modules:** - **spec-kit** (library crate): Config, retry, types - **tui/src/chatwidget/spec\_kit** (55 modules): TUI integration - **command\_registry.rs**: Dynamic command dispatch - **pipeline\_coordinator.rs**: Workflow state machine - **agent\_orchestrator.rs**: Agent lifecycle - **consensus\_coordinator.rs**: Multi-agent consensus - **native\_\*.rs**: Zero-cost operations (FREE) - **consensus\_db.rs**: SQLite artifact storage

**File:** codex-rs/spec-kit/src/lib.rs, codex-rs/tui/src/chatwidget/spec\_kit/mod.rs

---

### 3. Core Services (**codex-core** crate)

**Purpose:** Backend services for conversation, MCP, database, config

**Key Modules:** - **ConversationManager**: Agent conversation lifecycle - **McpConnectionManager**: MCP server aggregation - **Config**: Configuration loading, validation - **Database**: SQLite connection pooling, transactions - **Protocol**: OpenAI API client, model providers

**Responsibilities:** - Agent spawning and orchestration - MCP tool invocation - Model provider communication (OpenAI, Anthropic, Google) - Configuration hot-reload - Consensus artifact storage

**File:** codex-rs/core/src/lib.rs

---

### 4. MCP Integration (**mcp-client**, **mcp-types** crates)

**Purpose:** Model Context Protocol client and server support

**Components:** - **McpClient**: Async client for stdio communication - **McpConnectionManager**: Central hub for all MCP servers - **mcp-types**: JSON-RPC protocol types

**Architecture:**

```

MCP Server (subprocess)
  ↓ stdin/stdout
McpClient
  └── Writer Task: outgoing_rx → stdin (JSON-RPC requests)
  └── Reader Task: stdout → pending HashMap (JSON-RPC responses)
    └── Dispatcher: Request ID → oneshot::Sender (pair
      requests/responses)

```

**Performance:** - **Native integration**: 5.3x faster than subprocess (8.7ms typical) - **Concurrent I/O**: Separate reader/writer tasks prevent deadlock - **1MB buffer**: Handles large tool responses


**File:** codex-rs/mcp-client/src/mcp\_client.rs:63-150

---

## 5. Database Layer (**consensus\_db** in spec-kit, db module in core)

**Purpose:** SQLite storage for consensus artifacts, config, telemetry

**Architecture:**



**Optimizations:** - **WAL mode:** 6.6× read speedup (allows concurrent reads) - **Connection pooling:** Eliminates connection overhead -

**Optimized pragmas:** 32MB cache, memory-mapped I/O (1GB) -

**Incremental auto-vacuum:** 99.95% size reduction after cleanup

**Schema:** - consensus\_runs: Workflow execution tracking - agent\_outputs: Individual agent responses - consensus\_artifacts: Synthesized consensus results

**File:** codex-rs/core/src/db/connection.rs:39-105, codex-rs/tui/src/chatwidget/spec\_kit/consensus\_db.rs

---

## 6. Configuration System (config module in spec-kit)

**Purpose:** Layered configuration with hot-reload

**5-Tier Precedence** (highest to lowest): 1. **CLI flags:** --model gpt-5, -config key=value 2. **Shell environment:** export OPENAI\_API\_KEY=... 3.

**Profile:** [profiles.premium] in config.toml 4. **Config file:**

~/.code/config.toml 5. **Defaults:** Built-in fallback values

**Hot-Reload:**

```
File Change → Debouncer (300ms) → Validate → Lock → Replace → Event
                                         ↓ Fail
                                         Preserve Old Config
```

**Performance:** - Reload latency: <100ms (p95) - Lock contention: <1ms write locks - CPU overhead: <0.5% idle

**File:** codex-rs/spec-kit/src/config/hot\_reload.rs:1-100

---

## 7. Model Providers (codex-protocol, codex-chatgpt crates)

**Purpose:** Communication with AI model APIs

**Providers:** - **OpenAI:** GPT-5, GPT-4o, o3 (Responses API) - **Anthropic:** Claude Sonnet, Haiku, Opus (via CLI) - **Google:** Gemini Pro, Flash (via CLI) - **Custom:** Any OpenAI-compatible endpoint

**Features:** - Streaming responses (SSE) - Retry logic (exponential backoff) - Rate limit handling - Zero Data Retention support (ZDR)

**File:** codex-rs/protocol/src/lib.rs, codex-rs/chatgpt/src/lib.rs

---

## Data Flow

### Conversation Flow

1. User Input
  - ↳ ChatWidget.handle\_key\_event()
    - ↳ ChatWidget.submit\_prompt()
      - ↳ codex\_op\_tx.send(Op::NewMessage)
2. Async Backend
  - ↳ ConversationManager.submit(op)
    - ↳ Conversation.process()
      - ↳ Model Provider API (OpenAI/Anthropic/Google)
3. Response Streaming
  - ↳ conversation.next\_event()
    - ↳ app\_event\_tx.send(AppEvent::CodexEvent)
      - ↳ ChatWidget.handle\_event()
        - ↳ Render response tokens

### Spec-Kit Automation Flow

1. User Command: /speckit.auto SPEC-ID
  - ↳ CommandRegistry.find("speckit.auto")
    - ↳ AutoCommand.execute(widget, args)
2. Pipeline Initialization
  - ↳ PipelineCoordinator.start\_pipeline()
    - ↳ Load SPEC from docs/SPEC-{ID}/\*/spec.md
3. Stage Execution Loop
  - ↳ For each stage (specify, plan, tasks, implement, validate, audit, unlock):
    - ↳ NativeQualityGate.check() [if native stage]
    - ↳ AgentOrchestrator.spawn\_agents() [if multi-agent]
    - ↳ ConsensusCoordinator.synthesize() [aggregate results]
    - ↳ QualityGateBroker.validate() [checkpoint]
    - ↳ EvidenceRepository.store() [artifacts]
4. Completion
  - ↳ PipelineCoordinator.complete()
    - ↳ Push results to ChatWidget history

### MCP Tool Invocation Flow

1. Agent requests tool
  - ↳ Model response: {"type": "tool\_use", "name": "filesystem\_read\_file"}
2. Tool dispatch
  - ↳ McpConnectionManager.invoke\_tool()

```

    ↳ Find MCP client by tool prefix ("filesystem")
        ↳ McpClient.call_tool()

3. Server communication
    ↳ outgoing_tx.send(JSONRPCRequest)
        ↳ Writer task → stdin (JSON)
            ↳ MCP server processes request
            ↳ stdout (JSON) → Reader task
                ↳ pending.get(request_id).send(response)

4. Result returned
    ↳ Tool result forwarded to model
        ↳ Model continues generation with tool output

```

---

## Technology Stack

### Core Technologies

**Language:** Rust (Edition 2024) - Memory safety without garbage collection - Zero-cost abstractions - Fearless concurrency

**Async Runtime:** Tokio - Multi-threaded work-stealing scheduler - Async I/O for network, file, subprocess - Channels for sync/async boundary

**Terminal UI:** Ratatui (v0.29.0, patched fork) - Immediate-mode rendering - Cross-platform terminal support - Widget composition

**Database:** SQLite (rusqlite crate) - Embedded database (no server) - ACID transactions - WAL mode for concurrency

**Serialization:** Serde (JSON, TOML, YAML) - Compile-time serialization - Type-safe deserialization - Schema validation

### Supporting Libraries

**Networking:** - reqwest: HTTP client (model providers) - eventsource-stream: Server-Sent Events (streaming)

**Configuration:** - toml: Config file parsing - notify: Filesystem watching (hot-reload)

**Database:** - rusqlite: SQLite bindings - r2d2: Connection pooling - r2d2\_sqlite: SQLite adapter for r2d2

**MCP:** - mcp-types: Protocol definitions (internal) - tokio-util: Codec for line-delimited JSON

**Error Handling:** - anyhow: Flexible error types - thiserror: Custom error derive

**CLI:** - clap: Command-line argument parsing - clap\_complete: Shell completion generation

---

### Fork-Specific Additions

## Isolation Strategy

**Goal:** Add fork features without upstream conflicts

**Implementation:**

### 1. Isolated Module Tree:

```
codex-rs/tui/src/chatwidget/
└── mod.rs (upstream)
    └── spec_kit/ (fork, 55 modules, 98.8% isolated)
        ├── mod.rs
        ├── command_registry.rs
        ├── pipeline_coordinator.rs
        └── ... (50+ modules)
```

### 2. Friend Module Pattern:

```
// In chatwidget/mod.rs (upstream file, minimal change)
pub mod spec_kit; // Single line addition

// spec_kit modules can access ChatWidget private fields
impl ChatWidget {
    fn internal_method(&mut self) { /* ... */ }
}
```

### 3. Dynamic Command Registry (avoids upstream enum):

```
// Upstream: SlashCommand enum
pub enum SlashCommand { New, Model, Reasoning, ... }

// Fork: Dynamic registry (no enum growth)
pub trait SpecKitCommand { /* ... */ }
SPEC_KIT_REGISTRY.register(Box::new(AutoCommand));
```

### 4. Context Trait (decouples from ChatWidget):

```
pub trait SpecKitContext {
    fn submit_operation(&self, op: Op);
    fn push_error(&mut self, message: String);
    // ... methods spec-kit needs
}

impl SpecKitContext for ChatWidget { /* ... */ }
```

**Result:** - **98.8% isolation:** 55 modules, 1,222 lines extracted - **Zero upstream conflicts:** Rebases require <10 lines of merge -

**Testability:** MockSpecKitContext for unit tests - **Maintainability:** Clear separation of concerns

**Evidence:** docs/spec-kit/REFACTORING\_COMPLETE\_SUMMARY.md

---

## New Crates

**spec-kit** (codex-rs/spec-kit/): - Configuration system (hot-reload, 5-tier precedence) - Retry logic (exponential backoff, jitter) - Types and error handling - Evidence management - Cost tracking

**Purpose:** Reusable library for spec-kit automation (can extract as standalone in future per MAINT-10)

---

## Integration Points

### 1. TUI $\leftrightarrow$ Core Services

**Boundary:** Async/sync channel boundary

**Direction:** Bidirectional - **TUI**  $\rightarrow$  **Core**: UnboundedSender<0p> - **Core**  $\rightarrow$  **TUI**: AppEventSender

**Pattern:** Message passing with typed events

---

### 2. Core $\leftrightarrow$ MCP Servers

**Boundary:** stdio subprocess communication

**Direction:** Bidirectional (JSON-RPC over stdin/stdout)

**Protocol:** Line-delimited JSON (Model Context Protocol)

**Lifecycle:** 1. Spawn subprocess (tokio::process::Command) 2. Initialize with initialize request 3. List tools with tools/list request 4. Invoke tools with tools/call request 5. Kill on app exit (kill\_on\_drop = true)

---

### 3. Core $\leftrightarrow$ Model Providers

**Boundary:** HTTPS API requests

**Direction:** Request/response with streaming

**Protocols:** - **OpenAI**: Responses API (streaming SSE) - **Anthropic**: Messages API (via CLI subprocess) - **Google**: Gemini API (via CLI subprocess)

**Features:** - Retry logic (exponential backoff) - Rate limit handling (429 response) - Streaming token delivery

---

### 4. Spec-Kit $\leftrightarrow$ Database

**Boundary:** R2D2 connection pool

**Direction:** Read/write consensus artifacts

**Operations:** - Store agent outputs (per-agent JSON blobs) - Store consensus synthesis (aggregated results) - Query historical runs (evidence retrieval) - Atomic transactions (ACID guarantees)

---

### 5. Configuration $\leftrightarrow$ Filesystem

**Boundary:** File watching (notify crate)

**Direction:** Read config.toml, watch for changes

**Hot-Reload:** 1. Filesystem change event 2. Debounce (300ms window) 3. Parse and validate config 4. Atomic update via Arc 5. Emit reload event to UI

---

## Summary

### Architecture Highlights:

1. **Clean Layering:** TUI → Core → Infrastructure
2. **Async/Sync Hybrid:** Tokio backend, Ratatui UI
3. **Modular Design:** 24-crate workspace
4. **Fork Isolation:** 98.8% isolation via friend modules
5. **Performance-First:** 6.6x DB speedup, 5.3x MCP speedup
6. **Fault Tolerance:** Retry logic, graceful degradation

**Next Steps:** - [Cargo Workspace Guide](#) - Detailed crate documentation  
- [TUI Architecture](#) - Ratatui and async/sync patterns - [MCP Integration](#) - Native client details - [Database Layer](#) - SQLite optimization

---

**File References:** - Workspace: codex-rs/Cargo.toml - TUI: codex-rs/tui/src/app.rs, codex-rs/tui/src/chatwidget/mod.rs - Spec-Kit: codex-rs/spec-kit/src/lib.rs, codex-rs/tui/src/chatwidget/spec\_kit/mod.rs - Core: codex-rs/core/src/lib.rs - MCP: codex-rs/mcp-client/src/mcp\_client.rs - DB: codex-rs/core/src/db/connection.rs - Config: codex-rs/spec-kit/src/config/hot\_reload.rs

---

## TUI Architecture

Detailed architecture of the Terminal User Interface layer.

---

## Overview

The TUI layer uses **Ratatui** (v0.29.0, patched fork) with a hybrid **async/sync** architecture:

- **Sync Layer:** Ratatui event loop (blocking terminal operations)
- **Async Layer:** Tokio tasks (network I/O, subprocess management)
- **Bridge:** Channels (UnboundedSender, AppEventSender)

**File:** codex-rs/tui/src/

---

## Component Hierarchy

```
App (top-level state)
  └── ChatWidget (conversation interface)
    ├── BottomPane (input + status)
    ├── HistoryCell[] (message history)
    └── StreamController (token streaming)
```

```

    └── InterruptManager (Ctrl+C handling)
    └── SpecKitState (fork features)
└── TerminalInfo (size, capabilities)
└── AppEventReceiver (backend events)

```

---

## ChatWidget Structure

**Location:** codex-rs/tui/src/chatwidget/mod.rs:373+

```

pub(crate) struct ChatWidget<'a> {
    // === Backend Communication ===
    app_event_tx: AppEventSender,                                // Send events
    to App
    codex_op_tx: UnboundedSender<Op>,                           // Submit
    operations to backend

    // === UI Components ===
    bottom_pane: BottomPane<'a>,                                 // Input
    composer + status bar
    history_cells: Vec<Box<dyn HistoryCell>>,                // Conversation
    history

    // === State ===
    config: Config,                                              // Configuration
    auth_manager: Arc<AuthManager>,                               //
    Authentication state

    // === Spec-Kit (Fork) ===
    spec_auto_state: Option<SpecAutoState>,                     // Pipeline
    orchestration
    cost_tracker: Arc<spec_kit::cost_tracker::CostTracker>, // Cost
    tracking

    // === Agent Tracking ===
    active_agents: Vec<AgentInfo>,                                // Running
    agents
    agent_runtime: HashMap<String, AgentRuntime>,               // Agent
    metadata

    // === Execution State ===
    exec: ExecState,                                               // Command
    execution tracking
    terminal: TerminalState,                                       // Tmux session
    state

    // === Rendering ===
    stream: StreamController,                                      // Token
    streaming
    interrupts: InterruptManager,                                  // Interrupt
    handling
    cached_cell_size: OnceCell<(u16, u16)>,                   // Terminal
    dimensions
}

```

**Key Insights:** - **912K LOC in mod.rs** (large monolithic file) - **Friend module access:** spec\_kit modules can access private fields - **Hybrid ownership:** Arc<T> for shared state, direct ownership for UI

---

# Async/Sync Boundary Pattern

## The Problem

Ratatui requires synchronous event loop:

```
loop {
    terminal.draw(|f| ui(f, &app))?; // Blocking draw
    let event = event::read()?;
    app.handle_event(event); // Sync handler
}
```

Backend requires async I/O:


```
async fn conversation_loop() {
    let response = model_provider.chat(request).await?; // Async
    network I/O
    while let Some(token) = response.next().await? { // Async
        streaming
        // ...
    }
}
```

**Constraint:** Can't .await in sync event loop ✗

---

## The Solution: Channel Bridge

**Pattern:**



**Implementation:** codex-rs/tui/src/chatwidget/agent.rs:16-62

```
pub(crate) fn spawn_agent(
    config: Config,
    app_event_tx: AppEventSender,
    server: Arc<ConversationManager>,
) -> UnboundedSender<Op> {
    // Create channel for UI → Backend
    let (codex_op_tx, mut codex_op_rx) = unbounded_channel::<Op>();

    // Spawn async task
    tokio::spawn(async move {
        // Create conversation
        let NewConversation { conversation, .. } =
            server.new_conversation(config).await?;

        // Forward operations to conversation
        tokio::spawn(async move {
            while let Some(op) = codex_op_rx.recv().await {
                conversation.submit(op).await;
            }
        })
    })
}
```

```

    });

    // Forward events back to UI
    while let Ok(event) = conversation.next_event().await {
        app_event_tx.send(AppEvent::CodexEvent(event))?;
    }
});

// Return sync sender to UI
codex_op_tx
}

```

**Key Points:** - `UnboundedSender<Op>`: Sync side can send without `.await` - `app_event_tx.send()`: Backend sends events to UI queue - **Tokio runtime**: Spawns on separate thread pool - **No blocking**: UI thread never blocks on network I/O

---

## Operation Types (Op enum)

```

pub enum Op {
    NewMessage(String),                      // User prompt
    ToolResponse(ToolCallId, Result<String>), // Tool execution
    result
    RegenerateLastMessage,                  // Retry last response
    Interrupt,                            // Cancel current operation
    // ... 10+ variants
}

```

**Flow:** 1. User types message → ChatWidget.submit\_prompt() 2. Create `Op::NewMessage` → `codex_op_tx.send(op)` 3. Backend receives → `conversation.submit(op).await` 4. Model responds → `app_event_tx.send(Event::Token)` 5. UI receives → `ChatWidget.handle_event() → Render`

---

## Event Loop

**Location:** codex-rs/tui/src/app.rs

```

impl App {
    pub fn run(mut self) -> Result<()> {
        loop {
            // Draw UI
            self.terminal.draw(|f| {
                self.chat_widget.render(f, f.size());
            })?;

            // Handle events (non-blocking with timeout)
            if event::poll(Duration::from_millis(16))? {
                match event::read()? {
                    Event::Key(key) => self.handle_key(key)?,
                    Event::Resize(w, h) => self.handle_resize(w,
h)?,
                    Event::Mouse(mouse) =>
                    self.handle_mouse(mouse)?,
                    _ => {}
                }
            }
        }
    }
}

```

```

    // Process backend events
    while let Ok(app_event) = self.app_event_rx.try_recv() {
        self.handle_app_event(app_event)?;
    }

    // Check for exit
    if self.should_exit {
        break;
    }
}
Ok(())
}
}

```

**Loop Phases:** 1. **Draw:** Render UI to terminal buffer 2. **Poll:** Check for terminal events (16ms timeout = ~60 FPS) 3. **Handle Terminal Events:** Keyboard, mouse, resize 4. **Process Backend Events:** Tokens, completions, errors 5. **Check Exit:** Break if requested

---

## Rendering System

### Immediate Mode Rendering

**Ratatui uses immediate-mode:** - No retained UI tree - Full re-render every frame - Layout calculated on-the-fly

**Performance:** ~60 FPS for typical conversation UI

---

### Widget Composition

```

impl ChatWidget {
    pub fn render(&mut self, f: &mut Frame, area: Rect) {
        // Split layout
        let chunks = Layout::default()
            .direction(Direction::Vertical)
            .constraints([
                Constraint::Min(1),           // History
                Constraint::Length(3),        // Composer
                Constraint::Length(1),        // Status bar
            ])
            .split(area);

        // Render history
        self.render_history(f, chunks[0]);

        // Render input composer
        self.bottom_pane.render(f, chunks[1]);

        // Render status bar
        self.render_status(f, chunks[2]);
    }
}

```

---

### HistoryCell Trait

**Location:** codex-rs/tui/src/chatwidget/history/

```
pub trait HistoryCell: Send {
    fn height(&self, width: u16) -> u16; // Calculate cell
height
    fn render(&self, frame: &mut Frame, area: Rect); // Render to
area
}

// Implementations:
pub struct UserMessageCell { /* ... */ } // User prompt
pub struct AssistantMessageCell { /* ... */ } // AI response
pub struct ToolExecutionCell { /* ... */ } // Tool call/result
pub struct ExecCell { /* ... */ } // Command execution
pub struct BackgroundEventCell { /* ... */ } // System messages
```

**Dynamic Dispatch:** Vec<Box<dyn HistoryCell>> allows heterogeneous message types

---

## Spec-Kit Integration (Friend Module)

### Friend Module Pattern

Declared in chatwidget/mod.rs:

```
pub mod spec_kit; // Friend module - can access private fields
```

**Benefits:** - ✓ Spec-kit can read/write ChatWidget internals - ✓ No public API pollution - ✓ Clear encapsulation boundary - ✓ Easy to test (spec\_kit modules are independent)

---

### Context Trait Abstraction

**Location:** codex-rs/tui/src/chatwidget/spec\_kit/context.rs:1-140

```
pub trait SpecKitContext {
    // History operations
    fn history_push(&mut self, cell: impl HistoryCell + 'static);
    fn push_error(&mut self, message: String);
    fn push_background(&mut self, message: String, placement: BackgroundPlacement);

    // UI operations
    fn request_redraw(&mut self);

    // Agent/operation submission
    fn submit_operation(&self, op: Op);
    fn submit_prompt(&mut self, display: String, prompt: String);

    // Configuration access
    fn working_directory(&self) -> &Path;
    fn agent_config(&self) -> &[AgentConfig];

    // Spec auto state
    fn spec_auto_state_mut(&mut self) -> &mut Option<SpecAutoState>;
    fn spec_auto_state(&self) -> &Option<SpecAutoState>;
```

```

    // Guardrail operations
    fn collect_guardrail_outcome(&self, spec_id: &str, stage: SpecStage) -> Result<GuardrailOutcome>;
    fn run_spec_consensus(&mut self, spec_id: &str, stage: SpecStage)
        -> Result<(Vec<Line<'static>, bool)>;
}

impl SpecKitContext for ChatWidget {
    // Implementation delegates to ChatWidget methods
}

```

**Purpose:** Decouples spec-kit from ChatWidget implementation

**Testing:** Mock implementation in context::test\_mock

---

## Streaming & Interrupts

### StreamController

**Location:** codex-rs/tui/src/streaming/controller.rs

```

pub struct StreamController {
    active_stream: Option<StreamState>,
}

impl StreamController {
    pub fn start_stream(&mut self, request_id: String) {
        self.active_stream = Some(StreamState {
            request_id,
            tokens: Vec::new(),
            started_at: Instant::now(),
        });
    }

    pub fn append_token(&mut self, token: String) {
        if let Some(stream) = &mut self.active_stream {
            stream.tokens.push(token);
        }
    }

    pub fn finish_stream(&mut self) -> Option<StreamState> {
        self.active_stream.take()
    }
}

```

**Flow:** 1. Backend sends Event::StreamStart 2. StreamController.start\_stream() 3. Backend sends Event::Token(tok) repeatedly 4. StreamController.append\_token(tok) 5. UI renders partial response on each frame 6. Backend sends Event::StreamEnd 7. StreamController.finish\_stream()

---

### InterruptManager

**Location:** codex-rs/tui/src/chatwidget/interrupts.rs

```
pub struct InterruptManager {
```

```

        pending_interrupt: bool,
        last_interrupt_at: Option<Instant>,
    }

impl InterruptManager {
    pub fn request_interrupt(&mut self) {
        self.pending_interrupt = true;
        self.last_interrupt_at = Some(Instant::now());
    }

    pub fn consume_interrupt(&mut self) -> bool {
        std::mem::replace(&mut self.pending_interrupt, false)
    }
}

```

**Usage:** Ctrl+C → InterruptManager.request\_interrupt() → Send Op::Interrupt to backend → Backend cancels operation

---

## Input Handling

### BottomPane (Composer)

**Location:** codex-rs/tui/src/chatwidget/bottom\_pane.rs

```

pub struct BottomPane<'a> {
    input: String,                                // Current input buffer
    cursor_position: usize,                      // Cursor position
    history_index: Option<usize>,                // Command history
    navigation: FileSearch,                       // @ file search state
}

impl BottomPane {
    pub fn handle_key(&mut self, key: KeyEvent) -> InputAction {
        match key.code {
            KeyCode::Enter =>
                InputAction::Submit(std::mem::take(&mut self.input)),
            KeyCode::Char('@') if self.input.is_empty() => {
                self.navigation = Some(FileSearch::new());
                InputAction::None
            },
            KeyCode::Char(c) => {
                self.input.insert(self.cursor_position, c);
                self.cursor_position += 1;
                InputAction::None
            },
            KeyCode::Backspace => {
                if self.cursor_position > 0 {
                    self.input.remove(self.cursor_position - 1);
                    self.cursor_position -= 1;
                }
                InputAction::None
            },
            // ... more key handlers
        }
    }
}

```

**Features:** - Multi-line input (Shift+Enter) - Cursor movement (arrows, Home, End) - History navigation (Esc Esc) - File search (@ trigger) - Paste support (Ctrl+V with image detection)

---

## File Search (@-trigger)

**Location:** codex-rs/file-search/src/lib.rs

```
pub struct FileSearch {
    query: String,
    results: Vec<PathBuf>,
    selected_index: usize,
}

impl FileSearch {
    pub fn update_query(&mut self, query: String) {
        self.query = query;
        self.results = fuzzy_search(&query, max_results: 10);
        self.selected_index = 0;
    }

    fn fuzzy_search(query: &str, max_results: usize) -> Vec<PathBuf> {
        // Use nucleo-matcher for fuzzy matching
        // Search workspace for files matching query
    }
}
```

**UI Flow:** 1. User types @ → Activate file search 2. User types main → Update query, show results 3. User presses Up/Down → Navigate results 4. User presses Tab/Enter → Insert file path, exit search 5. User presses Esc → Cancel search

---

## Performance Considerations

### Rendering Optimizations

**Lazy Height Calculation:**

```
pub fn height(&self, width: u16) -> u16 {
    // Calculate height only when width changes
    if self.cached_width == Some(width) {
        return self.cached_height;
    }
    let height = self.calculate_height(width);
    self.cached_width = Some(width);
    self.cached_height = height;
    height
}
```

**Viewport Culling:** - Only render visible history cells - Skip cells outside viewport - Recalculate on scroll

---

### Event Processing

### Non-Blocking Event Poll:

```
if event::poll(Duration::from_millis(16))? {  
    // Process event  
}
```

- 16ms = ~60 FPS target
  - Non-blocking (returns immediately if no events)
  - Allows backend event processing every frame
- 

## Summary

### TUI Architecture Highlights:

1. **Async/Sync Hybrid:** Channels bridge sync UI and async backend
2. **Immediate-Mode Rendering:** Full re-render every frame (~60 FPS)
3. **Dynamic Dispatch:** Box<dyn HistoryCell> for heterogeneous messages
4. **Friend Module Pattern:** Spec-kit access to ChatWidget internals
5. **Context Trait:** Decouples spec-kit from ChatWidget implementation
6. **Streaming:** Real-time token rendering
7. **Interrupts:** Ctrl+C gracefully cancels operations

**Next Steps:** - [Core Execution](#) - Agent orchestration - [MCP Integration](#)  
- Native client details - [Database Layer](#) - SQLite optimization

---

**File References:** - ChatWidget: codex-  
rs/tui/src/chatwidget/mod.rs:373+ - Agent spawner: codex-  
rs/tui/src/chatwidget/agent.rs:16-62 - Event loop: codex-  
rs/tui/src/app.rs - Context trait: codex-  
rs/tui/src/chatwidget/spec\_kit/context.rs:1-140 - Streaming: codex-  
rs/tui/src/streaming/controller.rs - Bottom pane: codex-  
rs/tui/src/chatwidget/bottom\_pane.rs

---