

SPEC-DOC-004: Testing & Quality Assurance Documentation

- Objectives
- Scope
  - In Scope
  - Out of Scope
- Deliverables
- Success Criteria
- Related SPECs
- CI/CD Integration
- Overview
- Testing Stages
  1. Pre-Commit (Local) - <5s
  2. Pre-Push (Local) - ~2-5 min
  3. CI/CD (GitHub Actions) - ~10-15 min
- GitHub Actions Workflows
  - Preflight Tests Job
  - Build Matrix Job
  - Release Job
- Pre-Commit Hook
  - Installation
  - What It Checks
  - Bypass Pre-Commit (Emergencies Only)
- CI Test Script
  - Location
  - What It Tests
  - Why Curated Tests?
- Local Testing Before Push
  - Recommended Workflow
  - Fast Iteration Loop
- Code Coverage Integration
  - Local Coverage Measurement
  - CI Coverage (Future)
- Best Practices
  - DO
  - DON'T
- Troubleshooting
  - Pre-Commit Hook Not Running
  - CI Timeout
  - Flaky Tests
  - Build Cache Corruption
- Summary

Pattern 2: Login Flow E2E

E2E Test Setup Patterns

- Pattern 1: Temp Git Repository
- Pattern 2: Mock HTTP Server
- Pattern 3: Fixture Loading

Best Practices

- DO
- DON'T

Running E2E Tests

- Run All E2E Tests
- Run Specific E2E Test
- Run with Verbose Output

Summary

Integration Testing Guide

Overview

Integration Test Categories

- W01-W15: Workflow Integration Tests
- E01-E15: Error Recovery Integration Tests
- S01-S10: State Persistence Integration Tests
- Q01-Q10: Quality Gate Integration Tests
- C01-C10: Concurrent Operations Integration Tests

Test Structure

- Standard Integration Test Pattern

Workflow Integration Tests

- Pattern 1: Individual Stage Workflow
- Pattern 2: Multi-Stage Pipeline
- Pattern 3: Quality Gate Integration

Error Recovery Integration Tests

- Pattern 1: Consensus Failure → Retry → Recovery
- Pattern 2: MCP Failure → Fallback → Recovery

State Persistence Integration Tests

- Pattern 1: State Serialization → Load → Reconstruct
- Pattern 2: Pipeline Interrupt → Resume from Checkpoint

Evidence Verification Patterns

- Pattern 1: Comprehensive Evidence Verification
- Pattern 2: Degraded Consensus Detection

Best Practices

- DO
- DON'T

Running Integration Tests

- Run All Integration Tests
- Run Specific Category
- Run Specific Test
- Run with Output

Summary

Performance Testing Guide

Overview

Benchmarking with Criterion

- What is Criterion?
- Setup
- Basic Benchmark
- Database Performance Benchmark
- Throughput Benchmarks
- Running Benchmarks

Profiling

- Flamegraphs with cargo-flamegraph
- perf (Linux only)
- cargo-bloat (Binary Size Analysis)

Command-Line Benchmarking

- hyperfine

Benchmarking /speckit.auto

Performance Metrics

- Database Performance
- MCP Performance
- Config Hot-Reload

Regression Testing

- Baseline Comparison
- Continuous Performance Monitoring

Best Practices

- DO
- DON'T

Summary

Property-Based Testing Guide

Overview

What is Property-Based Testing?

- Traditional Example-Based Testing
- Property-Based Testing

Getting Started

- Add proptest Dependency
- Basic Property Test

Generators

- Built-in Generators
- Custom Generators

Testing Invariants

- Invariant 1: State Index Always Valid
- Invariant 2: Current Stage Mapping
- Invariant 3: Retry Count Never Negative

Testing Evidence Integrity

- Property 1: Written Evidence Always Parseable JSON
- Property 2: Evidence File Names Valid

Testing Collections

- Property 1: Filtering Never Increases Length
- Property 2: Sorting Preserves Length
- Property 3: Dedupe Length

Testing String Operations

- Property 1: Truncation Length
- Property 2: Regex Escape Safety

Shrinking

- What is Shrinking?
- Shrinking Example

Advanced Patterns

- Conditional Properties
- Composite Strategies
- Regression Testing

Configuration

- Adjust Test Cases
- Environment Variable
- Timeout

Best Practices

- DO
- DON'T

Running Property Tests

- Run All Property Tests
- Run with More Cases
- Debug Failing Test
- Re-run Regression Cases

Summary

Test Infrastructure

Overview

MockMcpManager

- Purpose
- API Reference
- Fixture Matching
- Usage Patterns
- Tests
- IntegrationTestContext
  - Purpose
  - API Reference
  - Usage Patterns
  - Tests
- StateBuilder
  - Purpose
  - API Reference
  - Usage Patterns
- EvidenceVerifier
  - Purpose
  - API Reference
  - Usage Patterns
- Fixture Library
  - Overview
  - File Naming Convention
  - Available Fixtures
  - Usage in Tests
  - Adding New Fixtures
- Coverage Tools
  - cargo-tarpaulin
  - cargo-llvm-cov
- Property-Based Testing
  - Overview
  - Proptest Basics
  - Test Categories
  - Custom Generators
  - Running Property Tests
- TestCodexBuilder
  - Purpose
  - API Reference
  - Custom Configuration
  - Usage with Wiremock
- Common Test Utilities
  - Test Module Structure
  - Shared Test Data
- Test Organization Best Practices
  - File Naming
  - Test Naming Conventions
  - Common Test Patterns
- Summary
- Testing Strategy
  - Overview
  - Coverage Goals
    - Overall Target: 40%+
    - Module-Specific Targets
- Testing Pyramid
  - Level 1: Unit Tests (~63%)
  - Level 2: Integration Tests (~33%)
  - Level 3: E2E Tests (~4%)
- Test Organization
  - Per-Module Tests
  - Workspace-Level Tests
- Coverage Measurement
  - Tools

Alternative: cargo-llvm-cov

Critical Path Coverage

- Priority 1: Spec-Kit Automation (70%+ target)
- Priority 2: MCP Client (70%+ target)
- Priority 3: Database Layer (50%+ target)

Test Execution Strategy

- Local Development
- Pre-Commit Hook
- Pre-Push Hook
- CI/CD Pipeline

Coverage Gaps

- Known Gaps (Acceptable)
- Prioritized Improvements

Testing Best Practices

- DO
- DON'T

Summary

Unit Testing Guide

Overview

Test Structure

- Arrange-Act-Assert Pattern
- Given-When-Then Pattern

Naming Conventions

- Test Function Names
- Test Module Organization

Testing Pure Functions

- What are Pure Functions?
- Example 1: Pattern Matching
- Example 2: Conditional Logic

Testing Error Handling

- Testing Error Cases
- Testing Error Messages
- Testing Panic Conditions

Testing with Test Data

- Inline Test Data
- External Test Fixtures
- Generated Test Data

Testing State Machines

- Example: SpecAutoState Transitions
- Testing Invalid Transitions
- Testing State Invariants

Testing Calculations

- Scoring Functions
- Penalty Calculations

Testing Collections

- Testing Filters
- Testing Sorting
- Testing Aggregations

Testing String Manipulation

- Regex Matching
- String Transformations
- Regex Escaping

Testing File Operations (with TempDir)

- Setup Pattern
- Testing Directory Creation

Testing with Mocks

- MockMcpManager Usage

Table-Driven Tests

- Pattern: Multiple Test Cases
- Parameterized Tests (with rtest)

Common Assertions  
    Equality  
    Boolean  
    Contains  
    Custom Messages  
    Floating Point  
Best Practices  
    DO  
    DON'T  
Running Tests  
    Run All Unit Tests  
    Run Specific Module  
    Run Specific Test  
    Run with Output  
    Run with Threads  
Test Coverage  
    Measure Coverage  
    Improve Coverage  
Summary

## SPEC-DOC-004: Testing & Quality Assurance Documentation

**Status:** Pending **Priority:** P1 (Medium) **Estimated Effort:** 12-16 hours **Target Audience:** Contributors, QA engineers **Created:** 2025-11-17

---

## Objectives

Document the complete testing and QA infrastructure: 1. Testing strategy (coverage goals: 40%+ target, currently 42-48%) 2. Test infrastructure (MockMcpManager, fixtures, tarpaulin) 3. Unit testing guide (patterns, examples, mocking) 4. Integration testing (workflow tests, cross-module) 5. E2E testing (pipeline validation, tmux automation) 6. Property-based testing (proptest, edge cases) 7. CI/CD integration (GitHub workflows, pre-commit hooks) 8. Performance testing (benchmarking, profiling)

---

## Scope

### In Scope

- Testing strategy and coverage targets (42-48% achieved, targeting 40%+)
- Test infrastructure (MockMcpManager implementation, fixtures)
- Unit testing patterns and examples
- Integration testing approach (604 tests total, 100% pass rate)
- E2E testing with tmux automation
- Property-based testing with proptest
- CI/CD workflows (.github/workflows/)
- Pre-commit/pre-push hooks
- Performance testing and benchmarking

- Test organization (per-module, integration tests)

## Out of Scope

- Writing new tests (implementation work)
  - Internal testing policy details (covered in testing-policy.md)
  - Spec-kit functional testing (covered in SPEC-DOC-003)
- 

## Deliverables

1. **content/testing-strategy.md** - Coverage goals, module targets
  2. **content/test-infrastructure.md** - MockMcpManager, fixtures, tools
  3. **content/unit-testing-guide.md** - Patterns, examples, mocking
  4. **content/integration-testing-guide.md** - Workflow tests, cross-module
  5. **content/e2e-testing-guide.md** - Pipeline validation, tmux
  6. **content/property-testing-guide.md** - Proptest usage, edge cases
  7. **content/ci-cd-integration.md** - GitHub workflows, hooks
  8. **content/performance-testing.md** - Benchmarking, profiling
- 

## Success Criteria

- Testing strategy clearly documented
  - MockMcpManager usage fully explained
  - Unit test patterns demonstrated with examples
  - Integration test approach documented
  - CI/CD workflow explained
  - All 604 existing tests referenced
- 

## Related SPECs

- SPEC-DOC-000 (Master)
  - SPEC-DOC-002 (Core Architecture - testing architecture)
  - SPEC-DOC-005 (Development - running tests locally)
- 

**Status:** Structure defined, content pending

---

## CI/CD Integration

Comprehensive guide to CI/CD integration and automated testing.

---

## Overview

**CI/CD Testing Strategy:** Automated testing at every stage (pre-commit, pre-push, CI, release)

**Goals:** - Fast feedback (<5s pre-commit, <2min pre-push) -  
Comprehensive coverage (all tests in CI) - Prevent regressions -  
Maintain code quality

**Current Status:** - Pre-commit hooks: 100% adoption - CI pipeline:  
GitHub Actions - Test execution: ~10-15 minutes - Pass rate: 100%

---

## Testing Stages

### 1. Pre-Commit (Local) - <5s

**Purpose:** Fast policy checks before commit

**Location:** .githooks/pre-commit

**What it runs:**

```
# Only runs if spec_kit modules modified
# Check 1: Storage policy
bash scripts/validate_storage_policy.sh

# Check 2: Tag schema
bash scripts/validate_tag_schema.sh
```

**Execution Time:** <5s

**Bypass** (emergencies only):

```
git commit --no-verify
```

---

### 2. Pre-Push (Local) - ~2-5 min

**Purpose:** Compile and lint checks before push

**Triggered:** Before git push

**What it runs:**

```
# Format check
cargo fmt --all --check

# Linting
cargo clippy --workspace --all-targets --all-features -- -D warnings

# Build (all features)
cargo build --workspace --all-features

# Optional: Targeted test compilation
cargo test --workspace --no-run
```

**Execution Time:** ~2-5 minutes

**Bypass** (emergencies only):

```
PREPUSH_FAST=0 git push
```

---

### 3. CI/CD (GitHub Actions) - ~10-15 min

**Purpose:** Complete testing and release

**Triggered:** Push to main, pull requests

**Location:** .github/workflows/release.yml

**Jobs:** 1. **Preflight Tests** (Linux fast E2E) 2. **Determine Version** (semantic versioning) 3. **Build** (Linux, macOS, Windows) 4. **Test** (all tests, all platforms) 5. **Release** (npm publish)

---

## GitHub Actions Workflows

### Preflight Tests Job

**Purpose:** Fast integration tests before full build matrix

**Platform:** Ubuntu 24.04

**Steps:**

```
jobs:
  preflight-tests:
    name: Preflight Tests (Linux fast E2E)
    runs-on: ubuntu-24.04
    steps:
      - name: Checkout code
        uses: actions/checkout@v4

      - name: Install Rust (1.90)
        run: |
          rustup set profile minimal
          rustup toolchain install 1.90.0 --profile minimal
          rustup default 1.90.0

      - name: Setup Rust Cache
        uses: Swatinem/rust-cache@v2
        with:
          prefix-key: v5-rust
          shared-key: codex-preflight-1.90
          workspaces: codex-rs -> target
          cache-targets: true
          cache-on-failure: true

      - name: Build (fast profile)
        run: ./build-fast.sh

      - name: Curated tests + CLI smokes
        run: bash scripts/ci-tests.sh
```

**What it tests:**

```
# scripts/ci-tests.sh

# Curated integration tests
cargo test -p codex-login --test all -q
cargo test -p codex-chatgpt --test all -q
```

```

cargo test -p codex-apply-patch --test all -q
cargo test -p codex-excpolicy --tests -q
cargo test -p mcp-types --tests -q

# CLI smoke tests
./codex-rs/target/dev-fast/code --version
./codex-rs/target/dev-fast/code completion bash
./codex-rs/target/dev-fast/code doctor

```

**Execution Time:** ~3-5 minutes

**Benefits:** - ✓ Fast feedback (before full matrix) - ✓ Catches common errors early - ✓ Tests critical integration points - ✓ Validates CLI functionality

---

## Build Matrix Job

**Purpose:** Build and test on all platforms

**Platforms:** - Linux (Ubuntu 24.04, x64 + arm64) - macOS (latest, x64 + arm64) - Windows (latest, x64)

**Rust Versions:** - Stable (1.90) - Beta (optional)

**Steps:**

```

jobs:
  build:
    strategy:
      matrix:
        os: [ubuntu-24.04, macos-latest, windows-latest]
        rust: [1.90.0]
    runs-on: ${{ matrix.os }}
    steps:
      - name: Checkout
        uses: actions/checkout@v4

      - name: Install Rust
        run: rustup toolchain install ${{ matrix.rust }}

      - name: Build
        run: cargo build --workspace --all-features

      - name: Test
        run: cargo test --workspace --all-features

```

**Execution Time:** ~8-12 minutes per platform

---

## Release Job

**Purpose:** Publish to npm after successful tests

**Triggers:** - Push to main - All tests pass

**Steps:** 1. Determine version (semantic versioning) 2. Build binaries (all platforms) 3. Publish to npm (@just-every/code)

**Packages Published:** - @just-every/code (main package) - @just-every/code-darwin-arm64 - @just-every/code-darwin-x64 - @just-every/code-linux-x64-musl - @just-every/code-linux-arm64-musl - @just-every/code-win32-x64

---

## Pre-Commit Hook

### Installation

#### One-time setup:

```
bash scripts/setup-hooks.sh
```

#### Verifies:

```
git config core.hooksPath  
# Should output: .githooks
```

---

### What It Checks

**File:** .githooks/pre-commit

```
#!/bin/bash  
# Pre-commit hook for policy compliance

# Only run if spec_kit modules modified
SPEC_KIT_CHANGES=$(git diff --cached --name-only | grep "spec_kit"  
|| true)

if [ -z "$SPEC_KIT_CHANGES" ]; then  
    # No spec_kit changes, skip policy checks  
    exit 0
fi

echo "✖️ Running policy compliance checks (spec_kit modified)..."

# Check 1: Storage policy
if ! bash scripts/validate_storage_policy.sh; then  
    echo "✖️ Storage policy violation detected"  
    exit 1
fi

# Check 2: Tag schema
if ! bash scripts/validate_tag_schema.sh; then  
    echo "✖️ Tag schema violation detected"  
    exit 1
fi

echo "✔️ Policy compliance checks passed"
exit 0
```

**Checks:** 1. **Storage policy:** Ensures local-memory usage compliant (MEMORY-POLICY.md) 2. **Tag schema:** Validates tag namespacing and naming

**Performance:** <5s (only runs for spec\_kit changes)

---

## Bypass Pre-Commit (Emergencies Only)

```
# Skip hook (use sparingly)
git commit --no-verify -m "Emergency fix"
```

**When to bypass:** - Critical production hotfix - Hook infrastructure broken - Reviewing/reverting broken commits

**When NOT to bypass:** - Avoiding policy violations (fix the code instead) - Convenience (hooks are fast) - Regular workflow

---

## CI Test Script

### Location

**File:** scripts/ci-tests.sh

**Purpose:** Fast integration tests for CI

---

### What It Tests

```
#!/usr/bin/env bash
set -euo pipefail

echo "[ci-tests] Running curated integration tests..."
pushd codex-rs >/dev/null

# Login integration tests
cargo test -p codex-login --test all -q

# ChatGPT integration tests
cargo test -p codex-chatgpt --test all -q

# Apply patch integration tests
cargo test -p codex-apply-patch --test all -q

# Execution policy tests
cargo test -p codex-execpolicy --tests -q

# MCP types tests
cargo test -p mcp-types --tests -q

popd >/dev/null

echo "[ci-tests] CLI smokes with host binary..."
BIN=./codex-rs/target/dev-fast/code

# Smoke tests
"${BIN}" --version >/dev/null
"${BIN}" completion bash >/dev/null
"${BIN}" doctor >/dev/null || true

echo "[ci-tests] Done."
```

---

### Why Curated Tests?

**Full test suite:** 604 tests, ~15 minutes

**Curated subset:** ~150 tests, ~3-5 minutes

**Selection Criteria:** - ✓ Integration tests (cross-module) - ✓ E2E tests (complete workflows) - ✓ Critical paths (login, apply, MCP) - ✗ Unit tests (fast, covered by local dev) - ✗ Property tests (slow, covered by weekly runs)

**Benefits:** - ✓ Fast feedback (3-5 min vs 15 min) - ✓ High signal (integration tests find real bugs) - ✓ CI efficiency (parallel preflight + full tests)

---

## Local Testing Before Push

### Recommended Workflow

**Step 1: Run affected tests** (iterative development):

```
cd codex-rs

# Test specific module you changed
cargo test -p codex-tui --lib

# Test specific file
cargo test -p codex-tui spec_kit::clarify_native
```

**Step 2: Run full test suite** (before committing):

```
cd codex-rs
cargo test --workspace --all-features
```

**Step 3: Check format and lint** (before committing):

```
cd codex-rs
cargo fmt --all -- --check
cargo clippy --workspace --all-targets --all-features -- -D warnings
```

**Step 4: Commit** (pre-commit hook runs automatically):

```
git add .
git commit -m "feat(tui): add clarify native checks"
# Hook runs: storage policy, tag schema (<5s)
```

**Step 5: Push** (pre-push hook runs automatically, optional):

```
git push
# Hook runs: fmt, clippy, build (~2-5min)
```

---

## Fast Iteration Loop

**For rapid development:**

```
# 1. Make changes
vim codex-rs/tui/src/chatwidget/spec_kit/clarify_native.rs

# 2. Test just this module (fast)
cd codex-rs
```

```

cargo test -p codex-tui clarify_native -- --nocapture

# 3. If tests pass, run clippy on this crate
cargo clippy -p codex-tui -- -D warnings

# 4. Commit (hook runs policy checks)
git add codex-rs/tui
git commit -m "fix(clarify): improve ambiguity detection"

# 5. Push later after multiple commits
git push

```

**Execution Time:** - Module tests: ~5-10s - Clippy: ~15-30s - Commit: <5s (hook) - **Total:** ~30-50s per iteration

---

## Code Coverage Integration

### Local Coverage Measurement

**Tool:** cargo-tarpaulin or cargo-llvm-cov

**Install:**

```

cargo install cargo-tarpaulin
# or
cargo install cargo-llvm-cov

```

**Usage:**

```

cd codex-rs

# Generate coverage report
cargo tarpaulin --workspace --all-features --out Html

# Open report
open target/tarpaulin/index.html

```

---

### CI Coverage (Future)

**GitHub Actions** (not yet implemented):

```

jobs:
  coverage:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4

      - name: Install tarpaulin
        run: cargo install cargo-tarpaulin

      - name: Run coverage
        run: cargo tarpaulin --workspace --all-features --out Xml

      - name: Upload toCodecov
        uses: codecov/codecov-action@v3
        with:
          files: ./cobertura.xml

```

- name: Comment PR with coverage
- uses: codecov/codecov-action@v3

**Benefits** (when implemented): - ✓ Track coverage trends - ✓ Fail PR if coverage drops >5% - ✓ Visualize coverage in PRs

---

## Best Practices

### DO

#### ✓ Run tests locally before pushing:

```
# Always test before pushing
cargo test --workspace --all-features

# Push after tests pass
git push
```

---

#### ✓ Fix CI failures immediately:

```
# CI failed? Fix it now, not later
git pull
cargo test --workspace
# Fix failures
git commit -m "fix(ci): resolve test failures"
git push
```

---

#### ✓ Keep CI green: - Main branch should always pass tests - Revert breaking commits if fix takes >1 hour - Document known flaky tests

---

#### ✓ Use caching effectively:

```
# GitHub Actions caching
- uses: Swatinem/rust-cache@v2
  with:
    prefix-key: v5-rust
    shared-key: codex-preflight-1.90
```

---

#### ✓ Run curated tests in CI (fast feedback):

```
# Preflight: curated subset (3-5 min)
bash scripts/ci-tests.sh

# Full matrix: all tests (10-15 min)
cargo test --workspace --all-features
```

---

### DON'T

#### ✗ Skip pre-commit hooks routinely:

```
# Bad: Habitual bypassing
git commit --no-verify # ✗ Don't make this a habit
```

---

### **✗ Push without testing:**

```
# Bad: Push untested code
git commit -m "quick fix"
git push # ✗ No local testing
```

---

### **✗ Ignore CI failures:**

```
# Bad: "CI is always red anyway"
# ✗ Fix CI or revert
```

---

### **✗ Commit broken tests:**

```
# Bad: Disable failing tests instead of fixing
#[test]
#[ignore] // ✗ Don't ignore, fix!
fn test_that_fails() {}
```

---

### **✗ Let coverage drop:**

```
# Bad: Coverage drops from 45% to 30%
# ✗ Add tests, don't delete them
```

---

## Troubleshooting

### Pre-Commit Hook Not Running

**Symptom:** Commits succeed without running hook

**Fix:**

```
# Check git config
git config core.hooksPath
# Should output: .githooks

# If not set, run setup
bash scripts/setup-hooks.sh
```

---

### CI Timeout

**Symptom:** CI job times out after 60 minutes

**Causes:** - Infinite loop in test - Deadlock in concurrent test - Slow property test (PROPTEST\_CASES too high)

**Fix:**

```
# Find slow tests locally
cargo test --workspace --nocapture --test-threads=1

# Reduce property test cases
PROPTEST_CASES=100 cargo test --test property_based_tests
```

---

## Flaky Tests

**Symptom:** Test passes locally, fails in CI (or vice versa)

**Common Causes:** - Race conditions (concurrent tests) - Hardcoded paths (use TempDir) - Network dependencies (use mocks) - Time-dependent tests (use fixed timestamps)

**Fix:**

```
# Run test multiple times locally
for i in {1..100}; do
    cargo test test_flaky_name || break
done

# If it fails, debug with single thread
cargo test test_flaky_name -- --test-threads=1 --nocapture
```

---

## Build Cache Corruption

**Symptom:** Build fails in CI with cryptic errors, passes locally

**Fix (GitHub Actions):**

```
# Clear cache by changing cache key
- uses: Swatinem/rust-cache@v2
  with:
    prefix-key: v6-rust # Increment version
```

---

## Summary

**CI/CD Testing Stages:**

1. **Pre-Commit** (<5s): Policy checks (storage, tags)
2. **Pre-Push** (2-5min): Format, clippy, build
3. **Preflight Tests** (3-5min): Curated integration tests
4. **Full CI** (10-15min): All tests, all platforms
5. **Release** (auto): Publish on main after tests pass

**Tools:** - ✓ GitHub Actions (CI/CD) - ✓ Rust Cache (faster builds) - ✓ Git Hooks (pre-commit, pre-push) - ✓ cargo-tarpaulin (coverage)

**Best Practices:** - ✓ Test locally before pushing - ✓ Keep CI green (100% pass rate) - ✓ Fast feedback (curated tests in preflight) - ✓ Fix failures immediately - ✓ Use caching (Rust Cache)

**Next Steps:** - Performance Testing - Benchmarks and profiling - Testing Strategy - Overall testing approach - Test Infrastructure - MockMcpManager, fixtures

---

**References:** - GitHub Actions: .github/workflows/release.yml - Pre-commit hook: .githooks/pre-commit - CI test script: scripts/ci-tests.sh - Setup hooks: scripts/setup-hooks.sh

---

# End-to-End Testing Guide

Comprehensive guide to end-to-end testing of complete user workflows.

---

## Overview

**End-to-End (E2E) Testing Philosophy:** Test complete user workflows from start to finish, simulating real-world usage

**Goals:** - Validate critical user journeys - Test system integration (TUI + backend + database + MCP) - Verify error recovery and degradation - Ensure configuration hot-reload works

**Current Status:** - ~24 E2E tests (4% of total) - 100% pass rate - Average execution time: ~10-60s per test - Categories: Pipeline automation, quality checkpoints, tmux sessions, config hot-reload

---

## E2E Test Categories

### Pipeline Automation Tests

**Purpose:** Test complete /speckit.auto pipeline (Plan → Tasks → Implement → Validate → Audit → Unlock)

**Location:** codex-rs/tui/tests/spec\_auto\_e2e.rs

**Coverage:** - Pipeline state machine (initialization, transitions, resume) - Quality checkpoint integration (PrePlanning, PostPlan, PostTasks) - Stage progression (all 6 stages) - Error handling and recovery

---

### Quality Checkpoint Tests

**Purpose:** Test quality gates at critical pipeline points

**Checkpoints:** - **PrePlanning** (BeforeSpecify): Clarify ambiguities before plan - **PostPlan** (AfterSpecify): Checklist quality scoring after plan - **PostTasks** (AfterTasks): Analyze consistency after tasks

**Coverage:** - Checkpoint triggering - Modification tracking - Escalation logic - Human intervention

---

### Tmux Session Tests

**Purpose:** Test tmux integration for long-running operations

**Location:** codex-rs/evidence/tmux-automation/

**Coverage:** - Session creation and lifecycle - Agent spawning in background - Session termination - Evidence collection

---

## Config Hot-Reload Tests

**Purpose:** Test configuration changes without restart

**Location:** codex-rs/tui/tests/config\_reload\_integration\_tests.rs

**Coverage:** - Config file watching - Hot-reload triggers - Provider switching - <100ms latency (p95)

---

## Pipeline E2E Tests

### Test Structure

**Standard Pattern:**

```
#[test]
fn test_spec_auto_state_initialization() {
    // 1. Create initial state
    let state = SpecAutoState::new(
        "SPEC-TEST-001".to_string(),
        "Test automation".to_string(),
        SpecStage::Plan,
        None, // HAL mode
    );

    // 2. Assert initial conditions
    assert_eq!(state.spec_id, "SPEC-TEST-001");
    assert_eq!(state.goal, "Test automation");
    assert_eq!(state.current_index, 0);
    assert_eq!(state.stages.len(), 6);
    assert_eq!(state.current_stage(), Some(SpecStage::Plan));
    assert!(state.quality_gates_enabled);
    assert!(state.completed_checkpoints.is_empty());
}
```

---

### Pattern 1: Pipeline Initialization

**Example:** spec\_auto\_e2e.rs:20

```
#[test]
fn test_spec_auto_state_initialization() {
    let state = SpecAutoState::new(
        "SPEC-TEST-001".to_string(),
        "Test automation".to_string(),
        SpecStage::Plan,
        None,
    );

    // Verify initial state
    assert_eq!(state.spec_id, "SPEC-TEST-001");
    assert_eq!(state.goal, "Test automation");
    assert_eq!(state.current_index, 0);
    assert_eq!(state.current_stage(), Some(SpecStage::Plan));

    // Verify stages
    assert_eq!(state.stages.len(), 6);
```

```

let expected = vec![
    SpecStage::Plan,
    SpecStage::Tasks,
    SpecStage::Implement,
    SpecStage::Validate,
    SpecStage::Audit,
    SpecStage::Unlock,
];
assert_eq!(state.stages, expected);

// Verify quality gates
assert!(state.quality_gates_enabled);
assert!(state.completed_checkpoints.is_empty());
}

```

**What This Tests:** - ✓ State initialization - ✓ Stage ordering (Plan → Tasks → Implement → Validate → Audit → Unlock) - ✓ Quality gates enabled by default - ✓ Checkpoint tracking initialized

---

## Pattern 2: Pipeline Stage Progression

```

#[test]
fn test_pipeline_full_progression() {
    let mut state = SpecAutoState::new(
        "SPEC-TEST-002".to_string(),
        "Full pipeline test".to_string(),
        SpecStage::Plan,
        None,
    );

    // ===== PLAN STAGE =====
    assert_eq!(state.current_stage(), Some(SpecStage::Plan));
    assert_eq!(state.current_index, 0);

    // Simulate plan completion
    state.current_index += 1;

    // ===== TASKS STAGE =====
    assert_eq!(state.current_stage(), Some(SpecStage::Tasks));
    assert_eq!(state.current_index, 1);

    state.current_index += 1;

    // ===== IMPLEMENT STAGE =====
    assert_eq!(state.current_stage(), Some(SpecStage::Implement));
    assert_eq!(state.current_index, 2);

    state.current_index += 1;

    // ===== VALIDATE STAGE =====
    assert_eq!(state.current_stage(), Some(SpecStage::Validate));
    assert_eq!(state.current_index, 3);

    state.current_index += 1;

    // ===== AUDIT STAGE =====
    assert_eq!(state.current_stage(), Some(SpecStage::Audit));
    assert_eq!(state.current_index, 4);
}

```

```

state.current_index += 1;

// ===== UNLOCK STAGE =====
assert_eq!(state.current_stage(), Some(SpecStage::Unlock));
assert_eq!(state.current_index, 5);

// ===== COMPLETION =====
state.current_index += 1;
assert_eq!(state.current_stage(), None); // Pipeline complete
}

```

**What This Tests:** - ✓ All 6 stages execute in order - ✓ Index advances correctly - ✓ State transitions deterministically - ✓ Pipeline completion (stage = None)

---

### Pattern 3: Resume from Middle Stage

```

#[test]
fn test_resume_from_tasks_stage() {
    // Start from Tasks (not Plan)
    let state = SpecAutoState::new(
        "SPEC-TEST-003".to_string(),
        "Resume test".to_string(),
        SpecStage::Tasks, // Resume from Tasks
        None,
    );

    // Verify resume point
    assert_eq!(state.current_index, 1); // Tasks is index 1
    assert_eq!(state.current_stage(), Some(SpecStage::Tasks));

    // Verify can still progress
    let mut state = state;
    state.current_index += 1;
    assert_eq!(state.current_stage(), Some(SpecStage::Implement));
}

```

**What This Tests:** - ✓ Pipeline can resume from any stage - ✓ Index calculated correctly for resume - ✓ Progression continues normally

---

## Quality Checkpoint E2E Tests

### Pattern 1: Checkpoint Tracking

```

#[test]
fn test_quality_checkpoints_track_completion() {
    let mut state = SpecAutoState::new(
        "SPEC-TEST-006".to_string(),
        "Checkpoint tracking".to_string(),
        SpecStage::Plan,
        None,
    );

    // Initially no checkpoints completed
    assert!(state.completed_checkpoints.is_empty());
}

```

```

// ===== PRE-PLANNING CHECKPOINT
=====

// Simulate PrePlanning checkpoint (Clarify)

state.completed_checkpoints.insert(QualityCheckpoint::PrePlanning);

    assert!
(state.completed_checkpoints.contains(&QualityCheckpoint::PrePlanning));

    assert!
(!state.completed_checkpoints.contains(&QualityCheckpoint::PostPlan));

    assert_eq!(state.completed_checkpoints.len(), 1);

// ===== POST-PLAN CHECKPOINT
=====

// Simulate PostPlan checkpoint (Checklist)
state.completed_checkpoints.insert(QualityCheckpoint::PostPlan);

    assert!
(state.completed_checkpoints.contains(&QualityCheckpoint::PrePlanning));

    assert!
(state.completed_checkpoints.contains(&QualityCheckpoint::PostPlan));

    assert_eq!(state.completed_checkpoints.len(), 2);

// ===== POST-TASKS CHECKPOINT
=====

// Simulate PostTasks checkpoint (Analyze)

state.completed_checkpoints.insert(QualityCheckpoint::PostTasks);

    assert_eq!(state.completed_checkpoints.len(), 3);
    assert!
(state.completed_checkpoints.contains(&QualityCheckpoint::PostTasks));

}

```

**What This Tests:** - ✓ Checkpoint completion tracked - ✓ Multiple checkpoints can coexist - ✓ No duplicate checkpoints (Set semantics)

---

## Pattern 2: Quality Modifications Tracking

```

#[test]
fn test_quality_modifications_tracked() {
    let mut state = SpecAutoState::new(
        "SPEC-TEST-007".to_string(),
        "Modification tracking".to_string(),
        SpecStage::Plan,
        None,
    );
}

// Initially no modifications
assert!(state.quality_modifications.is_empty());

```

```

// ===== PREPLANNING MODIFICATIONS
=====

// User fixes ambiguities in spec.md
state.quality_modifications.push("spec.md".to_string());

assert_eq!(state.quality_modifications.len(), 1);
assert!(
(state.quality_modifications.contains(&"spec.md".to_string()));

// ===== POSTPLAN MODIFICATIONS
=====

// User improves plan.md after checklist
state.quality_modifications.push("plan.md".to_string());

assert_eq!(state.quality_modifications.len(), 2);
assert!(
(state.quality_modifications.contains(&"plan.md".to_string()));

// ===== POSTTASKS MODIFICATIONS
=====

// User fixes tasks.md after analyze
state.quality_modifications.push("tasks.md".to_string());

assert_eq!(state.quality_modifications.len(), 3);
}

```

**What This Tests:** - ✓ Modifications tracked across checkpoints - ✓ Multiple files can be modified - ✓ Modification history preserved

---

### Pattern 3: Quality Gates Can Be Disabled

```

#[test]
fn test_quality_gates_can_be_disabled() {
    let state = SpecAutoState::with_quality_gates(
        "SPEC-TEST-008".to_string(),
        "No quality gates".to_string(),
        SpecStage::Plan,
        None,
        false, // Disable quality gates
    );

    // Verify quality gates disabled
    assert!(!state.quality_gates_enabled);

    // Pipeline should skip all checkpoints
    // (checkpoint logic would check quality_gates_enabled flag)
}

```

**What This Tests:** - ✓ Quality gates can be disabled - ✓ Flag persists in state - ✓ Pipeline can run without checkpoints

---

## Real-World E2E Tests

## Pattern 1: Apply Command E2E

Example: `apply_command_e2e.rs:78`

```
#[tokio::test]
async fn test_apply_command_creates_fibonacci_file() {
    // ===== SETUP: TEMP GIT REPO =====
    let temp_repo = create_temp_git_repo()
        .await
        .expect("Failed to create temp git repo");
    let repo_path = temp_repo.path();

    // ===== LOAD TASK FIXTURE =====
    let task_response = mock_get_task_with_fixture()
        .await
        .expect("Failed to load fixture");

    // ===== EXECUTE: APPLY DIFF =====
    apply_diff_from_task(task_response,
        Some(repo_path.to_path_buf()))
        .await
        .expect("Failed to apply diff from task");

    // ===== VERIFY: FILE CREATED =====
    let fibonacci_path = repo_path.join("scripts/fibonacci.js");
    assert!(fibonacci_path.exists(), "fibonacci.js was not created");

    // ===== VERIFY: FILE CONTENTS =====
    let contents = std::fs::read_to_string(&fibonacci_path)
        .expect("Failed to read fibonacci.js");

    assert!(
        contents.contains("function fibonacci(n)"),
        "fibonacci.js doesn't contain expected function"
    );
}
```

Helper: Create Temp Git Repo:

```
async fn create_temp_git_repo() -> anyhow::Result<TempDir> {
    let temp_dir = TempDir::new()?;
    let repo_path = temp_dir.path();
    let envs = vec![
        ("GIT_CONFIG_GLOBAL", "/dev/null"),
        ("GIT_CONFIG_NOSYSTEM", "1"),
    ];

    // Initialize git repo
    Command::new("git")
        .envs(envs.clone())
        .args(["init"])
        .current_dir(repo_path)
```

```

        .output()
        .await?;

    // Configure user
    Command::new("git")
        .envs(envs.clone())
        .args(["config", "user.email", "test@example.com"])
        .current_dir(repo_path)
        .output()
        .await?;

    Command::new("git")
        .envs(envs.clone())
        .args(["config", "user.name", "Test User"])
        .current_dir(repo_path)
        .output()
        .await?;

    // Create initial commit
    std::fs::write(repo_path.join("README.md"), "# Test Repo\n")?;

    Command::new("git")
        .envs(envs.clone())
        .args(["add", "README.md"])
        .current_dir(repo_path)
        .output()
        .await?;

    Command::new("git")
        .envs(envs.clone())
        .args(["commit", "-m", "Initial commit"])
        .current_dir(repo_path)
        .output()
        .await?;

    Ok(temp_dir)
}

```

**What This Tests:** - ✓ Complete apply command workflow - ✓ Git integration (temp repo, commits) - ✓ File creation and modification - ✓ Task fixture loading

---

## Pattern 2: Login Flow E2E

**Example: login\_server\_e2e.rs:79**

```

#[tokio::test]
async fn end_to_end_login_flow_persists_auth_json() -> Result<()> {
    // ===== SETUP: MOCK OAuth ISSUER
    =====

    let (issuer_addr, issuer_handle) = start_mock_issuer();
    let issuer = format!("http://{}:{}",
        issuer_addr.ip(),
        issuer_addr.port());

    // ===== SETUP: TEMP CODEX HOME
    =====

    let tmp = tempdir()?;

```

```

let codex_home = tmp.path().to_path_buf();

// Seed auth.json with stale data (should be overwritten)
let stale_auth = serde_json::json!({
    "OPENAI_API_KEY": "sk-stale",
    "tokens": {
        "id_token": "stale.header.payload",
        "access_token": "stale-access",
        "refresh_token": "stale-refresh",
    }
});
std::fs::write(
    codex_home.join("auth.json"),
    serde_json::to_string_pretty(&stale_auth)?,
)?;

// ===== EXECUTE: LOGIN FLOW =====

let options = ServerOptions {
    issuer: issuer.clone(),
    redirect_uri: "http://localhost:8080/callback".to_string(),
    codex_home: codex_home.clone(),
    // ... other options
};

run_login_server(options).await?;

// ===== VERIFY: AUTH.JSON UPDATED =====

let updated_auth =
std::fs::read_to_string(codex_home.join("auth.json"))?;
let auth_data: serde_json::Value =
serde_json::from_str(&updated_auth)?;

// Verify tokens refreshed
assert_ne!(auth_data["tokens"]["access_token"], "stale-access");
assert_eq!(auth_data["tokens"]["access_token"], "access-123");

// ===== CLEANUP: SHUTDOWN MOCK =====

drop(issuer_handle);

Ok(())
}

```

### Helper: Start Mock OAuth Issuer:

```

fn start_mock_issuer() -> (SocketAddr, thread::JoinHandle<()>) {
    let listener = TcpListener::bind(("127.0.0.1", 0)).unwrap();
    let addr = listener.local_addr().unwrap();
    let server = tiny_http::Server::from_listener(listener,
None).unwrap();

    let handle = thread::spawn(move || {
        while let Ok(mut req) = server.recv() {
            let url = req.url().to_string();
            if url.starts_with("/oauth/token") {
                // Build minimal JWT
                let payload = serde_json::json!({

```

```

        "email": "user@example.com",
        "https://api.openai.com/auth": {
            "chatgpt_plan_type": "pro",
        }
    });

    let id_token = create_jwt(&payload);

    let tokens = serde_json::json!({
        "id_token": id_token,
        "access_token": "access-123",
        "refresh_token": "refresh-123",
    });

    let resp = tiny_http::Response::from_data(
        serde_json::to_vec(&tokens).unwrap()
    );
    let _ = req.respond(resp);
}
}

});

(addr, handle)
}

```

**What This Tests:** - ✓ Complete login flow - ✓ OAuth integration (mock issuer) - ✓ Token persistence (auth.json) - ✓ Stale token replacement

---

## E2E Test Setup Patterns

### Pattern 1: Temp Git Repository

```

async fn create_temp_git_repo() -> anyhow::Result<TempDir> {
    let temp_dir = TempDir::new()?;
    let repo_path = temp_dir.path();

    // Disable global git config (isolation)
    let envs = vec![
        ("GIT_CONFIG_GLOBAL", "/dev/null"),
        ("GIT_CONFIG_NOSYSTEM", "1"),
    ];

    // Initialize repo
    run_git_command(repo_path, &envs, &["init"]).await?;

    // Configure user (required for commits)
    run_git_command(repo_path, &envs, &["config", "user.email", "test@example.com"]).await?;
    run_git_command(repo_path, &envs, &["config", "user.name", "Test User"]).await?;

    // Create initial commit
    std::fs::write(repo_path.join("README.md"), "# Test\n")?;
    run_git_command(repo_path, &envs, &["add", "."]).await?;
    run_git_command(repo_path, &envs, &["commit", "-m", "Initial commit"]).await?;
}

```

```

        Ok(temp_dir)
    }

async fn run_git_command(
    repo_path: &Path,
    envs: &[(&str, &str)],
    args: &[&str],
) -> anyhow::Result<()> {
    let output = Command::new("git")
        .envs(envs.iter().copied())
        .args(args)
        .current_dir(repo_path)
        .output()
        .await?;

    if !output.status.success() {
        anyhow::bail!(
            "Git command failed: {}",
            String::from_utf8_lossy(&output.stderr)
        );
    }
}

Ok(())
}

```

**Benefits:** - ✓ Isolated from global git config - ✓ Auto-cleanup (TempDir) - ✓ Reusable helper functions

---

## Pattern 2: Mock HTTP Server

```

fn start_mock_server() -> (SocketAddr, thread::JoinHandle<()>) {
    // Bind to random port
    let listener = TcpListener::bind(("127.0.0.1", 0)).unwrap();
    let addr = listener.local_addr().unwrap();
    let server = tiny_http::Server::from_listener(listener,
None).unwrap();

    let handle = thread::spawn(move || {
        while let Ok(req) = server.recv() {
            let url = req.url().to_string();

            let response = match url.as_str() {
                "/api/v1/endpoint" => {
                    serde_json::json!({"status": "ok"})
                }
                _ => {
                    serde_json::json!({"error": "not found"})
                }
            };

            let resp = tiny_http::Response::from_data(
                serde_json::to_vec(&response).unwrap()
            );
            let _ = req.respond(resp);
        }
    });

    (addr, handle)
}

```

```

#[tokio::test]
async fn test_with_mock_server() {
    let (addr, _handle) = start_mock_server();
    let base_url = format!("http://{}:{}",
        addr.ip(), addr.port());

    // Test code using base_url...
}

```

**Benefits:** - ✓ No external dependencies - ✓ Deterministic responses -  
✓ Fast (no network)

---

### Pattern 3: Fixture Loading

```

async fn load_fixture<T: serde::de::DeserializeOwned>(name: &str) ->
anyhow::Result<T> {
    let fixture_path = Path::new(env!("CARGO_MANIFEST_DIR"))
        .join("tests/fixtures")
        .join(format!("{}.json", name));

    let contents = std::fs::read_to_string(fixture_path)?;
    let data: T = serde_json::from_str(&contents)?;

    Ok(data)
}

#[tokio::test]
async fn test_with_fixture() {
    let task: GetTaskResponse = load_fixture("task_turn_fixture")
        .await
        .expect("Failed to load fixture");

    // Use task...
}

```

**Benefits:** - ✓ Realistic test data - ✓ Reusable across tests - ✓ Version controlled

---

## Best Practices

### DO

✓ **Test complete user workflows:**

```

// Good: Tests entire pipeline
#[test]
fn test_speckit_auto_full_pipeline() {
    // Create state
    // Run plan
    // Run tasks
    // ... all 6 stages
    // Verify completion
}

```

✓ **Use realistic test data:**

```
// Good: Load from fixture
let task = load_fixture("real_task_response").await?;

// Bad: Minimal mock data
let task = GetTaskResponse { id: "1", content: "test" };
```

---

#### ✓ Verify side effects:

```
// Verify file created
assert!(fibonacci_path.exists());

// Verify contents correct
let contents = std::fs::read_to_string(&fibonacci_path)?;
assert!(contents.contains("function fibonacci"));

// Verify git commit
let log = run_git(&["log", "--oneline"]).await?;
assert!(log.contains("Add fibonacci.js"));
```

---

#### ✓ Test error recovery:

```
#[tokio::test]
async fn test_pipeline_recovers_from_mcp_failure() {
    // Simulate MCP failure
    mock_mcp.fail_next_request();

    // Run pipeline
    let result = run_pipeline().await;

    // Verify fallback succeeded
    assert!(result.is_ok());
    assert!(result.unwrap().degraded);
}
```

---

#### ✓ Clean up resources:

```
#[tokio::test]
async fn test_with_cleanup() {
    let temp_dir = TempDir::new()?;
    let (_addr, handle) = start_mock_server();

    // Test logic...

    // Cleanup
    drop(handle); // Shutdown mock server
    drop(temp_dir); // Delete temp files

    Ok(())
}
```

---

## DON'T

#### ✗ Test too many workflows in one test:

```
// Bad: Tests multiple workflows (hard to debug)
#[test]
fn test_all_commands() {
```

```
    test_apply_command();
    test_login_flow();
    test_config_reload();
    test_tmux_session();
    // ... 500 lines
}
```

---

### ✗ Rely on external services:

```
// Bad: Depends on real OpenAI API
#[tokio::test]
async fn test_real_api() {
    let response =
reqwest::get("https://api.openai.com/v1/models").await?;
    // ✗ Flaky, slow, costs money
}

// Good: Use mock server
#[tokio::test]
async fn test_with_mock() {
    let (addr, _handle) = start_mock_server();
    let base_url = format!("http://{}", addr);
    // ✓ Fast, deterministic, free
}
```

---

### ✗ Skip verification:

```
// Bad: No assertions
#[tokio::test]
async fn test_pipeline() {
    run_pipeline().await?;
    // ✗ No verification
}

// Good: Verify outcomes
#[tokio::test]
async fn test_pipeline() {
    let result = run_pipeline().await?;
    assert_eq!(result.stages_completed, 6);
    assert!(result.plan_file.exists());
}
```

---

## Running E2E Tests

### Run All E2E Tests

```
cd codex-rs
cargo test --test '*_e2e'
```

**Runs:** - spec\_auto\_e2e.rs - apply\_command\_e2e.rs - login\_server\_e2e.rs

---

### Run Specific E2E Test

```
cargo test --test spec_auto_e2e test_spec_auto_state_initialization
```

---

## Run with Verbose Output

```
cargo test --test spec_auto_e2e -- --nocapture --test-threads=1
```

**Why --test-threads=1:** - E2E tests may conflict (ports, files) - Single-threaded ensures isolation

---

## Summary

### E2E Testing Best Practices:

1. **Complete Workflows:** Test from start to finish
2. **Realistic Data:** Use fixtures from real usage
3. **Isolation:** Temp dirs, mock servers, disable global config
4. **Verification:** Check files, state, side effects
5. **Error Recovery:** Test fallback and degradation
6. **Cleanup:** Auto-cleanup with TempDir, handle drops

**Test Categories:** - ✓ Pipeline automation (/speckit.auto, 6 stages) - ✓ Quality checkpoints (PrePlanning, PostPlan, PostTasks) - ✓ Real-world workflows (apply command, login flow) - ✓ Configuration hot-reload

**Key Patterns:** - ✓ Temp git repositories (isolated, auto-cleanup) - ✓ Mock HTTP servers (tiny\_http, deterministic) - ✓ Fixture loading (realistic test data) - ✓ State machine validation (initialization, progression, resume)

**Next Steps:** - [Property Testing Guide](#) - Generative invariant testing - [CI/CD Integration](#) - Automated testing pipeline - [Performance Testing](#) - Benchmarks and profiling

---

**References:** - Pipeline E2E: [codex-rs/tui/tests/spec\\_auto\\_e2e.rs](#) - Apply command: [codex-rs/chatgpt/tests/suite/apply\\_command\\_e2e.rs](#) - Login flow: [codex-rs/login/tests/suite/login\\_server\\_e2e.rs](#)

---

## Integration Testing Guide

Comprehensive guide to integration testing across modules.

---

## Overview

**Integration Testing Philosophy:** Test multiple modules working together in realistic workflows

**Goals:** - Verify module interactions - Test cross-cutting concerns (error recovery, state persistence) - Validate end-to-end workflows - Ensure evidence integrity

**Current Status:** - ~200 integration tests (33% of total) - 100% pass rate - Average execution time: ~3-5s per test - Categories: W01-W15 (workflows), E01-E15 (errors), S01-S10 (state), Q01-Q10 (quality), C01-C10 (concurrent)

---

## Integration Test Categories

### W01-W15: Workflow Integration Tests

**Purpose:** Test complete stage workflows across modules

**Flow:** Handler → Consensus → Evidence → Guardrail → State

**Location:** codex-rs/tui/tests/workflow\_integration\_tests.rs

**Coverage:** - W01-W05: Individual stage workflows (Plan, Tasks, Implement, Validate, Audit) - W06-W10: Multi-stage pipelines - W11-W15: Quality gate integration

---

### E01-E15: Error Recovery Integration Tests

**Purpose:** Test error propagation and recovery across modules

**Flow:** Error → Retry → Fallback → Recovery → Evidence

**Location:** codex-rs/tui/tests/error\_recovery\_integration\_tests.rs

**Coverage:** - E01-E05: Consensus and MCP failures - E06-E10: Guardrail validation errors - E11-E15: State corruption and recovery

---

### S01-S10: State Persistence Integration Tests

**Purpose:** Test state coordination with evidence storage

**Flow:** State Change → Evidence Write → Load from Disk → Reconstruct

**Location:** codex-rs/tui/tests/state\_persistence\_integration\_tests.rs

**Coverage:** - S01-S05: State serialization and reconstruction - S06-S10: Pipeline interrupt and resume

---

### Q01-Q10: Quality Gate Integration Tests

**Purpose:** Test quality gate orchestration across modules

**Flow:** Quality Gate → Native Checks → Consensus → Escalation → Guardrail

**Location:** codex-rs/tui/tests/quality\_flow\_integration\_tests.rs

**Coverage:** - Q01-Q05: BeforeSpecify and AfterSpecify gates - Q06-Q10: AfterTasks gate and consensus validation

---

### C01-C10: Concurrent Operations Integration Tests

**Purpose:** Test concurrent stage execution and evidence locking

**Flow:** Parallel Stages → Lock Acquisition → Evidence Writes → Lock Release

**Location:** codex-  
rs/tui/tests/concurrent\_operations\_integration\_tests.rs

**Coverage:** - C01-C05: Parallel consensus collection - C06-C10:  
Evidence write contention

---

## Test Structure

### Standard Integration Test Pattern

```
#[test]
fn w01_plan_stage_complete_workflow() {
    // 1. Setup: Create test context
    let ctx = IntegrationTestContext::new("SPEC-W01-001").unwrap();

    // 2. Arrange: Prepare filesystem (PRD, spec files)
    ctx.write_prd("test-feature", "# Test Feature\nBuild a test feature")
        .unwrap();
    ctx.write_spec("test-feature", "# Specification\nDetailed spec")
        .unwrap();

    // 3. Arrange: Create initial state
    let mut state = StateBuilder::new("SPEC-W01-001")
        .with_goal("Build test feature")
        .starting_at(SpecStage::Plan)
        .build();

    // 4. Act: Simulate module interactions
    // Write mock consensus artifacts (simulating consensus module output)
    let consensus_file = ctx
        .consensus_dir()
        .join("spec-plan_2025-10-19T12_00_00Z_gemini.json");
    std::fs::write(
        &consensus_file,
        json!({
            "agent": "gemini",
            "content": "Plan consensus output",
            "timestamp": "2025-10-19T12:00:00Z"
        })
        .to_string(),
    )
    .unwrap();

    // Write mock guardrail telemetry (simulating guardrail module output)
    let guardrail_file = ctx
        .commands_dir()
        .join("spec-plan_2025-10-19T12_00_00Z.json");
    std::fs::write(
        &guardrail_file,
        json!({
            "schemaVersion": 1,
            "baseline": {"status": "passed"},
```

```

        "tool": {"status": "passed"} ,
    })
    .to_string(),
)
.unwrap();

// 5. Assert: Verify evidence
let verifier = EvidenceVerifier::new(&ctx);
assert!(verifier.assert_structure_valid());
assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));
assert!(ctx.assert_guardrail_telemetry_exists(SpecStage::Plan));

// 6. Assert: Verify state transitions
state.current_index += 1;
assert_eq!(state.current_stage(), Some(SpecStage::Tasks));

// 7. Assert: Verify artifact counts
assert_eq!(ctx.count_consensus_files(), 1);
assert_eq!(ctx.count_guardrail_files(), 1);
}

```

---

## Workflow Integration Tests

### Pattern 1: Individual Stage Workflow

#### Example: W01 - Plan Stage Complete Workflow

**Test** (workflow\_integration\_tests.rs:22):

```

#[test]
fn w01_plan_stage_complete_workflow() {
    let ctx = IntegrationTestContext::new("SPEC-W01-001").unwrap();

    // Arrange: Create PRD and spec
    ctx.write_prd("test-feature", "# Test Feature\nBuild a test
feature")
        .unwrap();
    ctx.write_spec("test-feature", "# Specification\nDetailed spec")
        .unwrap();

    // Arrange: Initial state
    let mut state = StateBuilder::new("SPEC-W01-001")
        .with_goal("Build test feature")
        .starting_at(SpecStage::Plan)
        .build();

    assert_eq!(state.current_stage(), Some(SpecStage::Plan));

    // Act: Simulate consensus module output
    let consensus_file = ctx
        .consensus_dir()
        .join("spec-plan_2025-10-19T12_00_00Z_gemini.json");
    std::fs::write(
        &consensus_file,
        json!({
            "agent": "gemini",
            "content": "Plan consensus output",
        })
    )
}

```

```

        .to_string(),
    )
.unwrap();

// Act: Simulate guardrail module output
let guardrail_file = ctx
    .commands_dir()
    .join("spec-plan_2025-10-19T12_00_00Z.json");
std::fs::write(
    &guardrail_file,
    json!({"schemaVersion": 1, "baseline": {"status": "passed"})})
        .to_string(),
)
.unwrap();

// Assert: Verify evidence
assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));
assert!(ctx.assert_guardrail_telemetry_exists(SpecStage::Plan));

// Assert: Verify state advancement
state.current_index += 1;
assert_eq!(state.current_stage(), Some(SpecStage::Tasks));
}

```

---

## Pattern 2: Multi-Stage Pipeline

### Example: W06 - Plan → Tasks Pipeline

```

#[test]
fn w06_plan_tasks_pipeline() {
    let ctx = IntegrationTestContext::new("SPEC-W06-001").unwrap();

    // Arrange: Initial setup
    ctx.write_prd("multi-stage", "# Multi-stage Test").unwrap();
    let mut state = StateBuilder::new("SPEC-W06-001")
        .starting_at(SpecStage::Plan)
        .build();

    // ===== PLAN STAGE =====

    // Act: Plan stage consensus
    let plan_consensus = ctx
        .consensus_dir()
        .join("spec-plan_2025-10-19T10_00_00Z_gemini.json");
    std::fs::write(
        &plan_consensus,
        json!({"agent": "gemini", "stage": "plan", "content": "Plan
output"})
            .to_string(),
    )
    .unwrap();

    // Assert: Plan evidence exists
    assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));

    // Advance to Tasks
    state.current_index += 1;
    assert_eq!(state.current_stage(), Some(SpecStage::Tasks));
}

```

```

// ===== TASKS STAGE =====

// Act: Tasks stage consensus
let tasks_consensus = ctx
    .consensus_dir()
    .join("spec-tasks_2025-10-19T10_05_00Z_claude.json");
std::fs::write(
    &tasks_consensus,
    json!({"agent": "claude", "stage": "tasks", "content": "Task
list"})
        .to_string(),
)
.unwrap();

// Assert: Tasks evidence exists (accumulated, not replaced)
assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));
assert!(ctx.assert_consensus_exists(SpecStage::Tasks,
"claude"));
assert_eq!(ctx.count_consensus_files(), 2);

// Advance to Implement
state.current_index += 1;
assert_eq!(state.current_stage(), Some(SpecStage::Implement));
}

```

**Key Points:** - ✓ Evidence accumulates across stages (not replaced) -  
✓ State advances sequentially - ✓ Each stage verified independently

---

### Pattern 3: Quality Gate Integration

#### Example: W11 - BeforeSpecify Quality Gate

```

#[test]
fn w11_before_specify_quality_gate_workflow() {
    let ctx = IntegrationTestContext::new("SPEC-W11-001").unwrap();

    // Arrange: Create PRD with known ambiguities
    ctx.write_prd(
        "test",
        r#"
# PRD
## Requirements
- R1: System should be fast
- R2: Must handle TBD authentication
    "#,
    )
    .unwrap();

    let mut state = StateBuilder::new("SPEC-W11-001")
        .quality_gates(true)
        .starting_at(SpecStage::Plan)
        .build();

    // Act: Simulate quality gate execution (Clarify)
    let quality_gate_result = ctx
        .commands_dir()
        .join("quality-gate-clarify_2025-10-19T10_00_00Z.json");
    std::fs::write(

```

```

    &quality_gate_result,
    json!({
        "gate": "BeforeSpecify",
        "checks": ["clarify"],
        "results": {
            "ambiguities": [
                {"pattern": "should", "severity": "Important"},
                {"pattern": "TBD", "severity": "Critical"}
            ]
        },
        "verdict": "escalate", // Critical issues found
        "escalation_reason": "2 ambiguities found (1 critical)"
    })
    .to_string(),
)
.unwrap();

// Assert: Quality gate escalated
let content =
std::fs::read_to_string(&quality_gate_result).unwrap();
let data: serde_json::Value =
serde_json::from_str(&content).unwrap();
assert_eq!(data["verdict"], "escalate");
assert!(data["results"]["ambiguities"]
.as_array()
.unwrap()
.len() > 0);

// State remains at Plan (doesn't advance on escalation)
assert_eq!(state.current_stage(), Some(SpecStage::Plan));
}

```

**Key Points:** - ✓ Quality gate runs before stage - ✓ Escalation prevents advancement - ✓ Evidence records escalation reason

---

## Error Recovery Integration Tests

### Pattern 1: Consensus Failure → Retry → Recovery

#### Example: E01 - Consensus Failure with Retry

**Test** (error\_recovery\_integration\_tests.rs:23):

```

#[test]
fn
e01_consensus_failure_handler_retry_evidence_cleanup_state_reset() {
    let ctx = IntegrationTestContext::new("SPEC-E01-001").unwrap();

    let mut state = StateBuilder::new("SPEC-E01-001")
        .starting_at(SpecStage::Plan)
        .build();

    // ===== ATTEMPT 1: FAILURE =====

    // Act: Write failed consensus (empty result)
    let failed_consensus = ctx
        .consensus_dir()
        .join("spec-plan_2025-10-

```

```

19T10_00_00Z_gemini_attempt1.json");
    std::fs::write(
        &failed_consensus,
        json!({
            "agent": "gemini",
            "stage": "plan",
            "status": "failed",
            "error": "Empty consensus result",
            "attempt": 1,
        })
        .to_string(),
    )
    .unwrap();
}

// Assert: Failed attempt recorded
assert!(failed_consensus.exists());

// ===== RETRY: CLEANUP =====

// Simulate retry: cleanup failed evidence
std::fs::remove_file(&failed_consensus).unwrap();
assert!(!failed_consensus.exists());

// ===== ATTEMPT 2: SUCCESS =====

// Act: Retry with enhanced prompt
let success_consensus = ctx
    .consensus_dir()
    .join("spec-plan_2025-10-
19T10_05_00Z_gemini_attempt2.json");
    std::fs::write(
        &success_consensus,
        json!({
            "agent": "gemini",
            "stage": "plan",
            "status": "success",
            "content": "Enhanced prompt successful",
            "attempt": 2,
            "retry_reason": "empty_result",
        })
        .to_string(),
    )
    .unwrap();

// Assert: Retry succeeded
assert!(success_consensus.exists());
assert_eq!(ctx.count_consensus_files(), 1); // Only successful
attempt remains

// Assert: State advances after successful retry
state.current_index += 1;
assert_eq!(state.current_stage(), Some(SpecStage::Tasks));

// Assert: Evidence shows retry metadata
let content =
    std::fs::read_to_string(&success_consensus).unwrap();
    assert!(content.contains("retry_reason"));
    assert!(content.contains("attempt"));
}

```

**Key Points:** - ✅ Failed attempt recorded as evidence - ✅ Retry cleanup removes failed attempt - ✅ Success includes retry metadata - ✅ State advances only on success

---

## Pattern 2: MCP Failure → Fallback → Recovery

### Example: E02 - MCP Timeout with File Fallback

```
#[test]
fn e02_mcp_failure_fallback_to_file_evidence_recordsFallback() {
    let ctx = IntegrationTestContext::new("SPEC-E02-001").unwrap();

    // ===== MCP FAILURE =====

    // Write fallback marker evidence (MCP failed, using file
    fallback)
    let fallback_evidence = ctx
        .consensus_dir()
        .join("spec-plan_mcp_failure_2025-10-19T10_00_00Z.json");
    std::fs::write(
        &fallback_evidence,
        json!({
            "fallback_mode": "file_based",
            "mcp_error": "Timeout after 60s",
            "fallback_timestamp": "2025-10-19T10:00:00Z"
        })
        .to_string(),
    )
    .unwrap();

    // Assert: Fallback recorded
    assert!(fallback_evidence.exists());

    // ===== FILE-BASED CONSENSUS
    =====

    // Act: Write consensus from file-based fallback
    let file_consensus = ctx
        .consensus_dir()
        .join("spec-plan_2025-10-19T10_00_00Z_file_based.json");
    std::fs::write(
        &file_consensus,
        json!({
            "source": "file_based_fallback",
            "content": "Consensus from local files",
            "degraded": true
        })
        .to_string(),
    )
    .unwrap();

    // Assert: File-based consensus succeeded
    assert!(file_consensus.exists());
    assert_eq!(ctx.count_consensus_files(), 2); // Fallback marker +
    consensus

    // Assert: Degraded flag set
    let content = std::fs::read_to_string(&file_consensus).unwrap();
    assert!(content.contains("\\"degraded\\":true"));
```

```
}
```

**Key Points:** - ✓ MCP failure recorded as fallback evidence - ✓ File-based fallback produces consensus - ✓ Degraded flag indicates fallback mode - ✓ Multiple evidence files coexist

---

## State Persistence Integration Tests

### Pattern 1: State Serialization → Load → Reconstruct

#### Example: S01 - State Persistence and Reconstruction

Test (state\_persistence\_integration\_tests.rs:18):

```
#[test]
fn s01_state_change_evidence_write_load_from_disk_reconstruct() {
    let ctx = IntegrationTestContext::new("SPEC-S01-001").unwrap();
    let state = StateBuilder::new("SPEC-S01-001")
        .starting_at(SpecStage::Plan)
        .build();

    // ===== SERIALIZE STATE =====

    // Act: Write state to evidence
    let state_file =
        ctx.commands_dir().join("spec_auto_state.json");
    std::fs::write(
        &state_file,
        json!({
            "spec_id": state.spec_id,
            "current_index": state.current_index,
            "quality_gates_enabled": state.quality_gates_enabled,
        })
        .to_string(),
    )
    .unwrap();

    // ===== LOAD AND RECONSTRUCT
    =====

    // Act: Load from disk and verify reconstruction
    let loaded = std::fs::read_to_string(&state_file).unwrap();
    let data: serde_json::Value =
        serde_json::from_str(&loaded).unwrap();

    // Assert: All fields preserved
    assert_eq!(data["spec_id"], "SPEC-S01-001");
    assert_eq!(data["current_index"], 0);
    assert_eq!(data["quality_gates_enabled"], true);

    // Reconstruct state from loaded data
    let reconstructed =
        StateBuilder::new(data["spec_id"].as_str().unwrap())
            .starting_at(SpecStage::Plan)

            .quality_gates(data["quality_gates_enabled"].as_bool().unwrap())
            .build();
}
```

```
        assert_eq!(reconstructed.spec_id, state.spec_id);
        assert_eq!(reconstructed.current_index, state.current_index);
    }
```

---

## Pattern 2: Pipeline Interrupt → Resume from Checkpoint

### Example: S02 - Pipeline Interrupt and Resume

Test (state\_persistence\_integration\_tests.rs:45):

```
#[test]
fn s02_pipeline_interrupt_state_saved_resume_from_checkpoint() {
    let ctx = IntegrationTestContext::new("SPEC-S02-001").unwrap();
    let mut state = StateBuilder::new("SPEC-S02-001")
        .starting_at(SpecStage::Tasks)
        .build();

    // ===== SAVE CHECKPOINT =====

    // Act: Save checkpoint before interrupt
    let checkpoint = ctx.commands_dir().join("checkpoint.json");
    std::fs::write(
        &checkpoint,
        json!({
            "spec_id": state.spec_id,
            "checkpoint_index": state.current_index,
            "timestamp": "2025-10-19T10:00:00Z"
        })
        .to_string(),
    )
    .unwrap();

    assert_eq!(state.current_index, 1); // Tasks = index 1

    // ===== INTERRUPT =====

    // Simulate interrupt (state dropped)
    drop(state);

    // ===== RESUME =====

    // Act: Resume from checkpoint
    let loaded = std::fs::read_to_string(&checkpoint).unwrap();
    let data: serde_json::Value =
        serde_json::from_str(&loaded).unwrap();

    let resumed_state = StateBuilder::new("SPEC-S02-001")
        .starting_at(SpecStage::Plan)
        .build();

    // Assert: Checkpoint index preserved
    assert_eq!(data["checkpoint_index"], 1);
    assert_eq!(data["spec_id"], "SPEC-S02-001");

    // Resume would set current_index from checkpoint
    // (not shown: actual resume logic would apply checkpoint)
}
```

---

# Evidence Verification Patterns

## Pattern 1: Comprehensive Evidence Verification

```
#[test]
fn verify_complete_stage_evidence() {
    let ctx = IntegrationTestContext::new("SPEC-TEST").unwrap();

    // Simulate complete stage execution
    // ... (write consensus and guardrail artifacts)

    // ===== VERIFY STRUCTURE =====

    let verifier = EvidenceVerifier::new(&ctx);

    // Directory structure
    assert!(verifier.assert_structure_valid());

    // ===== VERIFY CONSENSUS =====

    // All agents present
    assert!(verifier.assert_consensus_complete(
        SpecStage::Plan,
        &["gemini", "claude", "gpt_pro"]
    ));

    // Individual agents
    assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));
    assert!(ctx.assert_consensus_exists(SpecStage::Plan, "claude"));
    assert!(ctx.assert_consensus_exists(SpecStage::Plan,
        "gpt_pro"));

    // ===== VERIFY GUARDRAIL =====

    assert!(
        verifier.assert_guardrail_valid(SpecStage::Plan).is_ok()
    );

    // ===== VERIFY COUNTS =====

    assert_eq!(ctx.count_consensus_files(), 3);
    assert_eq!(ctx.count_guardrail_files(), 1);
}
```

---

## Pattern 2: Degraded Consensus Detection

```
#[test]
fn verify_degraded_consensus() {
    let ctx = IntegrationTestContext::new("SPEC-TEST").unwrap();

    // Simulate degraded consensus (only 2/3 agents)
    // ... (write only gemini and claude consensus)

    let verifier = EvidenceVerifier::new(&ctx);

    // Should NOT be complete (missing gpt_pro)
    assert!(!verifier.assert_consensus_complete(
        SpecStage::Plan,
        &["gemini", "claude", "gpt_pro"]
```

```

));
// But 2/3 is still valid
assert!(verifier.assert_consensus_complete(
    SpecStage::Plan,
    &["gemini", "claude"]
));

// Verify degraded flag
let consensus = ctx
    .consensus_dir()
    .join("spec-plan_2025-10-19T10_00_00Z_synthesis.json");
std::fs::write(
    &consensus,
    json!({ "consensus_ok": true, "degraded": true,
"missing_agents": ["gpt_pro"] })
        .to_string(),
)
.unwrap();

let content = std::fs::read_to_string(&consensus).unwrap();
assert!(content.contains("\"degraded\":true"));
}

```

---

## Best Practices

### DO

✓ **Use IntegrationTestContext for isolation:**

```

#[test]
fn test_workflow() {
    // Each test gets isolated filesystem
    let ctx = IntegrationTestContext::new("SPEC-TEST-001").unwrap();
    // ... test logic
}

```

---

✓ **Verify evidence at each step:**

```

// After consensus
assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));

// After guardrail
assert!(ctx.assert_guardrail_telemetry_exists(SpecStage::Plan));

// After completion
assert_eq!(ctx.count_consensus_files(), 3);

```

---

✓ **Test both success and failure paths:**

```

#[test]
fn test_success_path() {
    // Happy path
}

#[test]

```

```
fn test_failure_path_with_retry() {
    // Error → Retry → Success
}

#[test]
fn test_failure_path_exhausted_retries() {
    // Error → Retry → Retry → Fail
}
```

---

✓ **Simulate realistic timing:**

```
let timestamp_attempt1 = "2025-10-19T10:00:00Z";
let timestamp_retry = "2025-10-19T10:05:00Z"; // 5 minutes later

// Evidence shows temporal sequence
```

---

✓ **Verify state transitions:**

```
assert_eq!(state.current_stage(), Some(SpecStage::Plan));

// Execute stage...

state.current_index += 1;
assert_eq!(state.current_stage(), Some(SpecStage::Tasks));
```

---

## DON'T

✗ **Share IntegrationTestContext across tests:**

```
// Bad: Shared context (tests interfere)
static mut CTX: Option<IntegrationTestContext> = None;

#[test]
fn test_a() {
    unsafe { CTX =
Some(IntegrationTestContext::new("SHARED").unwrap()); }
}

#[test]
fn test_b() {
    unsafe { /* use CTX */ } // ✗ Flaky (depends on test_a)
}
```

---

✗ **Test too many stages in one test:**

```
// Bad: Tests entire pipeline (hard to debug failures)
#[test]
fn test_entire_pipeline() {
    // Plan
    // Tasks
    // Implement
    // Validate
    // Audit
    // Unlock
    // → 200 lines, hard to maintain
}
```

```
// Good: Split into focused tests
#[test]
fn w01_plan_stage_workflow() { /* ... */ }

#[test]
fn w02_tasks_stage_workflow() { /* ... */ }
```

---

### ✗ Skip evidence verification:

```
// Bad: No verification
#[test]
fn test_workflow() {
    // Run workflow...
    // No assertions ✗
}

// Good: Verify evidence
#[test]
fn test_workflow() {
    // Run workflow...
    assert!(ctx.assert_consensus_exists(...));
    assert!(ctx.assert_guardrail_telemetry_exists(...));
}
```

---

### ✗ Use hard-coded paths:

```
// Bad: Hard-coded paths (breaks on other machines)
let consensus = Path::new("/tmp/consensus/SPEC-TEST/plan.json");

// Good: Use IntegrationTestContext
let consensus = ctx.consensus_dir().join("plan.json");
```

---

## Running Integration Tests

### Run All Integration Tests

```
cd codex-rs
cargo test --test '*_integration_tests'
```

**Runs:** - workflow\_integration\_tests.rs -  
error\_recovery\_integration\_tests.rs -  
state\_persistence\_integration\_tests.rs -  
quality\_flow\_integration\_tests.rs -  
concurrent\_operations\_integration\_tests.rs

---

### Run Specific Category

```
# Workflow tests only
cargo test --test workflow_integration_tests

# Error recovery tests only
cargo test --test error_recovery_integration_tests
```

---

## Run Specific Test

```
cargo test --test workflow_integration_tests  
w01_plan_stage_complete_workflow
```

---

## Run with Output

```
cargo test --test workflow_integration_tests -- --nocapture
```

Shows `println!()` output for debugging.

---

## Summary

### Integration Testing Best Practices:

1. **Isolation:** Use `IntegrationTestContext` for each test
2. **Evidence:** Verify evidence at each step
3. **Coverage:** Test success and failure paths
4. **Clarity:** One workflow per test
5. **Timing:** Simulate realistic sequences
6. **State:** Verify state transitions
7. **Cleanup:** Automatic (`TempDir` drops)

**Test Categories:** - ✓ W01-W15: Workflow integration (stage workflows, pipelines) - ✓ E01-E15: Error recovery (retry, fallback, degradation) - ✓ S01-S10: State persistence (serialize, resume, checkpoint) - ✓ Q01-Q10: Quality gates (BeforeSpecify, AfterSpecify, AfterTasks) - ✓ C01-C10: Concurrent operations (parallel, locking)

**Key Patterns:** - ✓ Multi-module workflows (Handler → Consensus → Evidence → Guardrail → State) - ✓ Error propagation (Failure → Retry → Recovery → Evidence) - ✓ State persistence (Serialize → Load → Reconstruct) - ✓ Evidence verification (EvidenceVerifier, counts, structure)

**Next Steps:** - [E2E Testing Guide](#) - Complete user workflows - [Property Testing Guide](#) - Generative invariant testing - [Test Infrastructure](#) - MockMcpManager, fixtures

---

**References:** - Workflow tests: `codex-rs/tui/tests/workflow_integration_tests.rs` - Error recovery: `codex-rs/tui/tests/error_recovery_integration_tests.rs` - State persistence: `codex-rs/tui/tests/state_persistence_integration_tests.rs` - `IntegrationTestContext`: `codex-rs/tui/tests/common/integration_harness.rs`

---

## Performance Testing Guide

Comprehensive guide to performance testing, benchmarking, and profiling.

---

# Overview

**Performance Testing Philosophy:** Measure, don't guess. Validate optimizations with data.

**Goals:** - Measure baseline performance - Validate optimizations - Detect regressions - Identify bottlenecks

**Tools:** - **criterion**: Statistical benchmarking - **cargo-flamegraph**: Profiling - **cargo-bloat**: Binary size analysis - **hyperfine**: Command-line benchmarking

**Current Benchmarks:** - Database performance ( $6.6\times$  read speedup validated) - MCP client ( $5.3\times$  faster than subprocess validated) - Connection pooling (R2D2)

---

## Benchmarking with Criterion

### What is Criterion?

**Criterion** is a statistical benchmarking tool for Rust that provides: - Accurate measurements (micro/nanosecond precision) - Statistical analysis (mean, stddev, outliers) - Regression detection (compare to baseline) - HTML reports with charts

**Website:** <https://bheisler.github.io/criterion.rs/>

---

### Setup

#### Add to Cargo.toml:

```
[dev-dependencies]
criterion = { version = "0.5", features = ["html_reports"] }

[[bench]]
name = "my_benchmark"
harness = false
```

---

### Basic Benchmark

**File:** benches/simple\_benchmark.rs

```
use criterion::{Criterion, black_box, criterion_group,
criterion_main};

fn fibonacci(n: u64) -> u64 {
    match n {
        0 => 1,
        1 => 1,
        n => fibonacci(n - 1) + fibonacci(n - 2),
    }
}

fn benchmark_fibonacci(c: &mut Criterion) {
    c.bench_function("fib 20", |b| {
```

```

        b.iter(|| fibonacci(black_box(20)));
    });
}

criterion_group!(benches, benchmark_fibonacci);
criterion_main!(benches);

```

### Run:

```
cargo bench --bench simple_benchmark
```

### Output:

```

fib 20           time:  [26.029 µs 26.251 µs 26.509 µs]
Found 11 outliers among 100 measurements (11.00%)
  6 (6.00%) high mild
  5 (5.00%) high severe

```

---

## Database Performance Benchmark

**Example:** codex-rs/core/benches/db\_performance.rs

**Performance Targets:** - Before: 850µs/read, 2.1ms/write, 78ms/100-read batch  
- After: 129µs/read, 0.9ms/write, 12ms/100-read batch  
- Improvement: 6.6× read, 2.3× write, 6.5× batch

---

### Benchmark Setup

```

use criterion::{Criterion, Throughput, black_box, criterion_group,
criterion_main};
use codex_core::db::initialize_pool;
use tempfile::TempDir;

/// Create temporary database with schema
fn setup_temp_db() -> (TempDir, PathBuf) {
    let temp_dir = TempDir::new().expect("Failed to create temp
dir");
    let db_path = temp_dir.path().join("test.db");

    let conn = Connection::open(&db_path).expect("Failed to open
connection");
    conn.execute_batch(
        "CREATE TABLE IF NOT EXISTS consensus_runs (
            id INTEGER PRIMARY KEY,
            spec_id TEXT NOT NULL,
            stage TEXT NOT NULL,
            consensus_ok INTEGER NOT NULL,
            created_at INTEGER NOT NULL
        );
        CREATE INDEX IF NOT EXISTS idx_spec_stage ON
consensus_runs(spec_id, stage);"
    )
    .expect("Failed to create schema");

    (temp_dir, db_path)
}

/// Create connection pool with WAL mode

```

```

fn setup_pool(db_path: &PathBuf) -> Pool<SqliteConnectionManager> {
    initialize_pool(db_path, 10).expect("Failed to initialize pool")
}

```

---

## Benchmark #1: Connection Pool vs Single Connection

```

fn benchmark_connection_pool_vs_single(c: &mut Criterion) {
    let mut group = c.benchmark_group("connection_pool_vs_single");

    // Setup: Create database with test data
    let (_temp_dir, db_path) = setup_temp_db();
    let pool = setup_pool(&db_path);

    // Insert 1000 test records
    {
        let conn = pool.get().expect("Failed to get connection");
        insert_test_data(&conn, 1000);
    }

    // Benchmark: Pooled connection reads
    group.bench_function("pooled_connection_read", |b| {
        b.iter(|| {
            let conn = pool.get().expect("Failed to get connection");
            let mut stmt = conn
                .prepare("SELECT * FROM consensus_runs WHERE spec_id
= ?1")
                .expect("Failed to prepare statement");
            let _count = stmt
                .query_map(["SPEC-TEST-050"], |_row| Ok(()))
                .expect("Failed to query")
                .count();
            black_box(_count);
        });
    });

    // Benchmark: Single connection reads (reused connection)
    group.bench_function("single_connection_read", |b| {
        let conn = setup_single_connection_wal(&db_path);
        b.iter(|| {
            let mut stmt = conn
                .prepare("SELECT * FROM consensus_runs WHERE spec_id
= ?1")
                .expect("Failed to prepare statement");
            let _count = stmt
                .query_map(["SPEC-TEST-050"], |_row| Ok(()))
                .expect("Failed to query")
                .count();
            black_box(_count);
        });
    });

    group.finish();
}

```

### Results:

```

connection_pool_vs_single/pooled_connection_read
    time: [129.45 µs 130.12 µs 130.89 µs]

```

```
connection_pool_vs_single/single_connection_read
    time: [127.89 µs 128.45 µs 129.12 µs]
```

**Analysis:** - ✓ Pool overhead minimal (~1-2µs) - ✓ Both achieve target (<150µs vs 850µs before) - ✓ 6.6× improvement validated

---

## Benchmark #2: WAL Mode Impact

```
fn benchmark_wal_mode_impact(c: &mut Criterion) {
    let mut group = c.benchmark_group("wal_mode_impact");

    let (_temp_dir, db_path) = setup_temp_db();

    // Setup: Connection with WAL mode
    let conn_wal = setup_single_connection_wal(&db_path);
    insert_test_data(&conn_wal, 1000);

    // Setup: Connection with DELETE mode (no WAL)
    let (_temp_dir2, db_path2) = setup_temp_db();
    let conn_delete = setup_single_connection_delete(&db_path2);
    insert_test_data(&conn_delete, 1000);

    // Benchmark: Read with WAL
    group.bench_function("read_wal", |b| {
        b.iter(|| {
            let mut stmt = conn_wal
                .prepare("SELECT * FROM consensus_runs WHERE spec_id
= ?1")
                .unwrap();
            black_box(stmt.query_map(["SPEC-TEST-050"], |_
Ok(())).unwrap().count());
        });
    });

    // Benchmark: Read with DELETE mode
    group.bench_function("read_delete", |b| {
        b.iter(|| {
            let mut stmt = conn_delete
                .prepare("SELECT * FROM consensus_runs WHERE spec_id
= ?1")
                .unwrap();
            black_box(stmt.query_map(["SPEC-TEST-050"], |_
Ok(())).unwrap().count());
        });
    });

    group.finish();
}
```

### Results:

```
wal_mode_impact/read_wal
    time: [129.12 µs 130.45 µs 131.89 µs]
```

```
wal_mode_impact/read_delete
    time: [847.34 µs 851.23 µs 856.78 µs]
```

Improvement: 6.58× faster with WAL ✓

---

## Throughput Benchmarks

**Pattern:** Measure operations per second

```
fn benchmark_batch_reads(c: &mut Criterion) {
    let mut group = c.benchmark_group("batch_reads");

    let (_temp_dir, db_path) = setup_temp_db();
    let pool = setup_pool(&db_path);
    let conn = pool.get().unwrap();
    insert_test_data(&conn, 1000);

    // Benchmark 100 reads (measure throughput)
    group.throughput(Throughput::Elements(100));
    group.bench_function("read_100", |b| {
        b.iter(|| {
            for i in 0..100 {
                let conn = pool.get().unwrap();
                let mut stmt = conn.prepare("SELECT * FROM
consensus_runs WHERE spec_id = ?1").unwrap();
                let _count = stmt.query_map([format!("SPEC-TEST-
{:03}", i % 100)], |_| Ok(())).unwrap().count();
                black_box(_count);
            }
        });
    });

    group.finish();
}
```

### Results:

```
batch_reads/read_100      time: [12.234 ms 12.456 ms 12.689 ms]
                           thrpt: [7.88 Kelem/s 8.03 Kelem/s 8.17
Kelem/s]
```

```
Before optimization: 78ms/100 reads → 1.28 Kelem/s
After optimization: 12ms/100 reads → 8.03 Kelem/s
Improvement: 6.27x faster ✓
```

---

## Running Benchmarks

### Run all benchmarks:

```
cd codex-rs
cargo bench
```

### Run specific benchmark:

```
cargo bench --bench db_performance
```

### Run specific function:

```
cargo bench --bench db_performance -- connection_pool
```

### Generate baseline (for regression detection):

```
cargo bench -- --save-baseline baseline_2025_11_17
```

### Compare to baseline:

```
cargo bench --baseline baseline_2025_11_17
```

#### View HTML reports:

```
open target/criterion/report/index.html
```

---

## Profiling

### Flamegraphs with cargo-flamegraph

**What are Flamegraphs?**: - Visual representation of stack traces - Shows where CPU time is spent - Width = time spent in function - Height = call stack depth

#### Install:

```
cargo install flamegraph
```

#### Usage:

```
# Profile specific benchmark
cargo flamegraph --bench db_performance --baseline baseline_2025_11_17

# Profile specific test
cargo flamegraph --test integration_test

# Profile binary
cargo flamegraph --bin code
```

**Output:** flamegraph.svg (interactive SVG)

**Interpretation:** - **Wide bars**: Hot paths (optimize these) - **Narrow bars**: Not worth optimizing - **Tall stacks**: Deep call chains

---

### perf (Linux only)

#### Install:

```
sudo apt install linux-tools-generic
```

#### Record:

```
cargo build --release
perf record --call-graph=dwarf ./target/release/code
```

#### Analyze:

```
perf report
```

#### Generate Flamegraph:

```
perf script | stackcollapse-perf.pl | flamegraph.pl > perf.svg
```

---

### cargo-bloat (Binary Size Analysis)

**Purpose:** Find large dependencies

### Install:

```
cargo install cargo-bloat
```

### Usage:

```
cd codex-rs
cargo bloat --release
```

### Output:

File	.text	Size	Crate
0.7%	1.2%	24.5KiB	regex
0.6%	1.0%	20.1KiB	serde_json
0.5%	0.9%	18.7KiB	tokio
...			

### Optimize (if needed):

```
# Cargo.toml
[profile.release]
lto = true                                # Link-time optimization
codegen-units = 1                           # Better optimization, slower build
strip = true                               # Strip symbols
opt-level = "z"                            # Optimize for size
```

---

## Command-Line Benchmarking

### hyperfine

**Purpose:** Benchmark CLI commands

### Install:

```
cargo install hyperfine
```

### Usage:

```
# Benchmark single command
hyperfine './codex-rs/target/release/code --version'

# Compare commands
hyperfine \
    './codex-rs/target/release/code doctor' \
    './codex-rs/target/dev-fast/code doctor'

# Warmup runs
hyperfine --warmup 3 'cargo test'

# Multiple runs
hyperfine --runs 100 './codex-rs/target/release/code --help'
```

### Example Output:

```
Benchmark 1: ./target/release/code --version
Time (mean ± σ):     12.3 ms ±   0.5 ms    [User: 8.2 ms, System:
3.1 ms]
Range (min ... max): 11.5 ms ... 14.2 ms    100 runs
```

---

## Benchmarking /speckit.auto

### Example:

```
hyperfine --warmup 1 --runs 5 \
    './codex-rs/target/release/code run "/speckit.auto SPEC-TEST-001"'
```

### Expected:

```
Time (mean ± σ):     45.2 s ±  2.1 s    [User: 38.1 s, System: 3.2
s]
Range (min ... max): 42.8 s ... 48.5 s    5 runs
```

---

## Performance Metrics

### Database Performance

**Measured Metrics:** - Read latency (μs): 850 → 129 (6.6× improvement)  
- Write latency (ms): 2.1 → 0.9 (2.3× improvement)  
- Batch reads (ms/100): 78 → 12 (6.5× improvement)

### How Measured:

```
// codex-rs/core/benches/db_performance.rs
criterion_group!(benches,
    benchmark_connection_pool_vs_single,
    benchmark_wal_mode_impact,
    benchmark_batch_reads,
);
```

---

### MCP Performance

**Measured Metrics:** - Native MCP client: 8.7ms typical - Subprocess MCP: 46ms typical - Improvement: 5.3× faster

### How Measured:

```
// Integration test timing
let start = std::time::Instant::now();
let result = mcp_client.call_tool(...).await?;
let elapsed = start.elapsed();
assert!(elapsed < Duration::from_millis(10)); // <10ms
```

---

### Config Hot-Reload

**Measured Metrics:** - Reload latency (p95): <100ms - File watch overhead: <1% CPU

### How Measured:

```
// Integration test
let start = std::time::Instant::now();
// Modify config file
std::fs::write(&config_path, new_content)?;
// Wait for reload
```

```
tokio::time::sleep(Duration::from_millis(50)).await;
// Verify reload
assert_eq!(app.current_model(), "gpt-5-medium");
let elapsed = start.elapsed();
assert!(elapsed < Duration::from_millis(100));
```

---

## Regression Testing

### Baseline Comparison

#### Save baseline:

```
cargo bench -- --save-baseline v1.0.0
```

#### Compare:

```
# After changes
cargo bench -- --baseline v1.0.0
```

#### Output:

```
connection_pool_vs_single/pooled_connection_read
    time: [129.45 µs 130.12 µs 130.89 µs]
    change: [-0.5% +0.2% +1.1%] (p = 0.23 >
0.05)
No change in performance detected.
```

**Interpretation:** - Change <5%: No regression - Change >5%:  
Investigate - Change >10%: **Regression detected** (fix before merge)

---

## Continuous Performance Monitoring

### CI Integration (future):

```
# .github/workflows/performance.yml
jobs:
  benchmark:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4

      - name: Run benchmarks
        run: cargo bench -- --save-baseline ci-baseline

      - name: Compare to previous
        run: cargo bench -- --baseline ci-baseline-previous

      - name: Fail if regression >10%
        run:
          if grep "change:[+][1-9][0-9]"
target/criterion/**/new/estimates.txt; then
            echo "Performance regression detected!"
            exit 1
        fi
```

---

# Best Practices

## DO

### ✓ Measure before optimizing:

```
# Before: Measure baseline
cargo bench -- --save-baseline before_optimization

# Optimize code...

# After: Measure improvement
cargo bench -- --baseline before_optimization
```

---

### ✓ Use black\_box() to prevent optimization:

```
// Good: Prevents compiler from optimizing away
b.iter(|| {
    black_box(expensive_function(black_box(input)));
});

// Bad: Compiler might optimize this away
b.iter(|| {
    expensive_function(input);
});
```

---

### ✓ Benchmark realistic workloads:

```
// Good: Real-world data
let data = load_fixture("real_prd.md");
b.iter(|| detect_ambiguities(black_box(&data)));

// Bad: Trivial input
let data = "test";
b.iter(|| detect_ambiguities(black_box(&data)));
```

---

### ✓ Run benchmarks on consistent hardware:

- Same machine (or CI runner)
- Close other programs
- Disable CPU frequency scaling (if possible)

---

### ✓ Set performance targets:

```
// Document targets in benchmark comments
/// Target: <150µs (was 850µs before optimization)
group.bench_function("pooled_read", |b| { ... });
```

---

## DON'T

### ✗ Optimize without measuring:

```
// Bad: Premature optimization
// "This looks slow, let me rewrite it"

// Good: Measure first
// cargo bench → identify hot path → optimize
```

---

### **✗ Trust microbenchmarks for macro performance:**

```
// Bad: Optimizing single function
fn fast_function() { /* 1μs faster */ }

// Better: Benchmark complete workflow
fn complete_pipeline() { /* Does 1μs matter here? */ }
```

---

### **✗ Ignore variance:**

```
# Bad: "It ran in 10ms once"
# Good: "Mean: 10.2ms ± 0.3ms (100 runs)"
```

---

### **✗ Benchmark in debug mode:**

```
# Bad: Debug mode (100x slower)
cargo bench

# Good: Release mode (default for benches)
cargo bench --release
```

---

## **Summary**

### **Performance Testing Best Practices:**

1. **Measure:** Use criterion for accurate benchmarks
2. **Profile:** Use flamegraphs to find hot paths
3. **Validate:** Confirm optimizations with data
4. **Regress:** Detect performance regressions
5. **Target:** Set clear performance goals

**Tools:** - ✓ criterion (statistical benchmarking) - ✓ cargo-flamegraph (profiling) - ✓ cargo-bloat (binary size) - ✓ hyperfine (CLI benchmarking) - ✓ perf (Linux profiling)

**Validated Improvements:** - ✓ Database: 6.6× read, 2.3× write - ✓ MCP: 5.3× faster (8.7ms vs 46ms) - ✓ Config reload: <100ms (p95)

**Key Metrics:** - ✓ Latency (μs, ms, s) - ✓ Throughput (ops/sec, elem/sec) - ✓ Percentiles (p50, p95, p99) - ✓ Variance (stddev, outliers)

**Next Steps:** - Testing Strategy - Overall testing approach - CI/CD Integration - Automated testing - Test Infrastructure - MockMcpManager, fixtures

---

**References:** - criterion: <https://bheisler.github.io/criterion.rs/> - Database benchmarks: [codex-rs.core/benches/db\\_performance.rs](https://codex-rs.core/benches/db_performance.rs) - Profiling guide: <https://nnethercote.github.io/perf-book/> - hyperfine: <https://github.com/sharkdp/hyperfine>

---

## **Property-Based Testing Guide**

Comprehensive guide to property-based testing with proptest.

---

## Overview

**Property-Based Testing Philosophy:** Generate random inputs to verify invariants hold across all possible values

**Tool:** `proptest` (Rust equivalent of QuickCheck/Hypothesis)

**Goals:** - Test invariants (properties that always hold) - Find edge cases automatically - Verify mathematical properties - Reduce test boilerplate

**Current Status:** - ~30 property-based tests - 100% pass rate - 100 test cases per property (default) - Integrated with standard test suite

---

## What is Property-Based Testing?

### Traditional Example-Based Testing

```
#![test]
fn test_reverse_twice_is_identity() {
    let vec = vec![1, 2, 3];
    let reversed = reverse(reverse(vec.clone()));
    assert_eq!(reversed, vec);
}
```

**Limitations:** - Only tests one input ([1, 2, 3]) - May miss edge cases (empty, single element, duplicates) - Requires manual case selection

---

### Property-Based Testing

```
use proptest::prelude::*;

proptest! {
    #[test]
    fn test_reverse_twice_is_identity(vec in any::<Vec<i32>>()) {
        let reversed = reverse(reverse(vec.clone()));
        prop_assert_eq!(reversed, vec);
    }
}
```

**Benefits:** - ✓ Tests 100 random inputs automatically - ✓ Finds edge cases (empty, single, large, etc.) - ✓ Shrinks failing input to minimal case - ✓ Focuses on **properties** not **examples**

---

## Getting Started

### Add proptest Dependency

**Cargo.toml:**

```
[dev-dependencies]
proptest = "1.3"
```

---

## Basic Property Test

```
use proptest::prelude::*;

proptest! {
    #[test]
    fn test_addition_commutative(a in any::<i32>(), b in any::<i32>())
    {
        // Property: a + b == b + a
        prop_assert_eq!(a + b, b + a);
    }
}
```

**How it works:** 1. Generate 100 random pairs of (a, b) 2. Run test with each pair 3. If any fails, shrink to minimal failing case 4. Report failure with minimal input

---

## Generators

### Built-in Generators

**Primitive Types:**

```
proptest! {
    #[test]
    fn test_primitives(
        n in any::<i32>(),
        s in any::<String>(),
        b in any::<bool>(),
    ) {
        // Test with random primitives
    }
}
```

**Collections:**

```
proptest! {
    #[test]
    fn test_collections(
        vec in any::<Vec<i32>>(),
        set in any::<HashSet<String>>(),
        map in any::<HashMap<i32, String>>(),
    ) {
        // Test with random collections
    }
}
```

**Ranges:**

```
proptest! {
    #[test]
    fn test_ranges(
```

```

        index in 0usize..10,           // 0-9
        score in 0.0..100.0,          // 0.0-99.999...
        percentage in 0..=100,         // 0-100 (inclusive)
    ) {
        prop_assert!(index < 10);
        prop_assert!(score < 100.0);
        prop_assert!(percentage <= 100);
    }
}

```

---

## Custom Generators

### Regex Patterns:

```

proptest! {
    #[test]
    fn test_spec_id_format(
        spec_id in "[A-Z]{4}-[A-Z]{3}-[0-9]{3}"
    ) {
        // Generates: "SPEC-KIT-001", "ABCD-XYZ-999", etc.
        prop_assert!(is_valid_spec_id(&spec_id));
    }
}

```

---

### Custom Strategies:

```

fn spec_stage_strategy() -> impl Strategy<Value = SpecStage> {
    prop_oneof![
        Just(SpecStage::Plan),
        Just(SpecStage::Tasks),
        Just(SpecStage::Implement),
        Just(SpecStage::Validate),
        Just(SpecStage::Audit),
        Just(SpecStage::Unlock),
    ]
}

proptest! {
    #[test]
    fn test_stage_valid(stage in spec_stage_strategy()) {
        // Tests all 6 stages
        prop_assert!(is_valid_stage(&stage));
    }
}

```

---

## Testing Invariants

### Invariant 1: State Index Always Valid

**Property:** State index  $\in [0, 5] \rightarrow \text{current\_stage}() \text{ returns } \text{Some}(\_), \text{ else None}$

**Test** (property\_based\_tests.rs:21):

```

proptest! {
    #[test]

```

```

fn pb01_state_index_always_in_valid_range(index in 0usize..20) {
    let mut state = StateBuilder::new("SPEC-PB01-TEST")
        .starting_at(SpecStage::Plan)
        .build();

    state.current_index = index;

    // Invariant: index ∈ [0, 5] → Some(_), else None
    if index < 6 {
        prop_assert!(state.current_stage().is_some());
    } else {
        prop_assert_eq!(state.current_stage(), None);
    }
}
}

```

**What This Tests:** - ✓ All indices 0-19 handled correctly - ✓ Valid indices (0-5) return Some - ✓ Invalid indices (6+) return None - ✓ No panics or crashes

---

## Invariant 2: Current Stage Mapping

**Property:** For index ∈ [0, 5], current\_stage() returns correct stage

**Test** (property\_based\_tests.rs:38):

```

proptest! {
    #[test]
    fn pb02_current_stage_always_some_when_index_under_six(
        index in 0usize..6
    ) {
        let mut state = StateBuilder::new("SPEC-PB02-TEST").build();
        state.current_index = index;

        prop_assert!(state.current_stage().is_some());

        // Verify correct stage mapping
        let expected_stages = vec![
            SpecStage::Plan,
            SpecStage::Tasks,
            SpecStage::Implement,
            SpecStage::Validate,
            SpecStage::Audit,
            SpecStage::Unlock,
        ];

        prop_assert_eq!(
            state.current_stage(),
            Some(expected_stages[index])
        );
    }
}

```

**What This Tests:** - ✓ All valid indices (0-5) return Some - ✓ Correct stage for each index - ✓ Consistent mapping

---

## Invariant 3: Retry Count Never Negative

**Property:** Retry count  $\leq$  max\_retries (capped at max)

**Test** (property\_based\_tests.rs:62):

```
proptest! {
    #[test]
    fn pb03_retry_count_never_negative(retries in 0usize..100) {
        let ctx = IntegrationTestContext::new("SPEC-PB03-TEST").unwrap();

        let max_retries = 3;
        let capped_retries = retries.min(max_retries);

        let retry_file = ctx.commands_dir().join("retry.json");
        std::fs::write(&retry_file, json!({
            "retry_count": capped_retries,
            "max_retries": max_retries,
            "within_limit": capped_retries <= max_retries
        }).to_string()).unwrap();

        let content = std::fs::read_to_string(&retry_file).unwrap();
        let data: serde_json::Value =
            serde_json::from_str(&content).unwrap();

        prop_assert!(data["retry_count"].as_u64().unwrap() <=
            max_retries as u64);
        prop_assert_eq!(data["within_limit"].as_bool(), Some(true));
    }
}
```

**What This Tests:** - ✓ Retry counts 0-99 all capped correctly - ✓ No retry count exceeds max - ✓ within\_limit flag always true

---

## Testing Evidence Integrity

### Property 1: Written Evidence Always Parseable JSON

**Property:** Any evidence written is valid JSON

**Test** (property\_based\_tests.rs:90):

```
proptest! {
    #[test]
    fn pb04_written_evidence_always_parseable_json(
        agent in "[a-z]{3,10}",
        content in ".*"
    ) {
        let ctx = IntegrationTestContext::new("SPEC-PB04-TEST").unwrap();

        let evidence = json!({
            "agent": agent,
            "content": content,
            "timestamp": "2025-10-19T00:00:00Z"
        });

        let file = ctx.consensus_dir().join("test.json");
        std::fs::write(&file, evidence.to_string()).unwrap();
    }
}
```

```

    // Invariant: File is valid JSON
    let content = std::fs::read_to_string(&file).unwrap();
    let parsed: Value = serde_json::from_str(&content).unwrap();

    prop_assert_eq!(parsed["agent"].as_str(),
Some(agent.as_str()));
}
}

```

**What This Tests:** - ✓ Random agent names (3-10 lowercase letters) - ✓ Random content (any string) - ✓ Always produces valid JSON - ✓ Round-trip serialization works

---

## Property 2: Evidence File Names Valid

**Property:** Generated filenames are valid filesystem paths

```

proptest! {
    #[test]
    fn pb05_evidence_filenames_always_valid(
        spec_id in "[A-Z]{4}-[A-Z]{3}-[0-9]{3}",
        stage in spec_stage_strategy(),
        agent in "[a-z]{5,10}",
    ) {
        let filename = format!(
            "spec-{:?}_{:}_{:}_{:}.json",
            stage,
            spec_id,
            "2025-10-19T10_00_00Z",
            agent
        );

        // Invariant: Filename contains no invalid characters
        prop_assert!(!filename.contains('/'));
        prop_assert!(!filename.contains('\\'));
        prop_assert!(!filename.contains('\0'));

        // Invariant: Filename is not empty
        prop_assert!(!filename.is_empty());

        // Invariant: Filename has .json extension
        prop_assert!(filename.ends_with(".json"));
    }
}

```

**What This Tests:** - ✓ Random SPEC IDs - ✓ All 6 stages - ✓ Random agent names - ✓ Filenames always valid (no /, \, null bytes) - ✓ Always has .json extension

---

## Testing Collections

### Property 1: Filtering Never Increases Length

**Property:** Filtered collection ≤ original length

```
proptest! {
```

```

#[test]
fn test_filter_never_increases_length(
    vec in any::<Vec<i32>>()
) {
    let filtered: Vec<_> = vec.iter()
        .filter(|&x| x > 0)
        .collect();

    prop_assert!(filtered.len() <= vec.len());
}

```

---

## Property 2: Sorting Preserves Length

**Property:** Sorted collection has same length as original

```

proptest! {
    #[test]
    fn test_sort_preserves_length(
        mut vec in any::<Vec<i32>>()
    ) {
        let original_len = vec.len();

        vec.sort();

        prop_assert_eq!(vec.len(), original_len);
    }
}

```

---

## Property 3: Dedupe Length

**Property:** Deduplicated length  $\leq$  original length

```

proptest! {
    #[test]
    fn test_dedupe_length(
        mut vec in any::<Vec<i32>>()
    ) {
        let original_len = vec.len();

        vec.sort();
        vec.dedup();

        prop_assert!(vec.len() <= original_len);
    }
}

```

---

# Testing String Operations

## Property 1: Truncation Length

**Property:** Truncated string  $\leq$  max length (plus ellipsis)

```

proptest! {
    #[test]

```

```

fn test_truncate_length(
    text in any:<String>(),
    max_len in lusize..100,
) {
    let truncated = truncate_context(&text, max_len);

    if text.len() <= max_len {
        // No truncation
        prop_assert_eq!(truncated.len(), text.len());
    } else {
        // Truncated with ...
        prop_assert_eq!(truncated.len(), max_len + 3);
    }
}

```

---

## Property 2: Regex Escape Safety

**Property:** Escaped string never causes regex parse error

```

proptest! {
    #[test]
    fn test_regex_escape_never_panics(s in ".*") {
        let escaped = regex_escape(&s);

        // Invariant: Escaped string is valid regex literal
        let pattern = format!("^{}$", escaped);
        let re = Regex::new(&pattern);

        prop_assert!(re.is_ok());
    }
}

```

---

## Shrinking

### What is Shrinking?

When a property test fails, proptest **shrinks** the failing input to the **minimal** failing case.

**Example:**

```

proptest! {
    #[test]
    fn test_all_positive(vec in any:<Vec<i32>>()) {
        prop_assert!(vec.iter().all(|&x| x > 0));
    }
}

```

**Failure:**

```

Test failed for input: [1, 2, 3, 0, 5, 6, 7, 8, 9]
Shrinking...
Minimal failing input: [0]

```

---

## Shrinking Example

**Original failure:** - Input: vec = [42, -17, 0, 99, -3, 100, 256, -1, 7] - Failed because: -17, -3, -1 are negative

**After shrinking:** - Input: vec = [-1] - Still fails, but minimal

**Benefits:** - ✓ Easier to debug - ✓ Clear failure reason - ✓ No noise from extra elements

---

## Advanced Patterns

### Conditional Properties

**Pattern:** Property holds only under certain conditions

```
proptest! {
    #[test]
    fn test_division_inverse(
        a in any::<f64>(),
        b in any::<f64>()
    ) {
        // Property only holds when b ≠ 0
        prop_assume!(b != 0.0);

        let result = a / b * b;
        prop_assert!((result - a).abs() < 0.0001);
    }
}
```

**prop\_assume!(condition):** - Skips test case if condition false - Generates new random input - Useful for preconditions

---

### Composite Strategies

**Pattern:** Combine multiple generators

```
fn state_and_index_strategy() -> impl Strategy<Value =
(SpecAutoState, usize)> {
    (spec_id_strategy(), 0usize..20)
        .prop_map(|(spec_id, index)| {
            let mut state = StateBuilder::new(&spec_id).build();
            state.current_index = index;
            (state, index)
        })
}

proptest! {
    #[test]
    fn test_with_composite(
        (state, index) in state_and_index_strategy()
    ) {
        if index < 6 {
            prop_assert!(state.current_stage().is_some());
        }
    }
}
```

```
}
```

---

## Regression Testing

**Pattern:** Save failing inputs, re-test on every run

**File:** proptest-regressions/property\_based\_tests.txt

```
# Seeds for failure cases
xs 1234567890
xs 9876543210
```

**Usage:** 1. Test fails with input xs = 1234567890 2. proptest saves seed to regression file 3. Next run always tests that seed first 4. Ensures bug doesn't resurface

---

## Configuration

### Adjust Test Cases

**Default:** 100 test cases per property

**Custom:**

```
proptest! {
    #![proptest_config(ProptestConfig::with_cases(1000))]

    #[test]
    fn test_with_more_cases(n in any::<i32>()) {
        // Runs 1000 times instead of 100
    }
}
```

---

### Environment Variable

```
# Run 10,000 test cases
PROPTEST_CASES=10000 cargo test --test property_based_tests
```

---

### Timeout

```
proptest! {
    #![proptest_config(ProptestConfig {
        cases: 100,
        max_shrink_iters: 10000,
        timeout: 5000, // 5 seconds
        .. ProptestConfig::default()
    })]

    #[test]
    fn test_with_timeout(vec in any::<Vec<i32>>()) {
        // Timeout if takes >5s
    }
}
```

---

# Best Practices

## DO

### ✓ Test invariants, not examples:

```
// Good: Tests property
proptest! {
    #[test]
    fn test_reverse_twice_identity(vec in any::<Vec<i32>>()) {
        prop_assert_eq!(reverse(reverse(vec.clone())), vec);
    }
}

// Bad: Tests specific example (use regular #[test])
proptest! {
    #[test]
    fn test_specific_case() {
        let vec = vec![1, 2, 3];
        prop_assert_eq!(reverse(reverse(vec.clone())), vec);
    }
}
```

---

### ✓ Use `prop_assume!` for preconditions:

```
proptest! {
    #[test]
    fn test_with_precondition(
        index in 0usize..100,
        vec in any::<Vec<i32>>()
    ) {
        prop_assume!(index < vec.len());

        let elem = vec[index];
        // Test with valid index
    }
}
```

---

### ✓ Test mathematical properties:

```
proptest! {
    #[test]
    fn test_addition_associative(a in any::<i32>(), b in any::<i32>(),
        c in any::<i32>()) {
        prop_assert_eq!((a + b) + c, a + (b + c));
    }

    #[test]
    fn test_multiplication_distributive(a in any::<i32>(), b in
        any::<i32>(), c in any::<i32>()) {
        prop_assert_eq!(a * (b + c), a * b + a * c);
    }
}
```

---

### ✓ Test round-trip properties:

```
proptest! {
```

```
#![test]
fn test_serialize_deserialize(state in any::<SpecAutoState>()) {
    let json = serde_json::to_string(&state).unwrap();
    let deserialized: SpecAutoState =
        serde_json::from_str(&json).unwrap();

        prop_assert_eq!(deserialized, state);
}
}
```

---

## DON'T

### ✗ Test concrete outputs:

```
// Bad: Property tests shouldn't check specific outputs
proptest! {
    #[test]
    fn test_bad(n in any::<i32>()) {
        prop_assert_eq!(add_one(n), n + 1); // ✗ This is just
example-based
    }
}
```

---

### ✗ Generate invalid inputs:

```
// Bad: Generates many invalid cases (slow)
proptest! {
    #[test]
    fn test_with_many_assumes(
        a in any::<i32>(),
        b in any::<i32>(),
    ) {
        prop_assume!(a > 0);
        prop_assume!(b > 0);
        prop_assume!(a < b);
        prop_assume!(b % 2 == 0);
        // ... many assumes = slow
    }
}

// Good: Use constrained generator
fn even_positive_pair_strategy() -> impl Strategy<Value = (i32,
i32)> {
    (i132..1000, i132..1000)
        .prop_filter("a < b and b even", |(a, b)| a < b && b % 2 ==
0)
}
```

---

## Running Property Tests

### Run All Property Tests

```
cd codex-rs
cargo test --test property_based_tests
```

---

## Run with More Cases

```
PROPTEST_CASES=1000 cargo test --test property_based_tests
```

---

## Debug Failing Test

```
# Run specific property test
cargo test --test property_based_tests pb01_state_index

# With verbose output
cargo test --test property_based_tests pb01_state_index -- --
nocapture
```

---

## Re-run Regression Cases

```
# Automatically runs saved regression cases from proptest-regressions/
cargo test --test property_based_tests
```

---

## Summary

### Property-Based Testing Best Practices:

1. **Invariants**: Test properties that always hold
2. **Generators**: Use appropriate generators (ranges, regex, custom)
3. **Shrinking**: Let proptest find minimal failing case
4. **Preconditions**: Use prop\_assume!() for preconditions
5. **Configuration**: Adjust test cases with PROPTEST\_CASES
6. **Regression**: Save failing cases automatically

**Common Properties to Test:** - ✓ Invariants (index bounds, retry limits) - ✓ Round-trip (serialize → deserialize) - ✓ Mathematical (associativity, commutativity, distributivity) - ✓ Collection operations (filter length, sort preserves length) - ✓ String operations (truncate length, regex escape safety) - ✓ Evidence integrity (valid JSON, valid filenames)

**Key Concepts:** - ✓ Generators create random inputs - ✓ Shrinking finds minimal failing case - ✓ Regression tests prevent regressions - ✓ 100 test cases per property (default)

**Next Steps:** - [CI/CD Integration](#) - Automated testing pipeline - [Performance Testing](#) - Benchmarks and profiling - [Test Infrastructure](#) - MockMcpManager, fixtures

---

**References:** - proptest docs: <https://docs.rs/proptest> - Property tests: [codex-rs/tui/tests/property\\_based\\_tests.rs](#) - Regression files: [proptest-regressions/](#)

---

## Test Infrastructure

Comprehensive testing infrastructure for the codebase.

---

## Overview

**Test Infrastructure Components:** - **MockMcpManager:** Mock MCP server for isolated testing - **IntegrationTestContext:** Multi-module test harness - **StateBuilder:** Test state configuration - **EvidenceVerifier:** Artifact validation helpers - **Fixture Library:** Real production data (20 files, 96 KB) - **Coverage Tools:** cargo-tarpaulin, cargo-llvm-cov - **Property Testing:** proptest for generative testing

**Location:** codex-rs/tui/tests/common/ (shared test utilities)

**Purpose:** Enable comprehensive testing without external dependencies

---

## MockMcpManager

### Purpose

Mock implementation of McpConnectionManager for testing MCP-dependent code without requiring a live local-memory server.

**Location:** codex-rs/tui/tests/common/mock\_mcp.rs (272 LOC)

**Use Cases:** - Test consensus logic without spawning agents - Verify MCP tool calls in isolation - Fast unit tests (<1ms vs 8.7ms real MCP) - Deterministic fixture responses

---

## API Reference

### Creating a Mock

```
use codex_tui::tests::common::MockMcpManager;

let mut mock = MockMcpManager::new();
```

**Methods:** - new() → Create empty mock - default() → Same as new() (implements Default)

---

### Adding Fixtures

#### Single Fixture:

```
mock.add_fixture(
    "local-memory",           // server name
    "search",                 // tool name
    Some("SPEC-TEST plan"),   // query pattern (or None for
wildcard)
    json!({
        "memory": {
            "id": "test-1",
            "content": "Test content"
    })
);
```

```

        }
    })
);

```

### **Multiple Fixtures:**

```

mock.add_fixtures(
    "local-memory",
    "search",
    Some("SPEC-TEST plan"),
    vec![
        json!({"memory": {"id": "test-1", "content": "Agent 1"}},),
        json!({"memory": {"id": "test-2", "content": "Agent 2"}},),
    ]
);

```

### **From File:**

```

mock.load_fixture_file(
    "local-memory",
    "search",
    Some("SPEC-KIT-DEMO plan"),
    "tests/fixtures/consensus/demo-plan-gemini.json"
)?;

```

---

## **Calling Tools**

### **Signature:**

```

pub async fn call_tool(
    &self,
    server: &str,
    tool: &str,
    arguments: Option<Value>,
    timeout: Option<Duration>,
) -> Result<CallToolResult>

```

### **Example:**

```

let args = json!({"query": "SPEC-TEST plan"});
let result = mock.call_tool(
    "local-memory",
    "search",
    Some(args),
    None // timeout
).await?;

// Extract response
if let ContentBlock::TextContent(text) = &result.content[0] {
    let data: Value = serde_json::from_str(&text.text)?;
    println!("{}: {:?}", data);
}

```

---

## **Call Logging**

### **Get Call History:**

```

let log = mock.call_log();
for entry in log {

```

```

    println!("Called: {} / {}", entry.server, entry.tool);
    println!("  Args: {:?}", entry.arguments);
}

```

### Clear Log:

```
mock.clear_log();
```

**Use Case:** Verify expected tool calls were made

```

assert_eq!(log.len(), 3);
assert_eq!(log[0].tool, "search");
assert_eq!(log[1].tool, "search");
assert_eq!(log[2].tool, "search");

```

---

## Fixture Matching

**Priority Order:** 1. **Exact query match:** query\_pattern = Some("SPEC-TEST plan") 2. **Wildcard match:** query\_pattern = None 3. **No match:** Returns error

### Example:

```

// Add wildcard fixture
mock.add_fixture("local-memory", "search", None, json!({ "default": true }));

// Add specific fixture
mock.add_fixture(
    "local-memory",
    "search",
    Some("SPEC-DEMO plan"),
    json!({ "specific": true })
);

// Query "SPEC-DEMO plan" → Returns {"specific": true}
// Query "anything else"   → Returns {"default": true}
// Query with no fixture  → Error

```

---

## Usage Patterns

### Pattern 1: Unit Testing Consensus

```

#[tokio::test]
async fn test_consensus_high_confidence() {
    let mut mock = MockMcpManager::new();

    // Load real production fixtures
    mock.load_fixture_file(
        "local-memory",
        "search",
        Some("SPEC-TEST plan"),
        "tests/fixtures/consensus/demo-plan-gemini.json"
    )?;
    mock.load_fixture_file(
        "local-memory",
        "search",
        Some("SPEC-TEST plan"),

```

```

        "tests/fixtures/consensus/demo-plan-claude.json"
    )?;

    // Test consensus collection
    let (results, degraded) = fetch_memory_entries(
        "SPEC-TEST",
        SpecStage::Plan,
        &mock
    ).await?;

    assert_eq!(results.len(), 2);
    assert!(!degraded, "Should have both agents");
}

```

---

## Pattern 2: Verifying Tool Calls

```

#[tokio::test]
async fn test_quality_gate_calls_all_tools() {
    let mut mock = MockMcpManager::new();
    mock.add_fixture("local-memory", "search", None, json!({}));

    // Run quality gate
    run_quality_gate("SPEC-TEST", &mock).await?;

    // Verify calls
    let log = mock.call_log();
    assert!(log.iter().any(|e| e.tool == "search"));

    // Verify call arguments
    let search_call = log.iter().find(|e| e.tool ==
"search").unwrap();
    assert!(search_call.arguments.is_some());
}

```

---

## Pattern 3: Testing Error Handling

```

#[tokio::test]
async fn test_consensus_degradation_on_missing_agent() {
    let mut mock = MockMcpManager::new();

    // Only add 2 of 3 agents
    mock.add_fixture("local-memory", "search", None, json!({"agent": "gemini"}));
    mock.add_fixture("local-memory", "search", None, json!({"agent": "claude"}));
    // gpt_pro deliberately missing

    let (results, degraded) = fetch_memory_entries(
        "SPEC-TEST",
        SpecStage::Plan,
        &mock
    ).await?;

    assert_eq!(results.len(), 2);
    assert!(degraded, "Should be degraded (missing 1 agent)");
}

```

---

## Tests

**Location:** codex-rs/tui/tests/mock\_mcp\_tests.rs (7 tests)

**Coverage:**

test_mock_mcp_returns_fixture	✓
test_mock_mcp_logs_calls	✓
test_mock_mcp_wildcard_matches	✓
test_mock_mcp_exact_query_precedence	✓
test_mock_mcp_multiple_fixtures_return_array	✓
test_mock_mcp_load_from_file	✓
test_mock_mcp_error_on_no_fixture	✓

**Run Tests:**

```
cd codex-rs
cargo test --test mock_mcp_tests
```

---

## IntegrationTestContext

### Purpose

Multi-module test harness for integration tests with isolated filesystem and evidence verification.

**Location:** codex-rs/tui/tests/common/integration\_harness.rs (254 LOC)

**Use Cases:** - Cross-module workflow tests - Evidence verification - Filesystem isolation (temp directories) - SPEC directory structure setup

---

### API Reference

#### Creating a Context

```
use codex_tui::tests::common::IntegrationTestContext;

let ctx = IntegrationTestContext::new("SPEC-TEST-001")?;
```

**Fields:**

```
pub struct IntegrationTestContext {
    pub temp_dir: TempDir,           // Auto-cleaned on drop
    pub spec_id: String,             // "SPEC-TEST-001"
    pub cwd: PathBuf,                // temp_dir path
    pub evidence_dir: PathBuf,       // docs/SPEC-OPS-004.../evidence
}
```

**Auto-Created Directories:** - docs/SPEC-OPS-004-integrated-coder-hooks/evidence/ - docs/SPEC-OPS-004.../evidence/consensus/{spec\_id}/ - docs/SPEC-OPS-004.../evidence/commands/{spec\_id}/

---

#### Directory Helpers

## Get Evidence Directories:

```
let consensus_dir = ctx.consensus_dir();
// → .../evidence/consensus/SPEC-TEST-001/  
  
let commands_dir = ctx.commands_dir();
// → .../evidence/commands/SPEC-TEST-001/
```

## Create SPEC Directory:

```
let spec_dir = ctx.create_spec_dirs("test-feature")?;
// → .../docs/SPEC-TEST-001-test-feature/
```

---

## File Helpers

### Write PRD:

```
ctx.write_prd("test-feature", "# PRD\n\nTest product
requirements")?;
// Creates: docs/SPEC-TEST-001-test-feature/PRD.md
```

### Write Spec:

```
ctx.write_spec("test-feature", "# SPEC-TEST-001\n\n## Goal\nTest");
// Creates: docs/SPEC-TEST-001-test-feature/spec.md
```

---

## Evidence Verification

### Check Consensus Artifacts:

```
// Single agent
let exists = ctx.assert_consensus_exists(SpecStage::Plan, "gemini");
assert!(exists);

// All agents (via EvidenceVerifier)
let verifier = EvidenceVerifier::new(&ctx);
assert!(verifier.assert_consensus_complete(
    SpecStage::Plan,
    &["gemini", "claude", "gpt_pro"]
));
```

### Check Guardrail Telemetry:

```
let exists = ctx.assert_guardrail_telemetry_exists(SpecStage::Plan);
assert!(exists);
```

### Count Files:

```
let count = ctx.count_consensus_files();
assert_eq!(count, 3, "Should have 3 agent outputs");

let guardrail_count = ctx.count_guardrail_files();
assert_eq!(guardrail_count, 1, "Should have 1 telemetry file");
```

---

## Usage Patterns

### Pattern 1: Workflow Integration Test

```

#[tokio::test]
async fn test_full_plan_stage_workflow() -> Result<()> {
    // Setup
    let ctx = IntegrationTestContext::new("SPEC-INT-001")?;
    ctx.write_prd("test-feature", "# Test PRD\n\n## Goal\nTest");

    // Run plan stage
    run_plan_stage(&ctx.spec_id, &ctx.cwd).await?;

    // Verify evidence
    assert!(ctx.assert_consensus_exists(SpecStage::Plan, "gemini"));
    assert!(ctx.assert_consensus_exists(SpecStage::Plan, "claude"));
    assert!(ctx.assert_consensus_exists(SpecStage::Plan,
    "gpt_pro"));
    assert!(ctx.assert_guardrail_telemetry_exists(SpecStage::Plan));

    // Verify file count
    assert_eq!(ctx.count_consensus_files(), 3);

    Ok(())
}

```

---

## Pattern 2: Error Recovery Test

```

#[tokio::test]
async fn test_error_recoveryCreatesEvidence() -> Result<()> {
    let ctx = IntegrationTestContext::new("SPEC-INT-002")?;

    // Simulate error (missing PRD)
    let result = run_plan_stage(&ctx.spec_id, &ctx.cwd).await;
    assert!(result.is_err());

    // Verify error evidence still created
    let verifier = EvidenceVerifier::new(&ctx);
    assert!
    (verifier.assert_guardrail_valid(SpecStage::Plan).is_ok());

    Ok(())
}

```

---

## Pattern 3: State Persistence Test

```

#[tokio::test]
async fn test_state_persists_across_stages() -> Result<()> {
    let ctx = IntegrationTestContext::new("SPEC-INT-003")?;
    ctx.write_prd("test", "# PRD")?;

    // Run plan
    run_plan_stage(&ctx.spec_id, &ctx.cwd).await?;
    assert_eq!(ctx.count_consensus_files(), 3);

    // Run tasks (should accumulate, not replace)
    run_tasks_stage(&ctx.spec_id, &ctx.cwd).await?;
    assert!(ctx.count_consensus_files() > 3, "Should accumulate
evidence");

    Ok(())
}

```

---

## Tests

**Location:** codex-rs/tui/tests/common/integration\_harness.rs (4 tests in mod tests)

**Coverage:**

test_integration_context_creation	✓
test_state_builder	✓
test_spec_dirs_creation	✓
test_evidence_verifier	✓

---

## StateBuilder

### Purpose

Builder pattern for creating SpecAutoState instances in tests with custom configuration.

**Location:** codex-rs/tui/tests/common/integration\_harness.rs

**Use Cases:** - Configure test automation state - Test different starting stages - Test HAL mode variations - Test quality gate configurations

---

### API Reference

#### Basic Usage

```
use codex_tui::tests::common::StateBuilder;

let state = StateBuilder::new("SPEC-TEST-001").build();
```

**Default Configuration:** - goal: "Integration test" - start\_stage: Plan - hal\_mode: None - quality\_gates\_enabled: true

---

#### Builder Methods

##### Custom Goal:

```
let state = StateBuilder::new("SPEC-TEST-001")
    .with_goal("Implement user authentication")
    .build();
```

##### Start at Different Stage:

```
let state = StateBuilder::new("SPEC-TEST-002")
    .starting_at(SpecStage::Implement)
    .build();
```

##### HAL Mode Configuration:

```
let state = StateBuilder::new("SPEC-TEST-003")
    .with_hal_mode(HalMode::Analyze)
    .build();
```

## Quality Gates Control:

```
let state = StateBuilder::new("SPEC-TEST-004")
    .quality_gates(false) // Disable quality gates
    .build();
```

## Chained Configuration:

```
let state = StateBuilder::new("SPEC-TEST-005")
    .with_goal("Test refactoring")
    .starting_at(SpecStage::Validate)
    .with_hal_mode(HalMode::TestOnly)
    .quality_gates(true)
    .build();
```

---

## Usage Patterns

### Pattern 1: Testing Stage Transitions

```
#[test]
fn test_stage_advancement() {
    let mut state = StateBuilder::new("SPEC-TEST-001")
        .starting_at(SpecStage::Plan)
        .build();

    assert_eq!(state.current_stage(), Some(SpecStage::Plan));

    state.advance_stage();
    assert_eq!(state.current_stage(), Some(SpecStage::Tasks));

    state.advance_stage();
    assert_eq!(state.current_stage(), Some(SpecStage::Implement));
}
```

---

### Pattern 2: Testing Quality Gate Behavior

```
#[test]
fn test_quality_gates_disabled() {
    let state = StateBuilder::new("SPEC-TEST-002")
        .quality_gates(false)
        .build();

    assert!(!state.quality_gates_enabled);

    // Quality gates should not run
    assert!(should_skip_quality_gate(&state));
}

#[test]
fn test_quality_gates_enabled() {
    let state = StateBuilder::new("SPEC-TEST-003")
        .quality_gates(true)
        .build();

    assert!(state.quality_gates_enabled);
}
```

---

### Pattern 3: Testing HAL Integration

```
#[test]
fn test_hal_mode_analyze() {
    let state = StateBuilder::new("SPEC-TEST-004")
        .with_hal_mode(HalMode::Analyze)
        .build();

    assert_eq!(state.hal_mode, Some(HalMode::Analyze));
}

#[test]
fn test_hal_mode_none() {
    let state = StateBuilder::new("SPEC-TEST-005")
        .build();

    assert_eq!(state.hal_mode, None);
}
```

---

## EvidenceVerifier

### Purpose

Helper for verifying evidence artifacts in integration tests.

**Location:** codex-rs/tui/tests/common/integration\_harness.rs

**Use Cases:** - Assert consensus artifacts exist - Validate guardrail telemetry - Verify directory structure - Check multi-agent completion

---

### API Reference

#### Creating a Verifier

```
use codex_tui::tests::common::EvidenceVerifier;

let ctx = IntegrationTestContext::new("SPEC-TEST-001")?;
let verifier = EvidenceVerifier::new(&ctx);
```

---

#### Verification Methods

**Consensus Complete** (all agents present):

```
let complete = verifier.assert_consensus_complete(
    SpecStage::Plan,
    &["gemini", "claude", "gpt_pro"]
);
assert!(complete);
```

**Guardrail Valid** (telemetry exists and parseable):

```
let result = verifier.assert_guardrail_valid(SpecStage::Plan);
assert!(result.is_ok());
```

**Structure Valid** (directories exist):

```
let valid = verifier.assert_structure_valid();
assert!(valid);
```

---

## Usage Patterns

### Pattern 1: Post-Workflow Verification

```
#[tokio::test]
async fn test_planCreatesCompleteEvidence() -> Result<()> {
    let ctx = IntegrationTestContext::new("SPEC-VER-001")?;
    ctx.write_prd("test", "# PRD")?;

    run_plan_stage(&ctx.spec_id, &ctx.cwd).await?;

    let verifier = EvidenceVerifier::new(&ctx);

    // Verify all artifacts
    assert!(verifier.assert_structure_valid());
    assert!(verifier.assert_consensus_complete(
        SpecStage::Plan,
        &["gemini", "claude", "gpt_pro"]
    ));
    assert!(
        verifier.assert_guardrail_valid(SpecStage::Plan).is_ok()
    );

    Ok(())
}
```

---

### Pattern 2: Degraded Consensus Detection

```
#[tokio::test]
async fn test_degraded_consensus_still_valid() -> Result<()> {
    let ctx = IntegrationTestContext::new("SPEC-VER-002")?;

    // Simulate degraded consensus (only 2/3 agents)
    simulate_agent_failure("gpt_pro")?;
    run_plan_stage(&ctx.spec_id, &ctx.cwd).await?;

    let verifier = EvidenceVerifier::new(&ctx);

    // Should NOT be complete (missing 1 agent)
    assert!(!verifier.assert_consensus_complete(
        SpecStage::Plan,
        &["gemini", "claude", "gpt_pro"]
    ));

    // But 2/3 is still valid
    assert!(verifier.assert_consensus_complete(
        SpecStage::Plan,
        &["gemini", "claude"]
    ));

    Ok(())
}
```

---

# Fixture Library

## Overview

**Location:** codex-rs/tui/tests/fixtures/consensus/ (20 files, 96 KB)

**Source:** Real production artifacts from docs/SPEC-OPS-004.../evidence/consensus/

**Coverage:** - Plan stage: 13 fixtures (DEMO, 025, 045) - Tasks stage: 3 fixtures (025) - Implement stage: 4 fixtures (025)

---

## File Naming Convention

**Format:** {spec\_id}-{stage}-{agent}.json

**Examples:** - demo-plan-gemini.json — SPEC-KIT-DEMO plan stage (Gemini output) - 025-implement-gpt\_codex.json — SPEC-KIT-025 implement stage (Codex output) - 045-plan-claude.json — SPEC-KIT-045 plan stage (Claude output)

---

## Available Fixtures

### Plan Stage (13 files)

**SPEC-KIT-DEMO:** - demo-plan-gemini.json (14 KB) - demo-plan-claude.json (12 KB) - demo-plan-gpt\_pro.json (15 KB)

**SPEC-KIT-025** (Native SPEC-ID generation): - 025-plan-gemini.json (16 KB) - 025-plan-claude.json (14 KB) - 025-plan-gpt\_pro.json (18 KB)

**SPEC-KIT-045** (Quality gate handler): - 045-plan-gemini.json (13 KB) - 045-plan-claude.json (11 KB) - 045-plan-gpt\_pro.json (17 KB)

---

### Tasks Stage (3 files)

**SPEC-KIT-025:** - 025-tasks-gemini.json (8 KB) - 025-tasks-claude.json (7 KB)

---

### Implement Stage (4 files)

**SPEC-KIT-025:** - 025-implement-gemini.json (9 KB) - 025-implement-claude.json (8 KB) - 025-implement-gpt\_codex.json (22 KB) — Code implementation - 025-implement-gpt\_pro.json (11 KB)

---

## Usage in Tests

### Loading Single Fixture:

```
let mut mock = MockMcpManager::new();
mock.load_fixture_file(
    "local-memory",
```

```

    "search",
    Some("SPEC-KIT-DEMO plan"),
    "tests/fixtures/consensus/demo-plan-gemini.json"
)?;

```

**Loading All Agents** (simulate 3-agent consensus):

```

let mut mock = MockMcpManager::new();
let agents = vec!["gemini", "claude", "gpt_pro"];

for agent in agents {
    mock.load_fixture_file(
        "local-memory",
        "search",
        Some("SPEC-KIT-DEMO plan"),
        &format!("tests/fixtures/consensus/demo-plan-{}.json",
agent)
    )?;
}

```

**Loading Different Stages:**

```

// Plan stage
mock.load_fixture_file("local-memory", "search", Some("SPEC-KIT-025
plan"),
    "tests/fixtures/consensus/025-plan-gemini.json")?;

// Tasks stage
mock.load_fixture_file("local-memory", "search", Some("SPEC-KIT-025
tasks"),
    "tests/fixtures/consensus/025-tasks-gemini.json")?;

// Implement stage
mock.load_fixture_file("local-memory", "search", Some("SPEC-KIT-025
implement"),
    "tests/fixtures/consensus/025-implement-gpt_coder.json")?;

```

---

## Adding New Fixtures

**Manual Creation:**

```

cd codex-rs/tui/tests/fixtures/consensus

# Copy from production evidence
cp ../../docs/SPEC-0PS-004.../evidence/consensus/SPEC-KIT-
070/spec-plan_*.json \
./070-plan-gemini.json

```

**Automated Extraction** (future):

```

# Extract fixtures from evidence repository
./scripts/extract_test_fixtures.sh SPEC-KIT-070

```

**Size Guidelines:** - Keep individual fixtures < 30 KB - Total fixture directory < 200 KB - Compress if needed (not implemented yet)

---

## Coverage Tools

## cargo-tarpaulin

**Purpose:** Line coverage measurement for Rust code

**Installation:**

```
cargo install cargo-tarpaulin
```

**Configuration:** codex-rs/tarpaulin.toml

---

### Configuration Details

```
[config]
# Only measure spec-kit coverage (fork-specific code)
run-types = ["Lib", "Tests"]

# Include patterns (spec-kit only)
include-pattern = "tui/src/chatwidget/spec_kit/.*\\".rs"

# Exclude test files and generated code
exclude-files = [
    "tui/src/chatwidget/spec_kit/*/tests/*",
    "tui/tests/*",
]

# Output formats
out = ["Html", "Stdout"]
output-dir = "target/tarpaulin"

# Timeout per test (integration tests are slow)
timeout = 120

# Verbose output
verbose = true
```

---

### Usage

**Full Coverage Report:**

```
cd codex-rs
cargo tarpaulin
```

**Output:**

```
|| Tested/Total Lines:
|| tui/src/chatwidget/spec_kit/handler.rs: 145/961
|| tui/src/chatwidget/spec_kit/consensus.rs: 120/992
|| tui/src/chatwidget/spec_kit/quality.rs: 178/807
|| ...
|| Coverage: 42.3%
```

**Specific Module:**

```
cargo tarpaulin -p codex-tui
```

**HTML Report:**

```
cargo tarpaulin --out Html
```

```
open target/tarpaulin/index.html
```

#### **XML for CI (Codecov):**

```
cargo tarpaulin --out Xml
```

---

### **Troubleshooting**

#### **Issue:** Timeout on slow tests

```
# Increase timeout  
cargo tarpaulin --timeout 300
```

#### **Issue:** Out of memory

```
# Reduce parallelism  
cargo tarpaulin --jobs 2
```

#### **Issue:** Incorrect coverage (too low)

```
# Ensure all features enabled  
cargo tarpaulin --all-features
```

---

## **cargo-llvm-cov**

#### **Purpose:** Alternative coverage tool using LLVM instrumentation

**Advantages:** - More accurate than tarpaulin - Faster execution - Better integration with IDEs

#### **Installation:**

```
cargo install cargo-llvm-cov
```

---

### **Usage**

#### **Generate Coverage:**

```
cd codex-rs  
cargo llvm-cov --workspace --all-features --html
```

#### **Open Report:**

```
open target/llvm-cov/html/index.html
```

#### **JSON Output** (for parsing):

```
cargo llvm-cov --workspace --all-features --json --output-path  
coverage.json
```

#### **Integration with VS Code:**

```
# Install Coverage Gutters extension  
# Run:  
cargo llvm-cov --workspace --all-features --lcov --output-path  
lcov.info  
  
# VS Code will show coverage inline
```

---

## Comparison: Tarpaulin vs llvm-cov

Feature	Tarpaulin	llvm-cov
<b>Accuracy</b>	~95%	~99%
<b>Speed</b>	Baseline	1.5-2× faster
<b>HTML Report</b>	✓ Good	✓ Excellent
<b>IDE Integration</b>	✗ Limited	✓ VS Code, IntelliJ
<b>CI Support</b>	✓ Codecov, Coveralls	✓ All platforms
<b>Install Size</b>	50 MB	150 MB (LLVM)

**Recommendation:** Use llvm-cov for local development, tarpaulin for CI (smaller install).

---

## Property-Based Testing

### Overview

**Purpose:** Generative testing with random inputs to verify invariants

**Tool:** `proptest` (Rust equivalent of Hypothesis/QuickCheck)

**Location:** `codex-rs/tui/tests/property_based_tests.rs`

**Use Cases:** - State machine invariants - Evidence integrity - Consensus edge cases - Input validation

---

### Proptest Basics

#### Simple Property Test:

```
use proptest::prelude::*;

proptest! {
    #[test]
    fn test_state_index_never_negative(index in 0usize..20) {
        // Property: State always handles any index gracefully
        let mut state = SpecAutoState::new(...);
        state.current_index = index;

        // Should never panic
        let _ = state.current_stage();
    }
}
```

**How It Works:** 1. Generate 100 random values for `index` (0-19) 2. Run test with each value 3. If any fails, shrink to minimal failing case 4. Report failure with minimal input

---

### Test Categories

#### PB01-PB03: State Invariants

### PB01: Index always in valid range

```
proptest! {
    #[test]
    fn pb01_state_index_always_in_valid_range(index in 0usize..20) {
        let mut state = StateBuilder::new("SPEC-PB01-TEST")
            .starting_at(SpecStage::Plan)
            .build();

        state.current_index = index;

        // Invariant: index ∈ [0, 5] → Some(_), else None
        if index < 6 {
            prop_assert!(state.current_stage().is_some());
        } else {
            prop_assert_eq!(state.current_stage(), None);
        }
    }
}
```

### PB02: Current stage always Some when index < 6

```
proptest! {
    #[test]
    fn pb02_current_stage_always_some_when_index_under_six(
        index in 0usize..6
    ) {
        let mut state = StateBuilder::new("SPEC-PB02-TEST").build();
        state.current_index = index;

        prop_assert!(state.current_stage().is_some());
    }
}
```

### PB03: Retry count never exceeds max

```
proptest! {
    #[test]
    fn pb03_retry_count_never_negative(retries in 0usize..100) {
        let max_retries = 3;
        let capped_retries = retries.min(max_retries);

        // Write retry file
        let retry_data = json!({
            "retry_count": capped_retries,
            "max_retries": max_retries,
        });

        // Invariant: retry_count ≤ max_retries
        prop_assert!(capped_retries <= max_retries);
    }
}
```

---

## PB04-PB06: Evidence Integrity

### PB04: Written evidence always parseable JSON

```
proptest! {
    #[test]
    fn pb04_written_evidence_always_parseable_json(
```

```

        agent in "[a-z]{3,10}",
        content in ".*"
    ) {
        let ctx = IntegrationTestContext::new("SPEC-PB04-TEST")?;

        let evidence = json!({
            "agent": agent,
            "content": content,
            "timestamp": "2025-10-19T00:00:00Z"
        });

        let file = ctx.consensus_dir().join("test.json");
        std::fs::write(&file, evidence.to_string())?;

        // Invariant: File is valid JSON
        let content = std::fs::read_to_string(&file)?;
        let parsed: Value = serde_json::from_str(&content)?;

        prop_assert_eq!(parsed["agent"].as_str(),
Some(agent.as_str()));
    }
}

```

---

## Custom Generators

### Generate SPEC IDs:

```

fn spec_id_strategy() -> impl Strategy<Value = String> {
    "[A-Z]{4}-[A-Z]{3}-[0-9]{3}"
        .prop_map(|s| s.to_string())
}

proptest! {
    #[test]
    fn test_spec_id_parsing(spec_id in spec_id_strategy()) {
        // Test SPEC ID validation
        assert!(is_valid_spec_id(&spec_id));
    }
}

```

### Generate Stages:

```

fn stage_strategy() -> impl Strategy<Value = SpecStage> {
    prop_oneof![
        Just(SpecStage::Plan),
        Just(SpecStage::Tasks),
        Just(SpecStage::Implement),
        Just(SpecStage::Validate),
        Just(SpecStage::Audit),
        Just(SpecStage::Unlock),
    ]
}

```

---

## Running Property Tests

### Run All Property Tests:

```
cd codex-rs
```

```
cargo test --test property_based_tests
```

#### Run Specific Test:

```
cargo test --test property_based_tests pb01_state_index
```

#### Adjust Iteration Count (default 100):

```
PROPTEST_CASES=1000 cargo test --test property_based_tests
```

#### Debug Failing Case:

```
# proptest creates a regression file
cat proptest-regressions/property_based_tests.txt

# Re-run with that specific input
cargo test --test property_based_tests -- --exact pb01_state_index
```

---

## TestCodexBuilder

### Purpose

Builder for creating test instances of CodexConversation with mock servers.

**Location:** codex-rs/core/tests/common/test\_codex.rs (76 LOC)

**Use Cases:** - Test agent spawning - Test conversation lifecycle - Test configuration variations - Integration with wiremock

---

### API Reference

#### Basic Usage:

```
use codex_core::tests::common::test_codex;

let server = wiremock::MockServer::start().await;
let codex = test_codex()
    .build(&server)
    .await?;
```

#### Fields:

```
pub struct TestCodex {
    pub home: TempDir,                                // Isolated home
    directory
    pub cwd: TempDir,                                 // Isolated working
    directory
    pub codex: Arc<CodexConversation>,           // Conversation
    instance
    pub session_configured: SessionConfiguredEvent, // Initial
    event
}
```

---

### Custom Configuration

## Modify Config:

```
let codex = test_codex()
    .with_config(|config| {
        config.model = "gpt-5-low".to_string();
        config.max_tokens = 4096;
    })
    .build(&server)
    .await?;
```

## Multiple Mutations:

```
let codex = test_codex()
    .with_config(|config| config.model = "gpt-5-low".to_string())
    .with_config(|config| config.max_tokens = 8192)
    .with_config(|config| config.temperature = 0.7)
    .build(&server)
    .await?;
```

---

## Usage with Wiremock

### Mock API Responses:

```
use wiremock::{MockServer, Mock, ResponseTemplate};
use wiremock::matchers::{method, path};

#[tokio::test]
async fn test_conversation_with_mock() -> Result<()> {
    let server = MockServer::start().await;

    // Mock /v1/chat/completions
    Mock::given(method("POST"))
        .and(path("/v1/chat/completions"))
        .respond_with(ResponseTemplate::new(200).set_body_json(json!({
            "id": "chatcmpl-test",
            "object": "chat.completion",
            "created": 1234567890,
            "model": "gpt-4",
            "choices": [
                {
                    "index": 0,
                    "message": {
                        "role": "assistant",
                        "content": "Test response"
                    },
                    "finish_reason": "stop"
                }
            ]
        }))
        .mount(&server)
        .await;

    let codex = test_codex().build(&server).await?;

    // Test conversation
    let response = codex.codex.send_message("Test").await?;
    assert_eq!(response.content, "Test response");

    Ok(())
}
```

---

---

## Common Test Utilities

### Test Module Structure

**Location:** codex-rs/tui/tests/common/mod.rs

```
//! Common test utilities for spec-kit

pub mod integration_harness;
pub mod mock_mcp;

pub use integration_harness::{
    EvidenceVerifier,
    IntegrationTestContext,
    StateBuilder,
};
pub use mock_mcp::MockMcpManager;
```

### Usage in Tests:

```
mod common;

use common::{
    MockMcpManager,
    IntegrationTestContext,
    StateBuilder,
    EvidenceVerifier,
};
```

---

## Shared Test Data

### Constants:

```
// tests/common/mod.rs

pub const TEST_SPEC_ID: &str = "SPEC-TEST-001";
pub const TEST_GOAL: &str = "Integration test";

pub fn default_test_prd() -> &'static str {
    r#"
# Product Requirements Document

## Goal
Test feature implementation

## Requirements
- R1: Feature should work
- R2: Feature should be tested
    "#}
}
```

### Usage:

```
use common::{TEST_SPEC_ID, default_test_prd};

#[tokio::test]
async fn test_with_shared_data() {
```

```
let ctx = IntegrationTestContext::new(TEST_SPEC_ID)?;
ctx.write_prd("test-feature", default_test_prd())?;
// ...
}
```

---

## Test Organization Best Practices

### File Naming

**Unit Tests** (in source files):

```
// src/chatwidget/spec_kit/handler.rs

#[cfg(test)]
mod tests {
    use super::*;

    #[test]
    fn test_handler_orchestration() { }
}
```

**Integration Tests** (separate files):

```
codex-rs/tui/tests/
├── workflow_integration_tests.rs
├── error_recovery_integration_tests.rs
├── state_persistence_integration_tests.rs
├── concurrent_operations_integration_tests.rs
└── quality_flow_integration_tests.rs
```

**Property Tests:**

```
codex-rs/tui/tests/property_based_tests.rs
```

---

### Test Naming Conventions

**Pattern:** test\_{what}\_{condition}\_{expected}

**Examples:**

```
#[test]
fn test_state_advance_increments_index() { }

#[test]
fn test_consensus_degraded_when_missing_agent() { }

#[test]
fn test_evidence_created_on_error() { }

#[tokio::test]
async fn test_quality_gate_passes_when_score_above_80() { }
```

**Avoid:**

```
#[test]
fn test1() { } // ✗ Meaningless
```

```
#[test]
fn it_works() { } // ✕ Too vague
```

---

## Common Test Patterns

### Pattern: Arrange-Act-Assert

```
#[test]
fn test_example() {
    // Arrange: Setup
    let ctx = IntegrationTestContext::new("SPEC-TEST")?;
    let state = StateBuilder::new("SPEC-TEST").build();

    // Act: Execute
    let result = do_something(&ctx, &state)?;

    // Assert: Verify
    assert_eq!(result, expected);
}
```

---

### Pattern: Given-When-Then

```
#[tokio::test]
async fn test_consensus_with_degradation() {
    // Given: 3-agent consensus with 1 agent failing
    let mut mock = MockMcpManager::new();
    mock.add_fixture("local-memory", "search", None, json!({"agent": "gemini"}));
    mock.add_fixture("local-memory", "search", None, json!({"agent": "claude"}));
    // gpt_pro missing (simulates failure)

    // When: Fetch consensus
    let (results, degraded) = fetch_memory_entries(
        "SPEC-TEST",
        SpecStage::Plan,
        &mock
    ).await?;

    // Then: Should have 2/3 agents and be degraded
    assert_eq!(results.len(), 2);
    assert!(degraded);
}
```

---

### Pattern: Table-Driven Tests

```
#[test]
fn test_stage_index_mapping() {
    let test_cases = vec![
        (0, Some(SpecStage::Plan)),
        (1, Some(SpecStage::Tasks)),
        (2, Some(SpecStage::Implement)),
        (3, Some(SpecStage::Validate)),
        (4, Some(SpecStage::Audit)),
        (5, Some(SpecStage::Unlock)),
        (6, None),
    ]
```

```
];
for (index, expected) in test_cases {
    let mut state = StateBuilder::new("SPEC-TEST").build();
    state.current_index = index;
    assert_eq!(state.current_stage(), expected);
}
}
```

---

## Summary

### Test Infrastructure Highlights:

1. **MockMcpManager**: Fixture-based MCP testing (272 LOC, 7 tests)
2. **IntegrationTestContext**: Isolated filesystem, evidence verification
3. **StateBuilder**: Test state configuration with fluent API
4. **EvidenceVerifier**: Artifact validation helpers
5. **Fixture Library**: 20 real production artifacts (96 KB)
6. **Coverage Tools**: cargo-tarpaulin (CI), cargo-llvm-cov (local)
7. **Property Testing**: proptest for generative invariant testing
8. **TestCodexBuilder**: Conversation mocking with wiremock

**Benefits:** - ✓ Fast tests (no external dependencies) - ✓ Deterministic (fixture-based) - ✓ Isolated (temp directories) - ✓ Comprehensive (unit, integration, property) - ✓ Measurable (coverage tools)

**Next Steps:** - [Unit Testing Guide](#) - Writing effective unit tests - [Integration Testing Guide](#) - Cross-module tests - [Property Testing Guide](#) - Generative testing patterns

---

**References:** - MockMcpManager: codex-  
rs/tui/tests/common/mock\_mcp.rs - IntegrationTestContext: codex-  
rs/tui/tests/common/integration\_harness.rs - Tarpaulin config: codex-  
rs/tarpaulin.toml - Property tests: codex-  
rs/tui/tests/property\_based\_tests.rs - TestCodexBuilder: codex-  
rs/core/tests/common/test\_codex.rs

---

## Testing Strategy

Comprehensive testing approach for the codebase.

---

## Overview

**Testing Philosophy:** Balance coverage, confidence, and development velocity

**Current Metrics** (as of 2025-11-17): - **Total Tests**: 604 tests across all modules - **Pass Rate**: 100% (all tests passing) - **Coverage**: 42-48% (estimated, varies by module) - **Target**: 40%+ coverage minimum

**Test Distribution:** - **Unit Tests:** ~380 tests (63%) - **Integration Tests:** ~200 tests (33%) - **E2E Tests:** ~24 tests (4%)

**Location:** Tests located alongside source in tests/ directories per module

---

## Coverage Goals

**Overall Target: 40%+**

**Rationale:** - Industry standard for Rust projects: 60-80% - Our target: 40%+ given complexity and time constraints - Current achievement: 42-48% ✓ **Target Met**

**Coverage by Priority:** - **Critical paths:** 70-80% (Spec-Kit automation, MCP client) - **Core functionality:** 50-60% (TUI, database, config) - **Supporting code:** 30-40% (utilities, helpers) - **Legacy code:** 20-30% (minimal coverage acceptable)

---

## Module-Specific Targets

Module	Priority	Target Coverage	Current Est.	Status
<b>codex-tui/spec_kit</b>	Critical	70%	~75%	✓ Exceeded
<b>codex-mcp-client</b>	Critical	70%	~65%	⌚ Near target
<b>codex-tui</b>	High	50%	~45%	⌚ Near target
<b>codex-core</b>	High	50%	~50%	✓ Met
<b>codex-db</b>	High	50%	~60%	✓ Exceeded
<b>config-loader</b>	Medium	40%	~55%	✓ Exceeded
<b>file-search</b>	Medium	40%	~40%	✓ Met
<b>utilities</b>	Low	30%	~35%	✓ Met

**Overall Status:** ✓ **42-48% coverage achieved** (exceeds 40% target)

---

## Testing Pyramid

### Level 1: Unit Tests (~63%)

**Purpose:** Test individual functions/components in isolation

**Characteristics:** - Fast execution (<1s for all unit tests) - No external dependencies (mocked) - High volume (~380 tests)

**What to Unit Test:** - ✓ Pure functions (input → output, no side effects) - ✓ Business logic (validation, parsing, calculations) - ✓ Data structures (serialization, deserialization) - ✓ Error handling (edge cases, invalid inputs)

**What NOT to Unit Test:** - ✗ Integration points (use integration tests) - ✗ UI rendering (hard to test, low ROI) - ✗ External APIs (mock in integration tests)

#### **Example Coverage:**

```
spec_kit/clarify_native.rs: 85% (pattern matching logic)  
spec_kit/checklist_native.rs: 90% (scoring algorithms)  
mcp-client/protocol.rs: 75% (JSON-RPC parsing)
```

---

## **Level 2: Integration Tests (~33%)**

**Purpose:** Test multiple modules working together

**Characteristics:** - Moderate execution time (1-10s per test) - Real module interactions (no mocks between modules) - Medium volume (~200 tests)

**What to Integration Test:** - ✓ Workflow orchestration (plan → tasks → implement) - ✓ Cross-module communication (TUI ↔ MCP client) - ✓ State persistence (database writes/reads) - ✓ Error propagation across modules

#### **Example Coverage:**

```
spec_kit/workflow_integration_tests.rs: 60 tests  
mcp_client/integration_tests.rs: 45 tests  
database/integration_tests.rs: 40 tests
```

---

## **Level 3: E2E Tests (~4%)**

**Purpose:** Test complete user workflows end-to-end

**Characteristics:** - Slow execution (10-60s per test) - Full stack (TUI + backend + database + MCP) - Low volume (~24 tests, high value)

**What to E2E Test:** - ✓ Critical user journeys (/speckit.auto full pipeline) - ✓ Error recovery (retry logic, degradation) - ✓ Tmux session management - ✓ Configuration hot-reload

#### **Example Coverage:**

```
spec_kit/e2e_tests.rs: 12 tests (full automation)  
tmux/e2e_tests.rs: 8 tests (session lifecycle)  
config/e2e_tests.rs: 4 tests (hot-reload)
```

---

# **Test Organization**

## **Per-Module Tests**

### **Structure:**

```
codex-rs/
└── tui/
    ├── src/
    │   └── chatwidget/
    │       └── spec_kit/
    │           ├── clarify_native.rs
    │           └── mod.rs
    └── tests/
        └── spec_kit/
            ├── clarify_native_tests.rs      (unit)
            ├── workflow_integration_tests.rs (integration)
            └── e2e_tests.rs                 (E2E)
```

**Naming Conventions:** - Unit tests: {module}\_tests.rs or #[cfg(test)]  
mod tests in source - Integration tests: {feature}\_integration\_tests.rs  
- E2E tests: e2e\_tests.rs or {workflow}\_e2e.rs

---

## Workspace-Level Tests

**Location:** codex-rs/tests/ (workspace root)

**Purpose:** Cross-crate integration tests

**Example:**

```
codex-rs/tests/
└── tui_mcp_integration.rs      # TUI ↔ MCP client integration
    └── full_pipeline_e2e.rs      # Complete /speckit.auto workflow
        └── hot_reload_integration.rs # Config changes across crates
```

---

## Coverage Measurement

### Tools

**Primary:** cargo-tarpaulin

**Installation:**

```
cargo install cargo-tarpaulin
```

**Usage:**

```
# All modules
cargo tarpaulin --workspace --all-features --timeout 300

# Specific module
cargo tarpaulin -p codex-tui --all-features

# HTML report
cargo tarpaulin --workspace --all-features --out Html
```

**Configuration** (.tarpaulin.toml):

```
[tarpaulin]
timeout = "300s"
exclude-files = [
    "target/*",
```

```
    "*/tests/*",
    "*/benches/*"
]
```

---

## Alternative: cargo-llvm-cov

### Installation:

```
cargo install cargo-llvm-cov
```

### Usage:

```
# Generate coverage
cargo llvm-cov --workspace --all-features --html

# Open report
open target/llvm-cov/html/index.html
```

**Advantage:** More accurate than tarpaulin, faster execution

---

## Critical Path Coverage

### Priority 1: Spec-Kit Automation (70%+ target)

**Critical Flows:** 1. ./speckit.new → SPEC creation 2. ./speckit.auto → Full 6-stage pipeline 3. Quality gates → Checkpoint validation 4. Consensus → Multi-agent synthesis

**Current Coverage:** ~75% ✓

#### Key Test Files:

```
tui/tests/spec_kit/
├── new_native_tests.rs          (95 tests)
├── pipeline_coordinator_tests.rs (85 tests)
├── quality_gate_handler_tests.rs (75 tests)
├── consensus_coordinator_tests.rs (45 tests)
└── workflow_integration_tests.rs (60 tests)
```

---

### Priority 2: MCP Client (70%+ target)

**Critical Flows:** 1. JSON-RPC protocol → Serialization/deserialization 2. Connection lifecycle → Connect, request, disconnect 3. Tool invocation → MCP tool calls 4. Error handling → Retry logic, timeouts

**Current Coverage:** ~65% ⚡

#### Key Test Files:

```
mcp-client/tests/
├── protocol_tests.rs          (40 tests)
├── connection_tests.rs        (30 tests)
├── tool_invocation_tests.rs   (25 tests)
└── integration_tests.rs       (45 tests)
```

---

## Priority 3: Database Layer (50%+ target)

**Critical Flows:** 1. Schema migrations → Up/down migrations  
2. CRUD operations → Insert, query, update, delete  
3. Connection pooling → R2D2 integration  
4. Transaction handling → Rollback on error

**Current Coverage:** ~60% ✓

### Key Test Files:

```
db/tests/
├── schema_tests.rs          (20 tests)
├── crud_tests.rs            (35 tests)
├── pool_tests.rs            (15 tests)
└── transaction_tests.rs     (10 tests)
```

---

## Test Execution Strategy

### Local Development

#### Run all tests:

```
cd codex-rs
cargo test --workspace --all-features
```

#### Run specific module:

```
cargo test -p codex-tui --all-features
```

#### Run specific test:

```
cargo test -p codex-tui
spec_kit::clarify_native::tests::detect_vague_language
```

#### Run with output:

```
cargo test -- --nocapture
```

---

## Pre-Commit Hook

**Location:** .githooks/pre-commit

#### What it runs:

```
# Format check
cargo fmt --all --check

# Linting
cargo clippy --workspace --all-targets --all-features -- -D warnings

# Quick test (compilation only, no execution)
cargo test --workspace --no-run
```

**Time:** ~30 seconds (fast feedback)

**Skip** (if needed):

```
PRECOMMIT_FAST_TEST=0 git commit -m "..."
```

---

## Pre-Push Hook

**Location:** .githooks/pre-push

**What it runs:**

```
# Format check
cargo fmt --all --check

# Linting
cargo clippy --workspace --all-targets --all-features -- -D warnings

# Build
cargo build --workspace --all-features

# Optional: Full test suite (slow)
# cargo test --workspace --all-features
```

**Time:** ~2-5 minutes

**Skip** (if needed):

```
PREPUSH_FAST=0 git push
```

---

## CI/CD Pipeline

**Location:** .github/workflows/rust.yml

**Triggers:** - Push to main - Pull requests - Manual workflow dispatch

**Jobs:** 1. **Test** (parallel matrix): - OS: Ubuntu, macOS, Windows - Rust: stable, beta - Features: all, default

2. **Coverage** (Ubuntu only):
  - Run cargo-tarpaulin
  - Upload to Codecov
  - Comment PR with coverage delta
3. **Lint**:
  - cargo fmt --check
  - cargo clippy -- -D warnings

**Time:** ~10-15 minutes total

---

## Coverage Gaps

### Known Gaps (Acceptable)

**UI Rendering** (~10% coverage): - **Reason:** Ratatui rendering hard to test - **Mitigation:** Manual testing, visual inspection

**Error Handling Paths** (~30% coverage): - **Reason:** Hard to trigger rare errors - **Mitigation:** Property-based testing (proptest)

**Legacy Code** (~20% coverage): - **Reason:** Technical debt, low ROI - **Mitigation:** Refactor on touch, add tests incrementally

---

## Prioritized Improvements

**Phase 1 (Completed):** 40%+ coverage - ✓ Spec-Kit core functionality (360 tests added) - ✓ MCP client protocol (140 tests added) - ✓ Database layer (80 tests added)

**Phase 2 (Optional):** 50%+ coverage - ✘ Error recovery scenarios - ✘ Concurrent operation tests - ✘ Edge case property testing

**Phase 3 (Future):** 60%+ coverage - ✘ UI interaction tests - ✘ Performance regression tests - ✘ Chaos engineering tests

---

## Testing Best Practices

### DO

#### ✓ Test behavior, not implementation:

```
// Good: Test behavior
#[test]
fn clarify_detects_vague_language() {
    let result = clarify("System should be fast");
    assert!(result.has_ambiguities());
    assert_eq!(result.ambiguities[0].pattern, "vague_language");
}

// Bad: Test implementation details
#[test]
fn clarify_calls_regex_find() {
    // Don't test internal regex usage
}
```

#### ✓ Use descriptive test names:

```
#[test]
fn checklist_fails_when_score_below_80() { }

#[test]
fn consensus_degraded_when_only_2_of_3_agents() { }
```

#### ✓ Arrange-Act-Assert pattern:

```
#[test]
fn test_feature() {
    // Arrange: Setup
    let input = "test input";

    // Act: Execute
    let result = functionUnderTest(input);

    // Assert: Verify
    assert_eq!(result, expected);
}
```

---

### DON'T

### **✗ Test framework internals:**

```
// Don't test that Tokio works
#[test]
fn tokio_runtime_spawns_tasks() { }
```

### **✗ Rely on test execution order:**

```
// Tests should be independent
#[test]
fn test_a() { /* modifies global state */ }

#[test]
fn test_b() { /* depends on test_a */ } // ✗ Bad
```

### **✗ Use magic numbers:**

```
// Bad
assert_eq!(result.len(), 42);

// Good
const EXPECTED_ITEM_COUNT: usize = 42;
assert_eq!(result.len(), EXPECTED_ITEM_COUNT);
```

---

## **Summary**

### **Testing Strategy Highlights:**

1. **Coverage Target:** 40%+ (achieved: 42-48%)
2. **Test Pyramid:** 63% unit, 33% integration, 4% E2E
3. **Critical Path Focus:** Spec-Kit (75%), MCP (65%), DB (60%)
4. **Tools:** cargo-tarpaulin, cargo-llvm-cov
5. **CI/CD:** GitHub Actions, pre-commit/pre-push hooks
6. **604 Tests Total:** 100% pass rate

**Next Steps:** - [Test Infrastructure](#) - MockMcpManager, fixtures - [Unit Testing Guide](#) - Patterns and examples - [Integration Testing](#) - Cross-module tests

---

**References:** - Rust testing guide: <https://doc.rust-lang.org/book/ch11-00-testing.html> - Tarpaulin docs: <https://github.com/xd009642/tarpaulin> - Test organization: [codex-rs/\\*/tests/](#) directories

---

## **Unit Testing Guide**

Comprehensive guide to writing effective unit tests.

---

## **Overview**

**Unit Testing Philosophy:** Test individual functions/components in isolation with no external dependencies

**Goals:** - Fast execution (<1s for all unit tests) - High coverage of business logic (70-80% for critical paths) - Deterministic and isolated - Easy to maintain

**Current Status:** - ~380 unit tests (63% of total) - 100% pass rate - Average execution time: ~800ms

---

## Test Structure

### Arrange-Act-Assert Pattern

**Standard Pattern** for all unit tests:

```
#[test]
fn test_feature_behavior() {
    // Arrange: Setup test data
    let input = "test input";
    let expected = "expected output";

    // Act: Execute function under test
    let result = function_under_test(input);

    // Assert: Verify expectations
    assert_eq!(result, expected);
}
```

**Example from codebase** (clarify\_native.rs:365):

```
#[test]
fn test_vague_language_detection() {
    // Arrange
    let detector = PatternDetector::default();
    let mut issues = Vec::new();

    // Act
    detector.check_vague_language("The system should be fast", 1,
&mut issues);

    // Assert
    assert_eq!(issues.len(), 1);
    assert!(issues[0].question.contains("should"));
}
```

---

### Given-When-Then Pattern

**Alternative Pattern** for behavior-driven tests:

```
#[test]
fn test_example() {
    // Given: Initial state
    let state = StateBuilder::new("TEST").build();

    // When: Action occurs
    state.advance_stage();

    // Then: Expected outcome
    assert_eq!(state.current_stage(), Some(SpecStage::Tasks));
```

```
}
```

---

## Naming Conventions

### Test Function Names

**Format:** test\_{what}\_{condition}\_{expected}

**Good Examples:**

```
#[test]
fn test_vague_language_detection() { }

#[test]
fn test_incomplete_markers_flagged_as_critical() { }

#[test]
fn test_quantifier_with_metrics_not_flagged() { }

#[test]
fn test_version_drift_detected_when_prd_newer() { }
```

**Bad Examples:**

```
#[test]
fn test1() { } // ✗ Meaningless

#[test]
fn it_works() { } // ✗ Too vague

#[test]
fn test_the_function() { } // ✗ Not descriptive
```

---

## Test Module Organization

**In-Source Tests** (preferred for unit tests):

```
// src/chatwidget/spec_kit/clarify_native.rs

pub fn detect_ambiguities(prd_content: &str) ->
Result<Vec<AmbiguityIssue>> {
    // Implementation...
}

#[cfg(test)]
mod tests {
    use super::*;

    #[test]
    fn test_vague_language_detection() {
        // Test implementation...
    }

    #[test]
    fn test_incomplete_markers() {
        // Test implementation...
    }
}
```

```
}
```

**Benefits:** - ✓ Tests live next to code - ✓ Private function access - ✓ Excluded from release builds

---

## Testing Pure Functions

### What are Pure Functions?

**Definition:** Functions that: 1. Always return same output for same input 2. Have no side effects (no I/O, no mutations) 3. Don't depend on external state

**Why Test Them:** Easiest to test, highest value per test

---

### Example 1: Pattern Matching

**Function** (clarify\_native.rs):

```
/// Check for vague language
fn check_vague_language(
    &self,
    line: &str,
    line_num: usize,
    issues: &mut Vec<AmbiguityIssue>,
) {
    for (pattern, severity, question, suggestion) in
&self.vague_patterns {
        if let Some(mat) = Regex::new(pattern).unwrap().find(line) {
            issues.push(AmbiguityIssue {
                id: format!("AMB-{:03}", issues.len() + 1),
                severity: *severity,
                pattern_name: "vague_language".to_string(),
                question: question.to_string(),
                suggestion: suggestion.to_string(),
                // ...
            });
        }
    }
}
```

**Unit Test** (clarify\_native.rs:365):

```
#[test]
fn test_vague_language_detection() {
    let detector = PatternDetector::default();
    let mut issues = Vec::new();

    detector.check_vague_language("The system should be fast", 1,
&mut issues);

    assert_eq!(issues.len(), 1);
    assert!(issues[0].question.contains("should"));
    assert_eq!(issues[0].pattern_name, "vague_language");
}
```

**What Makes This a Good Test:** - ✓ Tests one specific pattern (vague language) - ✓ Verifies both detection and message content - ✓ No external dependencies - ✓ Fast (<1ms)

---

## Example 2: Conditional Logic

**Function** (clarify\_native.rs:385):

```
fn check_quantifier_ambiguity(
    &self,
    line: &str,
    line_num: usize,
    issues: &mut Vec<AmbiguityIssue>,
) {
    for (pattern, question, suggestion) in &self.quantifier_patterns
    {
        if Regex::new(pattern).unwrap().is_match(line) {
            // Only flag if NO metrics present
            if !has_metrics(line) {
                issues.push(...);
            }
        }
    }
}
```

**Unit Tests** (clarify\_native.rs:385):

```
#[test]
fn test_quantifier_ambiguity() {
    let detector = PatternDetector::default();
    let mut issues = Vec::new();

    // Should flag: no metrics
    detector.check_quantifier_ambiguity("Must be fast", 1, &mut
issues);
    assert_eq!(issues.len(), 1);

    // Should NOT flag: has metrics
    issues.clear();
    detector.check_quantifier_ambiguity("Must be fast (<100ms)", 1,
&mut issues);
    assert_eq!(issues.len(), 0);
}
```

**What Makes This a Good Test:** - ✓ Tests both branches (with/without metrics) - ✓ Clear positive and negative cases - ✓ Reuses same detector (efficient)

---

## Testing Error Handling

### Testing Error Cases

**Pattern:** Verify function returns Err with expected error type

**Example 1: Missing File:**

```
#[test]
```

```
fn test_analyze_fails_when_prd_missing() {
    let temp_dir = TempDir::new().unwrap();
    let result = check_consistency("SPEC-TEST", temp_dir.path());

    assert!(result.is_err());
    let err = result.unwrap_err();
    assert!(err.to_string().contains("PRD.md not found"));
}
```

---

## Testing Error Messages

**Pattern:** Verify error messages are helpful

**Example:**

```
#[test]
fn test_error_message_includes_spec_id() {
    let result = find_spec_directory("SPEC-INVALID");

    assert!(result.is_err());
    let err = result.unwrap_err();
    assert!(err.to_string().contains("SPEC-INVALID"));
    assert!(err.to_string().contains("not found"));
}
```

---

## Testing Panic Conditions

**Use `should_panic` for panic tests:**

```
#[test]
#[should_panic(expected = "index out of bounds")]
fn test_invalid_index_panics() {
    let stages = vec![SpecStage::Plan];
    let _ = stages[10]; // Should panic
}
```

**Prefer `Result<()>` over panics:**

```
// Good: Returns error
fn validate_index(idx: usize) -> Result<()> {
    if idx >= 6 {
        return Err(anyhow!("Index {} out of range [0, 5]", idx));
    }
    Ok(())
}

// Bad: Panics
fn validate_index(idx: usize) {
    assert!(idx < 6, "Index out of range");
}
```

---

## Testing with Test Data

### Inline Test Data

**Pattern:** Small data inline in test

```
#[test]
fn test_requirement_extraction() {
    let prd_content = r#"
# PRD

## Requirements

- **R1**: User can log in
- **R2**: User can log out
"#;

    let requirements = extract_requirements(prd_content);

    assert_eq!(requirements.len(), 2);
    assert_eq!(requirements[0].id, "R1");
    assert_eq!(requirements[1].id, "R2");
}
```

---

## External Test Fixtures

**Pattern:** Large data from files (see test-infrastructure.md)

```
#[test]
fn test_with_real_prd() -> Result<()> {
    let prd_path = "tests/fixtures/prds/SPEC-DEMO-prd.md";
    let content = std::fs::read_to_string(prd_path)?;

    let ambiguities = detect_ambiguities(&content)?;

    // Real PRD should have known ambiguities
    assert!(ambiguities.len() > 0);
    assert!(ambiguities.iter().any(|a| a.severity == Severity::Critical));

    Ok(())
}
```

---

## Generated Test Data

**Pattern:** Use proptest for fuzz testing (see property-testing-guide.md)

```
use proptest::prelude::*;

proptest! {
    #[test]
    fn test_regex_escape_never_panics(s in ".*") {
        // Should handle any string
        let escaped = regex_escape(&s);
        assert!(escaped.len() >= s.len());
    }
}
```

---

## Testing State Machines

## Example: SpecAutoState Transitions

**State Machine:** - Plan → Tasks → Implement → Validate → Audit → Unlock

**Test Pattern:** Verify transitions

```
#[test]
fn test_stage_advancement() {
    let mut state = StateBuilder::new("SPEC-TEST")
        .starting_at(SpecStage::Plan)
        .build();

    // Initial state
    assert_eq!(state.current_stage(), Some(SpecStage::Plan));
    assert_eq!(state.current_index, 0);

    // Advance to Tasks
    state.advance_stage();
    assert_eq!(state.current_stage(), Some(SpecStage::Tasks));
    assert_eq!(state.current_index, 1);

    // Advance to Implement
    state.advance_stage();
    assert_eq!(state.current_stage(), Some(SpecStage::Implement));
    assert_eq!(state.current_index, 2);
}
```

---

## Testing Invalid Transitions

```
#[test]
fn test_CANNOT_advance_past_unlock() {
    let mut state = StateBuilder::new("SPEC-TEST")
        .starting_at(SpecStage::Unlock)
        .build();

    state.current_index = 5; // Unlock (last stage)

    // Advancing should be no-op or return None
    state.advance_stage();
    assert_eq!(state.current_stage(), None);
}
```

---

## Testing State Invariants

```
#[test]
fn test_state_index_never_negative() {
    let state = StateBuilder::new("SPEC-TEST").build();

    // Type system prevents negative (usize)
    assert!(state.current_index >= 0);

    // But ensure index is valid
    assert!(state.current_index < 6);
}
```

---

# Testing Calculations

## Scoring Functions

**Example** (checklist\_native.rs:350):

```
fn score_testability(prd_content: &str, issues: &mut  
Vec<QualityIssue>) -> f32 {  
    let mut score = 0.0;  
  
    // Check for acceptance criteria (40%)  
    let ac_re = Regex::new(r"(?mi)^###\s+Acceptance  
Criteria").unwrap();  
    if ac_re.is_match(prd_content) {  
        score += 40.0;  
    }  
  
    // Check for test scenarios (20%)  
    let test_re = Regex::new(r"(?mi)^##\s+Test  
(Strategy|Scenarios)").unwrap();  
    if test_re.is_match(prd_content) {  
        score += 20.0;  
    }  
  
    score.max(0.0)  
}
```

## Unit Tests:

```
#[test]  
fn test_score_testability_perfect() {  
    let prd = r#"  
### Acceptance Criteria  
- AC1: Test  
  
## Test Strategy  
- Test  
"#;  
  
    let mut issues = Vec::new();  
    let score = score_testability(prd, &mut issues);  
  
    assert_eq!(score, 60.0); // 40 + 20  
    assert_eq!(issues.len(), 0);  
}  
  
#[test]  
fn test_score_testability_missing_tests() {  
    let prd = r#"  
### Acceptance Criteria  
- AC1: Test  
"#;  
  
    let mut issues = Vec::new();  
    let score = score_testability(prd, &mut issues);  
  
    assert_eq!(score, 40.0); // 40 (AC) + 0 (no tests)  
    assert!(issues.iter().any(|i| i.category == "testability"));  
}
```

```

#[test]
fn test_score_testability_zero() {
    let prd = "# PRD\n\nNo structure";

    let mut issues = Vec::new();
    let score = score_testability(prd, &mut issues);

    assert_eq!(score, 0.0);
    assert!(issues.len() > 0);
}

```

---

## Penalty Calculations

**Example** (checklist\_native.rs:408):

```

fn score_consistency(issues: &[InconsistencyIssue]) -> f32 {
    let critical_count = issues.iter()
        .filter(|i| matches!(i.severity, Severity::Critical))
        .count();
    let important_count = issues.iter()
        .filter(|i| matches!(i.severity, Severity::Important))
        .count();

    let penalty = (critical_count as f32 * 20.0)
        + (important_count as f32 * 10.0);

    (100.0 - penalty).max(0.0)
}

```

**Unit Tests:**

```

#[test]
fn test_score_consistency_perfect() {
    let issues = vec![];
    let score = score_consistency(&issues);
    assert_eq!(score, 100.0);
}

#[test]
fn test_score_consistency_one_critical() {
    let issues = vec![{
        InconsistencyIssue {
            severity: Severity::Critical,
            // ...
        }
    ];
    let score = score_consistency(&issues);
    assert_eq!(score, 80.0); // 100 - 20
}

#[test]
fn test_score_consistency_multiple_issues() {
    let issues = vec![
        InconsistencyIssue { severity: Severity::Critical, /* ... */ },
        InconsistencyIssue { severity: Severity::Critical, /* ... */ },
        InconsistencyIssue { severity: Severity::Important, /* ... */ },
    ],
}

```

```

];
let score = score_consistency(&issues);
assert_eq!(score, 50.0); // 100 - (2*20 + 1*10)
}

#[test]
fn test_score_consistency_floor_at_zero() {
    let issues = vec![
        InconsistencyIssue { severity: Severity::Critical, /* ... */ };
        10
    ];
    let score = score_consistency(&issues);
    assert_eq!(score, 0.0); // Floor at 0 (would be -100)
}

```

---

## Testing Collections

### Testing Filters

```

#[test]
fn test_filter_critical_issues() {
    let issues = vec![
        AmbiguityIssue { severity: Severity::Critical, /* ... */ },
        AmbiguityIssue { severity: Severity::Important, /* ... */ },
        AmbiguityIssue { severity: Severity::Minor, /* ... */ },
    ];

    let critical: Vec<_> = issues.iter()
        .filter(|i| matches!(i.severity, Severity::Critical))
        .collect();

    assert_eq!(critical.len(), 1);
}

```

---

### Testing Sorting

**Example** (clarify\_native.rs:313):

```

fn sort_by_severity(issues: &mut Vec<AmbiguityIssue>) {
    issues.sort_by(|a, b| match (&a.severity, &b.severity) {
        (Severity::Critical, Severity::Critical) => Ordering::Equal,
        (Severity::Critical, _) => Ordering::Less,
        (_, Severity::Critical) => Ordering::Greater,
        // ...
    });
}

```

### Unit Test:

```

#[test]
fn test_sort_by_severity() {
    let mut issues = vec![
        AmbiguityIssue { severity: Severity::Minor, id: "1".into(),
        /* ... */ },
        AmbiguityIssue { severity: Severity::Critical, id:
        "2".into(), /* ... */ },
    ];
}

```

```

        AmbiguityIssue { severity: Severity::Important, id:
"3".into(), /* ... */ },
];

sort_by_severity(&mut issues);

assert_eq!(issues[0].severity, Severity::Critical);
assert_eq!(issues[1].severity, Severity::Important);
assert_eq!(issues[2].severity, Severity::Minor);
}

```

---

## Testing Aggregations

```

#[test]
fn test_count_by_severity() {
    let issues = vec![
        AmbiguityIssue { severity: Severity::Critical, /* ... */ },
        AmbiguityIssue { severity: Severity::Critical, /* ... */ },
        AmbiguityIssue { severity: Severity::Important, /* ... */ },
    ];

    let counts = count_by_severity(&issues);

    assert_eq!(counts.critical, 2);
    assert_eq!(counts.important, 1);
    assert_eq!(counts.minor, 0);
}

```

---

## Testing String Manipulation

### Regex Matching

```

#[test]
fn test_requirement_id_extraction() {
    let line = "- **R42**: User can authenticate";
    let re = Regex::new(r"\*\*(\d+)\*\*").unwrap();

    let cap = re.captures(line).unwrap();
    let id = &cap[1];

    assert_eq!(id, "42");
}

```

---

### String Transformations

**Example** (clarify\_native.rs:334):

```

fn truncate_context(text: &str, max_len: usize) -> String {
    if text.len() <= max_len {
        text.to_string()
    } else {
        format!("{}...", &text[..max_len])
    }
}

```

## Unit Tests:

```
#[test]
fn test_truncate_short_text() {
    let text = "Short";
    let result = truncate_context(text, 10);
    assert_eq!(result, "Short");
}

#[test]
fn test_truncate_long_text() {
    let text = "This is a very long text that should be truncated";
    let result = truncate_context(text, 10);
    assert_eq!(result, "This is a ...");
    assert_eq!(result.len(), 13); // 10 + ...
}

#[test]
fn test_truncate_exact_length() {
    let text = "Exactly10!";
    let result = truncate_context(text, 10);
    assert_eq!(result, "Exactly10!");
}
```

---

## Regex Escaping

### Function (clarify\_native.rs:349):

```
fn regex_escape(s: &str) -> String {
    s.chars()
        .map(|c| match c {
            '\\' | '.' | '+' | '*' | '?' | '(' | ')' | '['
            | ']' | '{' | '}' | '^' | '$' => {
                format!("\\{}", c)
            }
            _ => c.to_string(),
        })
        .collect()
}
```

### Unit Tests:

```
#[test]
fn test_regex_escape_special_chars() {
    assert_eq!(regex_escape("a.b"), "a\\.b");
    assert_eq!(regex_escape("a*b"), "a\\*b");
    assert_eq!(regex_escape("a?b"), "a\\?b");
    assert_eq!(regex_escape("a(b)", "a\\(b\\)");
}

#[test]
fn test_regex_escape_normal_chars() {
    assert_eq!(regex_escape("abc"), "abc");
    assert_eq!(regex_escape("123"), "123");
}

#[test]
fn test_regex_escape_multiple_special() {
    assert_eq!(regex_escape("a.b*c?"), "a\\.b\\*c\\?");
```

```
}
```

---

## Testing File Operations (with TempDir)

### Setup Pattern

```
use tempfile::TempDir;

#[test]
fn test_write_and_read_prd() -> Result<()> {
    // Arrange: Create temp directory
    let temp_dir = TempDir::new()?;
    let spec_dir = temp_dir.path().join("docs/SPEC-TEST-test");
    std::fs::create_dir_all(&spec_dir)?;

    let prd_path = spec_dir.join("PRD.md");
    let content = "# PRD\n\n## Goal\nTest";

    // Act: Write file
    std::fs::write(&prd_path, content)?;

    // Assert: Read and verify
    let read_content = std::fs::read_to_string(&prd_path)?;
    assert_eq!(read_content, content);

    Ok(())
    // TempDir auto-cleaned on drop
}
```

---

### Testing Directory Creation

```
#[test]
fn test_create_spec_directory() -> Result<()> {
    let temp_dir = TempDir::new()?;
    let spec_id = "SPEC-TEST-001";

    let spec_dir = create_spec_directory(temp_dir.path(), spec_id)?;

    assert!(spec_dir.exists());
    assert!(spec_dir.is_dir());
    assert!(spec_dir.ends_with("SPEC-TEST-001-test"));

    Ok(())
}
```

---

## Testing with Mocks

### MockMcpManager Usage

**Pattern:** Replace real MCP with mock

```
#[tokio::test]
async fn test_consensus_fetch() -> Result<()> {
```

```

// Arrange: Setup mock
let mut mock = MockMcpManager::new();
mock.add_fixture(
    "local-memory",
    "search",
    Some("SPEC-TEST plan"),
    json!({"memory": {"content": "Agent response"}})
);

// Act: Call function that uses MCP
let results = fetch_consensus("SPEC-TEST", SpecStage::Plan,
&mock).await?;

// Assert: Verify results
assert_eq!(results.len(), 1);

Ok()
}

```

See [test-infrastructure.md](#) for details.

---

## Table-Driven Tests

### Pattern: Multiple Test Cases

```

#[test]
fn test_stage_index_mapping() {
    let test_cases = vec![
        (0, Some(SpecStage::Plan)),
        (1, Some(SpecStage::Tasks)),
        (2, Some(SpecStage::Implement)),
        (3, Some(SpecStage::Validate)),
        (4, Some(SpecStage::Audit)),
        (5, Some(SpecStage::Unlock)),
        (6, None),
        (100, None),
    ];
}

for (index, expected) in test_cases {
    let mut state = StateBuilder::new("SPEC-TEST").build();
    state.current_index = index;

    assert_eq!(
        state.current_stage(),
        expected,
        "Failed for index {}",
        index
    );
}
}

```

**Benefits:** - ✓ Compact (many cases in one test) - ✓ Easy to add new cases - ✓ Clear failure messages

---

### Parameterized Tests (with rstest)

**Add to Cargo.toml:**

```
[dev-dependencies]
rtest = "0.18"
```

**Usage:**

```
use rtest::rtest;

#[rstest]
#[case("should", Severity::Important)]
#[case("must", Severity::Critical)]
#[case("TBD", Severity::Critical)]
#[case("TODO", Severity::Important)]
fn test_vague_language_severity(#[case] pattern: &str, #[case]
expected: Severity) {
    let detector = PatternDetector::default();
    let mut issues = Vec::new();

    detector.check_vague_language(
        &format!("The system {} work", pattern),
        1,
        &mut issues
    );

    assert_eq!(issues.len(), 1);
    assert_eq!(issues[0].severity, expected);
}
```

---

## Common Assertions

### Equality

```
assert_eq!(actual, expected);
assert_ne!(actual, unexpected);
```

---

### Boolean

```
assert!(condition);
assert!(!condition);
```

---

### Contains

```
assert!(vec.contains(&item));
assert!(string.contains("substring"));
```

---

### Custom Messages

```
assert_eq!(
    actual,
    expected,
    "Expected {}, got {} (context: {})",
    expected,
    actual,
```

```
    context  
);
```

---

## Floating Point

```
// Don't use assert_eq! for floats  
// Use approx crate instead  
  
use approx::assert_relative_eq;  
  
assert_relative_eq!(actual, expected, epsilon = 0.001);
```

---

## Best Practices

### DO

#### ✓ Test one thing per test:

```
#[test]  
fn test_vague_language_detection() {  
    // Only tests vague language, nothing else  
}  
  
#[test]  
fn test_incomplete_markers() {  
    // Only tests incomplete markers  
}
```

---

#### ✓ Use descriptive names:

```
#[test]  
fn test_quantifier_with_metrics_not_flagged() {  
    // Clear what's being tested  
}
```

---

#### ✓ Test edge cases:

```
#[test]  
fn test_truncate_empty_string() {  
    assert_eq!(truncate_context("", 10), "");  
}  
  
#[test]  
fn test_score_consistency_floor_at_zero() {  
    // Test penalty doesn't go negative  
}
```

---

#### ✓ Keep tests independent:

```
#[test]  
fn test_a() {  
    let state = StateBuilder::new("TEST-A").build();  
    // Uses own state, doesn't affect other tests  
}
```

```
#[test]
fn test_b() {
    let state = StateBuilder::new("TEST-B").build();
    // Independent
}
```

---

#### ✓ Use setup functions for common data:

```
fn create_test_prd() -> String {
    r#"
# PRD
## Requirements
- **R1**: Test
    "#.to_string()
}

#[test]
fn test_with_prd() {
    let prd = create_test_prd();
    // Use prd...
}
```

---

## DON'T

#### ✗ Test implementation details:

```
// Bad: Tests internal regex pattern
#[test]
fn test_regex_pattern_is_correct() {
    assert_eq!(VAGUE_PATTERN, r"(should|could|might)");
}

// Good: Tests behavior
#[test]
fn test_vague_language_detected() {
    // Tests that "should" is flagged
}
```

---

#### ✗ Rely on test execution order:

```
// Bad: test_b depends on test_a running first
static mut SHARED_STATE: i32 = 0;

#[test]
fn test_a() {
    unsafe { SHARED_STATE = 42; }
}

#[test]
fn test_b() {
    unsafe { assert_eq!(SHARED_STATE, 42); } // ✗ Flaky
}
```

---

#### ✗ Use magic numbers:

```
// Bad
```

---

```

assert_eq!(score, 42.0);

// Good
const EXPECTED_SCORE: f32 = 42.0;
assert_eq!(score, EXPECTED_SCORE);

// Or explain inline
assert_eq!(score, 60.0); // 40 (AC) + 20 (test strategy)

```

---

### **✗ Test too much in one test:**

```

// Bad: Tests everything at once
#[test]
fn test_entire_quality_system() {
    // 100 lines of setup
    // Tests clarify, analyze, checklist
    // Hard to debug when fails
}

// Good: Split into focused tests
#[test]
fn test_clarify_detects_vague_language() {}

#[test]
fn test_analyze_finds_missing_requirements() {}

#[test]
fn test_checklist_scores_completeness() {}

```

---

### **✗ Skip cleanup (use TempDir):**

```

// Bad: Leaves files behind
#[test]
fn test_write_file() {
    std::fs::write("/tmp/test.txt", "data")?;
    // File persists after test
}

// Good: Auto-cleanup
#[test]
fn test_write_file() -> Result<()> {
    let temp_dir = TempDir::new()?;
    std::fs::write(temp_dir.path().join("test.txt"), "data")?;
    Ok(())
    // temp_dir dropped, files deleted
}

```

---

## **Running Tests**

### **Run All Unit Tests**

```

cd codex-rs
cargo test --lib

```

**Explanation:** `--lib`: Only library tests (no integration tests) - Runs all `#[cfg(test)] mod tests { }` blocks

---

## Run Specific Module

```
cargo test -p codex-tui --lib clarify_native
```

**Breakdown:** -p codex-tui: Package ---lib: Unit tests only - clarify\_native: Module filter

---

## Run Specific Test

```
cargo test -p codex-tui test_vague_language_detection
```

**Output:**

```
running 1 test
test
chatwidget::spec_kit::clarify_native::tests::test_vague_language_detection
... ok

test result: ok. 1 passed; 0 failed; 0 ignored; 0 measured
```

---

## Run with Output

```
cargo test -- --nocapture
```

**Shows** `println!()` output even for passing tests.

---

## Run with Threads

```
# Single-threaded (for debugging)
cargo test -- --test-threads=1

# Parallel (default)
cargo test -- --test-threads=8
```

---

## Test Coverage

### Measure Coverage

**Using tarpaulin:**

```
cargo tarpaulin -p codex-tui --lib
```

**Output:**

```
|| Tested/Total Lines:
|| tui/src/chatwidget/spec_kit/clarify_native.rs: 89/120
||
|| Coverage: 74.2%
```

---

### Improve Coverage

**Identify Untested Lines:**

```
cargo tarpaulin -p codex-tui --lib --out Html  
open target/tarpaulin/index.html
```

**HTML Report** shows: - ✓ Green: Covered - ✗ Red: Not covered - △  
Yellow: Partially covered

---

## Summary

### Unit Testing Best Practices:

1. **Structure:** Use Arrange-Act-Assert pattern
2. **Naming:** test\_{what}\_{condition}\_{expected}
3. **Scope:** One thing per test
4. **Independence:** No shared state
5. **Speed:** Fast (<1ms typical)
6. **Coverage:** 70-80% for critical paths
7. **Cleanup:** Use TempDir for filesystem tests
8. **Mocks:** Use MockMcpManager for MCP

**Test Types Covered:** - ✓ Pure functions (pattern matching, calculations) - ✓ Error handling (missing files, invalid input) - ✓ State machines (transitions, invariants) - ✓ Collections (filtering, sorting, aggregation) - ✓ String manipulation (regex, truncation, escaping) - ✓ File operations (with TempDir)

**Next Steps:** - [Integration Testing Guide](#) - Cross-module tests - [Property Testing Guide](#) - Generative testing - [Test Infrastructure](#) - MockMcpManager, fixtures

---

**References:** - Rust testing guide: <https://doc.rust-lang.org/book/ch11-00-testing.html> - Example tests: codex-rs/tui/src/chatwidget/spec\_kit/\*/tests.rs - Test infrastructure: codex-rs/tui/tests/common/

---