

SPEC-DOC-006: Configuration & Customization Guide

- Objectives
- Scope
 - In Scope
 - Out of Scope
- Deliverables
- Success Criteria
- Related SPECs
- Agent Configuration
 - Overview
 - Agent Configuration Schema
 - Agent Fields
 - Default Agent Configuration
 - 5-Agent Setup
 - Agent Properties
 - name vs canonical_name
 - read_only Flag
 - enabled Flag
 - Advanced Agent Configuration
 - args_read_only vs args_write
 - Environment Variables
 - Per-Agent Instructions
- Subagent Commands
 - Default Commands
 - Subagent Command Configuration
 - Custom Subagent Command
- Quality Gate Integration
 - Agent Selection for Quality Gates
- Agent Cost Tiers
 - Cost Comparison
 - Strategic Agent Routing
- Example Configurations
 - Minimal (Single Agent)
 - Balanced (3 Agents)
 - Premium (5 Agents + Specialist)
- Debugging Agent Configuration
 - List Configured Agents
 - Validate Agent Commands
- Best Practices
 - 1. Use Canonical Names Consistently
 - 2. Enable Minimum Required Agents
 - 3. Use args_read_only for Cost Savings
 - 4. Leverage Per-Agent Instructions
- Summary

- Profiles
- Validation Rules
 - Required Fields
 - Type Validation
 - Semantic Validation
- Error Handling
- Summary
- Environment Variables
 - Overview
 - Core Environment Variables
 - CODEX_HOME / CODE_HOME
 - API Keys
 - OPENAI_API_KEY
 - ANTHROPIC_API_KEY
 - GOOGLE_API_KEY
 - AZURE_OPENAI_API_KEY
 - Custom Provider API Keys
 - Model Configuration Overrides
 - CODEX_MODEL
 - CODEX_PROVIDER
 - OPENAI_BASE_URL
 - OPENAI_WIRE_API
 - Spec-Kit Environment Variables
 - SPEC_OPS_CARGO_MANIFEST
 - SPEC_OPS_ALLOW_DIRTY
 - SPEC_OPS_TELEMETRY_HAL
 - SPEC_OPS_HAL_SKIP
 - SPECKIT_QUALITY_GATES_*
 - Logging and Debugging
 - RUST_LOG
 - RUST_BACKTRACE
 - Sandbox and Security
 - CODEX_SANDBOX_NETWORK_DISABLED
 - CI/CD Environment Variables
 - CI
 - GITHUB_ACTIONS
 - CODEX_AUTO_UPGRADE
 - Shell Environment Policy
 - Shell Environment Inheritance
 - Override Shell Environment Policy
 - MCP Server Environment Variables
 - MCP-Specific Variables
 - Global MCP Environment
 - HAL Secret Environment Variables
 - HAL_SECRET_KAVEDARR_API_KEY
 - Testing Environment Variables
 - PRECOMMIT_FAST_TEST
 - PREPUSH_FAST
 - Complete Environment Variable Reference
 - Core Variables
 - API Keys
 - Model Overrides
 - Spec-Kit Variables
 - Sandbox and Security
 - CI/CD Variables
 - Testing Variables
 - Best Practices
 - 1. Store Secrets in Environment Variables
 - 2. Use .env Files (Local Development)
 - 3. Use Profiles for Environment-Specific Config

4. Document Required Environment Variables

Debugging Environment Variables

- List Active Environment Variables
- Check Effective Configuration

Summary

Hot-Reload

- Overview
- Architecture
 - Reload Flow
- Configuration
 - Enable Hot-Reload
 - Configuration Fields
 - Debouncing
 - Debounce Tuning
 - Watch Additional Files
- Reload Events
 - Event Types
 - Event Flow
 - TUI Notifications
- Reload Performance
 - Latency Breakdown
 - Lock Performance
 - CPU Overhead
- Validation
 - Schema Validation
 - Semantic Validation
 - Validation Failure Behavior
- Deferring Reloads
 - When to Defer
 - Deferred Reload Behavior
- Change Detection
 - Detecting Config Changes
 - TUI Notification with Changes
- Debugging Hot-Reload
 - Enable Debug Logging
 - Test Hot-Reload
 - Disable Hot-Reload (Troubleshooting)

Best Practices

- 1. Use Default Debounce (2000ms)
- 2. Validate Config Before Saving
- 3. Monitor Reload Notifications
- 4. Test Config Changes Incrementally

Summary

MCP Servers

- Overview
- MCP Server Configuration
 - Basic Configuration
 - Configuration Fields
- Built-in MCP Servers
 - local-memory (Knowledge Persistence)
 - git-status (Repository Inspection)
 - HAL (Policy Validation)
- Custom MCP Servers
 - Creating a Custom Server
 - Custom Server Example (Node.js)

MCP Server Lifecycle

- Startup Process
- Lazy Loading
- Startup Optimization
- Shutdown Process

- Environment Variables
 - Server-Specific Environment
 - Global Environment Variables
- Timeouts and Retries
 - Startup Timeout
 - Tool Call Timeout
 - Retry Logic
- Debugging MCP Servers
 - Enable MCP Logging
 - Test MCP Server Manually
 - Check MCP Server Status
 - Force Restart MCP Server
- Common MCP Servers
 - Filesystem Server
 - HTTP Server
 - Database Servers
 - Custom API Server
- Security Considerations
 - 1. Validate Command Paths
 - 2. Avoid Secrets in Config
 - 3. Restrict Network Access
- Best Practices
 - 1. Use Lazy Loading
 - 2. Set Appropriate Timeouts
 - 3. Monitor MCP Server Logs
 - 4. Test Servers with MCP Inspector
- Summary
- Model Configuration
 - Overview
 - Basic Model Configuration
 - Minimal Setup
 - Model Selection
 - Provider Configuration
 - OpenAI (Default)
 - Anthropic (Claude)
 - Google (Gemini)
 - Ollama (Local)
 - Azure OpenAI
 - Custom Provider
 - Reasoning Configuration
 - Reasoning Effort
 - Reasoning Summary
 - Model Verbosity (GPT-5 Only)
 - Context Window Configuration
 - Context Window Size
 - Max Output Tokens
 - Network Tuning
 - Request Retries
 - Stream Retries
 - Stream Idle Timeout
 - Wire API Selection
 - Responses API (Default for GPT-5/o3)
 - Chat Completions API (Legacy)
 - Advanced Configuration
 - Force Reasoning Support
 - Disable Response Storage (ZDR Accounts)
 - Configuration Examples
 - Premium Quality Setup
 - Fast Iteration Setup
 - Local Development Setup

- Multi-Provider Setup
 - Debugging Model Configuration
 - Check Effective Configuration
 - Test Provider Connection
 - Summary
- Precedence System
 - Overview
 - Precedence Order
 - Tier 1: CLI Flags (Highest)
 - Tier 2: Shell Environment
 - Tier 3: Profile
 - Tier 4: Config File
 - Tier 5: Defaults (Lowest)
 - Precedence Examples
 - Example 1: Model Selection
 - Example 2: API Key
 - Example 3: Approval Policy
 - Example 4: Complex Nested Config
 - Special Cases
 - Shell Environment Policy Override
 - Profile Selection Precedence
 - Precedence Table
 - Debugging Precedence
 - Check Effective Configuration
 - Trace Configuration Source
- Best Practices
 - 1. Use Environment Variables for Secrets
 - 2. Use Profiles for Workflows
 - 3. Use CLI Flags for One-Off Overrides
 - 4. Keep config.toml for Persistent Preferences
- Summary

- Quality Gate Customization
- Overview
- Quality Gate Configuration
 - Basic Configuration
 - Field Reference
- Agent Selection Strategy
 - Multi-Agent Consensus (Plan, Validate)
 - Single-Agent Deterministic (Tasks)
 - Premium Consensus (Audit, Unlock)
- Custom Configurations
 - Cost-Optimized Setup
 - Premium Quality Setup
 - Specialist Configuration
- Per-Checkpoint Overrides
 - Override at Runtime
 - Environment Variable Overrides
- Consensus Thresholds
 - Minimum Consensus
 - Strict Consensus
 - Relaxed Consensus
- Degradation Handling
 - Agent Failure Behavior
 - Empty Consensus Handling
- Quality Gate Validation
 - Startup Validation
 - Runtime Validation
- Example Configurations
 - Balanced (Default)
 - Cost-Optimized

- Premium Quality
- Single-Agent (Development)
- Debugging Quality Gates
 - Check Quality Gate Configuration
 - Validate Agent Availability
- Best Practices
 - 1. Use 3 Agents for Consensus
 - 2. Use 1 Agent for Deterministic Tasks
 - 3. Reserve Premium Agents for Critical Gates
 - 4. Test Quality Gate Configuration
- Summary

Template Customization

- Overview
- Template Structure
 - Template Format
- Installing Templates
 - Method 1: Manual Installation
 - Method 2: Install from URL
 - Method 3: Install from Git Repository
- Using Templates
 - Apply Template Once
 - Set Default Template
 - Template Precedence
- Creating Custom Templates
 - Step 1: Define Template Metadata
 - Step 2: Define Configuration
 - Step 3: Test Template
 - Step 4: Version and Document
- Template Examples
 - Cost-Optimized Template
 - CI/CD Template
 - Team Standard Template
- Template Versioning
 - Version Schema
 - Version Compatibility
 - Automatic Template Updates
- Template Repositories
 - Official Template Repository
 - Install from Repository
 - Create Your Own Repository
- Debugging Templates
 - Validate Template
 - Dry-Run Template
 - Compare Templates
- Best Practices
 - 1. Version Templates Semantically
 - 2. Document Template Usage
 - 3. Test Templates Before Distribution
 - 4. Use Templates for Team Consistency
- Summary

Theme System

- Overview
- Theme Configuration
 - Basic Configuration
 - Built-in Themes
 - Theme Selection
 - Theme Previews
- Custom Color Overrides
 - Override Individual Colors
 - Available Color Fields

- Complete Custom Theme
- Syntax Highlighting
 - Highlight Configuration
 - Custom Syntax Theme
- Terminal Background Detection
 - Auto-Detection Process
 - Cached Terminal Background
 - Force Re-Detection
- Accessibility Options
 - High Contrast Mode
 - Large Text (Terminal Setting)
 - Color Blindness Support
- Theme Customization Examples
 - Solarized Dark
 - Gruvbox Dark
 - Dracula
- Debugging Themes
 - Test Theme
 - Preview All Themes
 - Dump Current Theme
 - Validate Custom Theme
- Hot-Reload Support
 - Live Theme Changes
 - Live Color Tweaking
- Spinner Customization
 - Built-in Spinners
 - Spinner Configuration
 - Custom Spinner
- Stream Animation
 - Stream Configuration
 - Responsive Preset
- Best Practices
 - 1. Use Built-in Themes When Possible
 - 2. Override Colors Sparingly
 - 3. Test Themes in Different Scenarios
 - 4. Consider Accessibility
- Summary

SPEC-DOC-006: Configuration & Customization Guide

Status: Pending **Priority:** P1 (Medium) **Estimated Effort:** 8-12 hours
Target Audience: Power users, advanced users **Created:** 2025-11-17

Objectives

Comprehensive guide to configuring and customizing the codex CLI:
1. Configuration file structure (config.toml complete schema)
2. 5-tier precedence (CLI, shell, profile, TOML, defaults)
3. Model configuration (providers, reasoning, profiles)
4. Agent configuration (5 agents, subagent commands, quality gates)
5. Quality gate customization (per-checkpoint agent selection)
6. Hot-reload configuration (config_reload.rs, 300ms debounce)
7. MCP server configuration (server definitions, lifecycle)
8. Environment variables

(CODEX_HOME, API keys, overrides) 9. Templates (installation, customization, versioning) 10. Theme system (TUI themes, accessibility)

Scope

In Scope

- Complete config.toml reference (all sections)
- 5-tier precedence system (with examples)
- Model provider configuration (OpenAI, Anthropic, Google, Ollama)
- Agent configuration (gemini, claude, code, gpt_pro, gpt_codex)
- Quality gate customization (per-checkpoint overrides)
- Hot-reload mechanism (300ms debounce, watch system)
- MCP server definitions (local-memory, git-status, hal, custom)
- Environment variables (CODEX_HOME, API_KEY, SPEC_OPS)
- Template customization (installing, modifying, versioning)
- TUI theme customization (colors, accessibility)

Out of Scope

- Architecture of config system (see SPEC-DOC-002)
 - Installation and setup (see SPEC-DOC-001)
 - Security of secrets (see SPEC-DOC-007)
-

Deliverables

1. **content/config-reference.md** - Complete config.toml schema
 2. **content/precedence-system.md** - 5-tier precedence with examples
 3. **content/model-configuration.md** - Provider setup, reasoning effort
 4. **content/agent-configuration.md** - 5 agents, subagent commands
 5. **content/quality-gate-customization.md** - Per-checkpoint overrides
 6. **content/hot-reload.md** - Config reload mechanism, debouncing
 7. **content/mcp-servers.md** - MCP server definitions, custom servers
 8. **content/environment-variables.md** - All env vars, overrides
 9. **content/template-customization.md** - Installing, modifying templates
 10. **content/theme-system.md** - TUI themes, accessibility options
-

Success Criteria

- Complete config.toml schema documented
 - 5-tier precedence clearly explained with examples
 - All agent configurations documented
 - Quality gate customization guide complete
 - Environment variables comprehensive list
 - Template customization tutorial complete
-

Related SPECs

- SPEC-DOC-000 (Master)
 - SPEC-DOC-001 (User Onboarding - basic config)
 - SPEC-DOC-002 (Core Architecture - config system internals)
 - SPEC-DOC-003 (Spec-Kit - quality gate config)
-

Status: Structure defined, content pending

Agent Configuration

Multi-agent setup, subagent commands, and agent profiles.

Overview

The **multi-agent system** enables consensus-driven decision-making through parallel execution of multiple AI agents.

Use Cases: - **Consensus Planning** - 3+ agents agree on architecture decisions - **Quality Gates** - Multiple agents validate test strategies - **Diverse Perspectives** - Combine strengths of different models

Configuration: `[[agents]]` array in `config.toml`

Agent Configuration Schema

Agent Fields

```
[[agents]]
  name = "gemini"                                # Display name
  canonical_name = "gemini"                      # Canonical identifier (for
quality gates)
  command = "gemini"                            # Executable command
  args = []                                     # Command arguments
  read_only = false                             # Force read-only mode
  enabled = true                               # Enable/disable agent
  description = "Google Gemini"                # Human-readable description
  env = {}                                      # Environment variables
  args_read_only = []                           # Args for read-only mode
(optional)
  args_write = []                             # Args for write mode (optional)
  instructions = ""                          # Per-agent instructions
(optional)
```

Default Agent Configuration

5-Agent Setup

```

# ~/.code/config.toml

#
=====

# Agent 1: Gemini (Fast, Cheap Consensus)
#
=====

[[agents]]
name = "gemini"
canonical_name = "gemini"
command = "gemini"
args = []
read_only = false
enabled = true
description = "Google Gemini Flash - Fast consensus agent (12.5x
cheaper than GPT-5)"

#
=====

# Agent 2: Claude (Balanced Reasoning)
#
=====

[[agents]]
name = "claude"
canonical_name = "claude"
command = "claude"
args = []
read_only = false
enabled = true
description = "Anthropic Claude Haiku - Balanced reasoning (12x
cheaper than GPT-5)"

#
=====

# Agent 3: Code (Strategic Planning)
#
=====

[[agents]]
name = "code"
canonical_name = "gpt_pro"
command = "code"
args = ["--model", "gpt-5"]
read_only = false
enabled = true
description = "OpenAI GPT-5 - Strategic planning and complex
reasoning"

#
=====

# Agent 4: GPT-Codex (Code Generation)

```

```

#
=====
[[agents]]
name = "gpt_codex"
canonical_name = "gpt_codex"
command = "code"
args = ["--model", "gpt-5-codex"]
read_only = false
enabled = true
description = "OpenAI GPT-5-Codex - Specialized code generation"

#
=====
# Agent 5: GPT-Pro (Premium Reasoning)
#
=====

[[agents]]
name = "gpt_pro"
canonical_name = "gpt_pro"
command = "code"
args = ["--model", "o3", "--config", "model_reasoning_effort=high"]
read_only = false
enabled = false # Disabled by default (premium cost)
description = "OpenAI o3 - Premium reasoning for critical decisions"

```

Agent Properties

name vs canonical_name

name: Display name, can change

canonical_name: Stable identifier used in quality gates

Example:

```

[[agents]]
name = "claude-sonnet"          # Display name (can evolve)
canonical_name = "claude"        # Canonical name (stable)
command = "anthropic"

```

Quality gate reference:

```

[quality_gates]
plan = ["claude"] # Uses canonical_name, not name

```

Benefit: Can rename display names without breaking quality gate configs

read_only Flag

Purpose: Force agent to run in read-only mode

Use Case: Agents that should never write files

Example:

```
[[agents]]
name = "readonly-advisor"
canonical_name = "advisor"
command = "gemini"
read_only = true # Never allow writes
enabled = true
```

enabled Flag

Purpose: Temporarily disable agent without removing config

Use Case: Testing, cost control, debugging

Example:

```
[[agents]]
name = "gpt_pro"
canonical_name = "gpt_pro"
command = "code"
args = ["--model", "o3"]
enabled = false # Disable premium agent to save cost
```

Advanced Agent Configuration

args_read_only vs args_write

Purpose: Different arguments for read vs write modes

Example:

```
[[agents]]
name = "claude"
canonical_name = "claude"
command = "anthropic"
args = [] # Default args

# Read-only mode: Use faster, cheaper model
args_read_only = ["--model", "claude-3-haiku"]

# Write mode: Use more capable model
args_write = ["--model", "claude-3-5-sonnet"]
```

Behavior: Automatically selects appropriate args based on operation mode

Environment Variables

Purpose: Pass environment variables to agent process

Example:

```
[[agents]]
```

```
name = "custom-agent"
canonical_name = "custom"
command = "/path/to/agent"
args = []
env = {
    LOG_LEVEL = "debug",
    CUSTOM_CONFIG = "/path/to/config.json",
    FEATURE_FLAGS = "experimental"
}
```

Use Case: Custom agents, debugging, feature flags

Per-Agent Instructions

Purpose: Prepend instructions to every prompt sent to this agent

Example:

```
[[agents]]
name = "security-focused"
canonical_name = "security"
command = "claude"
args = []
instructions = """
You are a security-focused code reviewer. Always prioritize:
1. Input validation and sanitization
2. Authentication and authorization checks
3. Secure cryptographic practices
4. Protection against OWASP Top 10 vulnerabilities

Flag any potential security issues with HIGH severity.
"""
```

Subagent Commands

Default Commands

The spec-kit framework provides **13 slash commands** that use agents:

Native (Tier 0 - Zero agents, FREE): - /speckit.new - SPEC creation (template-based, no agents) - /speckit.clarify - Ambiguity detection (heuristics) - /speckit.analyze - Consistency checking (structural diff) - /speckit.checklist - Quality scoring (rubric) - /speckit.status - Status dashboard (native)

Single-Agent (Tier 1 - 1 agent, ~\$0.10): - /speckit.specify - PRD drafting (gpt5-low) - /speckit.tasks - Task decomposition (gpt5-low)

Multi-Agent (Tier 2 - 2-3 agents, ~\$0.35): - /speckit.plan - Architectural planning (gemini-flash, claude-haiku, gpt5-medium) - /speckit.validate - Test strategy (gemini-flash, claude-haiku, gpt5-medium) - /speckit.implement - Code generation (gpt_codex HIGH, claude-haiku validator)

Premium (Tier 3 - 3 premium agents, ~\$0.80): - /speckit.audit - Compliance/security (gemini-pro, claude-sonnet, gpt5-high) - /speckit.unlock - Ship decision (gemini-pro, claude-sonnet, gpt5-high)

Full Pipeline (Tier 4 - Strategic routing, ~\$2.70): - /speckit.auto - Full 6-stage pipeline

Subagent Command Configuration

Table Format: [[subagents.commands]]

```
[[subagents.commands]]
name = "plan"                      # Command name (/speckit.plan)
read_only = true                     # Force read-only mode
agents = ["gemini", "claude", "code"] # Agents to use
orchestrator_instructions = "Focus on architectural decisions and
trade-offs."
agent_instructions = "Provide detailed reasoning for all
recommendations."
```

Fields: - name (string): Command name (matches /speckit.<name>) - read_only (boolean): Force read-only mode (default: command-specific) - agents (array): Agent names to enable (default: all enabled agents) - orchestrator_instructions (string): Extra instructions for orchestrator - agent_instructions (string): Instructions appended to each agent prompt

Custom Subagent Command

Example: Add custom consensus command

```
# ~/.code/config.toml

[[subagents.commands]]
name = "review" # Creates /speckit.review command
read_only = true
agents = ["claude", "gpt_pro"]
orchestrator_instructions = """
Focus on:
1. Code quality and maintainability
2. Performance implications
3. Security concerns
4. Test coverage adequacy
"""
agent_instructions = """
Provide specific, actionable feedback with code examples.
"""
```

Usage:

```
/speckit.review SPEC-KIT-065
```

Quality Gate Integration

Agent Selection for Quality Gates

Quality gates reference agents by **canonical_name**:

```
# Agents configuration
[[agents]]
name = "gemini-flash"
canonical_name = "gemini" # ← Used in quality gates
# ...

[[agents]]
name = "claude-haiku"
canonical_name = "claude" # ← Used in quality gates
# ...

# Quality gates configuration
[quality_gates]
plan = ["gemini", "claude", "code"] # Uses canonical_name
tasks = ["gemini"]
validate = ["gemini", "claude", "code"]
```

Validation: Config loader checks that all quality gate agents exist

Agent Cost Tiers

Cost Comparison

Based on OpenAI GPT-5 baseline (1.0x):

Agent	Model	Cost per 1M tokens	Relative Cost
gemini	Gemini Flash	\$0.40	12.5x cheaper
claude	Claude Haiku	\$0.40	12x cheaper
code	GPT-5	\$5.00	1.0x (baseline)
gpt_codex	GPT-5-Codex	\$5.00	1.0x
gpt_pro	o3 (high effort)	\$20.00	4x more expensive

Strategic Agent Routing

SPEC-KIT-070: Cost optimization via tiered agent selection

Principle: “Agents for reasoning, NOT transactions”

Tier 0 (Native): Pattern matching → FREE

```
# No agents needed for:
- /speckit.new (template-based SPEC-ID generation)
- /speckit.clarify (regex-based ambiguity detection)
- /speckit.analyze (structural consistency checking)
```

Tier 1 (Single Agent): Simple reasoning → \$0.10

```
[[subagents.commands]]
name = "specify"
agents = ["gpt5-low"] # Single cheap agent
```

Tier 2 (Multi-Agent): Complex decisions → \$0.35

```
[quality_gates]
plan = ["gemini", "claude", "gpt5-medium"] # 3 agents, diverse perspectives
```

Tier 3 (Premium): Critical decisions → \$0.80

```
[quality_gates]
unlock = ["gemini-pro", "claude-sonnet", "gpt5-high"] # Quality over cost
```

Example Configurations

Minimal (Single Agent)

```
[[agents]]
name = "gemini"
canonical_name = "gemini"
command = "gemini"
args = []
enabled = true
```

Use Case: Cost-conscious setup, simple tasks

Balanced (3 Agents)

```
# Cheap consensus
[[agents]]
name = "gemini"
canonical_name = "gemini"
command = "gemini"

# Balanced reasoning
[[agents]]
name = "claude"
canonical_name = "claude"
command = "claude"

# Strategic planning
[[agents]]
name = "code"
canonical_name = "gpt_pro"
command = "code"
args = ["--model", "gpt-5"]

[quality_gates]
plan = ["gemini", "claude", "gpt_pro"]
tasks = ["gemini"]
validate = ["gemini", "claude", "gpt_pro"]
```

Use Case: Most production workloads

Premium (5 Agents + Specialist)

```

# Full 5-agent setup with premium reasoning
[[agents]]
name = "gemini"
canonical_name = "gemini"
command = "gemini"
enabled = true

[[agents]]
name = "claude"
canonical_name = "claude"
command = "claude"
enabled = true

[[agents]]
name = "code"
canonical_name = "gpt_pro"
command = "code"
args = ["--model", "gpt-5"]
enabled = true

[[agents]]
name = "gpt_codex"
canonical_name = "gpt_codex"
command = "code"
args = ["--model", "o3", "--config", "model_reasoning_effort=high"]
enabled = true # Enable for critical decisions

[quality_gates]
plan = ["gemini", "claude", "gpt_pro"]
tasks = ["gemini"]
validate = ["gemini", "claude", "gpt_pro"]
audit = ["gemini", "claude", "gpt_codex", "gpt_pro"] # 4 agents for security
unlock = ["gemini", "claude", "gpt_pro"]

```

Use Case: Critical projects, maximum quality

Debugging Agent Configuration

List Configured Agents

```
code --agents-list
```

Output:

```
Configured Agents (5):
[✓] gemini      - Google Gemini Flash (enabled)
[✓] claude      - Anthropic Claude Haiku (enabled)
[✓] code         - OpenAI GPT-5 (enabled)
[✓] gpt_codex    - OpenAI GPT-5-Codex (enabled)
[✗] gpt_pro      - OpenAI o3 (disabled)
```

Validate Agent Commands

```
code --check-agents
```

Output:

```
Checking agent commands...
[✓] gemini: command 'gemini' found
[✓] claude: command 'claude' found
[✓] code: command 'code' found
[✓] gpt_codex: command 'code' found
[✗] gpt_pro: command 'code' found, but agent disabled
```

All enabled agents have valid commands.

Best Practices

1. Use Canonical Names Consistently

Good:

```
[[agents]]
canonical_name = "gemini" # Stable

[quality_gates]
plan = ["gemini"] # Matches canonical_name
```

Bad:

```
[[agents]]
name = "gemini-flash-2024" # Display name

[quality_gates]
plan = ["gemini-flash-2024"] # ✗ Breaks if name changes
```

2. Enable Minimum Required Agents

Good:

```
# Enable only what you need
[[agents]]
canonical_name = "gemini"
enabled = true

[[agents]]
canonical_name = "claude"
enabled = true

[[agents]]
canonical_name = "gpt_pro"
enabled = false # Disable premium agent unless needed
```

3. Use args_read_only for Cost Savings

Example:

```
[[agents]]
name = "claude"
canonical_name = "claude"
command = "anthropic"
args_read_only = ["--model", "claude-3-haiku"] # Cheap for read-only
args_write = ["--model", "claude-3-5-sonnet"] # Capable for writes
```

4. Leverage Per-Agent Instructions

Example:

```
[[agents]]
name = "security-agent"
canonical_name = "security"
command = "claude"
instructions = "Focus on security. Flag OWASP Top 10 vulnerabilities."
```

Summary

Agent Configuration covers: - 5-agent default setup (gemini, claude, code, gpt_codex, gpt_pro) - Agent properties (name, canonical_name, command, args, enabled) - Advanced features (args_read_only, env, instructions) - Subagent commands (13 built-in commands) - Quality gate integration - Cost tiers (Tier 0-4, \$0 to \$0.80 per stage)

Best Practices: - Use canonical_name for stability - Enable minimum required agents - Leverage args_read_only for cost savings - Use per-agent instructions for specialization

Next: [Quality Gate Customization](#)

Configuration Reference

Complete config.toml schema reference.

Overview

Location: `~/.code/config.toml`

Alternative: `~/.codex/config.toml` (legacy, read-only)

Format: TOML (Tom's Obvious, Minimal Language)

Validation: Schema validation on load, old config preserved on error

File Structure

Minimal Example

```
# ~/.code/config.toml (minimal)

model = "gpt-5"
model_provider = "openai"
approval_policy = "on-request"
```

Complete Example

```
# ~/.code/config.toml (comprehensive)

# =====

# Model Configuration
#
# =====

model = "gpt-5"
model_provider = "openai"
model_reasoning_effort = "medium" # minimal, low, medium, high
model_reasoning_summary = "auto" # auto, concise, detailed, none
model_verbosity = "medium" # low, medium, high (GPT-5 only)
model_context_window = 128000 # Override context window size
model_max_output_tokens = 16384 # Override max output tokens
model_supports_reasoning_summaries = false

#
# =====

# Model Providers
#
# =====

[model_providers.openai]
name = "OpenAI"
base_url = "https://api.openai.com/v1"
env_key = "OPENAI_API_KEY"
wire_api = "responses" # or "chat"
request_max_retries = 4
stream_max_retries = 10
stream_idle_timeout_ms = 300000 # 5 minutes

[model_providers.anthropic]
name = "Anthropic"
base_url = "https://api.anthropic.com"
env_key = "ANTHROPIC_API_KEY"
wire_api = "chat"

[model_providers.google]
name = "Google"
base_url = "https://generativelanguage.googleapis.com/v1beta"
env_key = "GOOGLE_API_KEY"
wire_api = "chat"

[model_providers.ollama]
```



```

validate = ["gemini", "claude", "code"]      # Multi-agent test
validation
audit = ["gemini", "claude", "gpt_codex"]  # Security/compliance
review
unlock = ["gemini", "claude", "gpt_codex"] # Ship decision

#
=====

# Hot-Reload Configuration
#
=====

[hot_reload]
enabled = true
debounce_ms = 2000 # Wait 2s after last change before reloading
watch_paths = ["config.toml"] # Additional paths to watch

#
=====

# Validation Configuration
#
=====

[validation]
check_api_keys = true          # Validate API keys on startup
check_commands = true          # Validate agent commands exist
strict_schema = true           # Enforce strict TOML schema
patch_harness = false          # Run patch validation harness

[validation.groups]
functional = true   # Functional checks (cargo, tsc, etc.)
stylistic = false    # Stylistic checks (prettier, shfmt)

[validation.tools]
shellcheck = true
cargo-check = true
# ... other tools (see Validation section below)

#
=====

# Approval Policy
#
=====

approval_policy = "on-request" # untrusted, on-failure, on-request,
never

#
=====

# Confirm Guard (Destructive Commands)
#
=====
```

```

[[confirm_guard.patterns]]
regex = "(?i)^\\s*git\\s+reset\\b"
message = "Blocked git reset. Reset rewrites the working
tree/index."

[[confirm_guard.patterns]]
regex = "(?i)^\\s*(?:sudo\\s+)?rm\\s+-[a-z-]*rf[a-z-]*\\s+"
message = "Blocked rm -rf. Force-recursive delete requires
confirmation."

#
=====

# Sandbox Configuration
#
=====

sandbox_mode = "workspace-write" # read-only, workspace-write,
danger-full-access

[sandbox_workspace_write]
exclude_tmpdir_env_var = false
exclude_slash_tmp = false
writable_roots = [] # Additional writable paths
network_access = false
allow_git_writes = true # Allow .git/ folder writes

#
=====

# Shell Environment Policy
#
=====

[shell_environment_policy]
inherit = "all" # all, core, none
ignore_default_excludes = false # If true, include *KEY*, *TOKEN*
vars
exclude = ["AWS_*", "AZURE_*"] # Additional exclusion patterns
set = { CI = "1" } # Force-set environment variables
include_only = [] # If non-empty, only these
patterns survive

#
=====

# MCP Servers
#
=====

[mcp_servers.local-memory]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-memory"]
startup_timeout_ms = 10000 # 10 seconds

[mcp_servers.git-status]

```

```

command = "npx"
args = ["-y", "@just-every/mcp-server-git"]
env = { LOG_LEVEL = "info" }

#
=====

# ACE (Agentic Context Engine)
#
=====

[ace]
enabled = true
mode = "auto" # auto, always, never
slice_size = 8 # Max 8 playbook bullets
db_path = "~/code/ace/playbooks_normalized.sqlite3"
use_for = ["speckit.constitution", "speckit.specify",
"speckit.tasks"]
complex_task_files_threshold = 4
rerun_window_minutes = 30

#
=====

# TUI Configuration
#
=====

[tui]
alternate_screen = true          # Use alternate screen mode
show_reasoning = false           # Show reasoning content by default

[tui.theme]
name = "dark-carbon-night"       # See Theme section for all themes
# Optional custom color overrides
colors = {}

[tui.highlight]
theme = "auto" # auto, or specific syntect theme

[tui.stream]
answer_header_immediate = false
show_answer_ellipsis = true
commit_tick_ms = 50
soft_commit_timeout_ms = 400
soft_commit_chars = 160
relax_list_holdback = false
relax_code_holdback = false
responsive = false # Enable snappier preset

[tui.spinner]
name = "diamond" # Spinner style from cli-spinners

[tui.notifications]
# false (disabled), true (all), or array of specific notifications
notifications = false

#

```

```
=====
# History Configuration
#
=====

[history]
persistence = "save-all" # save-all, none
max_bytes = 10485760      # 10 MB (not currently enforced)

#
=====

# Browser Configuration (Screenshot Tool)
#
=====

[browser]
enabled = false
fullpage = true
segments_max = 10
idle_timeout_ms = 30000
format = "png" # png, webp

[browser.viewport]
width = 1280
height = 720
device_scale_factor = 2.0
mobile = false

[browser.wait]
delay_ms = 1000 # Wait 1s before screenshot

#
=====

# GitHub Integration
#
=====

[github]
check_workflows_on_push = true
actionlint_on_patch = false
actionlint_strict = false

#
=====

# Project Hooks
#
=====

[[project_hooks]]
event = "session.start"
name = "install-deps"
command = ["npm", "install"]
```

```

timeout_ms = 60000

[[project_hooks]]
event = "file.after_write"
command = ["cargo", "fmt", "--all"]

#
=====

# Profiles
#
=====

profile = "default" # Active profile

[profiles.premium]
model = "o3"
model_reasoning_effort = "high"
model_reasoning_summary = "detailed"
approval_policy = "never"

[profiles.fast]
model = "gpt-4o-mini"
model_reasoning_effort = "low"
approval_policy = "never"

[profiles.ci]
model = "gpt-4o"
approval_policy = "never"
sandbox_mode = "read-only"
disable_response_storage = false

#
=====

# Miscellaneous
#
=====

disable_response_storage = false # Required for ZDR accounts
file_opener = "vscode" # vscode, vscode-insiders, cursor, windsurf,
none
hide_agent_reasoning = false
show_raw_agent_reasoning = false
project_doc_max_bytes = 32768 # 32 KiB
notify = [] # Command to execute for notifications

```

Configuration Sections

Model Configuration

Field	Type	Default	Description
model	string	"gpt-5"	Model name use Provider ID

model_provider	string	"openai"	from model_provider
model_reasoning_effort	string	"medium"	Reasoning effort: minimum, low, medium, high
model_reasoning_summary	string	"auto"	Summary mode: auto, concise, detailed, none
model_verbosity	string	"medium"	Verbosity level: low, medium, high (GPT-5 only)
model_context_window	integer	128000	Context window size in tokens
model_max_output_tokens	integer	16384	Max output tokens
model_supports_reasoning_summaries	boolean	false	Force reasoning support

Model Providers

Table Format: [model_providers.<id>]

Required Fields: - name (string): Display name - base_url (string): API base URL - env_key (string, optional): Environment variable for API key

Optional Fields: - wire_api (string): "chat" or "responses" (default: "chat") - query_params (table): Additional query parameters (e.g., Azure api-version) - http_headers (table): Static HTTP headers - env_http_headers (table): HTTP headers from environment variables - request_max_retries (integer): HTTP request retries (default: 4) - stream_max_retries (integer): SSE stream retries (default: 10) - stream_idle_timeout_ms (integer): Idle timeout in ms (default: 300000)

Example:

```
[model_providers.azure]
name = "Azure OpenAI"
base_url = "https://YOUR_PROJECT.openai.azure.com/openai"
env_key = "AZURE_OPENAI_API_KEY"
query_params = { api-version = "2025-04-01-preview" }
```

Agents

Array Format: [[agents]]

Field	Type	Required	Description
name	string	Yes	Agent name (display) Canonical identifier

canonical_name	string	No	(default: same as name)
command	string	Yes	Command to execute
args	array	No	Command arguments
read_only	boolean	No	Force read-only mode (default: false)
enabled	boolean	No	Enable agent (default: true)
description	string	No	Agent description
env	table	No	Environment variables
args_read_only	array	No	Args for read-only mode
args_write	array	No	Args for write mode
instructions	string	No	Per-agent instructions

Quality Gates

Table Format: [quality_gates]

Field	Type	Default	Description
plan	array	[]	Agent names for plan stage
tasks	array	[]	Agent names for tasks stage
validate	array	[]	Agent names for validate stage
audit	array	[]	Agent names for audit stage
unlock	array	[]	Agent names for unlock stage

Agent names must match canonical_name from [[agents]].

Hot-Reload

Table Format: [hot_reload]

Field	Type	Default	Description
enabled	boolean	true	Enable hot-reload
debounce_ms	integer	2000	Debounce window in ms (default: 2s)
watch_paths	array	[]	Additional paths to watch

Validation

Table Format: [validation]

Field	Type	Default	Description
check_api_keys	boolean	true	Validate API keys on startup
check_commands	boolean	true	Validate agent commands exist
strict_schema	boolean	true	Enforce strict TOML schema
patch_harness	boolean	false	Run patch validation harness

<code>tools_allowlist</code>	array	null	Restrict allowed tools
<code>timeout_seconds</code>	integer	null	Tool execution timeout

Groups ([validation.groups]): - functional (boolean): Functional checks (cargo, tsc, eslint) - stylistic (boolean): Stylistic checks (prettier, shfmt)

Tools ([validation.tools]): - shellcheck, markdownlint, hadolint, yamllint (stylistic) - cargo-check, tsc, eslint, mypy, pyright, golangci-lint (functional) - shfmt, prettier (stylistic)

Sandbox Configuration

Field	Type	Default	Description
<code>sandbox_mode</code>	string	"read-only"	Sandbox mode: read-only, workspace-write, danger-full-access

[`sandbox_workspace_write`] (only applies when `sandbox_mode = "workspace-write"`):

Field	Type	Default	Description
<code>exclude_tmpdir_env_var</code>	boolean	false	Exclude \$TMPDIR from writable roots
<code>exclude_slash_tmp</code>	boolean	false	Exclude /tmp from writable roots
<code>writable_roots</code>	array	[]	Additional writable paths
<code>network_access</code>	boolean	false	Allow network access
<code>allow_git_writes</code>	boolean	true	Allow .git/ folder writes

MCP Servers

Table Format: [mcp_servers.<name>]

Field	Type	Required	Description
<code>command</code>	string	Yes	Command to execute
<code>args</code>	array	No	Command arguments
<code>env</code>	table	No	Environment variables
<code>startup_timeout_ms</code>	integer	No	Startup timeout in ms (default: 10000)

Example:

```
[mcp_servers.custom-tool]
command = "/path/to/mcp-server"
args = ["--port", "8080"]
```

```
env = { API_KEY = "secret" }
startup_timeout_ms = 15000
```

TUI Configuration

Theme ([tui.theme]): - name (string): Theme name (see Theme System guide) - colors (table): Custom color overrides - label (string, optional): Custom theme label - is_dark (boolean, optional): Dark theme hint

Highlight ([tui.highlight]): - theme (string): Syntax highlighting theme (default: "auto")

Stream ([tui.stream]): - answer_header_immediate (boolean): Show header immediately - show_answer_ellipsis (boolean): Show ellipsis while waiting - commit_tick_ms (integer): Animation commit rate (default: 50ms) - soft_commit_timeout_ms (integer): Soft-commit timeout - soft_commit_chars (integer): Soft-commit character threshold - relax_list_holdback (boolean): Relax list marker hold-back - relax_code_holdback (boolean): Relax code block hold-back - responsive (boolean): Enable snappier preset

Profiles

Table Format: [profiles.<name>]

Profiles can override any top-level config field. See [Precedence System](#) for details.

Example:

```
[profiles.premium]
model = "o3"
model_reasoning_effort = "high"
approval_policy = "never"
```

Activation: Set profile = "premium" or use --profile premium flag.

Validation Rules

Required Fields

None - All fields have defaults

Type Validation

- Strings: Non-empty (whitespace trimmed)
- Integers: Must be positive (where applicable)
- Booleans: true or false
- Arrays: Can be empty unless semantically invalid

Semantic Validation

1. **Model provider must exist:** model_provider must be a key in

- ```
model_providers
```
2. **Quality gate agents must exist:** Agent names in quality\_gates.\* must match canonical\_name in [[agents]]
  3. **Evidence size must be reasonable:** evidence.max\_size\_mb ≤ 1000
  4. **Debounce must be reasonable:** hot\_reload.debounce\_ms ≥ 100
- 

## Error Handling

**On validation failure:** 1. Old config is **preserved** (no reload) 2. ReloadFailed event emitted with error message 3. TUI shows notification with error details

**Example error:**

```
Config validation failed: Agent 'unknown-agent' not found in
quality_gates.plan
Old config preserved.
```

---

## Summary

**Config File:** `~/.code/config.toml`

**Sections:** 20+ configuration sections covering: - Model/provider configuration - Multi-agent setup - Quality gates - Hot-reload settings - Validation rules - Sandbox policy - MCP servers - TUI customization - Profiles

**Validation:** Schema validation with old config preservation on error

**Next:** [Precedence System](#)

---

## Environment Variables

Complete reference for all environment variables and override behavior.

---

## Overview

Environment variables provide **Tier 2 precedence** (higher than config.toml, lower than CLI flags).

**Use Cases:** - API keys and secrets - Environment-specific overrides (dev, staging, production) - CI/CD configuration - Temporary configuration changes

---

## Core Environment Variables

**CODEX\_HOME / CODE\_HOME**

**Purpose:** Installation directory

**Default:** `~/code`

**Legacy:** `~/codex` (read-only, deprecated)

**Usage:**

```
export CODEX_HOME="/custom/path"
or
export CODE_HOME="/custom/path"
```

**Precedence:** `CODE_HOME > CODEX_HOME > ~/code`

**Files Stored:**

```
$CODEX_HOME/
├── config.toml # Configuration file
├── history.jsonl # Session history
├── debug.log # Debug logs
├── mcp-memory/ # MCP memory database
├── mcp-cache/ # MCP tool cache
└── ace/ # ACE playbook database
 └── playbooks_normalized.sqlite3
```

---

## API Keys

### OPENAI\_API\_KEY

**Purpose:** OpenAI API authentication

**Required:** When using `model_provider = "openai"`

**Usage:**

```
export OPENAI_API_KEY="sk-proj-..."
```

**Security:** Never commit to git, never store in `config.toml`

---

### ANTHROPIC\_API\_KEY

**Purpose:** Anthropic API authentication

**Required:** When using `model_provider = "anthropic"`

**Usage:**

```
export ANTHROPIC_API_KEY="sk-ant-..."
```

---

### GOOGLE\_API\_KEY

**Purpose:** Google Gemini API authentication

**Required:** When using `model_provider = "google"`

**Usage:**

```
export GOOGLE_API_KEY="..."
```

---

## AZURE\_OPENAI\_API\_KEY

**Purpose:** Azure OpenAI API authentication

**Required:** When using Azure model provider

**Usage:**

```
export AZURE_OPENAI_API_KEY="..."
```

**Alternative:** OPENAI\_API\_KEY also works for Azure

---

## Custom Provider API Keys

**Pattern:** <PROVIDER\_NAME>\_API\_KEY

**Example:**

```
[model_providers.custom]
env_key = "CUSTOM_API_KEY"

export CUSTOM_API_KEY="..."
```

---

## Model Configuration Overrides

### CODEX\_MODEL

**Purpose:** Override default model

**Precedence:** Env var > config.toml

**Usage:**

```
export CODEX_MODEL=o3
code "task"
```

**Equivalent:**

```
code --model o3 "task"
```

---

### CODEX\_PROVIDER

**Purpose:** Override model provider

**Usage:**

```
export CODEX_PROVIDER=anthropic
code "task"
```

**Equivalent:**

```
code --config model_provider=anthropic "task"
```

---

## **OPENAI\_BASE\_URL**

**Purpose:** Override OpenAI base URL

**Use Case:** Custom proxy, Azure, local endpoint

**Usage:**

```
export OPENAI_BASE_URL="https://custom.openai.com/v1"
```

**Overrides:** model\_providers.openai.base\_url

---

## **OPENAI\_WIRE\_API**

**Purpose:** Force OpenAI wire protocol

**Options:** "responses" or "chat"

**Usage:**

```
export OPENAI_WIRE_API="chat" # Force chat completions
```

**Overrides:** model\_providers.openai.wire\_api

---

## **Spec-Kit Environment Variables**

### **SPEC\_OPS\_CARGO\_MANIFEST**

**Purpose:** Override cargo manifest path for workspace commands

**Default:** Auto-detected (codex-rs/Cargo.toml)

**Usage:**

```
export SPEC_OPS_CARGO_MANIFEST="/path/to/Cargo.toml"
```

---

### **SPEC\_OPS\_ALLOW\_DIRTY**

**Purpose:** Allow guardrail commands with dirty git tree

**Default:** 0 (require clean tree)

**Usage:**

```
export SPEC_OPS_ALLOW_DIRTY=1
/guardrail.auto SPEC-KIT-065
```

**Use Case:** Testing, development iteration

---

### **SPEC\_OPS\_TELEMETRY\_HAL**

**Purpose:** Enable HAL telemetry collection

**Default:** 0 (disabled)

#### **Usage:**

```
export SPEC_OPS_TELEMETRY_HAL=1
/guardrail.plan SPEC-KIT-065
```

**Output:** Captures hal.summary.{status,failed\_checks,artifacts} in telemetry

---

#### **SPEC\_OPS\_HAL\_SKIP**

**Purpose:** Skip HAL validation (when secrets unavailable)

**Default:** 0 (run HAL validation)

#### **Usage:**

```
export SPEC_OPS_HAL_SKIP=1
/guardrail.audit SPEC-KIT-065
```

**Use Case:** Development without HAL secrets

---

#### **SPECKIT\_QUALITY\_GATES\_\***

**Purpose:** Override quality gate agent selection

**Pattern:** SPECKIT\_QUALITY\_GATES\_<STAGE>=agent1,agent2,agent3

#### **Usage:**

```
export SPECKIT_QUALITY_GATES_PLAN="gemini,claude,code,gpt_pro"
export SPECKIT_QUALITY_GATES_TASKS="code"
export SPECKIT_QUALITY_GATES_VALIDATE="gemini,claude,code"
export SPECKIT_QUALITY_GATES_AUDIT="gemini,claude,gpt_codex,gpt_pro"
export SPECKIT_QUALITY_GATES_UNLOCK="gemini,claude,gpt_pro"
```

**Precedence:** Env var > config.toml

---

## **Logging and Debugging**

#### **RUST\_LOG**

**Purpose:** Rust logging level

**Options:** error, warn, info, debug, trace

#### **Usage:**

```
export RUST_LOG=debug
code
```

#### **Module-Specific:**

```
export RUST_LOG=codex_tui::chatwidget::spec_kit=debug
code
```

#### **Multiple Modules:**

```
export RUST_LOG=codex_mcp_client=debug,codex_spec_kit=trace
code
```

---

## RUST\_BACKTRACE

**Purpose:** Enable backtraces on panic

**Usage:**

```
export RUST_BACKTRACE=1 # Short backtrace
export RUST_BACKTRACE=full # Full backtrace
code
```

**Use Case:** Debugging crashes

---

## Sandbox and Security

### CODEX\_SANDBOX\_NETWORK\_DISABLED

**Purpose:** Disable network access in sandbox

**Auto-Set:** When `sandbox_mode = "read-only"` or `sandbox_mode = "workspace-write"` with `network_access = false`

**Usage** (manual override):

```
export CODEX_SANDBOX_NETWORK_DISABLED=1
```

---

## CI/CD Environment Variables

### CI

**Purpose:** Detect CI environment

**Auto-Set:** By most CI systems (GitHub Actions, GitLab CI, CircleCI, etc.)

**Usage:**

```
[shell_environment_policy]
set = { CI = "1" }
```

**Effect:** Triggers CI-specific behavior (non-interactive mode, strict validation)

---

### GITHUB\_ACTIONS

**Purpose:** Detect GitHub Actions environment

**Auto-Set:** By GitHub Actions

**Usage:**

```
if ["$GITHUB_ACTIONS" = "true"]; then
 export CODEX_MODEL="gpt-4o" # Use cheaper model in CI
fi
```

---

## CODEX\_AUTO\_UPGRADE

**Purpose:** Enable/disable auto-upgrade

**Options:** true/false, 1/0, yes/no, on/off

**Usage:**

```
export CODEX_AUTO_UPGRADE=false # Disable auto-upgrade in CI
```

**Overrides:** auto\_upgrade\_enabled in config.toml

---

## Shell Environment Policy

### Shell Environment Inheritance

**Configuration:**

```
[shell_environment_policy]
inherit = "all" # all, core, none
ignore_default_excludes = false
exclude = ["AWS_*", "AZURE_*"]
set = { CI = "1" }
include_only = []
```

**Default Excludes** (when ignore\_default\_excludes = false): - \*KEY\*  
(case-insensitive) - \*TOKEN\* (case-insensitive) - \*SECRET\* (case-insensitive)

**Example:**

```
These are excluded by default:
export AWS_ACCESS_KEY="..." # Excluded (*KEY*)
export GITHUB_TOKEN="..." # Excluded (*TOKEN*)
export DB_SECRET="..." # Excluded (*SECRET*)

These are included (no KEY/TOKEN/SECRET):
export PATH="/usr/bin" # Included
export HOME="/home/user" # Included
```

---

### Override Shell Environment Policy

**Usage:**

```
export SHELL_ENV_INHERIT="core" # Override inherit mode
export SHELL_ENV_IGNORE_DEFAULT_EXCLUDES="1" # Include KEY/TOKEN
vars
```

---

## MCP Server Environment Variables

## MCP-Specific Variables

**Pattern:** Set in env field of [mcp\_servers.<name>]

**Example:**

```
[mcp_servers.database]
command = "/path/to/db-server"
env = {
 DB_HOST = "localhost",
 DB_PORT = "5432",
 DB_NAME = "mydb"
}
```

**Scope:** Only available to that specific MCP server

---

## Global MCP Environment

**Pattern:** MCP\_\* prefix

**Usage:**

```
export MCP_LOG_LEVEL="debug"
export MCP_TIMEOUT="30000"
```

**Scope:** Available to all MCP servers

---

## HAL Secret Environment Variables

### HAL\_SECRET\_KAVEDARR\_API\_KEY

**Purpose:** Kavedarr API key for HAL validation

**Required:** When running HAL smoke tests or policy validation

**Usage:**

```
export HAL_SECRET_KAVEDARR_API_KEY="..."
```

**Security:** Never commit, never store in config

---

## Testing Environment Variables

### PRECOMMIT\_FAST\_TEST

**Purpose:** Skip test compilation in pre-commit hook

**Default:** 1 (skip test compilation)

**Usage:**

```
export PRECOMMIT_FAST_TEST=0 # Run test compilation
git commit
```

---

## **PREPUSH\_FAST**

**Purpose:** Skip pre-push hooks

**Default:** 1 (run hooks)

**Usage:**

```
export PREPUSH_FAST=0 # Skip pre-push hooks
git push
```

**Warning:** Only use for emergencies

---

## **Complete Environment Variable Reference**

### **Core Variables**

| Variable       | Purpose                     | Default           | Example      |
|----------------|-----------------------------|-------------------|--------------|
| CODEX_HOME     | Installation directory      | ~/.code           | /custom/path |
| CODE_HOME      | Alt. installation directory | (uses CODEX_HOME) | /custom/path |
| RUST_LOG       | Logging level               | info              | debug        |
| RUST_BACKTRACE | Backtrace on panic          | 0                 | 1, full      |

---

### **API Keys**

| Variable             | Purpose                        | Required For                 |
|----------------------|--------------------------------|------------------------------|
| OPENAI_API_KEY       | OpenAI authentication          | model_provider = "openai"    |
| ANTHROPIC_API_KEY    | Anthropic authentication       | model_provider = "anthropic" |
| GOOGLE_API_KEY       | Google Gemini authentication   | model_provider = "google"    |
| AZURE_OPENAI_API_KEY | Azure OpenAI authentication    | Azure model provider         |
| <PROVIDER>_API_KEY   | Custom provider authentication | Custom providers             |

---

### **Model Overrides**

| Variable           | Overrides                       | Example         |
|--------------------|---------------------------------|-----------------|
| CODEX_MODEL        | model                           | o3              |
| CODEX_PROVIDER     | model_provider                  | anthropic       |
| OPENAI_BASE_URL    | model_providers.openai.base_url | https://custom. |
| OPENAI_WIRE_API    | model_providers.openai.wire_api | chat, responses |
| CODEX_AUTO_UPGRADE | auto_upgrade_enabled            | true, false     |

---

## Spec-Kit Variables

| Variable                | Purpose                      | Default       | Example              |
|-------------------------|------------------------------|---------------|----------------------|
| SPEC_OPS_CARGO_MANIFEST | Cargo manifest path          | Auto-detected | /path/to/Cargo.toml  |
| SPEC_OPS_ALLOW_DIRTY    | Allow dirty git tree         | 0             | 1                    |
| SPEC_OPS_TELEMETRY_HAL  | Enable HAL telemetry         | 0             | 1                    |
| SPEC_OPS_HAL_SKIP       | Skip HAL validation          | 0             | 1                    |
| SPECKIT_QUALITY_GATES_* | Override quality gate agents | (from config) | gemini, claude, code |

---

## Sandbox and Security

| Variable                       | Purpose                    | Auto-Set                          | Manual Override |
|--------------------------------|----------------------------|-----------------------------------|-----------------|
| CODEX_SANDBOX_NETWORK_DISABLED | Disable network in sandbox | Yes (when network_access = false) | 1               |

---

## CI/CD Variables

| Variable       | Purpose                  | Auto-Set By     | Example |
|----------------|--------------------------|-----------------|---------|
| CI             | CI environment detection | Most CI systems | 1, true |
| GITHUB_ACTIONS | GitHub Actions detection | GitHub Actions  | true    |

---

## Testing Variables

| Variable            | Purpose                             | Default | Example |
|---------------------|-------------------------------------|---------|---------|
| PRECOMMIT_FAST_TEST | Skip test compilation in pre-commit | 1       | 0       |
| PREPUSH_FAST        | Skip pre-push hooks                 | 1       | 0       |

---

## Best Practices

### 1. Store Secrets in Environment Variables

**Good:**

```
export OPENAI_API_KEY="sk-proj-..."
export ANTHROPIC_API_KEY="sk-ant-..."
```

**Bad:**

```
DON'T: Never store secrets in config.toml
[model_providers.openai]
api_key = "sk-proj-..." # ✗ Security risk!
```

---

### 2. Use .env Files (Local Development)

.env file (git-ignored):

```
.env
OPENAI_API_KEY=sk-proj-...
ANTHROPIC_API_KEY=sk-ant-...
GOOGLE_API_KEY=...
```

Load with direnv:

```
Install direnv
brew install direnv # macOS
apt install direnv # Linux

Enable for shell
echo 'eval "$(direnv hook bash)"' >> ~/.bashrc

Allow .envrc
echo 'dotenv' > .envrc
direnv allow
```

---

### 3. Use Profiles for Environment-Specific Config

config.toml:

```
[profiles.dev]
model = "gpt-4o-mini"
approval_policy = "never"

[profiles.staging]
model = "gpt-5"
approval_policy = "on-request"

[profiles.production]
model = "o3"
approval_policy = "on-failure"
model_reasoning_effort = "high"
```

Usage:

```
Development
code --profile dev "task"

Staging
code --profile staging "task"

Production
code --profile production "task"
```

---

## 4. Document Required Environment Variables

**README.md:**

```
Required Environment Variables

- `OPENAI_API_KEY` - OpenAI API key
- `ANTHROPIC_API_KEY` - Anthropic API key (optional)
- `CODEX_HOME` - Installation directory (optional, default: ~/.code)
```

---

## Debugging Environment Variables

### List Active Environment Variables

```
All CODEX_* and *_API_KEY variables
env | grep -E 'CODEX|API_KEY'
```

**Output:**

```
CODEX_HOME=/home/user/.code
CODEX_MODEL=gpt-5
OPENAI_API_KEY=sk-proj-***
ANTHROPIC_API_KEY=sk-ant-***
```

---

### Check Effective Configuration

```
code --config-dump | grep -A 5 "# Source:"
```

**Output:**

```
model = "o3" # Source: Environment variable (CODEX_MODEL)
model_provider = "openai" # Source: config.toml
approval_policy = "never" # Source: Profile 'premium'
```

---

## Summary

**Environment Variables** provide:

- Tier 2 precedence (env var > config.toml)
- API key storage (secure, never in config)
- Environment-specific overrides (dev, staging, production)
- CI/CD configuration
- Temporary configuration changes

**Categories:**

- Core (CODEX\_HOME, RUST\_LOG) - API Keys (OPENAI\_API\_KEY, ANTHROPIC\_API\_KEY, GOOGLE\_API\_KEY)
- Model Overrides (CODEX\_MODEL, CODEX\_PROVIDER)
- Spec-Kit

(SPEC\_OPS\_, SPECKIT\_) - Sandbox  
(CODEX\_SANDBOX\_NETWORK\_DISABLED) - CI/CD (CI,  
GITHUB\_ACTIONS) - Testing (PRECOMMIT\_FAST\_TEST,  
PREPUSH\_FAST)

**Best Practices:** - Store secrets in environment variables - Use .env files (git-ignored) for local development - Use profiles for environment-specific config - Document required variables in README

**Next:** [Template Customization](#)

---

## Hot-Reload

Config reload mechanism with 300ms debouncing.

---

### Overview

**Hot-reload** enables configuration changes to apply **without restarting** the application.

**Benefits:** - Instant config updates (<100ms latency) - No session interruption - Safe validation (old config preserved on error)

**Performance:** <0.5% CPU overhead, <336ms reload latency (p50)

---

## Architecture

### Reload Flow

```
File Change → notify crate → Debouncer (300ms) → Validate → Lock →
Replace → Event
 ↓ Fail
 Preserve Old Config
```

**Components:** 1. **File Watcher** (notify crate) - Detects filesystem changes  
2. **Debouncer** - Buffers events for 300ms to prevent storms  
3. **Validator** - Validates new config (schema, semantic)  
4. **Lock** - Atomic config replacement (RwLock)  
5. **Event** - Notification to TUI/app

---

## Configuration

### Enable Hot-Reload

**Default:** Enabled

```
~/.code/config.toml
[hot_reload]
```

```
enabled = true
debounce_ms = 2000 # Wait 2s after last change
watch_paths = ["config.toml"] # Additional files to watch
```

---

## Configuration Fields

| Field       | Type    | Default | Description                                                    |
|-------------|---------|---------|----------------------------------------------------------------|
| enabled     | boolean | true    | Enable/disable hot-reload                                      |
| debounce_ms | integer | 2000    | Debounce window in milliseconds                                |
| watch_paths | array   | []      | Additional paths to watch (relative to <code>~/.code/</code> ) |

---

## Debouncing

**Purpose:** Prevent reload storms from multiple filesystem events

**Example Scenario:**

```
t=0ms: File save event 1 (vim writes temp file)
t=50ms: File save event 2 (vim renames temp file)
t=100ms: File save event 3 (vim updates mtime)
t=2100ms: No events for 2000ms → Trigger reload
```

**Result:** Only **one reload** despite 3 filesystem events

---

## Debounce Tuning

**Fast Debounce** (impatient users):

```
[hot_reload]
debounce_ms = 500 # 500ms (more responsive, more reloads)
```

**Slow Debounce** (complex editors):

```
[hot_reload]
debounce_ms = 5000 # 5s (less responsive, fewer reloads)
```

**Recommended:** 2000ms (2 seconds)

---

## Watch Additional Files

**Example:** Watch model provider configs

```
[hot_reload]
watch_paths = [
 "config.toml", # Default
 "models/openai.toml", # Custom model config
 "models/anthropic.toml", # Custom model config
]
```

**Use Case:** Split configuration across multiple files

---

## Reload Events

### Event Types

```
pub enum ConfigReloadEvent {
 /// File change detected (before reload attempt)
 FileChanged(PathBuf),

 /// Config successfully reloaded
 ReloadSuccess,

 /// Reload failed (old config preserved)
 ReloadFailed(String),
}
```

---

### Event Flow

#### Successful Reload:

1. FileChanged(~/.code/config.toml) # File changed
2. [Debounce wait 2000ms]
3. [Parse TOML: OK]
4. [Validate: OK]
5. [Replace config]
6. ReloadSuccess # Notify TUI

#### Failed Reload:

1. FileChanged(~/.code/config.toml) # File changed
2. [Debounce wait 2000ms]
3. [Parse TOML: ERROR]
4. ReloadFailed("Invalid TOML: missing closing bracket")
5. [Old config preserved]

---

## TUI Notifications

#### Success:

- ✓ Config reloaded successfully
  - 2 model configs changed
  - Quality gates updated

#### Failure:

- ✗ Config reload failed: Invalid TOML syntax at line 42  
Old configuration preserved.

---

## Reload Performance

### Latency Breakdown

#### End-to-end reload latency:

```

File save → Filesystem event → Debounce wait → Parse TOML → Validate
→ Write lock → Event
 0ms ~10ms 2000ms ~20ms ~5ms
<1ms ~1ms

Total: ~2036ms (p50)
 ~2120ms (p95)

```

**Acceptable:** Sub-3-second reload for manual config edits

---

## Lock Performance

**Read Lock** (frequent, fast):

```
let config = watcher.get_config(); // Arc::clone
```

**Timing:**

```

Acquire read lock: <1μs
Clone Arc: <100ns
Release read lock: <100ns

```

Total: <1μs

**Concurrency:** Multiple readers allowed (RwLock)

---

**Write Lock** (rare, fast):

```
*config.write().unwrap() = new_config;
```

**Timing:**

```

Acquire write lock: <500μs (wait for readers to finish)
Replace config: <100ns
Release write lock: <100ns

```

Total: <1ms

**Blocking:** Briefly blocks readers (<1ms)

---

## CPU Overhead

**Idle** (file watching):

```

CPU usage: <0.5%
Memory: ~2 MB (notify crate + debouncer)

```

**During Reload:**

```

CPU spike: ~10-20% for ~50ms (parsing + validation)
Memory spike: ~1 MB (temporary during validation)

```

---

## Validation

### Schema Validation

**Checks:** 1. TOML syntax validity 2. Required fields present 3. Type correctness (string, int, bool, array) 4. Enum values valid

#### Example Errors:

```
x Invalid TOML: unexpected character ']' at line 42
x Missing required field: model_providers.openai.base_url
x Type mismatch: quality_gates.plan expected array, got string
x Invalid enum value: approval_policy="unknown" (expected:
untrusted, on-failure, on-request, never)
```

---

## Semantic Validation

**Checks:** 1. Model provider exists 2. Quality gate agents exist and are enabled 3. Evidence size limits reasonable 4. Debounce timing reasonable

#### Example Errors:

```
x Model provider 'unknown' not found in model_providers
x Quality gate agent 'gpt_pro' not found or disabled
x Evidence max_size_mb=5000 exceeds limit (1000 MB)
x Hot-reload debounce_ms=50 too low (minimum: 100ms)
```

---

## Validation Failure Behavior

**On validation failure:** 1. **Preserve old config** (no changes applied) 2. **Emit ReloadFailed event** with error message 3. **Show TUI notification** with error details 4. **Log error** to `~/code/debug.log`

**User Action:** Fix `config.toml` and save again (triggers new reload)

---

## Deferring Reloads

### When to Defer

**Defer reload if:** 1. Quality gate is active (don't interrupt validation) 2. Agents are running (don't interrupt execution) 3. Critical operation in progress (file write, git commit)

#### Implementation:

```
pub fn should_defer_reload(quality_gate_active: bool, agent_running: bool) -> bool {
 quality_gate_active || agent_running
}
```

---

## Deferred Reload Behavior

**Scenario:** User edits config while quality gate is running

#### Behavior:

1. FileChanged event received

2. Check if quality gate active: YES
3. Queue reload for later
4. Quality gate completes
5. Execute queued reload

**Result:** Config reloads after quality gate completes (no interruption)

---

## Change Detection

### Detecting Config Changes

**Purpose:** Show user what changed in TUI notification

#### Implementation:

```
pub fn detect_config_changes(old: & AppConfig, new: & AppConfig) ->
(usize, bool, bool) {
 let models_changed = count_model_changes(old, new);
 let quality_gates_changed = old.quality_gates != new.quality_gates;
 let cost_changed = old.cost != new.cost;

 (models_changed, quality_gates_changed, cost_changed)
}
```

**Returns:** (models\_changed, quality\_gates\_changed, cost\_changed)

---

### TUI Notification with Changes

#### Example:

- ✓ Config reloaded successfully
  - 3 model configs changed (openai, anthropic, google)
  - Quality gates updated (plan: 3→2 agents)
  - Cost limits changed (\$10/day → \$20/day)

## Debugging Hot-Reload

### Enable Debug Logging

```
export RUST_LOG=codex_spec_kit::config::hot_reload=debug
code
```

#### Log Output:

```
[DEBUG] HotReloadWatcher initialized
[DEBUG] Watching: ~/.code/config.toml
[DEBUG] Debounce window: 2000ms
[DEBUG] FileChanged event: ~/.code/config.toml
[DEBUG] Debouncing... (waiting 2000ms)
[DEBUG] Debounce complete, attempting reload
[DEBUG] Parsing TOML: OK
[DEBUG] Validating config: OK
[DEBUG] Acquiring write lock...
```

```
[DEBUG] Write lock acquired (<1ms)
[DEBUG] Config replaced
[DEBUG] ReloadSuccess event emitted
```

---

## Test Hot-Reload

### Manual Test:

```
Terminal 1: Run app with debug logging
export RUST_LOG=debug
code

Terminal 2: Edit config
vim ~/.code/config.toml
Make change and save

Terminal 1: Check logs
[DEBUG] FileChanged event: ~/.code/config.toml
[DEBUG] Debouncing...
[DEBUG] Config reloaded successfully
```

---

## Disable Hot-Rewload (Troubleshooting)

```
[hot_reload]
enabled = false # Disable hot-reload
```

**Use Case:** Debugging config loading issues, performance profiling

---

## Best Practices

### 1. Use Default Debounce (2000ms)

#### Recommended:

```
[hot_reload]
debounce_ms = 2000 # 2 seconds
```

**Reason:** Balances responsiveness with reload frequency

---

### 2. Validate Config Before Saving

#### Workflow:

```
Edit config
vim ~/.code/config.toml

Validate locally (optional tool)
toml-lint ~/.code/config.toml

Save (triggers hot-reload)
```

---

### 3. Monitor Reload Notifications

**Good Practice:** Check TUI notifications after config changes

**Example:**

- ✓ Config reloaded successfully
  - 2 agents enabled
  - Quality gates updated

**Bad Sign:**

- ✗ Config reload failed: Invalid agent name  
Old configuration preserved.

**Action:** Fix error and save again

---

## 4. Test Config Changes Incrementally

**Good:**

1. Change one section (e.g., model config)
2. Save and verify reload
3. Change next section (e.g., quality gates)
4. Save and verify reload

**Bad:**

1. Change 10 sections at once
  2. Save
  3. Error in section 7
  4. Hard to debug which change caused error
- 

## Summary

**Hot-Reload Features:** - Instant config updates (<100ms latency) - 300ms debouncing (prevents reload storms) - Safe validation (old config preserved on error) - TUI notifications (success/failure) - Deferred reload (don't interrupt operations) - Change detection (show what changed)

**Performance:** - <0.5% CPU overhead (idle) - ~2036ms reload latency (p50) - <1µs read locks - <1ms write locks

**Configuration:**

```
[hot_reload]
enabled = true
debounce_ms = 2000
watch_paths = ["config.toml"]
```

**Best Practices:** - Use default 2000ms debounce - Validate config before saving - Monitor TUI notifications - Test changes incrementally

**Next:** [MCP Servers](#)

---

## MCP Servers

MCP server configuration, custom servers, and lifecycle management.

---

## Overview

**MCP** (Model Context Protocol) enables AI agents to access external tools and resources through standardized servers.

**Use Cases:** - Memory systems (local-memory for knowledge persistence) - Git operations (git-status for repository inspection) - Custom tools (HAL for policy validation) - External services (databases, APIs, file systems)

**Configuration:** [mcp\_servers.<name>] sections in config.toml

---

## MCP Server Configuration

### Basic Configuration

```
~/.code/config.toml

[mcp_servers.local-memory]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-memory"]
startup_timeout_ms = 10000 # 10 seconds
```

---

### Configuration Fields

| Field              | Type    | Required | Description                        |
|--------------------|---------|----------|------------------------------------|
| command            | string  | Yes      | Executable command                 |
| args               | array   | No       | Command arguments                  |
| env                | table   | No       | Environment variables              |
| startup_timeout_ms | integer | No       | Startup timeout (default: 10000ms) |

---

## Built-in MCP Servers

### local-memory (Knowledge Persistence)

**Purpose:** Store and retrieve high-value knowledge (architecture decisions, patterns, bug fixes)

**Configuration:**

```
[mcp_servers.local-memory]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-memory"]
startup_timeout_ms = 10000
```

## Installation:

```
Auto-installed on first use via npx -y
Or install globally:
npm install -g @modelcontextprotocol/server-memory
```

**Tools Provided:** - mcp\_local-memory\_store\_memory - Store knowledge  
- mcp\_local-memory\_search - Search knowledge - mcp\_local-memory\_analysis - Analyze patterns

## Usage:

```
Use mcp_local-memory_store_memory:
- content: "Routing bug fixed: SpecKitCommand wasn't passing config..."
- domain: "debugging"
- tags: ["type:bug-fix", "spec:SPEC-KIT-066"]
- importance: 9
```

**Storage:** ~/.code/mcp-memory/ (SQLite database)

---

## git-status (Repository Inspection)

**Purpose:** Inspect Git repository state, history, changes

### Configuration:

```
[mcp_servers.git-status]
command = "npx"
args = ["-y", "@just-every/mcp-server-git"]
env = { LOG_LEVEL = "info" }
```

**Tools Provided:** - mcp\_git-status\_status - Get git status - mcp\_git-status\_diff - Get diff for files - mcp\_git-status\_log - Get commit history

**Use Case:** Automated commit message generation, change analysis

---

## HAL (Policy Validation)

**Purpose:** Validate spec-kit policies (storage policy, tag schema, quality gates)

### Configuration:

```
[mcp_servers.hal]
command = "/path/to/hal-server"
args = ["--mode", "strict"]
env = { HAL_SECRET_KAVEDARR_API_KEY = "..." }
startup_timeout_ms = 15000
```

**Tools Provided:** - mcp\_hal\_validate\_storage\_policy - Check local-memory usage - mcp\_hal\_validate\_tag\_schema - Check tag naming - mcp\_hal\_validate\_quality\_gates - Check consensus

**Note:** HAL server is project-specific (not publicly available)

---

# Custom MCP Servers

## Creating a Custom Server

**Example:** Database query server

```
[mcp_servers.database]
command = "/path/to/db-mcp-server"
args = ["--connection-string", "postgres://localhost/mydb"]
env = { DB_PASSWORD = "secret" }
startup_timeout_ms = 20000 # Longer timeout for DB connection
```

**Server Implementation:** See [MCP Server SDK](#)

---

## Custom Server Example (Node.js)

```
// db-mcp-server.js
const { MCPServer } = require('@modelcontextprotocol/sdk');
const { Pool } = require('pg');

const server = new MCPServer({
 name: 'database',
 version: '1.0.0',
});

const pool = new Pool({
 connectionString: process.argv[2],
});

server.tool({
 name: 'query',
 description: 'Execute SQL query',
 parameters: {
 sql: { type: 'string', description: 'SQL query to execute' },
 },
 async handler({ sql }) {
 const result = await pool.query(sql);
 return { rows: result.rows };
 },
});

server.start();
```

### Configuration:

```
[mcp_servers.database]
command = "node"
args = ["/path/to/db-mcp-server.js", "postgres://localhost/mydb"]
```

---

# MCP Server Lifecycle

## Startup Process

1. Config loaded → Parse [mcp\_servers.\*] sections
2. Spawn process → Execute command with args

3. Handshake → Initialize MCP protocol
4. List tools → Request tools/list from server
5. Cache tools → Store tool metadata
6. Ready → Server available for use

**Timeout:** startup\_timeout\_ms (default: 10000ms)

---

## Lazy Loading

**Default Behavior:** MCP servers are **not** started until first use

**Benefit:** Save resources by only starting needed servers

**Example:**

```
Configured but not started
[mcp_servers.database]
command = "node"
args = ["/path/to/db-server.js"]

Only started when tool is called:
Use mcp_database_query: "SELECT * FROM users"
```

---

## Startup Optimization

**Cache Tool List:**

First session:

1. Start MCP server (~500ms)
2. Request tools/list (~100ms)
3. Cache to ~/.code/mcp-cache/database.json
4. Use tools

Subsequent sessions:

1. Load cached tools from ~/.code/mcp-cache/database.json (~10ms)
2. Lazy-start server only when tool is called

**Benefit:** Faster session startup (no waiting for MCP servers)

---

## Shutdown Process

1. Session end → Send shutdown signal to all MCP servers
  2. Wait for clean shutdown (max 5s)
  3. Force kill if timeout
  4. Clean up temp files
- 

## Environment Variables

### Server-Specific Environment

```
[mcp_servers.custom]
command = "/path/to/server"
env = {
 API_KEY = "secret",
```

```
 LOG_LEVEL = "debug",
 FEATURE_FLAG = "experimental"
}
```

**Scope:** Only available to the MCP server process

---

## Global Environment Variables

```
Available to all MCP servers
export MCP_LOG_LEVEL="debug"
export MCP_TIMEOUT="30000"
```

**Use Case:** Global MCP debugging settings

---

## Timeouts and Retries

### Startup Timeout

**Default:** 10000ms (10 seconds)

**Configuration:**

```
[mcp_servers.slow-server]
command = "/path/to/slow-server"
startup_timeout_ms = 30000 # 30 seconds for slow startup
```

**Behavior:** If server doesn't respond within timeout, startup fails

---

### Tool Call Timeout

**Default:** Inherited from validation.timeout\_seconds

**Override:**

```
[validation]
timeout_seconds = 60 # 60 seconds for all MCP tool calls
```

---

### Retry Logic

**Startup Failures:** No automatic retry (manual restart required)

**Tool Call Failures:** Retry up to 3 times with exponential backoff

**Example:**

1. Tool call fails (network error)
  2. Wait 1s
  3. Retry (1/3)
  4. Wait 2s
  5. Retry (2/3)
  6. Wait 4s
  7. Retry (3/3)
  8. Give up, report error to agent
-

# Debugging MCP Servers

## Enable MCP Logging

```
code export RUST_LOG=codex_mcp_client=debug
```

### Log Output:

```
[DEBUG] Starting MCP server: local-memory
[DEBUG] Command: npx -y @modelcontextprotocol/server-memory
[DEBUG] Handshake complete
[DEBUG] Requesting tools/list...
[DEBUG] Received 3 tools: store_memory, search, analysis
[DEBUG] MCP server ready: local-memory
```

---

## Test MCP Server Manually

### MCP Inspector (official debugging tool):

```
code npm install -g @modelcontextprotocol/inspector

Test local-memory server
code npx @modelcontextprotocol/inspector npx -y
@modelcontextprotocol/server-memory
```

**Features:** - Test tool calls - Inspect responses - Debug connection issues

---

## Check MCP Server Status

```
code --mcp-status
```

### Output:

```
MCP Servers (3 configured):

local-memory:
 Status: Running (PID: 12345)
 Command: npx -y @modelcontextprotocol/server-memory
 Uptime: 2h 15m
 Tools: 3 (store_memory, search, analysis)

git-status:
 Status: Not started (lazy-load)
 Command: npx -y @just-every/mcp-server-git
 Tools: 3 (cached)

database:
 Status: Failed (startup timeout)
 Command: /path/to/db-server
 Error: Connection timeout after 20000ms
```

---

## Force Restart MCP Server

```
code --mcp-restart local-memory
```

**Use Case:** Server crashed, hung, or behaving incorrectly

---

## Common MCP Servers

### Filesystem Server

```
[mcp_servers.filesystem]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-filesystem",
"/allowed/path"]
```

**Tools:** Read/write files in allowed directory

---

### HTTP Server

```
[mcp_servers.http]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-http"]
```

**Tools:** Make HTTP requests

---

### Database Servers

#### PostgreSQL:

```
[mcp_servers.postgres]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-postgres",
"postgres://localhost/mydb"]
```

#### SQLite:

```
[mcp_servers.sqlite]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-sqlite",
"/path/to/db.sqlite"]
```

---

### Custom API Server

```
[mcp_servers.custom-api]
command = "/path/to/custom-mcp-server"
args = [--api-url, "https://api.example.com"]
env = { API_TOKEN = "secret" }
```

---

## Security Considerations

### 1. Validate Command Paths

#### Good:

```
[mcp_servers.trusted]
```

```
command = "npx" # Well-known command
args = ["-y", "@modelcontextprotocol/server-memory"]
```

**Bad:**

```
[mcp_servers.untrusted]
command = "/tmp/random-script.sh" # ✗ Untrusted source
```

---

## 2. Avoid Secrets in Config

**Good:**

```
[mcp_servers.database]
command = "/path/to/db-server"
env = { DB_PASSWORD = "secret" } # △ Still visible in config

Better: Use environment variable
env = { DB_PASSWORD = "$DB_PASSWORD_FROM_ENV" }
```

**Best:**

```
Store secret in environment
export DB_PASSWORD="secret"

[mcp_servers.database]
command = "/path/to/db-server"
Server reads $DB_PASSWORD from environment
```

---

## 3. Restrict Network Access

**Sandbox Mode:** MCP servers inherit sandbox restrictions

```
sandbox_mode = "read-only" # MCP servers also read-only

[mcp_servers.filesystem]
command = "npx"
args = ["-y", "@modelcontextprotocol/server-filesystem",
"/safe/path"]
```

---

## Best Practices

### 1. Use Lazy Loading

**Default behavior** (don't change): - Servers start on first use - Faster session startup - Lower resource usage

---

### 2. Set Appropriate Timeouts

**Fast servers** (in-memory):

```
[mcp_servers.memory]
startup_timeout_ms = 5000 # 5s
```

**Slow servers** (database, network):

```
[mcp_servers.database]
startup_timeout_ms = 30000 # 30s
```

---

### 3. Monitor MCP Server Logs

```
export RUST_LOG=debug
code

Check logs for MCP errors
tail -f ~/.code/debug.log | grep MCP
```

---

### 4. Test Servers with MCP Inspector

```
npx @modelcontextprotocol/inspector <command> <args>
```

**Benefit:** Catch configuration errors before using in production

---

## Summary

**MCP Servers** enable: - Knowledge persistence (local-memory) - Git operations (git-status) - Custom tools (database, API, filesystem) - External service integration

### Configuration:

```
[mcp_servers.<name>]
command = "executable"
args = ["arg1", "arg2"]
env = { KEY = "value" }
startup_timeout_ms = 10000
```

**Features:** - Lazy loading (start on first use) - Tool caching (faster startup) - Automatic retry (tool call failures) - Hot-reload support (config changes)

**Debugging:** - MCP Inspector (test servers) - --mcp-status (check status) - --mcp-restart (force restart) - Debug logging (RUST\_LOG=debug)

**Next:** [Environment Variables](#)

---

## Model Configuration

Provider setup, reasoning effort, and model tuning.

---

## Overview

Model configuration controls: 1. **Provider Selection** - Which AI service to use (OpenAI, Anthropic, Google, Ollama) 2. **Model Selection** - Which specific model (GPT-5, o3, Claude, Gemini) 3.

**Reasoning Configuration** - Effort level, summaries, verbosity 4.  
**Network Tuning** - Retries, timeouts, streaming

---

## Basic Model Configuration

### Minimal Setup

```
~/.code/config.toml

model = "gpt-5"
model_provider = "openai"
```

### Environment:

```
export OPENAI_API_KEY="sk-proj-..."
```

---

### Model Selection

**Available Models** (OpenAI): - gpt-5 - Default, balanced reasoning and cost - gpt-5-codex - Optimized for code generation - o3 - Maximum reasoning capability (premium) - o4-mini - Fast reasoning model - gpt-4o - Legacy model - gpt-4o-mini - Fast, cheap legacy model

### Configuration:

```
model = "o3" # Use premium reasoning model
```

### CLI Override:

```
code --model o3 "complex task"
```

---

## Provider Configuration

### OpenAI (Default)

```
model_provider = "openai"

[model_providers.openai]
name = "OpenAI"
base_url = "https://api.openai.com/v1"
env_key = "OPENAI_API_KEY"
wire_api = "responses" # or "chat"
request_max_retries = 4
stream_max_retries = 10
stream_idle_timeout_ms = 300000 # 5 minutes
```

### Environment Variables:

```
export OPENAI_API_KEY="sk-proj-..."

Optional overrides
export OPENAI_BASE_URL="https://custom.openai.com/v1"
export OPENAI_WIRE_API="chat" # Force chat completions API
```

---

## Anthropic (Claude)

```
model_provider = "anthropic"
model = "claude-3-5-sonnet"

[model_providers.anthropic]
name = "Anthropic"
base_url = "https://api.anthropic.com"
env_key = "ANTHROPIC_API_KEY"
wire_api = "chat"
```

### Environment:

```
export ANTHROPIC_API_KEY="sk-ant-..."
```

---

## Google (Gemini)

```
model_provider = "google"
model = "gemini-2.0-flash-001"

[model_providers.google]
name = "Google"
base_url = "https://generativelanguage.googleapis.com/v1beta"
env_key = "GOOGLE_API_KEY"
wire_api = "chat"
```

### Environment:

```
export GOOGLE_API_KEY="..."
```

---

## Ollama (Local)

```
model_provider = "ollama"
model = "mistral"

[model_providers.ollama]
name = "Ollama"
base_url = "http://localhost:11434/v1"
No env_key needed for local Ollama
```

### Setup:

```
Install Ollama
curl -fsSL https://ollama.com/install.sh | sh

Pull model
ollama pull mistral

Start server
ollama serve
```

---

## Azure OpenAI

```
model_provider = "azure"
model = "gpt-5"

[model_providers.azure]
```

```

name = "Azure OpenAI"
base_url = "https://YOUR_PROJECT.openai.azure.com/openai"
env_key = "AZURE_OPENAI_API_KEY"
wire_api = "chat"
query_params = { api-version = "2025-04-01-preview" }

```

### Environment:

```
export AZURE_OPENAI_API_KEY="..."
```

---

## Custom Provider

```

[model_providers.custom]
name = "Custom Provider"
base_url = "https://custom.api.com/v1"
env_key = "CUSTOM_API_KEY"
wire_api = "chat"

Optional: Static HTTP headers
http_headers = { "X-Custom-Header" = "value" }

Optional: Dynamic HTTP headers from environment
env_http_headers = { "X-Features" = "CUSTOM_FEATURES" }

Network tuning
request_max_retries = 3
stream_max_retries = 5
stream_idle_timeout_ms = 180000 # 3 minutes

```

---

## Reasoning Configuration

### Reasoning Effort

Controls how much computational effort the model uses for reasoning.

**Options:** - minimal - Fastest, least reasoning (previously “none”) - low - Light reasoning - medium - Balanced (default) - high - Maximum reasoning (premium cost)

### Configuration:

```
model_reasoning_effort = "high"
```

### Use Cases:

| Effort  | Use Case                          | Cost    | Speed    |
|---------|-----------------------------------|---------|----------|
| minimal | Simple formatting, trivial tasks  | Lowest  | Fastest  |
| low     | Straightforward code changes      | Low     | Fast     |
| medium  | Moderate complexity tasks         | Medium  | Moderate |
| high    | Complex refactoring, architecture | Highest | Slowest  |

### Example:

```
Premium profile for complex tasks
[profiles.premium]
```

```
model = "o3"
model_reasoning_effort = "high"

Fast profile for simple tasks
[profiles.fast]
model = "gpt-4o-mini"
model_reasoning_effort = "minimal"
```

---

## Reasoning Summary

Controls summarization of reasoning process.

**Options:** - auto - Model decides (default) - concise - Brief summary - detailed - Comprehensive summary - none - No summary

### Configuration:

```
model_reasoning_summary = "detailed"
```

### Example Output:

**auto:**

```
Reasoning: Analyzing code structure...
```

**concise:**

```
Reasoning: Identified 3 refactoring opportunities.
```

**detailed:**

```
Reasoning: Analyzed codebase structure. Identified 3 refactoring opportunities:
```

1. Extract duplicate validation logic into shared function
2. Replace switch statement with strategy pattern
3. Simplify nested conditionals with early returns

**none:**

```
(No reasoning summary shown)
```

---

## Model Verbosity (GPT-5 Only)

Controls output length/detail for GPT-5 family models.

**Options:** - low - Concise output - medium - Balanced (default) - high - Detailed explanations

### Configuration:

```
model = "gpt-5"
model_verbosity = "low"
```

### Example:

**low:**

```
Refactored validation logic. See main.rs:42.
```

**medium:**

Refactored validation logic into shared function `validate\_input()`  
in main.rs:42. Updated 3 call sites.

**high:**

Refactored validation logic to improve maintainability:  
1. Extracted duplicate validation code into new function  
`validate\_input()`  
- Location: main.rs:42-58  
- Parameters: &str input, bool strict\_mode  
- Returns: Result<(), ValidationError>  
2. Updated call sites: handler.rs:15, api.rs:33, cli.rs:67  
3. Added unit tests: tests/validation\_test.rs:10-45

---

## Context Window Configuration

### Context Window Size

**Default:** Auto-detected based on model

**Manual Override:**

```
model_context_window = 128000 # 128K tokens
```

**Use Case:** New models not yet recognized by Codex

---

### Max Output Tokens

**Default:** Auto-detected based on model

**Manual Override:**

```
model_max_output_tokens = 16384 # 16K tokens
```

**Use Case:** Limit output length for cost control

---

## Network Tuning

### Request Retries

**Default:** 4 retries

**Configuration:**

```
[model_providers.openai]
request_max_retries = 6 # Increase for unreliable networks
```

**Behavior:** Exponential backoff (1s, 2s, 4s, 8s, 16s, 32s)

---

### Stream Retries

**Default:** 10 retries

## **Configuration:**

```
[model_providers.openai]
stream_max_retries = 15 # Increase for flaky connections
```

**Use Case:** Unstable network, frequent disconnects

---

## **Stream Idle Timeout**

**Default:** 300,000 ms (5 minutes)

## **Configuration:**

```
[model_providers.openai]
stream_idle_timeout_ms = 600000 # 10 minutes for slow models
```

**Use Case:** Very slow models or complex tasks

---

# **Wire API Selection**

## **Responses API (Default for GPT-5/o3)**

**Features:** - Native reasoning support - Reasoning summaries - Verbosity control - Optimized for GPT-5 family

## **Configuration:**

```
[model_providers.openai]
wire_api = "responses"
```

---

## **Chat Completions API (Legacy)**

**Features:** - Compatible with all OpenAI models - Compatible with most third-party providers - Simpler protocol

## **Configuration:**

```
[model_providers.openai]
wire_api = "chat"
```

**Use Case:** Third-party providers, older models

---

# **Advanced Configuration**

## **Force Reasoning Support**

**Use Case:** Custom models that support reasoning but aren't auto-detected

## **Configuration:**

```
model_supports_reasoning_summaries = true
```

---

## Disable Response Storage (ZDR Accounts)

**Use Case:** Zero Data Retention accounts

**Configuration:**

```
disable_response_storage = true
```

**Effect:** Forces Chat Completions API instead of Responses API

---

## Configuration Examples

### Premium Quality Setup

```
Maximum reasoning quality
model = "o3"
model_provider = "openai"
model_reasoning_effort = "high"
model_reasoning_summary = "detailed"
model_verbosity = "high"

[model_providers.openai]
wire_api = "responses"
```

---

### Fast Iteration Setup

```
Speed over quality
model = "gpt-4o-mini"
model_provider = "openai"
model_reasoning_effort = "minimal"
model_reasoning_summary = "none"
model_verbosity = "low"

[model_providers.openai]
wire_api = "chat"
```

---

### Local Development Setup

```
Ollama for offline development
model = "mistral"
model_provider = "ollama"

[model_providers.ollama]
name = "Ollama"
base_url = "http://localhost:11434/v1"
```

---

### Multi-Provider Setup

```
Default to OpenAI
model = "gpt-5"
model_provider = "openai"

OpenAI provider
```

```

[model_providers.openai]
name = "OpenAI"
base_url = "https://api.openai.com/v1"
env_key = "OPENAI_API_KEY"
wire_api = "responses"

Anthropic provider
[model_providers.anthropic]
name = "Anthropic"
base_url = "https://api.anthropic.com"
env_key = "ANTHROPIC_API_KEY"
wire_api = "chat"

Ollama provider (local)
[model_providers.ollama]
name = "Ollama"
base_url = "http://localhost:11434/v1"

Profiles for quick switching
[profiles.openai]
model_provider = "openai"
model = "gpt-5"

[profiles.claude]
model_provider = "anthropic"
model = "claude-3-5-sonnet"

[profiles.local]
model_provider = "ollama"
model = "mistral"

```

#### Usage:

```

code --profile openai "task"
code --profile claude "task"
code --profile local "task"

```

---

## Debugging Model Configuration

### Check Effective Configuration

```
code --config-dump | grep -A 10 "model"
```

#### Output:

```

model = "o3" # From: CLI flag
model_provider = "openai" # From: config.toml
model_reasoning_effort = "high" # From: profile 'premium'
model_reasoning_summary = "detailed" # From: profile 'premium'

```

---

## Test Provider Connection

```

Enable debug logging
export RUST_LOG=debug
code "Hello world"

```

#### Log Output:

```
[DEBUG] Model provider: openai
[DEBUG] Base URL: https://api.openai.com/v1
[DEBUG] Wire API: responses
[DEBUG] Model: o3
[DEBUG] Reasoning effort: high
[INFO] Connection successful
```

---

## Summary

**Model Configuration** covers: - Provider selection (OpenAI, Anthropic, Google, Ollama, Azure, custom) - Model selection (GPT-5, o3, Claude, Gemini, etc.) - Reasoning effort (minimal, low, medium, high) - Reasoning summaries (auto, concise, detailed, none) - Model verbosity (low, medium, high) - Network tuning (retries, timeouts) - Wire API selection (responses, chat)

**Best Practices:** - Use profiles for different quality/speed tradeoffs - Store API keys in environment variables - Tune network settings for your connection quality - Use local providers (Ollama) for offline development

**Next:** [Agent Configuration](#)

---

## Precedence System

5-tier configuration precedence with examples.

---

### Overview

The configuration system implements **5-tier precedence** (highest to lowest):

1. **CLI Flags** (highest priority) - Command-line arguments
2. **Shell Environment** - Environment variables
3. **Profile** - Named configuration sets
4. **Config File** - `~/.code/config.toml`
5. **Defaults** (lowest priority) - Built-in fallback values

**Rule:** Higher tiers override lower tiers

---

### Precedence Order

#### Tier 1: CLI Flags (Highest)

**Priority:** Highest

**Usage:**

```
Specific model flags
code --model o3 "task description"
code --profile premium "task"
```

```

Generic config flag
code --config model="gpt-5"
code --config approval_policy=never
code -c model_reasoning_effort=high

Deep config paths (dot notation)
code --config model_providers.openai.wire_api="chat"
code --config shell_environment_policy.include_only='["PATH",
"HOME"]'

```

**Characteristics:** - Overrides all other tiers - Session-specific (not persisted) - Supports dot notation for nested values - Values in TOML format (not JSON)

#### Examples:

```

Override model
code --model o3

Override approval policy
code --config approval_policy=never

Override provider config
code --config
model_providers.openai.base_url="https://custom.api.com"

```

---

## Tier 2: Shell Environment

**Priority:** 2nd highest

**Patterns:** - CODEX\_HOME, CODE\_HOME - Installation directory - <PROVIDER>\_API\_KEY - API keys (e.g., OPENAI\_API\_KEY) - OPENAI\_BASE\_URL - Provider base URL override - OPENAI\_WIRE\_API - Wire protocol override ("responses" or "chat") - CODEX\_MODEL, CODEX\_PROVIDER - Model/provider overrides

#### Usage:

```

API keys (most common)
export OPENAI_API_KEY="sk-proj-..."
export ANTHROPIC_API_KEY="sk-ant-..."
export GOOGLE_API_KEY="..."

Home directory
export CODEX_HOME="/custom/path"

Provider overrides
export OPENAI_BASE_URL="https://custom.openai.com/v1"
export OPENAI_WIRE_API="responses"

Model overrides
export CODEX_MODEL="gpt-5"
export CODEX_PROVIDER="anthropic"

```

**Characteristics:** - Persistent for session duration - Useful for secrets (API keys) - Environment-specific overrides - Case-insensitive for most values

---

## Tier 3: Profile

**Priority:** 3rd highest

**Activation:**

```
Via CLI
code --profile premium "task"

Via config.toml
profile = "premium"
```

**Definition:**

```
~/.code/config.toml

[profiles.premium]
model = "o3"
model_reasoning_effort = "high"
model_reasoning_summary = "detailed"
approval_policy = "never"

[profiles.fast]
model = "gpt-4o-mini"
model_reasoning_effort = "low"
approval_policy = "never"

[profiles.ci]
model = "gpt-4o"
approval_policy = "never"
sandbox_mode = "read-only"
```

**Characteristics:** - Named configuration sets - Overrides config.toml base values - Can be selected per-session via CLI - Useful for different workflows

---

## Tier 4: Config File

**Priority:** 4th highest

**Location:** `~/.code/config.toml`

**Example:**

```
model = "gpt-5"
model_provider = "openai"
approval_policy = "on-request"
sandbox_mode = "workspace-write"

[quality_gates]
plan = ["gemini", "claude", "code"]
tasks = ["gemini"]
```

**Characteristics:** - Persistent across sessions - User-specific configuration - Hot-reloadable (changes apply without restart) - TOML format (human-readable)

---

## Tier 5: Defaults (Lowest)

**Priority:** Lowest

**Source:** Built-in code defaults

**Example:**

```
impl Default for AppConfig {
 fn default() -> Self {
 Self {
 model: "gpt-5-codex".to_string(),
 model_provider: "openai".to_string(),
 approval_policy: ApprovalPolicy::OnRequest,
 sandbox_mode: SandboxMode::ReadOnly,
 // ... 30+ more fields
 }
 }
}
```

**Characteristics:** - Fallback values when no other tier specifies - Hardcoded in Rust source - Guaranteed sensible defaults - Work out-of-the-box without configuration

---

## Precedence Examples

### Example 1: Model Selection

**Setup:**

```
~/.code/config.toml
model = "gpt-5"

[profiles.premium]
model = "o3"

export CODEX_MODEL="gpt-4o"
```

**Scenarios:**

| Command                                  | Effective Model | Why                                              |
|------------------------------------------|-----------------|--------------------------------------------------|
| code "task"                              | gpt-4o          | Env var<br>(Tier 2) ><br>config.toml<br>(Tier 4) |
| code --profile premium "task"            | o3              | Profile<br>(Tier 3) ><br>env var<br>(Tier 2)     |
| code --model o1 "task"                   | o1              | CLI flag<br>(Tier 1) ><br>all others             |
| code --profile premium --model o1 "task" | o1              | CLI flag<br>(Tier 1)<br>wins                     |

---

### Example 2: API Key

### Setup:

```
~/.code/config.toml
(no API key specified)

export OPENAI_API_KEY="sk-proj-env-key"
```

### Scenarios:

| Command                                                                   | Effective Key   | Why                                              |
|---------------------------------------------------------------------------|-----------------|--------------------------------------------------|
| code "task"                                                               | sk-proj-env-key | Env var<br>(Tier 2)<br>><br>defaults<br>(Tier 5) |
| code --config<br>model_providers.openai.env_key="sk- proj-cli-key" "task" | sk-proj-cli-key | CLI flag<br>(Tier 1)<br>> env<br>var (Tier<br>2) |

**Note:** API keys should **always** be stored in environment variables, never in config.toml.

## Example 3: Approval Policy

### Setup:

```
~/.code/config.toml
approval_policy = "on-request"

[profiles.ci]
approval_policy = "never"

No environment overrides
```

### Scenarios:

| Command                                                               | Effective Policy | Why                                               |
|-----------------------------------------------------------------------|------------------|---------------------------------------------------|
| code "task"                                                           | on-request       | config.toml<br>(Tier 4) ><br>defaults<br>(Tier 5) |
| code --profile ci "task"                                              | never            | Profile<br>(Tier 3) ><br>config.toml<br>(Tier 4)  |
| code --profile ci --<br>config<br>approval_policy=untrusted<br>"task" | untrusted        | CLI flag<br>(Tier 1) ><br>profile<br>(Tier 3)     |

## Example 4: Complex Nested Config

### Setup:

```

~/.code/config.toml
[model_providers.openai]
base_url = "https://api.openai.com/v1"
wire_api = "responses"

export OPENAI_BASE_URL="https://custom.openai.com/v1"

```

### Scenarios:

| Command                                                        | Effective URL                |
|----------------------------------------------------------------|------------------------------|
| code "task"                                                    | https://custom.openai.com/v1 |
| code --config<br>model_providers.openai.wire_api="chat" "task" | https://custom.openai.com/v1 |

## Special Cases

### Shell Environment Policy Override

**Warning:** shell\_environment\_policy.set can override config values at runtime.

#### Example:

```

~/.code/config.toml
approval_policy = "always"

[shell_environment_policy]
set = { APPROVAL_POLICY = "never" } # △ OVERIDES top-level
setting!

```

**Behavior:** APPROVAL\_POLICY=never wins at runtime (subprocess environment)

**Best Practice:** Avoid using shell\_environment\_policy.set for keys that exist as top-level config options.

---

### Profile Selection Precedence

**Priority:** CLI --profile > config.toml profile field > no profile

#### Example:

```

~/.code/config.toml
profile = "fast" # Default profile

```

```
[profiles.fast]
```

```
model = "gpt-4o-mini"
```

```
[profiles.premium]
```

```
model = "o3"
```

| Command                              | Effective Profile | Model       | Why                                 |
|--------------------------------------|-------------------|-------------|-------------------------------------|
| code "task" --profile premium "task" | fast              | gpt-4o-mini | config.toml profile field           |
|                                      | premium           | o3          | CLI --profile overrides config.toml |

## Precedence Table

**Summary:**

| Tier | Source      | Example            | Persistence                    | Override M                         |
|------|-------------|--------------------|--------------------------------|------------------------------------|
| 1    | CLI Flags   | --model o3         | Session-only                   | Command-line                       |
| 2    | Environment | OPENAI_API_KEY=... | Session/shell                  | export VAR=value                   |
| 3    | Profile     | [profiles.premium] | Persistent<br>(in config.toml) | --profile name<br>profile = "name" |
| 4    | Config File | model = "gpt-5"    | Persistent                     | Edit<br>~/.code/config             |
| 5    | Defaults    | "gpt-5-codex"      | Built-in                       | (Cannot over-                      |

## Debugging Precedence

### Check Effective Configuration

**Command:**

```
code --config-dump
```

**Output:**

```
Effective configuration (after precedence resolution)
model = "o3" # From: CLI flag (--model o3)
model_provider = "openai" # From: config.toml
approval_policy = "never" # From: profile 'premium'
... full effective config
```

## Trace Configuration Source

**Example:**

```
With verbose logging
```

```
export RUST_LOG=debug
code --model o3 "task"
```

#### Log Output:

```
[DEBUG] Config layer 5 (defaults): model=gpt-5-codex
[DEBUG] Config layer 4 (config.toml): model=gpt-5
[DEBUG] Config layer 3 (profile 'premium'): model=o3
[DEBUG] Config layer 1 (CLI flag): model=o3
[INFO] Effective model: o3 (source: CLI flag)
```

---

## Best Practices

### 1. Use Environment Variables for Secrets

#### Good:

```
export OPENAI_API_KEY="sk-proj-..."
```

#### Bad:

```
DON'T: API keys should NOT be in config.toml
[model_providers.openai]
api_key = "sk-proj-..." # ✗ Security risk!
```

---

### 2. Use Profiles for Workflows

#### Example:

```
Fast iteration
[profiles.fast]
model = "gpt-4o-mini"
approval_policy = "never"

Premium quality
[profiles.premium]
model = "o3"
model_reasoning_effort = "high"

CI/automation
[profiles.ci]
model = "gpt-4o"
approval_policy = "never"
sandbox_mode = "read-only"
```

#### Usage:

```
code --profile fast "quick formatting"
code --profile premium "complex refactor"
code --profile ci "generate report"
```

---

### 3. Use CLI Flags for One-Off Overrides

#### Example:

```
One-time model override
```

```
code --model o3 "complex task"

One-time approval policy override
code --config approval_policy=never "trusted script"
```

---

## 4. Keep config.toml for Persistent Preferences

**Example:**

```
~/.code/config.toml

Personal preferences (persistent)
model = "gpt-5"
approval_policy = "on-request"
sandbox_mode = "workspace-write"
file_opener = "vscode"

[tui.theme]
name = "dark-carbon-night"
```

---

## Summary

**5-Tier Precedence** (highest to lowest): 1. CLI Flags - Session-specific overrides 2. Environment Variables - Secrets and env-specific config 3. Profiles - Named configuration sets 4. Config File - Persistent user preferences 5. Defaults - Built-in fallback values

**Rule:** Higher tiers override lower tiers

**Best Practices:** - Secrets → Environment variables - Workflows → Profiles - One-off overrides → CLI flags - Persistent preferences → config.toml

**Next:** [Model Configuration](#)

---

## Quality Gate Customization

Per-checkpoint agent selection and override rules.

---

## Overview

**Quality Gates** are checkpoints in the spec-kit workflow that ensure standards are met before proceeding.

**5 Quality Gates:** 1. **Plan** - Architectural planning (multi-agent consensus) 2. **Tasks** - Task decomposition (single-agent) 3. **Validate** - Test strategy validation (multi-agent) 4. **Audit** - Security/compliance review (premium agents) 5. **Unlock** - Ship/no-ship decision (premium agents)

**Configuration:** [quality\_gates] section in config.toml

---

# Quality Gate Configuration

## Basic Configuration

```
~/.code/config.toml

[quality_gates]
plan = ["gemini", "claude", "code"] # Multi-agent planning
tasks = ["gemini"] # Single-agent task
breakdown
validate = ["gemini", "claude", "code"] # Multi-agent test
validation
audit = ["gemini", "claude", "gpt_codex"] # Security/compliance
unlock = ["gemini", "claude", "gpt_codex"] # Ship decision
```

---

## Field Reference

| Field    | Purpose                 | Recommended Agents | Cost Tier        |
|----------|-------------------------|--------------------|------------------|
| plan     | Architectural decisions | 3 agents (diverse) | Tier 2 (~\$0.35) |
| tasks    | Task breakdown          | 1 agent (cheap)    | Tier 1 (~\$0.10) |
| validate | Test strategy           | 3 agents (diverse) | Tier 2 (~\$0.35) |
| audit    | Security/compliance     | 3+ premium         | Tier 3 (~\$0.80) |
| unlock   | Ship decision           | 3 premium          | Tier 3 (~\$0.80) |

---

# Agent Selection Strategy

## Multi-Agent Consensus (Plan, Validate)

**Purpose:** Diverse perspectives on complex decisions

**Recommended Setup:**

```
[quality_gates]
plan = ["gemini", "claude", "code"] # Fast + Balanced + Strategic
```

**Agent Roles:** - gemini - Fast consensus, broad coverage (12.5x cheaper) - claude - Balanced reasoning, edge case detection (12x cheaper) - code (GPT-5) - Strategic planning, complex reasoning (baseline)

**Why 3 Agents:** - 2 agents: Risk of tie (no consensus) - 3 agents: Majority vote possible - 4+ agents: Diminishing returns, higher cost

---

## Single-Agent Deterministic (Tasks)

**Purpose:** Straightforward decomposition without opinion diversity

### **Recommended Setup:**

```
[quality_gates]
tasks = ["gemini"] # Single cheap agent
```

**Why Single Agent:** - Task breakdown is mechanical (not strategic) - No benefit from consensus - Cost savings (1 agent vs 3)

---

### **Premium Consensus (Audit, Unlock)**

**Purpose:** Critical decisions requiring maximum reasoning

### **Recommended Setup:**

```
[quality_gates]
audit = ["gemini", "claude", "gpt_codex"] # Security-focused
unlock = ["gemini", "claude", "gpt_codex"] # Ship decision
```

**Agent Selection:** - gemini - Broad vulnerability scanning - claude - Edge case security analysis - gpt\_codex - Code-specific security patterns

**Why Premium:** - Audit prevents security incidents (\$1000s in damages) - Unlock prevents production bugs (\$1000s in incidents) - \$0.80 cost per stage justifiable for critical gates

---

## **Custom Configurations**

### **Cost-Optimized Setup**

**Goal:** Minimize cost while maintaining quality

```
[quality_gates]
plan = ["gemini", "claude"] # 2 cheap agents (no GPT-5)
tasks = ["gemini"] # Single cheap agent
validate = ["gemini", "claude"] # 2 cheap agents
audit = ["gemini", "claude", "code"] # 2 cheap + 1 mid-tier
unlock = ["gemini", "claude", "code"] # 2 cheap + 1 mid-tier
```

**Cost Savings:** ~60% reduction (from \$2.70 to ~\$1.08 per full pipeline)

**Tradeoff:** Less strategic depth (no GPT-5 on plan/validate)

---

### **Premium Quality Setup**

**Goal:** Maximum quality, cost secondary

```
[quality_gates]
plan = ["gemini", "claude", "code", "gpt_pro"] # 4 agents (premium)
tasks = ["code"] # GPT-5 for task breakdown
validate = ["gemini", "claude", "code", "gpt_pro"] # 4 agents
audit = ["gemini", "claude", "code", "gpt_codex", "gpt_pro"] # 5 agents
unlock = ["gemini", "claude", "gpt_codex", "gpt_pro"] # 4 premium
```

**Cost:** ~\$4.50 per full pipeline (66% increase)

**Benefit:** Maximum reasoning, redundant validation

---

## Specialist Configuration

**Goal:** Assign specialists per gate

```
Define specialized agents
[[agents]]
name = "security-specialist"
canonical_name = "security"
command = "claude"
instructions = "Focus on OWASP Top 10, cryptography, auth/authz."

[[agents]]
name = "test-specialist"
canonical_name = "test"
command = "gemini"
instructions = "Focus on test coverage, edge cases, property-based
tests."

Quality gates with specialists
[quality_gates]
plan = ["gemini", "claude", "code"] # General agents
tasks = ["gemini"] # General agent
validate = ["test", "claude", "code"] # Test specialist for
validation
audit = ["security", "claude", "gpt_codex"] # Security specialist
for audit
unlock = ["gemini", "claude", "gpt_codex"] # General agents
```

---

## Per-Checkpoint Overrides

### Override at Runtime

Quality gates can be overridden per-command:

```
Override plan agents
/speckit.plan SPEC-KIT-065 --agents gemini,claude

Override validate agents (premium quality)
/speckit.validate SPEC-KIT-065 --agents gemini,claude,code,gpt_pro

Override audit agents (cost-optimized)
/speckit.audit SPEC-KIT-065 --agents gemini,claude
```

**Use Case:** One-off quality/cost tradeoffs

---

### Environment Variable Overrides

```
Override plan agents via env var
export SPECKIT_QUALITY_GATES_PLAN="gemini,claude,code,gpt_pro"
/speckit.plan SPEC-KIT-065
```

```
Override tasks agents
export SPECKIT_QUALITY_GATES_TASKS="code"
/speckit.tasks SPEC-KIT-065
```

**Precedence:** Env var > config.toml

---

## Consensus Thresholds

### Minimum Consensus

**Default:** 2/3 agents (66.7%)

**Configuration:**

```
[quality_gates]
plan = ["gemini", "claude", "code"]
consensus_threshold = 0.67 # 2/3 agents must agree
```

**Example:** - 3 agents, 2 agree → ✅ Pass (2/3 = 66.7%) - 3 agents, 1 agrees → ❌ Fail (1/3 = 33.3%)

---

### Strict Consensus

**Configuration:**

```
[quality_gates]
unlock = ["gemini", "claude", "gpt_codex"]
consensus_threshold = 1.0 # 100% agreement required
```

**Use Case:** Critical ship decisions (unlock gate)

**Behavior:** All agents must agree to pass

---

### Relaxed Consensus

**Configuration:**

```
[quality_gates]
plan = ["gemini", "claude", "code"]
consensus_threshold = 0.5 # 50% majority
```

**Use Case:** Exploratory planning (early stages)

**Behavior:** Simple majority sufficient

---

## Degradation Handling

### Agent Failure Behavior

**Scenario:** One agent fails (timeout, error)

**Default Behavior:** 1. Retry up to 3 times (AR-2) 2. If still fails, continue with remaining agents 3. Consensus valid if remaining agents  $\geq$  threshold

**Example:**

```
[quality_gates]
plan = ["gemini", "claude", "code"] # 3 agents
consensus_threshold = 0.67
```

**If code agent fails:** - Remaining: gemini, claude (2 agents) - If both agree:  $2/2 = 100\% \geq 67\%$   $\rightarrow \checkmark$  Pass - If disagree:  $1/2 = 50\% < 67\%$   $\rightarrow \times$  Fail

---

## Empty Consensus Handling

**Scenario:** All agents fail

**Behavior:** Fall back to degraded mode

**Example:**

```
All agents failed
✗ Quality gate failed: No agents returned valid consensus
△ Continuing in degraded mode (manual review required)
```

**User Action:** Manual PRD review and approval

---

## Quality Gate Validation

### Startup Validation

**Validation Rules:** 1. All agent names must exist in [[agents]] 2. Agent canonical\_name must match quality gate references 3. Agents must be enabled 4. Minimum 1 agent per gate

**Example Error:**

```
Config validation error:
 quality_gates.plan: Agent 'unknown-agent' not found
 quality_gates.audit: Agent 'gpt_pro' exists but is disabled
```

Fix: Check [[agents]] configuration

---

### Runtime Validation

#### Per-command validation:

```
/speckit.plan SPEC-KIT-065
```

**Validation:** 1. All specified agents are available 2. Agents can be spawned (commands exist) 3. Consensus threshold achievable

**Example Error:**

```
✗ Cannot execute /speckit.plan:
```

```
- Agent 'claude' command not found
- Consensus threshold 0.67 requires ≥2 agents, only 1 available
```

Fix: Install missing agent or adjust consensus\_threshold

---

## Example Configurations

### Balanced (Default)

```
[quality_gates]
plan = ["gemini", "claude", "code"] # 3 agents, diverse
tasks = ["gemini"] # 1 agent, cheap
validate = ["gemini", "claude", "code"] # 3 agents, diverse
audit = ["gemini", "claude", "gpt_codex"] # 3 agents, security-focused
unlock = ["gemini", "claude", "gpt_codex"] # 3 agents, ship decision
```

**Cost:** ~\$2.70 per full pipeline

---

### Cost-Optimized

```
[quality_gates]
plan = ["gemini", "claude"] # 2 agents (no GPT-5)
tasks = ["gemini"] # 1 agent
validate = ["gemini", "claude"] # 2 agents
audit = ["gemini", "claude"] # 2 agents (no premium)
unlock = ["gemini", "claude"] # 2 agents
```

**Cost:** ~\$0.80 per full pipeline (70% reduction)

---

### Premium Quality

```
[quality_gates]
plan = ["gemini", "claude", "code", "gpt_pro"] # 4 agents
tasks = ["code"] # GPT-5 for tasks
validate = ["gemini", "claude", "code", "gpt_pro"] # 4 agents
audit = ["gemini", "claude", "code", "gpt_codex", "gpt_pro"] # 5 agents
unlock = ["gemini", "claude", "gpt_codex", "gpt_pro"] # 4 agents
```

**Cost:** ~\$4.50 per full pipeline (66% increase)

---

### Single-Agent (Development)

```
[quality_gates]
plan = ["gemini"] # Fast iteration
tasks = ["gemini"]
validate = ["gemini"]
audit = ["gemini"]
unlock = ["gemini"]
```

**Cost:** ~\$0.20 per full pipeline (93% reduction)

**Use Case:** Rapid prototyping, development iteration

---

## Debugging Quality Gates

### Check Quality Gate Configuration

```
code --quality-gates-dump
```

#### Output:

```
[quality_gates]
plan = ["gemini", "claude", "code"] # 3 agents
tasks = ["gemini"] # 1 agent
validate = ["gemini", "claude", "code"] # 3 agents
audit = ["gemini", "claude", "gpt_codex"] # 3 agents
unlock = ["gemini", "claude", "gpt_codex"] # 3 agents

Consensus thresholds (effective)
plan.consensus_threshold = 0.67
validate.consensus_threshold = 0.67
unlock.consensus_threshold = 1.0 # Strict (100%)
```

---

## Validate Agent Availability

```
code --check-quality-gates
```

#### Output:

Validating quality gates...

```
plan:
 [✓] gemini (enabled, command found)
 [✓] claude (enabled, command found)
 [✓] code (enabled, command found)

tasks:
 [✓] gemini (enabled, command found)

validate:
 [✓] gemini (enabled, command found)
 [✓] claude (enabled, command found)
 [✓] code (enabled, command found)

audit:
 [✓] gemini (enabled, command found)
 [✓] claude (enabled, command found)
 [✗] gpt_codex (disabled)

unlock:
 [✓] gemini (enabled, command found)
 [✓] claude (enabled, command found)
 [✗] gpt_codex (disabled)

⚠ Warning: gpt_codex is disabled but referenced in audit, unlock
gates
```

---

## Best Practices

### 1. Use 3 Agents for Consensus

**Recommended:** 3 agents for plan, validate, audit, unlock

**Reason:** Allows majority vote, avoids ties

---

### 2. Use 1 Agent for Deterministic Tasks

**Recommended:** 1 agent for tasks

**Reason:** Task breakdown is mechanical, no consensus needed

---

### 3. Reserve Premium Agents for Critical Gates

**Good:**

```
[quality_gates]
plan = ["gemini", "claude", "code"] # Mid-tier for planning
audit = ["gemini", "claude", "gpt_pro"] # Premium for security
unlock = ["gemini", "claude", "gpt_pro"] # Premium for ship
decision
```

---

### 4. Test Quality Gate Configuration

```
Dry-run to validate config
/speckit.plan SPEC-TEST-001 --dry-run
```

---

## Summary

**Quality Gate Customization** covers: - 5 quality gates (plan, tasks, validate, audit, unlock) - Agent selection strategies (multi-agent, single-agent, premium) - Cost optimization (70-93% reduction possible) - Consensus thresholds (50-100%) - Degradation handling (agent failures) - Runtime overrides (CLI, env vars)

**Best Practices:** - 3 agents for consensus gates - 1 agent for deterministic gates - Premium agents for critical decisions - Test configuration with dry-run

**Next:** [Hot-Reload](#)

---

## Template Customization

Installing, modifying, and versioning custom templates.

---

## Overview

**Templates** provide pre-configured settings for common workflows.

**Use Cases:** - Team-specific default configurations - Project-specific quality gate settings - Environment-specific profiles (dev, staging, production) - Organization-wide standards

**Location:** `~/.code/templates/`

---

## Template Structure

### Template Format

```
~/.code/templates/premium-quality.toml

[template]
name = "Premium Quality"
version = "1.0.0"
description = "Premium quality configuration with maximum reasoning"
author = "Your Name"
created = "2025-11-17"

Template configuration (will be merged with config.toml)
[config]
model = "o3"
model_reasoning_effort = "high"
model_reasoning_summary = "detailed"
approval_policy = "never"

[config.quality_gates]
plan = ["gemini", "claude", "code", "gpt_pro"]
tasks = ["code"]
validate = ["gemini", "claude", "code", "gpt_pro"]
audit = ["gemini", "claude", "code", "gpt_coder", "gpt_pro"]
unlock = ["gemini", "claude", "gpt_coder", "gpt_pro"]

[config.hot_reload]
enabled = true
debounce_ms = 2000

[[config.agents]]
name = "gpt_pro"
canonical_name = "gpt_pro"
command = "code"
args = ["--model", "o3", "--config", "model_reasoning_effort=high"]
enabled = true
```

---

## Installing Templates

### Method 1: Manual Installation

**Steps:**

```
Create templates directory
mkdir -p ~/.code/templates
```

```
Copy template file
cp premium-quality.toml ~/.code/templates/

List installed templates
code --templates-list
```

---

## Method 2: Install from URL

```
code --template-install https://example.com/templates/premium-quality.toml
```

**Behavior:** 1. Download template file 2. Validate template structure 3. Save to `~/.code/templates/` 4. Confirm installation

---

## Method 3: Install from Git Repository

```
code --template-install github:theturtlecsz/code-templates/premium-quality.toml
```

**Behavior:** 1. Clone/fetch from GitHub 2. Extract template file 3. Validate and install

---

# Using Templates

## Apply Template Once

```
code --template premium-quality "task"
```

**Behavior:** Merges template config with `config.toml` for this session only

---

## Set Default Template

```
~/.code/config.toml

template = "premium-quality" # Apply on every session
```

**Behavior:** Template config merged on startup

---

## Template Precedence

**Precedence** (highest to lowest): 1. CLI flags (`--model o3`) 2. Environment variables (`CODEX_MODEL=o3`) 3. **Template config** (new tier) 4. Profile (`[profiles.premium]`) 5. `config.toml` 6. Defaults

**Example:**

```
~/.code/config.toml
model = "gpt-5"

~/.code/templates/premium.toml
[config]
```

```
model = "o3"

Usage:
code --template premium "task"
Effective model: "o3" (template > config.toml)

code --template premium --model gpt-4o "task"
Effective model: "gpt-4o" (CLI > template)
```

---

## Creating Custom Templates

### Step 1: Define Template Metadata

```
[template]
name = "My Custom Template"
version = "1.0.0"
description = "Custom configuration for my team"
author = "Team Lead"
created = "2025-11-17"
tags = ["team", "production"] # Optional
```

---

### Step 2: Define Configuration

```
[config]
Model configuration
model = "gpt-5"
model_provider = "openai"
approval_policy = "on-request"

Quality gates
[config.quality_gates]
plan = ["gemini", "claude", "code"]
tasks = ["gemini"]
validate = ["gemini", "claude", "code"]
audit = ["gemini", "claude", "gpt_codex"]
unlock = ["gemini", "claude", "gpt_codex"]

Agents
[[config.agents]]
name = "gemini"
canonical_name = "gemini"
command = "gemini"
enabled = true

... more configuration
```

---

### Step 3: Test Template

```
Save template
cp my-template.toml ~/.code/templates/

Test application
code --template my-template --dry-run "test task"

Check effective configuration
```

```
code --template my-template --config-dump
```

---

## Step 4: Version and Document

**Version Incrementing:** - Major: Breaking changes (agent names changed, quality gates restructured) - Minor: New features (new agents, new quality gates) - Patch: Bug fixes, clarifications

**Documentation:**

```
[template]
name = "My Template"
version = "2.1.0" # Incremented version
changelog = """
2.1.0 (2025-11-17):
- Added gpt_pro agent for premium reasoning
- Increased audit quality gate to 4 agents

2.0.0 (2025-11-10):
- BREAKING: Renamed 'code' agent to 'gpt_pro'
- Added cost optimization profile

1.0.0 (2025-11-01):
- Initial release
"""


```

---

## Template Examples

### Cost-Optimized Template

```
~/.code/templates/cost-optimized.toml

[template]
name = "Cost Optimized"
version = "1.0.0"
description = "Minimize cost while maintaining quality"

[config]
model = "gpt-4o-mini"
model_reasoning_effort = "low"
approval_policy = "never"

[config.quality_gates]
plan = ["gemini", "claude"] # 2 cheap agents
tasks = ["gemini"]
validate = ["gemini", "claude"]
audit = ["gemini", "claude"]
unlock = ["gemini", "claude"]

[[config.agents]]
name = "gemini"
canonical_name = "gemini"
command = "gemini"
enabled = true

[[config.agents]]
```

```
name = "claude"
canonical_name = "claude"
command = "claude"
enabled = true
```

#### Usage:

```
code --template cost-optimized "task"
```

---

## CI/CD Template

```
~/.code/templates/ci-cd.toml

[template]
name = "CI/CD"
version = "1.0.0"
description = "Configuration optimized for CI/CD pipelines"

[config]
model = "gpt-4o"
approval_policy = "never"
sandbox_mode = "read-only"
disable_response_storage = false

[config.quality_gates]
plan = ["gemini", "claude"]
tasks = ["gemini"]
validate = ["gemini", "claude"]
audit = ["gemini", "claude"]
unlock = ["gemini", "claude"]

[config.hot_reload]
enabled = false # No hot-reload in CI

[config.history]
persistence = "none" # Don't persist history in CI
```

#### Usage (in CI):

```
code --template ci-cd "generate report"
```

---

## Team Standard Template

```
~/.code/templates/team-standard.toml

[template]
name = "Team Standard"
version = "1.2.0"
description = "Standard configuration for our team"
author = "Engineering Team"
organization = "ACME Corp"

[config]
model = "gpt-5"
model_provider = "openai"
approval_policy = "on-request"

Custom quality gates for our workflow
```

```

[config.quality_gates]
plan = ["gemini", "claude", "code"]
tasks = ["gemini"]
validate = ["gemini", "claude", "code"]
audit = ["gemini", "claude", "gpt_codex"]
unlock = ["gemini", "claude", "gpt_codex"]

Team-specific agents
[[config.agents]]
name = "team-security"
canonical_name = "security"
command = "claude"
instructions = """
Focus on ACME Corp security standards:
- OWASP Top 10 compliance
- PCI-DSS requirements for payment processing
- GDPR compliance for user data
"""
enabled = true

Use team security agent for audit
[config.quality_gates]
audit = ["security", "gemini", "gpt_codex"]

```

---

## Template Versioning

### Version Schema

**Format:** MAJOR.MINOR.PATCH

**Versioning Rules:** - **MAJOR:** Breaking changes (incompatible with previous versions) - **MINOR:** New features (backward compatible) - **PATCH:** Bug fixes, documentation updates

---

### Version Compatibility

#### Check Template Version:

```
code --template-info premium-quality
```

#### Output:

```

Template: Premium Quality
Version: 2.1.0
Compatible with: codex-rs >= 0.5.0
Author: Your Name
Description: Premium quality configuration with maximum reasoning

```

#### Changelog:

```

2.1.0 (2025-11-17):
- Added gpt_pro agent
- Increased audit quality gate to 4 agents
2.0.0 (2025-11-10):
- BREAKING: Renamed agents
1.0.0 (2025-11-01):
- Initial release

```

---

## Automatic Template Updates

### Enable Auto-Update:

```
~/.code/config.toml

template_auto_update = true # Check for updates on startup
template_update_channel = "stable" # stable, beta, nightly
```

### Manual Update:

```
code --template-update premium-quality
```

### Output:

```
Checking for updates...
New version available: 2.2.0 (current: 2.1.0)
```

```
Changelog:
2.2.0 (2025-11-20):
- Added performance optimizations
- Fixed quality gate configuration bug
```

```
Update? [Y/n]: y
```

```
Downloading... ✓
Installing... ✓
Template updated successfully.
```

---

## Template Repositories

### Official Template Repository

**URL:** <https://github.com/theturtlecsz/code-templates>

**Templates:** - premium-quality.toml - Maximum quality, high cost - cost-optimized.toml - Minimum cost, acceptable quality - ci-cd.toml - CI/CD pipelines - team-standard.toml - Team collaboration - solo-developer.toml - Individual productivity

---

### Install from Repository

```
Install from official repository
code --template-install official:premium-quality

Install from GitHub
code --template-install github:theturtlecsz/code-templates/premium-
quality.toml

Install from URL
code --template-install
https://raw.githubusercontent.com/theturtlecsz/code-templates/main/template.toml
```

---

### Create Your Own Repository

### **Structure:**

```
my-templates/
├── README.md
└── templates.json # Template index
└── templates/
 ├── premium.toml
 ├── cost.toml
 └── ci.toml
```

### **templates.json:**

```
{
 "repository": "my-templates",
 "version": "1.0.0",
 "templates": [
 {
 "name": "premium",
 "file": "templates/premium.toml",
 "description": "Premium quality template",
 "version": "1.0.0"
 },
 {
 "name": "cost",
 "file": "templates/cost.toml",
 "description": "Cost-optimized template",
 "version": "1.0.0"
 }
]
}
```

---

## **Debugging Templates**

### **Validate Template**

```
code --template-validate ~/.code/templates/my-template.toml
```

#### **Output:**

```
Validating template...
```

```
Template Metadata:
```

- ✓ name: "My Template"
- ✓ version: "1.0.0"
- ✓ description: Present

```
Configuration:
```

- ✓ model: "gpt-5" (valid)
- ✓ quality\_gates.plan: 3 agents (valid)
- ✓ agents: 3 configured (all valid)

```
Template is valid ✓
```

---

### **Dry-Run Template**

```
code --template my-template --dry-run "task"
```

## **Output:**

Dry-run mode: No actions will be executed

```
Effective configuration (with template "my-template"):
 model: o3 (from template)
 model_reasoning_effort: high (from template)
 quality_gates.plan: ["gemini", "claude", "code", "gpt_pro"] (from
 template)
```

Would execute: [task description]

---

## **Compare Templates**

```
code --template-diff premium-quality cost-optimized
```

## **Output:**

```
Comparing templates:
 premium-quality v2.1.0
 cost-optimized v1.0.0
```

Differences:

```
model:
 - premium-quality: "o3"
 + cost-optimized: "gpt-4o-mini"

model_reasoning_effort:
 - premium-quality: "high"
 + cost-optimized: "low"

quality_gates.plan:
 - premium-quality: ["gemini", "claude", "code", "gpt_pro"] (4
 agents)
 + cost-optimized: ["gemini", "claude"] (2 agents)
```

---

## **Best Practices**

### **1. Version Templates Semantically**

**Good:**

```
[template]
version = "2.1.0"
changelog = ""
2.1.0: Added gpt_pro agent
2.0.0: BREAKING: Renamed agents
1.0.0: Initial release
"""
```

---

### **2. Document Template Usage**

**Include README:**

```
Premium Quality Template

Purpose: Maximum reasoning quality for critical projects

Cost: ~$4.50 per full pipeline (66% increase over default)

When to Use:
- Critical production features
- Security-sensitive code
- Architecture decisions

When NOT to Use:
- Simple formatting tasks
- Routine bug fixes
- Development iteration
```

---

### 3. Test Templates Before Distribution

```
Validate template
code --template-validate my-template.toml

Dry-run test
code --template my-template --dry-run "test task"

Full test with real task
code --template my-template "simple test task"
```

---

### 4. Use Templates for Team Consistency

#### Team workflow:

```
Install team template
code --template-install github:myorg/code-templates/team-
standard.toml

Set as default
Add to ~/.code/config.toml:
template = "team-standard"
```

---

## Summary

**Template Customization** provides:

- Pre-configured settings for common workflows
- Team-specific default configurations
- Environment-specific profiles (dev, staging, production)
- Organization-wide standards

**Features:**

- Template installation (URL, GitHub, local)
- Template versioning (semantic versioning)
- Template repositories (official + custom)
- Automatic updates
- Template validation and dry-run

#### Usage:

```
Install template
code --template-install official:premium-quality

Use template once
```

```
code --template premium-quality "task"

Set as default
Add to config.toml:
template = "premium-quality"
```

**Best Practices:** - Version templates semantically - Document template usage (purpose, cost, when to use) - Test templates before distribution - Use templates for team consistency

**Next:** [Theme System](#)

---

## Theme System

TUI themes, color customization, and accessibility options.

---

### Overview

The **theme system** provides visual customization for the TUI (Terminal User Interface).

**Features:** - 14 built-in themes (7 light + 7 dark) - Custom color overrides - Syntax highlighting themes - Accessibility options - Hot-reload support

**Configuration:** [tui.theme] section in config.toml

---

## Theme Configuration

### Basic Configuration

```
~/.code/config.toml

[tui.theme]
name = "dark-carbon-night" # Built-in theme
```

---

### Built-in Themes

**Light Themes:** 1. light-photon (default light) 2. light-prism-rainbow  
3. light-vivid-triad 4. light-porcelain 5. light-sandbar 6. light-glacier

**Dark Themes:** 7. dark-carbon-night (default dark) 8. dark-shinobi-dusk 9. dark-oled-black-pro 10. dark-amber-terminal 11. dark-aurora-flux 12. dark-charcoal-rainbow 13. dark-zen-garden 14. dark-paper-light-pro

---

### Theme Selection

**Auto-Detection** (default):

```
Omit theme name to auto-detect based on terminal background
[tui.theme]
name not specified - auto-detect
```

**Behavior:** Probes terminal background, selects appropriate light/dark theme

---

### Manual Selection:

```
[tui.theme]
name = "dark-carbon-night" # Explicitly select theme
```

---

## Theme Previews

**Light Photon** (default light): - Background: Light gray (#F5F5F5) - Foreground: Dark gray (#333333) - Primary: Blue (#007ACC) - Secondary: Purple (#8B008B) - Success: Green (#28A745) - Warning: Orange (#FFA500) - Error: Red (#DC3545)

---

**Dark Carbon Night** (default dark): - Background: Very dark gray (#1E1E1E) - Foreground: Light gray (#D4D4D4) - Primary: Cyan (#00D4FF) - Secondary: Magenta (#FF00D4) - Success: Green (#4EC9B0) - Warning: Yellow (#DCDCAA) - Error: Red (#F48771)

---

**Dark OLED Black Pro** (true black for OLED displays): - Background: Pure black (#000000) - Foreground: White (#FFFFFF) - Primary: Bright cyan (#00FFFF) - Secondary: Bright magenta (#FF00FF) - Success: Bright green (#00FF00) - Warning: Bright yellow (#FFFF00) - Error: Bright red (#FF0000)

---

## Custom Color Overrides

### Override Individual Colors

```
[tui.theme]
name = "dark-carbon-night"

[tui.theme.colors]
primary = "#00D4FF" # Override primary color
background = "#1A1A1A" # Slightly darker background
border_focused = "#FFD700" # Gold border for focused elements
```

---

### Available Color Fields

**Primary Colors:** - primary - Primary accent color - secondary - Secondary accent color - background - Background color - foreground - Foreground (text) color

**UI Elements:** - border - Default border color - border\_focused - Focused element border - selection - Selected item background - cursor - Cursor color

**Status Colors:** - success - Success messages (green) - warning - Warning messages (yellow/orange) - error - Error messages (red) - info - Info messages (blue)

**Text Colors:** - text - Primary text color - text\_dim - Dimmed/secondary text - text\_bright - Bright/emphasized text

**Syntax Colors:** - keyword - Syntax keywords (if, for, function) - string - String literals - comment - Code comments - function - Function names

**Animation Colors:** - spinner - Loading spinner color - progress - Progress bar color

---

## Complete Custom Theme

```
[tui.theme]
name = "custom" # Use 'custom' to define fully custom theme
label = "My Custom Theme" # Display name
is_dark = true # Dark theme hint

[tui.theme.colors]
Primary colors
primary = "#0080FF"
secondary = "#FF0080"
background = "#1C1C1C"
foreground = "#E0E0E0"

UI elements
border = "#444444"
border_focused = "#0080FF"
selection = "#2A2A2A"
cursor = "#FFFFFF"

Status colors
success = "#00FF00"
warning = "#FFAA00"
error = "#FF0000"
info = "#00AAFF"

Text colors
text = "#E0E0E0"
text_dim = "#808080"
text_bright = "#FFFFFF"

Syntax colors
keyword = "#569CD6"
string = "#CE9178"
comment = "#6A9955"
function = "#DCDCAA"

Animation colors
spinner = "#0080FF"
progress = "#00FF00"
```

---

## Syntax Highlighting

## Highlight Configuration

```
[tui.highlight]
theme = "auto" # Auto-select based on UI theme
```

**Options:** - "auto" - Auto-detect (default) - "<theme-name>" - Specific syntect theme

**Available Syntect Themes:** - base16-ocean.dark - base16-ocean.light - InspiredGitHub - Solarized (dark) - Solarized (light) - Monokai

---

## Custom Syntax Theme

```
[tui.highlight]
theme = "Monokai" # Use Monokai theme for code blocks
```

---

## Terminal Background Detection

### Auto-Detection Process

1. Query terminal environment variables
    - \$TERM (terminal type)
    - \$TERM\_PROGRAM (terminal program)
    - \$COLORFGBG (foreground/background color hint)
  2. Probe terminal background (if supported)
    - Send OSC 11 query
    - Parse RGB response
    - Determine if dark/light
  3. Select appropriate theme
    - Dark background → dark-carbon-night
    - Light background → light-photon
  4. Cache result
    - Store in ~/.code/config.toml
    - Skip probe on subsequent starts
- 

## Cached Terminal Background

### Auto-Cached:

```
[tui]
[tui.cached_terminal_background]
is_dark = true
term = "xterm-256color"
term_program = "iTerm.app"
source = "osc11-probe"
rgb = "#1E1E1E"
```

**Benefit:** Faster startup (no terminal probe)

---

## Force Re-Detection

```
Delete cached background
code --clear-terminal-cache

Or manually edit config.toml and remove
[tui.cached_terminal_background]
```

---

## Accessibility Options

### High Contrast Mode

**Enable via Custom Theme:**

```
[tui.theme]
name = "custom"
label = "High Contrast"

[tui.theme.colors]
background = "#000000" # Pure black
foreground = "#FFFFFF" # Pure white
primary = "#00FFFF" # Bright cyan
error = "#FF0000" # Bright red
success = "#00FF00" # Bright green
border_focused = "#FFFF00" # Bright yellow
```

---

### Large Text (Terminal Setting)

**Increase Terminal Font Size:**

Terminal Settings → Font Size → 16pt (or larger)

**Note:** TUI adapts to terminal font size automatically

---

### Color Blindness Support

**Protanopia/Deuteranopia** (red-green color blindness):

```
[tui.theme]
name = "custom"
label = "Color Blind Friendly"

[tui.theme.colors]
Avoid red/green distinction
success = "#0080FF" # Blue instead of green
error = "#FF8800" # Orange instead of red
warning = "#FFFF00" # Yellow (safe)
info = "#00FFFF" # Cyan (safe)
```

---

**Tritanopia** (blue-yellow color blindness):

```
[tui.theme.colors]
Avoid blue/yellow distinction
primary = "#FF00FF" # Magenta instead of blue
warning = "#FF8800" # Orange instead of yellow
```

---

# Theme Customization Examples

## Solarized Dark

```
[tui.theme]
name = "custom"
label = "Solarized Dark"
is_dark = true

[tui.theme.colors]
background = "#002B36" # base03
foreground = "#839496" # base0
primary = "#268BD2" # blue
secondary = "#D33682" # magenta
success = "#859900" # green
warning = "#B58900" # yellow
error = "#DC322F" # red
info = "#2AA198" # cyan
```

---

## Gruvbox Dark

```
[tui.theme]
name = "custom"
label = "Gruvbox Dark"
is_dark = true

[tui.theme.colors]
background = "#282828" # dark0
foreground = "#EBDBB2" # light1
primary = "#83A598" # blue
secondary = "#D3869B" # purple
success = "#B8BB26" # green
warning = "#FABD2F" # yellow
error = "#FB4934" # red
info = "#8EC07C" # aqua
```

---

## Dracula

```
[tui.theme]
name = "custom"
label = "Dracula"
is_dark = true

[tui.theme.colors]
background = "#282A36" # Background
foreground = "#F8F8F2" # Foreground
primary = "#BD93F9" # Purple
secondary = "#FF79C6" # Pink
success = "#50FA7B" # Green
warning = "#F1FA8C" # Yellow
error = "#FF5555" # Red
info = "#8BE9FD" # Cyan
```

---

## Debugging Themes

## Test Theme

```
Test theme without saving to config
code --theme dark-carbon-night
```

---

## Preview All Themes

```
code --themes-preview
```

**Output:** Opens TUI showing all themes side-by-side

---

## Dump Current Theme

```
code --theme-dump
```

**Output:**

```
[tui.theme]
name = "dark-carbon-night"

[tui.theme.colors]
primary = "#00D4FF"
background = "#1E1E1E"
foreground = "#D4D4D4"
... all effective colors
```

---

## Validate Custom Theme

```
code --theme-validate ~/.code/config.toml
```

**Output:**

Validating theme...

```
Theme: custom (My Custom Theme)
✓ primary: #0080FF (valid hex)
✓ background: #1C1C1C (valid hex)
✓ foreground: #E0E0E0 (valid hex)
✓ All 24 color fields valid
```

Theme is valid ✓

---

## Hot-Reload Support

### Live Theme Changes

**Edit config.toml:**

```
[tui.theme]
name = "dark-carbon-night" # Change to different theme
```

**Save:** TUI reloads theme within 2 seconds (debounced)

**Notification:**

- ✓ Config reloaded successfully
  - Theme changed: light-photon → dark-carbon-night

## Live Color Tweaking

## Edit config.toml:

```
[tui.theme.colors]
primary = "#FF0080" # Change primary color
```

**Save:** Color updates instantly (hot-reload)

**Use Case:** Iterative theme customization

## Spinner Customization

## Built-in Spinners

**Default:** "diamond"

**Available Spinners** (from cli-spinners): - dots, dots2, dots3 (simple dots) - line, line2 (horizontal line) - pipe, simpleDots (classic spinners) - star, star2 (star animation) - flip, hamburger (quirky animations) - growVertical, growHorizontal (growth animations) - balloon, balloon2 (balloon animations) - noise, bounce (dynamic animations) - boxBounce, boxBounce2 (box animations) - triangle, arc (geometric shapes) - circle, circleQuarters, circleHalves (circle animations) - squish, toggle (squish/toggle animations) - layer, betaWave (wave animations) - fingerDance, fistBump (emoji animations) - soccerHeader, mindblown (emoji animations) - speaker, orangePulse (pulse animations) - bluePulse, orangeBluePulse (multi-color pulses) - timeTravel, aesthetic (special effects) - dqpb, weather (themed spinners) - christmas, grenade (themed spinners) - point, layer (pointer animations) - betaWave, shark (wave/shark animations)

## Spinner Configuration

```
[tui.spinner]
name = "dots" # Simple dots spinner
```

## Custom Spinner

```
[tui.spinner]
name = "my-spinner"

[tui.spinner.custom.my-spinner]
interval = 80 # Milliseconds between frames
frames = ["\"", "\\"", "\\"", "\\"", "\\"", "\\"", "\\"", "\\"", "\\"", "\\"", "\\""]
label = "My Custom Spinner" # Optional display name
```

## Stream Animation

## Stream Configuration

```
[tui.stream]
answer_header_immediate = false # Show header before first text
show_answer_ellipsis = true # Show "..." while waiting
commit_tick_ms = 50 # Animation speed (50ms default)
soft_commit_timeout_ms = 400 # Commit after 400ms idle
soft_commit_chars = 160 # Commit after 160 chars
relax_list_holdback = false # Allow list line commits
relax_code_holdback = false # Allow code block commits
responsive = false # Enable snappier preset
```

---

## Responsive Preset

**Enable:**

```
[tui.stream]
responsive = true # Enable snappier preset
```

**Effect:** Overrides to: - commit\_tick\_ms = 30 (faster animation) - soft\_commit\_timeout\_ms = 400 - soft\_commit\_chars = 160

**Use Case:** Users who prefer instant response over smooth animation

---

## Best Practices

### 1. Use Built-in Themes When Possible

**Reason:** Pre-tested, well-balanced, maintained

**Example:**

```
[tui.theme]
name = "dark-carbon-night" # Built-in theme
```

---

### 2. Override Colors Sparingly

**Good** (1-2 color overrides):

```
[tui.theme]
name = "dark-carbon-night"

[tui.theme.colors]
primary = "#00FFAA" # Just change primary accent
```

**Bad** (override everything):

```
[tui.theme.colors]
Defining all 24 colors - hard to maintain
primary = "..."
secondary = "..."
... 22 more fields
```

---

### 3. Test Themes in Different Scenarios

**Test Cases:** - Success messages (green) - Error messages (red) - Warning messages (yellow) - Info messages (blue) - Code syntax highlighting - Spinner animations - Border focus states

---

## 4. Consider Accessibility

**Contrast Ratio:** WCAG AA requires 4.5:1 for normal text

**Check Contrast:**

```
Use online tool: https://webaim.org/resources/contrastchecker/
Background: #1E1E1E
Foreground: #D4D4D4
Contrast: 12.63:1 ✓ (WCAG AAA)
```

---

## Summary

**Theme System** provides: - 14 built-in themes (7 light + 7 dark) - Custom color overrides (24 color fields) - Syntax highlighting themes - Spinner customization (50+ built-in, custom support) - Stream animation tuning - Hot-reload support (live theme changes) - Accessibility options (high contrast, color blind support)

**Configuration:**

```
[tui.theme]
name = "dark-carbon-night" # Built-in theme

[tui.theme.colors]
primary = "#00D4FF" # Optional color override

[tui.highlight]
theme = "auto" # Syntax highlighting

[tui.spinner]
name = "dots" # Spinner style

[tui.stream]
responsive = false # Animation speed
```

**Best Practices:** - Use built-in themes when possible - Override colors sparingly (1-2 overrides) - Test themes in different scenarios - Consider accessibility (contrast, color blindness)

**Next:** [Configuration Reference](#) (for complete schema)

---