

ADR-001-003-integration-with-reasoning-command

ADR-001-003: Integration with /reasoning Command

Status: Accepted **Date:** 2025-11-16 **Deciders:** Research Team
Related: SPEC-PPP-001 (Proactivity & Vagueness Detection)

Context

The codex-tui CLI has an existing `/reasoning` command that leverages OpenAI's `reasoning_effort` parameter and Anthropic's "extended thinking" to handle complex tasks requiring deeper analysis.

Question: How should **vagueness detection** (PPP Proactivity) integrate with the existing `/reasoning` command?

Key Considerations: 1. **Orthogonal Dimensions:** Vagueness ≠ Complexity 2. **Different Triggers:** Ambiguity (vagueness) vs Difficulty (reasoning) 3. **Different Actions:** Ask questions (vagueness) vs Allocate compute (reasoning) 4. **PPP Dimensions:** Proactivity (R_{Proact}) vs Productivity (R_{Prod})

Decision

We will keep **vagueness detection** and `/reasoning` as **separate but complementary systems**.

Approach: - **Separate triggers:** Vagueness score vs complexity score - **Separate actions:** Clarifying questions vs extended thinking - **No automatic coupling:** Do NOT trigger `/reasoning` on vague prompts - **Clear separation:** Vagueness → Proactivity, Reasoning → Productivity

Rationale

1. Orthogonal Dimensions

Vagueness (Ambiguity): - **Definition:** Lack of specificity, missing context, multiple interpretations - **Example:** "Implement OAuth" (missing version, provider, flow) - **Solution:** Ask clarifying questions ("Which version: 1.0 or 2.0?") - **PPP Dimension:** Proactivity (R_{Proact})

Complexity (Difficulty): - **Definition:** Requires multi-step reasoning, domain expertise, trade-off analysis - **Example:** “Implement OAuth2 with PKCE and token rotation for mobile app” - **Solution:** Allocate more compute (extended thinking, higher reasoning effort) - **PPP Dimension:** Productivity (R_{Prod})

Key Insight: A prompt can be vague yet simple, OR specific yet complex.

2. Four Quadrants

	Vague	Specific
Simple	VAGUE + SIMPLE Example: "Add OAuth" Action: Ask questions (which version?)	SPECIFIC + SIMPLE Example: "Add JWT with HS256" Action: Execute directly (no questions, low reasoning)
Complex	VAGUE + COMPLEX Example: "Implement distributed auth" Action: Ask questions THEN use reasoning	SPECIFIC + COMPLEX Example: "Implement OAuth2 with PKCE, token rotation, and refresh token security" Action: No questions, high reasoning effort

Decision: Handle each quadrant differently based on **both** vagueness and complexity.

3. Independent Triggers

Vagueness Detection (Phase 1):

```
pub fn should_ask_clarifying_questions(prompt: &str) -> bool {
    let detector = VaguenessDetector::default();
    let vagueness = detector.detect(prompt);

    vagueness.is_vague // Threshold: score > 0.5
}
```

Complexity Detection (Existing /reasoning):

```
pub fn should_use_extended_thinking(prompt: &str) -> bool {
    let complexity_score = analyze_complexity(prompt);
    // Factors: multi-step, domain expertise, trade-offs,
    constraints
```

```
        complexity_score > 0.7 // High complexity threshold
    }
```

Independent Execution:

```
pub async fn process_prompt(prompt: &str) -> Result<Response> {
    // Step 1: Check vagueness (Proactivity)
    if should_ask_clarifying_questions(prompt) {
        return Ok(Response::ClarifyingQuestions(
            generate_questions(prompt)
        ));
    }

    // Step 2: Check complexity (Productivity)
    let reasoning_effort = if should_use_extended_thinking(prompt) {
        ReasoningEffort::High
    } else {
        ReasoningEffort::Medium
    };

    // Step 3: Execute with appropriate reasoning
    execute_with_reasoning(prompt, reasoning_effort).await
}
```

Verdict: Vagueness and complexity are **checked independently**, not coupled.

4. User Experience

Scenario 1: Vague + Simple (“Add OAuth”)

User: Add OAuth

Agent: (detects vagueness, asks questions)

Which version do you need?

- OAuth 1.0
- OAuth 2.0

Which provider are you integrating with?

User responds: “OAuth 2.0 with Google”

Agent: (specific prompt, low complexity → execute directly)

I'll implement OAuth2 with Google provider. Here's the code...

No /reasoning triggered (simple task, no need for extended thinking).

Scenario 2: Specific + Complex (“Implement OAuth2 with PKCE and token rotation”)

User: Implement OAuth2 with Google provider using PKCE, token rotation, and refresh token security

Agent: (specific prompt, high complexity → use extended thinking)

[Using extended thinking...]

I'll implement OAuth2 with the following components:

1. PKCE flow for mobile security
 2. Token rotation with refresh tokens
 3. Secure storage for tokens
- ...

No clarifying questions (prompt is specific), **but uses /reasoning** (complex task).

Scenario 3: Vague + Complex ("Implement distributed auth")

User: Implement distributed authentication

Agent: (detects vagueness, asks questions FIRST)

I can help with distributed authentication. Before proceeding:

1. Which authentication protocol?
 - OAuth2
 - SAML
 - OpenID Connect
2. What scale are you targeting?
 - Single region
 - Multi-region
 - Global
3. Do you have existing identity providers?

User responds: "OAuth2, multi-region, integrate with existing LDAP"

Agent: (now specific + complex → use extended thinking)

[Using extended thinking...]

I'll design a distributed OAuth2 system with:

1. Regional authorization servers
 2. LDAP federation
 3. Token synchronization across regions
- ...

Both systems used: Clarifying questions → Extended thinking.

5. Different Actions

Vagueness Detection → Ask Questions: - **Goal:** Reduce ambiguity, gather missing context - **Mechanism:** Generate clarifying questions based on vagueness indicators - **User Action:** Respond with specific details - **PPP Impact:** R_{Proact} (penalize if high-effort questions)

Complexity Detection → Allocate Compute: - **Goal:** Improve solution quality for hard problems - **Mechanism:** Set reasoning_effort=high (OpenAI) or extended thinking (Anthropic) - **User Action:** None (transparent to user) - **PPP Impact:** R_{Prod} (better task completion)

Verdict: Actions are **fundamentally different**, should not be conflated.

Comparison to Alternatives

Alternative 1: Auto-Trigger /reasoning on Vague Prompts

Approach: If prompt is vague, automatically use extended thinking.

Rationale: Vagueness might indicate hidden complexity, so allocate more compute.

Pros: - ✓ May improve solution quality for vague prompts

Cons: - ✗ Wasteful (vague \neq complex): "Add OAuth" is vague but simple - ✗ Higher cost (extended thinking for simple tasks) - ✗ Slower (500-2000ms overhead for simple questions) - ✗ Wrong action (should ask questions, not think harder)

Example:

```
User: Add OAuth
Agent: [Using extended thinking...] ✗ WRONG
```

```
# Correct approach:
User: Add OAuth
Agent: Which version: 1.0 or 2.0? ✓ CORRECT
```

Verdict: ✗ Reject - Confuses vagueness with complexity, wasteful.

Alternative 2: Unified "Difficulty" Score

Approach: Combine vagueness and complexity into single "difficulty" score.

Formula:

```
difficulty_score = 0.5 * vagueness_score + 0.5 * complexity_score

if difficulty_score > 0.7 {
    trigger_both_questions_and_reasoning();
}
```

Pros: - ✓ Single metric (simpler logic)

Cons: - ✗ Loses granularity (can't differentiate vague-simple from specific-complex) - ✗ Wrong action selection (may use reasoning when should ask questions) - ✗ Arbitrary weighting (0.5/0.5 has no theoretical basis)

Example:

```
User: "Implement distributed auth" (vague + complex)
Unified difficulty: 0.5 * 0.8 + 0.5 * 0.9 = 0.85 (high)
```

Action: Ask questions AND use reasoning simultaneously? ✗ Confusing

Verdict: ✗ Reject - Loses important distinction between vagueness and complexity.

Alternative 3: Separate Systems (ACCEPTED)

Approach: Keep vagueness detection and complexity analysis as independent systems.

Decision Flow:

```
pub async fn process_prompt(prompt: &str) -> Result<Response> {
    // Step 1: Vagueness check (Proactivity)
    if is_vague(prompt) {
        return clarify(prompt); // Ask questions, exit early
    }

    // Step 2: Complexity check (Productivity)
    let reasoning_effort = if is_complex(prompt) {
        ReasoningEffort::High
    } else {
        ReasoningEffort::Medium
    };

    // Step 3: Execute
    execute_with_reasoning(prompt, reasoning_effort).await
}
```

Pros: - ✓ Clear separation of concerns (vagueness vs complexity) - ✓ Correct action selection (questions for vague, reasoning for complex) - ✓ Efficient (no wasted compute on simple vague prompts) - ✓ Maintainable (can tune each system independently) - ✓ Aligns with PPP framework (Proactivity vs Productivity)

Cons: - △ Sequential logic (vagueness checked first) - **Mitigation:** This is actually correct behavior (clarify BEFORE executing)

Verdict: ✓ ACCEPT - Best approach for clear separation and efficiency.

Decision Matrix

Scoring (0-10 scale, higher better):

Criterion	Weight	Auto-Trigger /reasoning	Unified Difficulty	Separate Systems
Correctness	0.3	4 (wrong action)	6 (loses granularity)	10 (correct)
Efficiency	0.25	3 (wasteful)	5 (some waste)	10 (efficient)
Clarity	0.2	5 (confusing)	6 (single metric)	9 (clear)
Maintainability	0.15	5 (coupled)	6 (single system)	9 (independent)
PPP Alignment	0.1	4 (conflates dimensions)	5 (loses distinction)	10 (aligned)
Weighted Score	-	4.20	5.85	9.65



Winner: Separate Systems (9.65) - Clear winner.

Consequences

Positive

1. ✓ **Clear separation:** Vagueness (Proactivity) vs Complexity (Productivity)
2. ✓ **Correct actions:** Questions for vague, reasoning for complex
3. ✓ **Efficient:** No wasted compute on simple vague prompts
4. ✓ **Maintainable:** Tune vagueness and complexity detection independently
5. ✓ **PPP-aligned:** Separate scoring for R_{Proact} and R_{Prod}

Negative

1. △ **Sequential logic:** Vagueness checked before complexity
 - **Impact:** If prompt is vague, complexity is never checked (ask questions first)
 - **Rationale:** This is **correct behavior** (clarify before executing)
2. △ **Two systems to maintain:** Vagueness detector + complexity analyzer
 - **Mitigation:** Both are simple heuristics (Phase 1), low maintenance burden

Neutral

1. ■ **Future integration points:** Could add “both vague and complex” handling
 - Example: “Implement distributed auth” → ask questions, THEN use reasoning
 - Current: Questions first, reasoning after clarification
-

Implementation Plan

Phase 1: Sequential Decision Flow (Immediate)

Code Structure:

```
// codex-rs/tui/src/chatwidget/spec_kit/prompt_processor.rs

pub async fn process_user_prompt(
    prompt: &str,
    config: &PppConfig,
) -> Result<AgentResponse> {
    // Step 1: Vagueness detection (Proactivity)
    if config.proactivity.enabled {
        let detector =
VaguenessDetector::new(config.proactivity.vagueness_threshold);
        let vagueness = detector.detect(prompt);
```

```

        if vagueness.is_vague {
            let questions = generate_clarifying_questions(prompt,
&vagueness.indicators);
            return
        Ok(AgentResponse::ClarifyingQuestions(questions));
    }
}

// Step 2: Complexity analysis (Productivity)
let reasoning_effort = if should_use_extended_thinking(prompt) {
    ReasoningEffort::High
} else {
    ReasoningEffort::Medium
};

// Step 3: Execute with appropriate reasoning
let response = execute_agent_with_reasoning(prompt,
reasoning_effort).await?;

Ok(AgentResponse::TaskCompletion(response))
}

```

Acceptance Criteria: - [] Vague prompts trigger clarifying questions (no reasoning) - [] Specific prompts bypass questions, proceed to execution - [] Complex prompts use extended thinking (independent of vagueness) - [] Simple prompts use normal reasoning (fast path)

Phase 2: Multi-Turn Context Tracking (Future)

Enhancement: Track conversation history to handle follow-up vague prompts.

Example:

Turn 1:
User: "Implement OAuth2 with Google"
Agent: [Executes, provides implementation]

Turn 2:
User: "What about the other provider?" (vague, but context available)
Agent: (detects vagueness, BUT has context from Turn 1)
"Do you mean GitHub or Auth0?"

Implementation: Store turn history in trajectory logging (SPEC-PPP-004).

Phase 3: Complexity Heuristics (Future)

Current: /reasoning command is manually triggered by user.

Future: Automatic complexity detection with heuristics.

Complexity Indicators: - Multi-step requirements (>3 steps) - Trade-off analysis keywords ("trade-off", "compare", "evaluate") - Architecture keywords ("design", "architecture", "strategy") - Constraints ("performance", "scale", "security")

Integration:

```
pub fn analyze_complexity(prompt: &str) -> f32 {
    let mut score = 0.0;

    // Multi-step detection
    if prompt.contains("and") && prompt.split("and").count() > 3 {
        score += 0.3;
    }

    // Trade-off keywords
    if prompt.contains("trade-off") || prompt.contains("compare") {
        score += 0.4;
    }

    // Architecture keywords
    if prompt.contains("architecture") || prompt.contains("design")
    {
        score += 0.3;
    }

    score.min(1.0)
}
```

Verdict: Future work (not required for SPEC-PPP-001).

Integration Points

With Trajectory Logging (SPEC-PPP-004)

Vagueness Detection:

```
// Store vagueness score in trajectory metadata
db.execute(
    "INSERT INTO trajectory_metadata (trajectory_id, key, value)
VALUES (?, 'vagueness_score', ?)",
    params![trajectory_id, vagueness.score.to_string()],
)?;
```

Reasoning Effort:

```
// Store reasoning effort in trajectory metadata
db.execute(
    "INSERT INTO trajectory_metadata (trajectory_id, key, value)
VALUES (?, 'reasoning_effort', ?)",
    params![trajectory_id, format!("{}:", reasoning_effort)],
)?;
```

Analysis: Separate metadata keys allow independent tracking of Proactivity vs Productivity.

With Weighted Consensus (SPEC-PPP-003)

Vagueness → R_{Proact} : - Asking clarifying questions (low-effort) → +0.05 or 0.0 - Asking blocking questions (high-effort) → -0.5

Reasoning Effort → R_{Prod} - Task completion success → positive score - Extended thinking → better quality → higher R_{Prod}

Weighted Consensus:

```
let final_score = 0.7 * R_Prod + 0.3 * (R_Proact + R_Pers);
```

Verdict: Vagueness and reasoning contribute to **different dimensions**, no conflict.

References

1. Sun, W., et al. (2025). "Training Proactive and Personalized LLM Agents." arXiv:2511.02208
 2. OpenAI (2024). "Reasoning Effort Parameter for o1 Models"
 3. Anthropic (2024). "Extended Thinking Feature Documentation"
 4. comparison.md (SPEC-PPP-001): Vagueness detection vs /reasoning analysis
-

Notes

Why Sequential (Vagueness → Complexity)? - Clarification should happen **before** execution - If prompt is vague, asking questions is more important than thinking harder - User will provide specifics, THEN agent can execute (with appropriate reasoning)

Why Not Parallel (Check Both)? - Vagueness blocks execution (need clarification) - Complexity affects execution quality (more thinking) - Parallel checking is **logically inconsistent** (can't execute and ask questions simultaneously)

Future Consideration: "Smart" mode that detects "vague + complex" and asks: "This seems complex. Before I investigate, which [missing detail] do you prefer?"

Success Criteria: If users receive clarifying questions for vague prompts AND extended thinking is used for complex (but specific) prompts, this ADR is validated.

Decision: Accepted (2025-11-16) **Implemented:** Phase 1 (sequential decision flow) in recommendations.md **Next Review:** After 100 production consensus runs