

# comparison

## SPEC-PPP-002: Comparative Analysis

Last Updated: 2025-11-16

### Tool Comparison: AI Coding Assistants

Feature	Cursor	GitHub Copilot	Continue.dev	Aid
<b>Config Format</b>	.cursorrules (MD) + .mcp	JSON (settings.json)	YAML (config.yaml)	CLI flags
<b>Preference Count</b>	~10 (implicit)	~15 (toggles)	~20 (model/context)	0 (managed by prompt)
<b>Interaction Style</b>	Via rules text	✗ None	Via system message	✗ Manual
<b>Output Format</b>	✗ None	✗ None	✗ None	✗ None
<b>Language Support</b>	✗ English only	localeOverride (UI)	✗ English only	✗ English only
<b>Format Enforcement</b>	Prompt only (~75%)	Prompt only (~70%)	Prompt only (~75%)	N/A
<b>Conflict Detection</b>	✗ None	✗ None	✗ None	N/A
<b>Version Control</b>	✓ Project rules	✓ Workspace	✓ YAML file	N/A
<b>Scope</b>	Project + User	Workspace + User	Global	Per-session
<b>Validation</b>	✗ None	✗ None	✗ None	N/A
<b>Best For</b>	Team standards	VS Code users	Multi-model	CLI power users
<b>PPP Coverage</b>	10% (2/20)	5% (1/20)	15% (3/20)	0% (0/20)

**Winner:** Proposed (PPP) - Only solution targeting full framework compliance

### Translation Service Comparison

Service	Type	Languages	Quality (BLEU)	Latency	
<b>LibreTranslate (Self-Hosted)</b>	OSS	100+	35-40	~200-500ms	FR
<b>LibreTranslate Cloud</b>	Hosted	100+	35-40	~800-1200ms	\$5/mo
<b>DeepL API</b>	Commercial	32	45-50	~300-600ms	\$4.50/char
<b>LLM-Native (Claude/GPT)</b>	Inline	All	40-48 (varies)	~1-3s	~\$100/mo
<b>Google Translate API</b>	Commercial	130+	38-42	~400-700ms	\$20/mo

#### Recommendations by Use Case: - Self-hosted deployment:

LibreTranslate (self-hosted) - Privacy, no API costs -

**Prototype/testing:** LLM-native - Simplest, no external service -

**Production (high volume):** DeepL API - Best quality, cost-effective -

**Production (self-hosted):** LibreTranslate (self-hosted) - Privacy, control

## Format Enforcement Strategy Comparison

Approach	Compliance %	Quality Impact	Latency	Complexity	Cost
<b>Prompt Injection Only</b>	70-85%	✓ None	✓ 0ms	✓ LOW	✓ \$0
<b>Post-Processing</b>	100%	△ -10-20%	✓ <1ms	△ MEDIUM	✓ \$0
<b>Validation + Retry (1x)</b>	90-95%	✓ -0-5%	△ +50-100%	△ MEDIUM	△ +token cost
<b>Validation + Retry (2x)</b>	95-98%	✓ -0-5%	✗ +100-200%	△ MEDIUM	✗ +2x tokens
<b>Hybrid (Prompt + Post)</b>	100%	△ -5-10%	✓ <1ms	✗ HIGH	✓ \$0

**Recommended:** Validation + Retry (1x) for production - Balances compliance (90-95%) with quality (<5% degradation) - Acceptable latency increase for critical preferences - Fall back to post-processing if 2 retries fail

## Rust Crate Evaluation

### Configuration & Validation

Crate	Version	Maturity	Purpose	Pros	Cons
<b>serde</b>	1.0	✓ Very Mature	Serialization	Industry standard, compile-time	None
<b>toml</b>	0.8	✓ Very Mature	TOML parsing	Human-readable, well-supported	Limited nesting
<b>validator</b>	0.18	✓ Mature	Validation macros	Declarative, less boilerplate	Limited cross-field
<b>serde_json</b>	1.0	✓ Very Mature	JSON parsing	Fast, required for JSON validation	None

## Text Processing & Translation

Crate	Version	Maturity	Purpose	Pros
<b>regex</b>	1.10	✓ Very Mature	Pattern matching	Fast, cached compilation
<b>libretranslate-</b> <b>rs</b>	0.1	△ Young	LibreTranslate API	Simple API
<b>reqwest</b>	0.12	✓ Very Mature	HTTP client	Flexible, async

**Decision:** Use reqwest directly for both LibreTranslate and DeepL to avoid unmaintained dependencies

## PPP Preference Implementation Complexity

Preference	Category	Complexity	Strategy	Estimated Effort
<b>1. no_preference</b>	Baseline	✓ TRIVIAL	Default	1 hour
<b>2. concise_question</b>	Interaction	✓ LOW	Prompt injection	2 hours
<b>3. detail_question</b>	Interaction	✓ LOW	Prompt injection	2 hours
<b>4. answer_more</b>	Interaction	✓ LOW	Prompt injection	2 hours
<b>5. only_begin</b>	Interaction	△ MEDIUM	State tracking	4 hours
<b>6. no_ask</b>	Interaction	✓ LOW	Prompt injection	2 hours
<b>7. do_selection</b>	Interaction	△ MEDIUM	Prompt + validation	4 hours

<b>8. professional</b>	Interaction	✓ LOW	Prompt injection	2 hours
<b>9. amateur</b>	Interaction	✓ LOW	Prompt injection	2 hours
<b>10. ask_many</b>	Interaction	△ MEDIUM	State tracking	4 hours
<b>11. one_question</b>	Interaction	△ MEDIUM	State tracking	4 hours
<b>12. first_try</b>	Interaction	△ MEDIUM	State tracking	4 hours
<b>13. lang_ita</b>	Language	✗ HIGH	Translation service	8 hours
<b>14. lang_multi</b>	Language	✗ HIGH	Translation service	12 hours
<b>15. capital</b>	Format	✓ LOW	Post-processing	2 hours
<b>16. commas</b>	Format	✓ LOW	Post-processing	2 hours
<b>17. json</b>	Format	△ MEDIUM	Validation + retry	4 hours
<b>18. joke</b>	Content	△ MEDIUM	Prompt + detection	4 hours
<b>19. snippet</b>	Content	△ MEDIUM	Validation	4 hours
<b>20. length</b>	Format	✓ LOW	Post-processing	2 hours

**Total Effort:** ~68 hours (~2 weeks for 1 engineer)

**Phase 1 Target** (12 preferences, LOW-MEDIUM only): ~34 hours (~1 week)

## Configuration Schema Comparison

Approach	Format	Validation	Complexity	Flexibility
<b>TOML (Proposed)</b>	Structured	Pre-parse	△ MEDIUM	✓ HIGH
<b>JSON</b>	Structured	Pre-parse	△ MEDIUM	✓ HIGH
<b>YAML</b>	Structured	Pre-parse	✗ HIGH	✓ VERY HIGH
<b>Markdown (Cursor-style)</b>	Unstructured	✗ None	✓ LOW	△ LOW

**Decision:** **TOML** - Best balance of human-readability, validation, and Rust ecosystem support

# Mult-Agent Framework Scoring Comparison

Framework	Consensus Method	Interaction Scoring	Technical Scoring	Weighting
<b>theturtlecsz (Current)</b>	Binary (ok/degraded)	✗ None	✓ Completeness	N/A
<b>CrewAI</b>	Voting	✗ None	✓ Output quality	Equal
<b>AutoGen</b>	First-valid	✗ None	✓ Success/fail	N/A
<b>LangGraph</b>	Custom (user-defined)	△ Possible	✓ Custom	Custom
<b>Proposed (PPP)</b>	Weighted	✓ $R_{Proact} + R_{Pers}$	✓ Completeness	70/30 (tunable)

**Gap:** No existing multi-agent framework implements interaction-quality-based consensus

**Opportunity:** Novel contribution to multi-agent research

## Cost Analysis

### Translation Costs (1000 requests/month, avg 200 chars/request)

Service	Monthly Cost	Per-Request	Setup Cost	Infrastructure
LibreTranslate (Self-Hosted)	~\$5 (VPS)	\$0	~4 hours	Docker container
LibreTranslate Cloud	\$5 (paid tier)	\$0.005	0 hours	None
DeepL API	\$0.98 (200K chars)	\$0.00098	0 hours	None
LLM-Native (Claude)	~\$0.20 (token cost)	\$0.0002	0 hours	None

**Recommendation:** - **Low volume** (<100/mo): LLM-native (\$0.20/mo)  
- **Medium volume** (100-1000/mo): DeepL API (\$0.98/mo) - **High volume** (>1000/mo): Self-hosted LibreTranslate (\$5/mo VPS)

## Validation Overhead Costs

Strategy	Token Overhead	Monthly Cost (1000 req)	Latency Overhead
Prompt Injection	+50-100 tokens	+\$0.05-0.10	0ms

Validation + Retry (1x, 10% retry rate)	+200 tokens (10%)	+\$0.04	+500ms (10%)
Validation + Retry (2x, 1% retry rate)	+400 tokens (1%)	+\$0.008	+1000ms (1%)

**Total Monthly Cost Estimate** (1000 requests, all preferences enabled): - Base agent cost: ~\$10-20 (existing) - +Translation (DeepL): +\$0.98 - +Validation overhead: +\$0.10 - **Total:** ~\$11-21/mo (+5-10% increase)

---

## Performance Benchmarks (Estimated)

Component	Operation	Latency	Throughput	Bottleneck
<b>TOML Parse</b>	Load config	<1ms	N/A	Disk I/O
<b>Validation</b>	Check preferences	<0.1ms	10K/sec	CPU (regex)
<b>Prompt Injection</b>	Append constraints	<0.5ms	2K/sec	String concat
<b>JSON Validation</b>	Parse output	1-5ms	500/sec	JSON pars
<b>Regex (no_commas)</b>	Check commas	<0.1ms	10K/sec	Pattern match
<b>Sentence Count</b>	Count sentences	0.5-1ms	1K/sec	Regex + split
<b>Translation (LibreTranslate)</b>	Translate 200 chars	200-500ms	2-5/sec	Network + Model
<b>Translation (DeepL)</b>	Translate 200 chars	300-600ms	2-3/sec	Network
<b>Translation (LLM-native)</b>	Translate 200 chars	1-3s	0.3-1/sec	LLM call

**Total Overhead** (worst case, all preferences enabled): - Validation: ~10ms - Translation (if needed): ~500ms - Retry (10% of requests): +50% latency - **Expected:** <10% overhead for 90% of requests, +50-100% for 10% requiring translation/retry

---

## Recommendations Summary

Decision	Recommended Option	Alternative	Rationale
<b>Config Format</b>	TOML	JSON	Human-readable, Rust-native
<b>Validation</b>	Validation +	Post-	Best quality/compliance

<b>Strategy</b>	Retry (1x)	processing	balance
<b>Translation (Production)</b>	DeepL API	LibreTranslate self-hosted	Best quality, affordable
<b>Translation (Self-Hosted)</b>	LibreTranslate	LLM-native	Privacy, control
<b>Translation (Low-Volume)</b>	LLM-native	DeepL	Simplest, cheapest
<b>Enforcement Layers</b>	3-layer (Prompt + Validation + Post)	2-layer	Maximum compliance
<b>Phase 1 Scope</b>	12/20 preferences	8/20	60% coverage, avoid HIGH complexity
<b>Conflict Detection</b>	Pre-parse validation	Runtime warning	Fail-fast UX
<b>Storage</b>	Extend config_types.rs	New crate	Minimizes dependencies

## Next Steps

1. **Validate recommendations** with project maintainers
2. **Prototype** 5 preferences (json, no\_commas, one\_question, concise, professional)
3. **Benchmark** validation overhead with real agent workloads
4. **User testing** to identify most valuable preferences
5. **Phase 1 implementation** (12 preferences, ~1 week effort)