

recommendations

SPEC-PPP-002: Implementation Recommendations

Target: 95%+ PPP Framework Compliance (Personalization Dimension) **Phased Rollout:** 3 phases over 4-6 months **Total Effort:** 68 hours (~2 weeks for 1 engineer)

✓ Phase 1: Foundation (12/20 Preferences, 60% Compliance)

Duration: 1 week **Effort:** 34 hours **Compliance Target:** 60% (12/20 preferences) **Risk Level:** LOW

Scope

Include (LOW-MEDIUM complexity only): 1. ✓ no_preference (baseline) 2. ✓ concise_question (prompt injection) 3. ✓ detail_question (prompt injection) 4. ✓ answer_more (prompt injection) 5. ✓ no_ask (prompt injection) 6. ✓ professional (prompt injection) 7. ✓ amateur (prompt injection) 8. ✓ capital (post-processing) 9. ✓ commas (post-processing) 10. ✓ json (validation + retry) 11. ✓ length (post-processing) 12. ✓ snippet (validation)

Defer to Phase 2: - do_selection, only_begin, ask_many, one_question, first_try (require trajectory tracking) - lang_ita, lang_multi (require translation service) - joke (complex content generation)

Implementation Tasks

Task 1: Configuration Schema (8 hours)

```
// File: codex-rs/core/src/config_types.rs

#[derive(Deserialize, Serialize, Debug, Clone, PartialEq)]
pub struct UserPreferences {
    // Phase 1: Core preferences
    pub interaction_mode: InteractionMode,           // PPP #1, #4, #6
    pub question_detail: QuestionDetail,             // PPP #2, #3
    pub expertise_level: ExpertiseLevel,            // PPP #8, #9
    pub format_constraints: FormatConstraints,      // PPP #15, #16,
#17, #20
    pub content_requirements: ContentRequirements, // PPP #19
}

impl UserPreferences {
    pub fn validate(&self) -> Result<(), Vec<PreferenceConflict>>;
    pub fn apply_to_prompt(&self, base: &str) -> String;
```

```

    pub fn validate_output(&self, output: &str) -> ValidationResult;
}

```

Task 2: Prompt Injection Layer (4 hours)

```

// File: codex-rs/tui/src/chatwidget/spec_kit/consensus.rs

// In artifact processing (line ~220):
let personalized_prompt = if let Some(prefs) =
&config.user_preferences {
    prefs.apply_to_prompt(&base_prompt)
} else {
    base_prompt
};

```

Task 3: Output Validation Module (8 hours)

```

// File: codex-rs/tui/src/output_formatter.rs (NEW)

pub struct OutputFormatter {
    preferences: UserPreferences,
}

impl OutputFormatter {
    pub fn format_and_validate(&self, output: &str) -> FormattedOutput {
        // 1. Validate against preferences
        // 2. Apply transformations if needed
        // 3. Return formatted output + validation result
    }
}

```

Task 4: Post-Processing Transformations (6 hours)

```

impl OutputFormatter {
    fn enforce_no_commas(&self, text: &str) -> String {
        text.replace(',', " and").replace(" , ", " and ")
    }

    fn enforce_all_caps(&self, text: &str) -> String {
        text.to_uppercase()
    }

    fn enforce_sentence_count(&self, text: &str, target: usize) -> String {
        let sentences = split_sentences(text);
        if sentences.len() == target {
            text.to_string()
        } else if sentences.len() > target {
            sentences[..target].join(". ") + "."
        } else {
            // Pad with filler sentences (e.g., "More details available upon request.")
            pad_sentences(text, target)
        }
    }
}

```

Task 5: Validation + Retry Logic (6 hours)

```
// In consensus flow:
```

```

let mut attempts = 0;
let max_retries = 1;

loop {
    let output = agent.call(&prompt).await?;
    let validation = preferences.validate_output(&output);

    if validation.valid || attempts >= max_retries {
        break output;
    }

    // Strengthen constraints for retry
    prompt = format!(
        "{}\n\n⚠ CRITICAL: Previous response violated:\n{}",
        original_prompt,
        validation.violations.iter()
            .map(|v| format!("- {}", v.description))
            .collect::<Vec<_>>()
            .join("\n")
    );
    attempts += 1;
}

```

Task 6: Unit Tests (2 hours)

```

#[cfg(test)]
mod tests {
    #[test]
    fn test_no_commas_enforcement() {
        let prefs = UserPreferences {
            format_constraints: FormatConstraints {
                no_commas: true,
                ..Default::default()
            },
            ..Default::default()
        };

        let input = "Hello, world! How are you, today?";
        let result = prefs.validate_output(input);
        assert!(!result.valid);
        assert_eq!(result.violations.len(), 1);
    }

    #[test]
    fn test_json_validation() {
        let prefs = UserPreferences {
            format_constraints: FormatConstraints {
                require_json: true,
                ..Default::default()
            },
            ..Default::default()
        };

        let valid = r#"{"message": "Hello world"}"#;
        let invalid = "Hello world";

        assert!(prefs.validate_output(valid).valid);
        assert!(!prefs.validate_output(invalid).valid);
    }
}

```

Success Metrics (Phase 1)

- ☒ 12/20 preferences implemented (60%)
 - ☒ Validation + retry achieves 90%+ compliance
 - ☒ <10% latency overhead for non-retry cases
 - ☒ All unit tests pass
 - ☒ No breaking changes to existing consensus flow
-

☒ Phase 2: Advanced Features (18/20 Preferences, 90% Compliance)

Duration: 2-3 weeks **Effort:** 24 hours **Compliance Target:** 90%
(18/20 preferences) **Risk Level:** MEDIUM

Scope

Add (require trajectory tracking): 13. ✓ do_selection (A/B/C format validation) 14. ✓ only_begin (state: has_asked_before) 15. ✓ ask_many (state: questions_in_current_turn) 16. ✓ one_question (state: questions_in_current_turn) 17. ✓ first_try (state: attempt_count) 18. ✓ lang_ita (translation service - LibreTranslate)

Defer to Phase 3: - lang_multi (complex: 5+ languages) - joke (complex: humor detection)

Implementation Tasks

Task 7: Trajectory Tracking Integration (8 hours) - Depends on SPEC-PPP-004 (trajectory schema) - Track turn number, questions asked, timing - Store in SQLite trajectories table

Task 8: State-Based Preference Enforcement (8 hours)

```
impl UserPreferences {
    pub fn check_state_constraints(&self, trajectory: &AgentTrajectory) -> Result<()> {
        // Check only_begin: already asked?
        if self.question_timing == QuestionTiming::OnlyBegin
            && trajectory.turns.len() > 1
            && !trajectory.questions_asked.is_empty() {
            return Err("only_begin violated: agent already asked
questions".into());
        }

        // Check one_question: multiple questions in current turn?
        if self.question_format == QuestionFormat::OnePerTurn {
            let current_turn_questions = trajectory.turns.last()
                .map(|t| count_questions(&t.response))
                .unwrap_or(0);

            if current_turn_questions > 1 {
                return Err("one_question violated: multiple
questions in turn".into());
            }
        }
    }
}
```

```

        Ok(())
    }
}

```

Task 9: Translation Service Integration (6 hours)

```
// File: codex-rs/tui/src/translation_service.rs (NEW)

pub enum TranslationBackend {
    LibreTranslate { url: String },
    DeepL { api_key: String },
    LlmNative,
}

pub async fn translate(
    text: &str,
    target_lang: &str,
    backend: &TranslationBackend,
) -> Result<String> {
    match backend {
        TranslationBackend::LibreTranslate { url } => {
            // HTTP POST to LibreTranslate API
        },
        TranslationBackend::DeepL { api_key } => {
            // HTTP POST to DeepL API
        },
        TranslationBackend::LlmNative => {
            // Meta-prompt: "Translate to {lang}: {text}"
        },
    }
}
```

Task 10: Self-Hosted LibreTranslate Setup (2 hours)

```
# Docker deployment
docker run -d -p 5000:5000 \
    libretranslate/libretranslate:latest

# Test
curl -X POST http://localhost:5000/translate \
    -H "Content-Type: application/json" \
    -d '{"q": "Hello", "source": "en", "target": "it"}'
```

Success Metrics (Phase 2)

- ☒ 18/20 preferences implemented (90%)
 - ☒ Trajectory-dependent preferences validated
 - ☒ Translation service operational (LibreTranslate)
 - ☒ <500ms translation latency (self-hosted)
-

❖ Phase 3: Full Compliance (20/20 Preferences, 100% Compliance)

Duration: 1-2 weeks **Effort:** 10 hours **Compliance Target:** 100% (20/20 preferences) **Risk Level:** MEDIUM-HIGH

Scope

Add (complex features): 19. ✓ lang_multi (5+ languages in single response) 20. ✓ joke (humor detection and injection)

Implementation Tasks

Task 11: Multi-Language Support (6 hours)

```
async fn enforce_multilingual(
    text: &str,
    languages: &[String],
    translator: &TranslationBackend,
) -> Result<String> {
    // Strategy: Split response into sections, translate each to
    // different language
    let sentences = split_sentences(text);
    let mut result = Vec::new();

    for (i, sentence) in sentences.iter().enumerate() {
        let target_lang = &languages[i % languages.len()];
        let translated = translate(sentence, target_lang,
translator).await?;
        result.push(format!("[{}]\n{}", target_lang.to_uppercase(),
translated));
    }

    Ok(result.join(" "))
}
```

Task 12: Humor Detection (4 hours)

```
fn contains_joke(text: &str) -> bool {
    // Heuristics:
    // - Contains "haha", "lol", emoji 😂
    // - Has punchline structure (setup → punchline)
    // - Contains wordplay patterns

    let joke_markers = ["haha", "lol", "😂", "🤣", "joke", "pun"];
    joke_markers.iter().any(|m| text.to_lowercase().contains(m))

    // TODO: Improve with LLM-based joke detection
}

fn inject_joke(text: &str) -> String {
    // Append a programming joke if none detected
    let jokes = [
        "Why do programmers prefer dark mode? Because light attracts
bugs!",
        "There are 10 types of people: those who understand binary,
and those who don't.",
        "A SQL query walks into a bar, walks up to two tables and
asks... 'Can I join you?'",
    ];

    if contains_joke(text) {
        text.to_string()
    } else {
        format!("{}\n\nHere's a joke: {}", text,
jokes[rand::random::<usize>() % jokes.len()])
    }
}
```

```
        }  
    }
```

Success Metrics (Phase 3)

- ☒ 20/20 preferences implemented (100%)
 - ☒ Multi-language responses validated
 - ☒ Joke injection/detection functional
 - ☒ Full PPP framework compliance achieved
-

Technical Recommendations

Recommendation 1: Use TOML for Configuration

Rationale: - Human-readable (better UX than JSON) - Native Rust support (serde + toml crates) - Supports complex nesting (enums, structs, arrays) - Widely used in Rust ecosystem (Cargo.toml)

Example:

```
[user_preferences]  
interaction_mode = "no_preference"  
question_detail = "concise"  
expertise_level = "professional"  
  
[user_preferences.format_constraints]  
no_commas = true  
require_json = false  
exact_sentence_count = 3
```

Recommendation 2: Implement 3-Layer Enforcement

Layer 1: Prompt Injection (70-85% compliance, 0ms overhead)

```
fn apply_to_prompt(&self, base: &str) -> String {  
    format!("{}\n\nUSER PREFERENCES:\n{}", base,  
self.constraint_text())  
}
```

Layer 2: Validation + Retry (90-95% compliance, +20-30% latency)

```
if !prefs.validate_output(&output).valid {  
    retry_with_stronger_constraints();  
}
```

Layer 3: Post-Processing Fallback (100% compliance, may degrade quality)

```
if still_invalid {  
    force_transformation(&output); // e.g., strip commas, convert  
    to JSON  
}
```

Recommendation 3: Defer Translation to Phase 2

Rationale: - Translation adds significant complexity (service integration, latency) - Only 2/20 preferences require translation (lang_ita, lang_multi) - Can validate 90% of framework without translation - Phase 1 proves core concept, Phase 2 adds polish

Migration Path: 1. Phase 1: English-only, 12 preferences 2. Phase 2: Add LibreTranslate (self-hosted), 6 more preferences 3. Phase 3: Optimize with DeepL (optional premium), 2 final preferences

Recommendation 4: Prioritize Format Constraints

Reason: Format constraints (json, no_commas, capital, length) are: - Easy to implement (post-processing) - High user value (ensure output compatibility) - Deterministic (100% enforceable)

Implementation Order: 1. ✓ no_commas (simplest: regex replace) 2. ✓ capital (trivial: .to_uppercase()) 3. ✓ length (moderate: sentence counting) 4. ✓ json (moderate: validation + retry)

Recommendation 5: Validate Conflicts Early

Problem: Users may specify contradictory preferences - no_ask + do_selection (can't ask questions AND require A/B/C format) - all_caps + require_json (JSON requires lowercase keywords)

Solution: Pre-parse validation in config loading

```
impl UserPreferences {
    pub fn validate(&self) -> Result<(), Vec<PreferenceConflict>> {
        let mut conflicts = Vec::new();

        if self.interaction_mode == InteractionMode::NoQuestions
            && self.question_format != QuestionFormat::Freeform
        {
            conflicts.push(PreferenceConflict {
                description: "no_ask conflicts with
question_format".into(),
                severity: ConflictSeverity::Error,
            });
        }

        if conflicts.is_empty() { Ok(()) } else { Err(conflicts) }
    }
}
```

III Success Metrics & Validation

Quantitative Metrics

Metric	Phase 1 Target	Phase 2 Target	Phase 3 Target
Preference Coverage	12/20 (60%)	18/20 (90%)	20/20 (100%)
Validation Compliance	90%+	95%+	98%+

Latency Overhead	<10% (no retry)	<10% (no retry)	<10% (no retry)
Retry Rate	<15%	<10%	<5%
Translation Latency	N/A	<500ms (self-hosted)	<300ms (DeepL)
Unit Test Coverage	80%+	90%+	95%+

Qualitative Metrics

- User Acceptance:** 80%+ users find preferences useful (survey)
- Agent Compliance:** Agents respect preferences without quality degradation
- Error Messages:** Clear, actionable conflict warnings
- Documentation:** Comprehensive guide for config.toml setup

☒ Risk Mitigation

Risk 1: Agent Non-Compliance

Risk: Agents may ignore preference constraints despite prompt injection

Mitigation: - Use validation + retry (Layer 2) for critical preferences
- Fall back to post-processing (Layer 3) if 2 retries fail - Log violations to identify which preferences need stronger enforcement

Acceptance: Some preferences (e.g., joke) may have lower compliance (~80%) - acceptable trade-off

Risk 2: Translation Quality Degradation

Risk: LibreTranslate may produce poor translations for technical content

Mitigation: - Offer multiple backends (LibreTranslate, DeepL, LLM-native) - Default to LLM-native for low-volume use cases (better technical accuracy) - Allow users to configure translation backend in config.toml

Fallback: If translation fails, return English with warning

Risk 3: Performance Regression

Risk: Validation + translation adds latency to agent execution

Mitigation: - Measure baseline latency before implementation - Set performance budgets (<10% overhead for 90% of requests) - Optimize hot paths (cache regex compilation, reuse translators) - Make translation async (don't block agent execution)

Monitoring: Add telemetry to track P50/P95/P99 latencies

⌘ Implementation Checklist

Phase 1 (Week 1)

- Extend `config_types.rs` with `UserPreferences` struct
- Implement `validate()` for conflict detection
- Implement `apply_to_prompt()` for prompt injection
- Create `output_formatter.rs` module
- Implement post-processing transformations (no_commas, capital, length)
- Implement JSON validation + retry logic
- Write unit tests (80%+ coverage)
- Integration test with `/speckit.plan`
- Update `config.toml.example` with preferences section
- Document in CLAUDE.md

Phase 2 (Weeks 2-4)

- Integrate with trajectory tracking (SPEC-PPP-004)
- Implement state-based preferences (only_begin, one_question, etc.)
- Deploy LibreTranslate (self-hosted Docker)
- Implement translation service integration
- Add `lang_ita` support
- Write integration tests for trajectory-dependent preferences
- Benchmark translation latency
- Update documentation

Phase 3 (Weeks 5-6)

- Implement multi-language support (`lang_multi`)
 - Implement joke detection and injection
 - Final compliance validation (20/20 preferences)
 - Performance tuning
 - User acceptance testing
 - Production rollout (feature flag: `ppp.enabled = true`)
-

☷ References

1. **PPP Framework:** arXiv:2511.02208
 2. **Feasibility Analysis:** `docs/ppp-framework-feasibility-analysis.md`
 3. **SPEC-PPP-000:** Master coordination SPEC
 4. **Consensus System:** `codex-rs/tui/src/chatwidget/spec_kit/consensus.rs`
 5. **Config Types:** `codex-rs/core/src/config_types.rs`
 6. **LibreTranslate:** <https://github.com/LibreTranslate/LibreTranslate>
 7. **DeepL API:** <https://www.deepl.com/docs-api>
-

End of Recommendations ✓