

comparison

SPEC-PPP-003: Comparative Analysis

Last Updated: 2025-11-16

Multi-Agent Framework Comparison

Feature	CrewAI	AutoGen	LangGraph	PF (Prop)
Agent Selection	Task delegation	First-valid	Custom logic	Weight consent
Orchestration	Centralized (manager)	Decentralized (network)	Graph-based	Graph-based (existing)
Consensus Mechanism	None (single agent)	Implicit (discussion)	User-defined	70/30 weightage
Technical Scoring	✗ None	✗ None	△ User-defined	✓ Completeness + correctness
Interaction Scoring	✗ None	✗ None	✗ None	✓ R_Proficiency + R_Performance
Multi-Agent Voting	✗ No	✗ No	△ Custom	✓ Yes
Weight Configuration	N/A	N/A	N/A	✓ configurable
Stage-Specific Weights	N/A	N/A	△ Custom	✓ Yes (1, 2)
Human-in-Loop	△ Limited	✓ Yes	✓ Yes	✓ Yes (via config)
Parallelization	△ Sequential	✓ Network	✓ Graph	✓ Existence (MCP)
License	MIT	Apache 2.0	MIT	Project specific
Best For	Project workflows	Conversations	Custom routing	Coding
PPP Compliance	0%	5%	40%	100%

Winner: PPP (Proposed) - Only framework combining technical quality with interaction quality scoring.

Consensus Mechanism Comparison

Mechanism	Agreement	Speed	Quality	Robustness	PF
Majority Voting	>50%	✓ Fast	△ Medium	✓ Good	✗ 20'
Supermajority	>66%	△ Medium	✓ High	✓ Very good	✗ 20'
Unanimity	100%	✗ Slow	✓ Highest	△ Fragile	✗ 10'
Weighted Voting	Variable	✓ Fast	✓ High	✓ Good	△ 60'
Confidence-Weighted	Variable	✓ Fast	△ Medium	△ Medium	△ 50'
First-Valid	N/A	✓ Fastest	✗ Low	✗ Poor	✗ 10'
PPP Weighted	N/A (scoring)	✓ Fast	✓ Highest	✓ Excellent	✓ 10'

Details

1. Majority Voting - How it works: >50% of agents must agree on solution - **Strengths:** Simple, fast (single round), democratic - **Weaknesses:** Treats all agents equally (ignores expertise), binary (agree/disagree only) - **Research:** +13.2% improvement on reasoning tasks (arXiv:2502.19130)

Example (3 agents):

Agent 1: "Use OAuth2 with PKCE" (2 votes)
 Agent 2: "Use OAuth2 with PKCE" (2 votes)
 Agent 3: "Use basic auth" (1 vote)
 Winner: OAuth2 with PKCE (majority)

PPP Fit: ✗ 20% - Ignores quality differences (agent 3's solution might be higher quality but loses)

2. Supermajority (66%+) - **How it works:** >66% of agents must agree - **Strengths:** More robust than simple majority, filters out outliers - **Weaknesses:** May fail to converge (no clear winner), slower - **Research:** +2.8% improvement on knowledge tasks (arXiv:2502.19130)

Example (3 agents, need 2/3 = 67%):

Agent 1: "Use OAuth2 with PKCE" (1 vote = 33%)
 Agent 2: "Use basic auth" (1 vote = 33%)
 Agent 3: "Use JWT tokens" (1 vote = 33%)
 Winner: None (no supermajority, trigger discussion round)

PPP Fit: ✗ 20% - Same issue as majority (ignores quality), plus convergence problems

3. Unanimity (100%) - How it works: All agents must agree -
Strengths: Highest confidence, safety for critical tasks -
Weaknesses: Very slow (many discussion rounds), often fails to converge - **Research:** Used for critical tasks (safety, compliance)

Example (3 agents):

Round 1: Agent 1, 2, 3 propose different solutions → No unanimity
Round 2: Agent 1, 2 converge, Agent 3 still different → No unanimity
Round 3: Agent 3 convinced by Agent 1's argument → Unanimity achieved

PPP Fit: ✗ 10% - Too slow for coding (agents use different models, rarely 100% agree)

4. Weighted Voting (Domain Expertise) - How it works: Agents weighted by expertise, majority of *weighted* votes wins - **Strengths:** Prioritizes skilled agents, standard ML ensemble technique -
Weaknesses: Need to define expertise weights (which agent is "better"?) - **Research:** Commonly used in ML ensembles (inverse error weighting)

Example (3 agents with expertise weights):

Agent 1 (gemini-pro): "Use OAuth2" (weight: 0.5)
Agent 2 (claude-opus): "Use OAuth2" (weight: 0.3)
Agent 3 (gpt-4): "Use JWT" (weight: 0.2)

Weighted votes:
OAuth2: $0.5 + 0.3 = 0.8$
JWT: 0.2
Winner: OAuth2 ($0.8 > 0.2$)

PPP Fit: △ 60% - Good for technical quality weighting, but no interaction quality dimension

5. Confidence-Weighted Voting - How it works: Agents self-report confidence (0-1 scale), votes weighted by confidence - **Strengths:** Accounts for uncertainty, agents can express doubt - **Weaknesses:** Agents may overestimate confidence (calibration problem) - **Research:** Used in multi-agent debate systems (arXiv:2502.19130)

Example (3 agents with self-reported confidence):

Agent 1: "Use OAuth2" (confidence: 0.9)
Agent 2: "Use OAuth2" (confidence: 0.6)
Agent 3: "Use JWT" (confidence: 0.3)

Weighted votes:
OAuth2: $0.9 + 0.6 = 1.5$
JWT: 0.3
Winner: OAuth2 ($1.5 > 0.3$)

PPP Fit: △ 50% - Interesting, but confidence ≠ quality, and no interaction scoring

6. First-Valid Output - How it works: First agent to produce valid output wins (no voting) - **Strengths:** Fastest (terminates immediately), lowest cost - **Weaknesses:** Ignores quality (first \neq best), no ensemble benefit - **Used by:** CrewAI (task delegation model)

Example (3 agents, parallel execution):

Agent 1 (gemini-flash): Responds in 2s → "Use OAuth2 with PKCE"
Agent 2 (claude-haiku): Responds in 3s → "Use OAuth2 with Authorization Code"
Agent 3 (gpt-4): Responds in 5s → "Use OAuth2 with state parameter"

Winner: Agent 1 (first valid output, ignoring agents 2 & 3)

PPP Fit: ✗ 10% - Completely ignores quality and interaction preferences

7. PPP Weighted Consensus (Proposed) - How it works: Score each agent on technical quality (70%) + interaction quality (30%), select highest - **Strengths:** Balances correctness with UX, configurable weights, no voting rounds needed - **Weaknesses:** Novel (needs validation), requires trajectory logging (SPEC-PPP-004) -

Formula:

$$\text{score_i} = 0.7 \times \text{technical_i} + 0.3 \times \text{interaction_i}$$

where:

$$\begin{aligned}\text{technical_i} &= \text{completeness} + \text{correctness} \text{ (existing)} \\ \text{interaction_i} &= R_{\text{Proact}} + R_{\text{Pers}} \text{ (from trajectory)}\end{aligned}$$

Example (3 agents):

Agent 1 (gemini-flash):
Technical: 0.85 (good completeness, minor issues)
R_Proact: 0.05 (asked 2 low-effort questions)
R_Pers: 0.05 (no violations)
Interaction: 0.10
Final Score: $0.7 \times 0.85 + 0.3 \times 0.10 = 0.625$

Agent 2 (claude-opus):
Technical: 0.95 (excellent completeness)
R_Proact: -0.50 (asked 1 high-effort question)
R_Pers: 0.05 (no violations)
Interaction: -0.45
Final Score: $0.7 \times 0.95 + 0.3 \times (-0.45) = 0.530$

Agent 3 (gpt-4):
Technical: 0.80 (good, but less complete)
R_Proact: 0.05 (no questions)
R_Pers: -0.03 (1 major violation: didn't use JSON)
Interaction: 0.02
Final Score: $0.7 \times 0.80 + 0.3 \times 0.02 = 0.566$

Winner: Agent 1 (0.625) - Balanced technical quality + excellent interaction

Key Insight: Agent 2 has best technical score (0.95) but loses due to poor interaction (-0.45). PPP prefers Agent 1's balance.

PPP Fit: ✓ 100% - Designed specifically for PPP framework

Weight Selection Strategy Comparison

Strategy	Approach	Pros	Cons
Equal Weights	50/50 technical/interaction	Simple	Ignores relative importance
Grid Search	Try 0.0, 0.1, ..., 1.0	Exhaustive	Slow ($11^2 = 121$ trials for 2D)
Inverse Error	$w = 1/error$	Adaptive, proven (ML)	Needs validation dataset
Bayesian Optimization	Use previous trials to guide search	Sample-efficient	Complex implementation
Domain Expert	Manual selection (70/30)	Interpretable, fast	Subjective
User Configurable	User sets via config	Flexible, personalized	Requires user expertise
Stage-Specific	Different weights per stage	Adaptive to task criticality	More config complexity

Details

1. Equal Weights (50/50)

```
[ppp.weights]
technical = 0.5
interaction = 0.5
```

Rationale: Treat technical and interaction equally

Problem: Technical correctness more important than UX for coding tools (users prefer correct code with poor interaction over incorrect code with great interaction)

Verdict: ✗ Avoid - Doesn't reflect coding tool priorities

2. Grid Search

```
# Pseudocode
best_weights = None
best_score = -inf

for w_tech in [0.0, 0.1, 0.2, ..., 1.0]:
    w_interact = 1.0 - w_tech
    score = evaluate_on_validation_set(w_tech, w_interact)
    if score > best_score:
        best_score = score
        best_weights = (w_tech, w_interact)
```

```
    return best_weights
```

Pros: - Exhaustive (tries all combinations) - Guaranteed to find best in grid

Cons: - Expensive (11 trials for 1D, 121 trials for 2D if tuning per-stage) - Requires labeled validation set (which output is "best"?)

Verdict: Δ Use in Phase 3 for fine-tuning after initial deployment

3. Inverse Error Weighting

```
# Standard ML ensemble technique
technical_error = 1 - technical_score_avg # e.g., 0.15 (85% avg)
interaction_error = abs(interaction_score_avg) # e.g., 0.02 (assume
avg +0.02)

w_tech = (1 / technical_error) / ((1 / technical_error) + (1 /
interaction_error))
w_interact = (1 / interaction_error) / ((1 / technical_error) + (1 /
interaction_error))

# Example:
# technical_error = 0.15 → 1/0.15 = 6.67
# interaction_error = 0.02 → 1/0.02 = 50.0
# w_tech = 6.67 / (6.67 + 50.0) = 0.12
# w_interact = 50.0 / (6.67 + 50.0) = 0.88
```

Problem with naive application: Interaction scores are small (± 0.05), technical scores are large (0-1), so interaction gets over-weighted.

Solution: Normalize scores first:

```
technical_norm = (technical - 0.5) / 0.5 # Map [0,1] → [-1,1]
interaction_norm = interaction / 0.1 # Map [-0.5,0.05] →
[-5,0.5]

# Then apply inverse error
```

Verdict: Δ Phase 2 - Interesting but requires careful normalization

4. Domain Expert Selection (70/30)

```
[ppp.weights]
technical = 0.7 # Correctness is primary goal
interaction = 0.3 # UX is important but secondary
```

Rationale: - Coding tools: Correctness > UX (users tolerate questions if code is right) - ML ensemble literature: Stronger model gets 60-80% weight - Similar ratio to other multi-objective systems (e.g., Pareto optimization)

Validation (analogies): - **Google Search**: ~70% relevance, ~30% diversity (estimated from research) - **ML AutoML**: ~70% accuracy, ~30% interpretability (typical trade-off) - **Amazon**

Recommendations: ~70% purchase probability, ~30% diversity

Verdict: \checkmark Use for Phase 1 - Well-justified, interpretable, standard practice

5. User Configurable

```
# User can override defaults
[ppp.weights]
technical = 0.8    # User prefers correctness even more
interaction = 0.2

# Or emphasize UX
[ppp.weights]
technical = 0.6
interaction = 0.4
```

Use Cases: - **Prototyping:** User prefers fast iteration (low interaction weight, accept questions) - **Production:** User needs correct code (high technical weight) - **Beginner:** User needs good UX (balanced weights 60/40)

Verdict: ✓ Phase 2 - Empowers users, low implementation cost

6. Stage-Specific Weights

```
# Early stages: Exploration (lower technical weight)
[ppp.weights.plan]
technical = 0.6
interaction = 0.4

# Mid stages: Balanced
[ppp.weights.implement]
technical = 0.7
interaction = 0.3

# Late stages: Correctness critical (higher technical weight)
[ppp.weights.unlock]
technical = 0.8
interaction = 0.2
```

Rationale: - **Plan:** Ideas phase, interaction matters more (asking questions is OK) - **Implement:** Balance (need correct code, but UX still important) - **Unlock:** Final validation, correctness critical (no room for errors)

Verdict: ✓ Phase 2 - Adaptive, aligns with task criticality research

Interaction Quality Metric Comparison

Metric	Dimension	Granularity	PPP Alignment
CORE Score	Dialog quality	Turn-level	△ 30%
Communication Efficiency	Message count	Session-level	✓ 80%
Decision Synchronization	Agreement rate	Turn-level	△ 40%
Coordination Quality	Planning score	Session-level	△ 50%
Question Effort			

(PPP)	Low/Med/High	Turn-level	<input checked="" type="checkbox"/> 100%
Preference Violations (PPP)	Minor/Major/Critical	Turn-level	<input checked="" type="checkbox"/> 100%

Details

CORE Score:

$$\text{CORE} = \alpha \times \text{Entropy} + \beta \times (1 - \text{Repetition}) + \gamma \times \text{Similarity}$$

Pros: - Comprehensive dialog quality metric - Captures diversity, coherence, redundancy

Cons: - Designed for game theory scenarios (not coding) - Doesn't measure user impact (agent-to-agent focused)

PPP Alignment: 30% (interesting but not user-centric)

Communication Efficiency:

$$\text{Efficiency} = \text{Tasks Completed} / \text{Messages Sent}$$

Example: - Agent 1: Asks 3 questions, completes 1 task → Efficiency = $1/3 = 0.33$ - Agent 2: Asks 0 questions, completes 1 task → Efficiency = $1/0 = \infty$

Pros: - Simple, intuitive (fewer messages = better) - Similar to PPP's proactivity (fewer questions = bonus)

Cons: - Doesn't distinguish question types (low vs high effort)

PPP Alignment: 80% (close to R_{Proact} but less granular)

PPP Interaction Quality (Proposed):

$$\text{Interaction} = R_{\text{Proact}} + R_{\text{Pers}}$$

where:

```
R_Proact = f(question_effort)      # Penalizes high-effort
questions
R_Pers   = f(preference_violations) # Penalizes violations
```

Pros: - User-centric (measures impact on user) - Granular (distinguishes low/medium/high effort) - Actionable (agent can improve by reducing high-effort questions)

Cons: - Novel (no prior validation) - Requires trajectory logging infrastructure

PPP Alignment: 100% (designed for PPP framework)

Implementation Approach Comparison

Approach	Complexity	Flexibility	Performance	Recommendation
----------	------------	-------------	-------------	----------------

Refactor Existing	△ Medium	△ Medium	✓ Fast	✓ Phase 1
New Module	✓ Low	✓ High	✓ Fast	△ Phase 2 (if refactor too complex)
Separate Crate	✗ High	✓ Highest	△ Slower (IPC)	✗ Avoid (over-engineering)

Details

1. Refactor consensus.rs (Recommended)

Current (consensus.rs:681-958):

```
pub async fn run_spec_consensus(...) -> Result<ConsensusResult> {
    // ... existing logic
    // Select best artifact (currently: first with highest technical
    score)
    let best = artifacts.iter()
        .max_by_key(|a| calculate_technical_score(a))
        .unwrap();

    Ok(best)
}
```

Proposed (weighted):

```
pub async fn run_spec_consensus_weighted(
    artifacts: Vec<ConsensusArtifactData>,
    weights: (f32, f32), // (technical, interaction)
) -> Result<WeightedConsensus> {
    let (w_tech, w_interact) = weights;
    let db = open_consensus_db()?;

    let mut scores = Vec::new();
    for artifact in artifacts {
        // Technical score (existing)
        let technical = calculate_technical_score(&artifact)?;

        // Interaction score (new: from trajectory)
        let trajectory_id = get_trajectory_id(&db,
            &artifact.spec_id, &artifact.agent)?;
        let proact = calculate_r_proact(&db, trajectory_id)?;
        let pers = calculate_r_pers(&db, trajectory_id)?;
        let interaction = proact.r_proact + pers.r_pers;

        // Weighted combination
        let final_score = (w_tech * technical) + (w_interact *
            interaction);

        scores.push(AgentScore {
            agent_name: artifact.agent.clone(),
            technical_score: technical,
            interaction_score: interaction,
            final_score,
        });
    }
}
```

```

        scores.sort_by(|a, b|
b.final_score.partial_cmp(&a.final_score).unwrap()));

    Ok(WeightedConsensus {
        best_agent: scores[0].agent_name.clone(),
        confidence: scores[0].final_score,
        scores,
    })
}

```

Pros: - Extends existing function (minimal disruption) - Reuses existing infrastructure (consensus_db, scoring logic) - Backward compatible (can keep old function for comparison)

Cons: - Must refactor ~300 lines (medium effort) - Adds dependency on SPEC-PPP-004 (trajectory logging)

Verdict: ✓ Recommended for Phase 1

2. New Module (weighted_consensus.rs)

Structure:

```

codex-rs/tui/src/chatwidget/spec_kit/
└── consensus.rs (existing, unchanged)
    └── weighted_consensus.rs (new)

```

Implementation:

```

// weighted_consensus.rs
pub struct WeightedConsensusScorer {
    db: Arc<Connection>,
    weights: (f32, f32),
}

impl WeightedConsensusScorer {
    pub fn new(db_path: String, weights: (f32, f32)) -> Self { ... }

    pub fn score_agent(
        &self,
        artifact: &ConsensusArtifactData,
    ) -> Result<AgentScore> {
        // ... scoring logic
    }

    pub fn select_best(
        &self,
        artifacts: Vec<ConsensusArtifactData>,
    ) -> Result<WeightedConsensus> {
        // ... selection logic
    }
}

```

Pros: - Clean separation (doesn't touch existing code) - Easier to test (isolated) - Can swap implementations easily

Cons: - Duplicate logic (technical scoring copied from consensus.rs) - More files to maintain

Verdict: △ Use if refactoring consensus.rs proves too complex

Cost Analysis

Development Effort

Task	Complexity	Estimated Effort
Refactor consensus.rs	△ MEDIUM	6 hours
Calculate interaction scores	✓ LOW	2 hours (reuse SPEC-PPP-004)
Configuration (weights)	✓ LOW	2 hours
Integration tests	△ MEDIUM	4 hours
Documentation	✓ LOW	2 hours

Total: ~16 hours (~2 days)

Runtime Cost

Operation	Current	With Weighted Consensus	Overhead
Score Agents	~10ms (technical only)	~30ms (technical + interaction)	+20ms
Query Trajectories	N/A	~20ms (2 queries: R_Proact + R_Pers)	+20ms
Total Consensus	~10ms	~50ms	+40ms

Impact: +40ms per consensus run (negligible vs agent execution time ~10-30 seconds)

Verdict: Performance impact acceptable (<0.2% overhead)

Recommendations Summary

Decision	Recommended Option	Alternative	Rationale
Consensus Mechanism	PPP Weighted (70/30)	Weighted voting	Balances technical + interaction
Weight Selection	Domain expert (70/30)	User configurable (Phase 2)	Standard practice, interpretable
Weight Granularity	Global (same for all stages)	Stage-specific (Phase 2)	Simplicity first
Implementation	Refactor consensus.rs	New module	Reuses infrastructure
Scoring	Linear weighted	Custom	Standard ML

Formula	average	function	technique
Configuration	config.toml	Hardcoded	User flexibility
Fallback	Technical-only (if no trajectory)	Fail	Backward compatible

Next Steps

1. **Validate 70/30 weights** with sample agents (run retrospective analysis)
2. **Prototype** weighted consensus with existing consensus artifacts
3. **Benchmark** overhead (<50ms target)
4. **Phase 1 implementation:** Refactor consensus.rs with weighted scoring
5. **Phase 2 enhancement:** User-configurable weights, stage-specific tuning