

recommendations

SPEC-PPP-003: Implementation Recommendations

Last Updated: 2025-11-16 **Target:** Weighted consensus combining technical quality + interaction quality **Total Effort:** ~16 hours (2 days, 1 engineer)

Executive Summary

Implement PPP weighted consensus by refactoring existing consensus.rs to combine technical quality (70%) with interaction quality (30%) when selecting the best agent output. Use standard ML ensemble weighted averaging technique with configurable weights.

Key Decisions: 1. ✓ Refactor consensus.rs (NOT new module) - Extends existing infrastructure 2. ✓ 70/30 weights (technical/interaction) - Domain expert selection 3. ✓ Linear weighted average - Standard ML ensemble technique 4. ✓ User-configurable weights - Via config.toml (Phase 2) 5. ✓ Fallback to technical-only - If no trajectory available (backward compatible)

Phased Rollout

Phase 1: Core Weighted Consensus (8 hours, 1 day)

Goal: Basic weighted scoring with fixed 70/30 weights

Deliverables: - ✓ Refactor consensus.rs:681-958 to add weighted scoring - ✓ Integration with SPEC-PPP-004 trajectory logging - ✓ Unit tests (weighted formula, edge cases) - ✓ Backward compatibility (fallback to technical-only)

Capabilities: - Score agents on technical (70%) + interaction (30%) - Select best agent based on weighted score - Handle missing trajectories gracefully

NOT Included (defer to Phase 2): - User-configurable weights - Stage-specific weights - Weight optimization/tuning

Phase 2: Configurability & Tuning (6 hours, 0.75 days)

Goal: User control and stage-specific weights

Deliverables: - ✓ Configuration via config.toml - ✓ Stage-specific weights (plan: 60/40, unlock: 80/20) - ✓ Validation (weights must sum to 1.0) - ✓ Integration tests (different weight scenarios)

Capabilities: - Users can override default 70/30 weights - Different weights per spec-kit stage - Config validation with clear error messages

NOT Included (defer to Phase 3): - Dynamic weight adaptation - Weight optimization from historical data

Phase 3: Advanced Optimization (Optional, 12 hours, 1.5 days)

Goal: Data-driven weight tuning

Deliverables: - ✓ Historical regret tracking (selected vs best) - ✓ Grid search for optimal weights - ✓ Inverse error weighting (adaptive) - ✓ A/B testing framework

Capabilities: - Measure weight effectiveness - Automatic weight tuning based on outcomes - Per-user weight preferences

Detailed Task Breakdown

Task 1: Refactor consensus.rs for Weighted Scoring (4 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/consensus.rs:681-958

Current Implementation:

```
pub async fn run_spec_consensus(
    spec_id: &str,
    stage: &str,
    artifacts: Vec<ConsensusArtifactData>,
) -> Result<ConsensusResult> {
    // Current: Select based on technical score only
    let best = artifacts.iter()
        .max_by_key(|a| calculate_technical_score(a))
        .unwrap();

    Ok(ConsensusResult {
        selected_agent: best.agent.clone(),
        selected_content: best.content.clone(),
    })
}
```

Proposed Implementation:

```
use crate::chatwidget::spec_kit::ppp_scoring::{calculate_r_proact,
calculate_r_pers};
use crate::chatwidget::spec_kit::consensus_db::open_consensus_db;

pub struct AgentScore {
    pub agent_name: String,
```

```

    pub technical_score: f32,
    pub interaction_score: f32,
    pub final_score: f32,
    pub details: ScoreDetails,
}

pub struct ScoreDetails {
    pub proactivity: ProactivityScore, // From SPEC-PPP-004
    pub personalization: PersonalizationScore,
}

pub struct WeightedConsensus {
    pub best_agent: String,
    pub confidence: f32,
    pub scores: Vec<AgentScore>,
}

pub async fn run_spec_consensus_weighted(
    spec_id: &str,
    stage: &str,
    artifacts: Vec<ConsensusArtifactData>,
    weights: Option<(f32, f32)>, // (technical, interaction),
defaults to (0.7, 0.3)
) -> Result<WeightedConsensus> {
    let (w_tech, w_interact) = weights.unwrap_or((0.7, 0.3));

    // Validate weights
    if (w_tech + w_interact - 1.0).abs() > 0.001 {
        return Err(anyhow!("Weights must sum to 1.0, got {} + {} = {}",
                           w_tech, w_interact, w_tech + w_interact));
    }

    let db = open_consensus_db()?;
    let mut scores = Vec::new();

    for artifact in artifacts {
        // Technical score (existing logic)
        let technical = calculate_technical_score(&artifact)?;

        // Interaction score (new: from trajectory)
        let interaction = if let Ok(traj_id) =
get_trajectory_id(&db, spec_id, &artifact.agent) {
            let proact = calculate_r_proact(&db, traj_id)?;
            let pers = calculate_r_pers(&db, traj_id)?;

            let score = proact.r_proact + pers.r_pers;

            scores.push(AgentScore {
                agent_name: artifact.agent.clone(),
                technical_score: technical,
                interaction_score: score,
                final_score: (w_tech * technical) + (w_interact * score),
                details: ScoreDetails {
                    proactivity: proact,
                    personalization: pers,
                },
            });
        }
    }
}

```

```

        score
    } else {
        // Fallback: No trajectory available, use technical only
        warn!("No trajectory found for agent {}, using technical
score only", artifact.agent);

        scores.push(AgentScore {
            agent_name: artifact.agent.clone(),
            technical_score: technical,
            interaction_score: 0.0,
            final_score: technical, // 100% technical (backward
compatible)
            details: ScoreDetails {
                proactivity: ProactivityScore::default(),
                personalization:
PersonalizationScore::default(),
            },
        });
    }

    0.0
}
}

// Sort by final_score descending
scores.sort_by(|a, b|
b.final_score.partial_cmp(&a.final_score).unwrap());

Ok(WeightedConsensus {
    best_agent: scores[0].agent_name.clone(),
    confidence: scores[0].final_score,
    scores,
})
}
}

```

Integration Points:

```

// In consensus.rs, update public API
pub async fn run_spec_consensus(
    spec_id: &str,
    stage: &str,
    artifacts: Vec<ConsensusArtifactData>,
) -> Result<ConsensusResult> {
    // Delegate to weighted consensus
    let weighted = run_spec_consensus_weighted(spec_id, stage,
artifacts, None).await?;

    // Convert to old format for backward compatibility
    Ok(ConsensusResult {
        selected_agent: weighted.best_agent,
        selected_content: get_artifact_content(&weighted.best_agent,
&artifacts)?,
    })
}

```

Acceptance Criteria: - [] Weighted consensus selects best agent (validated with test cases) - [] Fallback to technical-only if no trajectory - [] Weights validated (must sum to 1.0) - [] Backward compatible (existing code still works)

Task 2: Configuration Support (2 hours)

File: config.toml.example

Add PPP Weights Section:

```
# PPP Framework - Weighted Consensus
[ppp.weights]
# Global default weights
technical = 0.7      # Technical quality weight (correctness,
completeness)
interaction = 0.3     # Interaction quality weight (proactivity +
personalization)

# Stage-specific weights (Phase 2)
[ppp.weights.plan]
technical = 0.6      # Planning: Exploration phase, interaction
matters more
interaction = 0.4

[ppp.weights.tasks]
technical = 0.7      # Tasks: Balanced
interaction = 0.3

[ppp.weights.implement]
technical = 0.7      # Implementation: Balanced
interaction = 0.3

[ppp.weights.validate]
technical = 0.75     # Validation: Correctness important
interaction = 0.25

[ppp.weights.audit]
technical = 0.8      # Audit: Security/compliance critical
interaction = 0.2

[ppp.weights.unlock]
technical = 0.8      # Unlock: Final validation, correctness
critical
interaction = 0.2
```

Configuration Loading:

```
// codex-rs/core/src/config_types.rs

#[derive(Deserialize, Serialize, Debug, Clone)]
pub struct PppWeights {
    #[serde(default = "default_technical_weight")]
    pub technical: f32,

    #[serde(default = "default_interaction_weight")]
    pub interaction: f32,

    // Stage-specific overrides (Phase 2)
    #[serde(default)]
    pub plan: Option<StageWeights>,

    #[serde(default)]
    pub tasks: Option<StageWeights>,

    #[serde(default)]
```

```

    pub implement: Option<StageWeights>,
#[serde(default)]
    pub validate: Option<StageWeights>,
#[serde(default)]
    pub audit: Option<StageWeights>,
#[serde(default)]
    pub unlock: Option<StageWeights>,
}

#[derive(Deserialize, Serialize, Debug, Clone)]
pub struct StageWeights {
    pub technical: f32,
    pub interaction: f32,
}

fn default_technical_weight() -> f32 { 0.7 }
fn default_interaction_weight() -> f32 { 0.3 }

impl PppWeights {
    pub fn validate(&self) -> Result<()> {
        let sum = self.technical + self.interaction;
        if (sum - 1.0).abs() > 0.001 {
            return Err(anyhow!(
                "PPP weights must sum to 1.0, got technical={} + interaction={} = {}",
                self.technical, self.interaction, sum
            ));
        }
    }

    // Validate stage-specific weights (Phase 2)
    for (stage, weights) in [
        ("plan", &self.plan),
        ("tasks", &self.tasks),
        ("implement", &self.implement),
        ("validate", &self.validate),
        ("audit", &self.audit),
        ("unlock", &self.unlock),
    ] {
        if let Some(w) = weights {
            let sum = w.technical + w.interaction;
            if (sum - 1.0).abs() > 0.001 {
                return Err(anyhow!(
                    "PPP weights.{} must sum to 1.0, got {}",
                    stage, sum
                ));
            }
        }
    }
    Ok(())
}

pub fn get_weights_for_stage(&self, stage: &str) -> (f32, f32) {
    let stage_weights = match stage {
        "plan" => &self.plan,
        "tasks" => &self.tasks,
        "implement" => &self.implement,
    }
}

```

```

        "validate" => &self.validate,
        "audit" => &self.audit,
        "unlock" => &self.unlock,
        _ => &None,
    };

    stage_weights
        .as_ref()
        .map(|w| (w.technical, w.interaction))
        .unwrap_or((self.technical, self.interaction))
    }
}

// Add to AgentConfig
#[derive(Deserialize, Serialize, Debug, Clone)]
pub struct AgentConfig {
    // ... existing fields

    #[serde(default)]
    pub ppp: Option<PppConfig>,
}

#[derive(Deserialize, Serialize, Debug, Clone)]
pub struct PppConfig {
    #[serde(default)]
    pub enabled: bool,

    #[serde(default)]
    pub weights: PppWeights,

    #[serde(default)]
    pub trajectory: TrajectoryConfig, // From SPEC-PPP-004
}

```

Usage in consensus.rs:

```

pub async fn run_spec_consensus_with_config(
    spec_id: &str,
    stage: &str,
    artifacts: Vec<ConsensusArtifactData>,
    config: &AgentConfig,
) -> Result<WeightedConsensus> {
    let weights = if let Some(ppp) = &config.ppp {
        ppp.weights.get_weights_for_stage(stage)
    } else {
        (0.7, 0.3) // Default
    };

    run_spec_consensus_weighted(spec_id, stage, artifacts,
        Some(weights)).await
}

```

Acceptance Criteria: - [] Config loads with defaults (70/30) - [] User can override via config.toml - [] Stage-specific weights work (Phase 2) - [] Validation catches invalid weights (sum ≠ 1.0)

Task 3: Integration Tests (4 hours)

File: codex-rs/tui/tests/ppp_weighted_consensus_test.rs

Test Scenarios:

```
use codex_tui::chatwidget::spec_kit::consensus::*;

#[tokio::test]
async fn test_weighted_consensus_basic() {
    // Setup: 3 agents with different technical + interaction scores
    let artifacts = vec![
        create_test_artifact("agent1", 0.85, 0.10), // Balanced
        create_test_artifact("agent2", 0.95, -0.45), // High tech,
        poor interaction
        create_test_artifact("agent3", 0.80, 0.02), // Medium
        tech, ok interaction
    ];

    let result = run_spec_consensus_weighted(
        "SPEC-TEST-001",
        "implement",
        artifacts,
        Some((0.7, 0.3)),
    ).await.unwrap();

    // Agent1 should win: 0.7*0.85 + 0.3*0.10 = 0.625
    assert_eq!(result.best_agent, "agent1");
    assert_eq!(result.scores.len(), 3);
    assert!((result.confidence - 0.625).abs() < 0.01);
}

#[tokio::test]
async fn test_fallback_no_trajectory() {
    // Setup: Agent with no trajectory (no interaction score)
    let artifacts = vec![
        create_test_artifact_no_trajectory("agent1", 0.90),
    ];

    let result = run_spec_consensus_weighted(
        "SPEC-TEST-002",
        "plan",
        artifacts,
        Some((0.7, 0.3)),
    ).await.unwrap();

    // Should use technical score only (100%)
    assert_eq!(result.best_agent, "agent1");
    assert_eq!(result.scores[0].final_score, 0.90);
    assert_eq!(result.scores[0].interaction_score, 0.0);
}

#[tokio::test]
async fn test_stage_specific_weights() {
    let config = load_test_config_with_stage_weights();
    let artifacts = vec![
        create_test_artifact("agent1", 0.80, 0.05),
        create_test_artifact("agent2", 0.70, 0.10),
    ];

    // Plan stage: 60/40 (interaction matters more)
    let result = run_spec_consensus_with_config(
        "SPEC-TEST-003",
        "plan",
    );
```

```

        artifacts.clone(),
        &config,
    ) .await.unwrap();

    // Agent2 should win in plan: 0.6*0.70 + 0.4*0.10 = 0.46
    // vs Agent1: 0.6*0.80 + 0.4*0.05 = 0.50
    assert_eq!(result.best_agent, "agent1");

    // Unlock stage: 80/20 (technical matters more)
    let result = run_spec_consensus_with_config(
        "SPEC-TEST-003",
        "unlock",
        artifacts.clone(),
        &config,
    ) .await.unwrap();

    // Agent1 should still win: 0.8*0.80 + 0.2*0.05 = 0.65
    assert_eq!(result.best_agent, "agent1");
}

#[tokio::test]
async fn test_weight_validation() {
    let artifacts = vec![create_test_artifact("agent1", 0.8, 0.1)];

    // Invalid: weights don't sum to 1.0
    let result = run_spec_consensus_weighted(
        "SPEC-TEST-004",
        "plan",
        artifacts,
        Some((0.5, 0.6)), // Sum = 1.1
    ) .await;

    assert!(result.is_err());
    assert!(result.unwrap_err().to_string().contains("must sum to
1.0"));
}

#[tokio::test]
async fn test_equal_scores_tiebreaker() {
    // Setup: 2 agents with identical final scores
    let artifacts = vec![
        create_test_artifact("agent1", 0.80, 0.05), // 0.7*0.8 +
0.3*0.05 = 0.575
        create_test_artifact("agent2", 0.75, 0.10), // 0.7*0.75 +
0.3*0.10 = 0.555
    ];

    let result = run_spec_consensus_weighted(
        "SPEC-TEST-005",
        "implement",
        artifacts,
        Some((0.7, 0.3)),
    ) .await.unwrap();

    // Agent1 should win (higher score)
    assert_eq!(result.best_agent, "agent1");
}

```

Helper Functions:

```
fn create_test_artifact(agent: &str, technical: f32, interaction:
```

```

f32) -> ConsensusArtifactData {
    // Create artifact with pre-computed scores
    // Insert trajectory with calculated R_Proact + R_Pers =
interaction
    // ...
}

fn create_test_artifact_no_trajectory(agent: &str, technical: f32) -
> ConsensusArtifactData {
    // Create artifact without trajectory (test fallback)
    // ...
}

```

Acceptance Criteria: - [] All test scenarios pass - [] Weighted formula validated (spot checks) - [] Fallback behavior tested - [] Stage-specific weights tested

Task 4: Documentation & Examples (2 hours)

File: docs/ppp-weighted-consensus-guide.md

Contents:

```

# PPP Weighted Consensus Guide

## Overview

PPP weighted consensus combines **technical quality** (correctness, completeness) with **interaction quality** (proactivity, personalization) when selecting the best agent output.

## Formula

final_score = (w_technical × technical_score) + (w_interaction × interaction_score)

Default weights: - w_technical = 0.7 (70%) - w_interaction = 0.3 (30%)

## Configuration

### Global Weights

```toml
[ppp.weights]
technical = 0.7
interaction = 0.3
```

```

Stage-Specific Weights

```

[ppp.weights.plan]
technical = 0.6    # Exploration phase
interaction = 0.4

[ppp.weights.unlock]
technical = 0.8    # Correctness critical
interaction = 0.2

```

Examples

Example 1: Balanced Agent Wins

Agent 1:

Technical: 0.85 (good code quality)
Interaction: 0.10 (asked 2 low-effort questions)
Final: $0.7 \times 0.85 + 0.3 \times 0.10 = 0.625 \leftarrow \text{Winner}$

Agent 2:

Technical: 0.95 (excellent code quality)
Interaction: -0.45 (asked 1 high-effort blocking question)
Final: $0.7 \times 0.95 + 0.3 \times (-0.45) = 0.530$

Winner: Agent 1 - Better overall balance

Example 2: Technical Quality Dominates

Weights: 80/20 (unlock stage)

Agent 1:

Technical: 0.95
Interaction: -0.20
Final: $0.8 \times 0.95 + 0.2 \times (-0.20) = 0.72 \leftarrow \text{Winner}$

Agent 2:

Technical: 0.85
Interaction: 0.05
Final: $0.8 \times 0.85 + 0.2 \times 0.05 = 0.69$

Winner: Agent 1 - Technical quality prioritized

Tuning Weights

When to Increase Technical Weight (>0.7)

- Critical stages (audit, unlock)
- Production deployments
- Security-sensitive code
- Complex algorithms

When to Increase Interaction Weight (>0.3)

- Early stages (plan, tasks)
- Prototyping/exploration
- Beginner users
- Interactive workflows

Finding Your Ideal Weights

1. Start with defaults (70/30)
2. Run 10-20 spec-kit executions
3. Review selected agents vs your preference
4. Adjust weights in 0.05 increments
5. Repeat until satisfied

```

**Acceptance Criteria**:
- [ ] Guide explains formula clearly
- [ ] Examples demonstrate use cases
- [ ] Tuning guidance provided
- [ ] Links to full SPEC-PPP-003 research

---

## Performance Validation

### Benchmark Targets

Operation	Target	Measurement
**Calculate Interaction Score**	<20ms	2 SQL queries (R_Proact + R_Pers)
**Calculate Final Score**	<1ms	Weighted average (arithmetic)
**Select Best Agent**	<5ms	Sort by score
**Total Consensus Overhead**	**<50ms**	All operations combined

### Test Queries

```sql
-- Test 1: R_Proact with 100 turns, 50 questions
SELECT ... FROM trajectory_turns t LEFT JOIN trajectory_questions q ...
;
-- Expected: <10ms

-- Test 2: R_Pers with 100 turns, 10 violations
SELECT ... FROM trajectory_turns t LEFT JOIN trajectoryViolations v ...
;
-- Expected: <10ms

-- Test 3: Sort 5 agents by final score
SELECT * FROM agent_scores ORDER BY final_score DESC;
-- Expected: <1ms

```

---

## Phase 2 Enhancements

### User-Configurable Weights

**UI Addition** (optional, codex-tui):

PPP Consensus Settings

Technical Weight:	[0.70] (0.0 - 1.0)
Interaction Weight:	[0.30] (auto-calc)

Presets:

- [ ] Balanced (70/30)
- [ ] Technical Focus (80/20)
- [ ] Interaction Focus (60/40)

[Apply] [Reset to Defaults]

## Stage-Specific Weight Recommendations

Based on task criticality research:

Stage	Technical	Interaction	Rationale
<b>plan</b>	0.6	0.4	Exploration, questions OK
<b>tasks</b>	0.7	0.3	Balanced
<b>implement</b>	0.7	0.3	Balanced
<b>validate</b>	0.75	0.25	Testing correctness important
<b>audit</b>	0.8	0.2	Security/compliance critical
<b>unlock</b>	0.8	0.2	Final validation

## Phase 3 Advanced Features

### Regret Tracking

**Metric:** Regret = score(best) - score(selected)

**Implementation:**

```
pub struct ConsensusRegret {
 pub selected_agent: String,
 pub selected_score: f32,
 pub best_agent: String, // In retrospect (user feedback)
 pub best_score: f32,
 pub regret: f32,
}

pub fn track_regret(
 consensus: &WeightedConsensus,
 user_preferred: &str,
) -> ConsensusRegret {
 let selected_score = consensus.confidence;
 let best_score = consensus.scores.iter()
 .find(|s| s.agent_name == user_preferred)
 .map(|s| s.final_score)
 .unwrap_or(0.0);

 ConsensusRegret {
 selected_agent: consensus.best_agent.clone(),
 selected_score,
 best_agent: user_preferred.to_string(),
 best_score,
 regret: best_score - selected_score,
 }
}
```

**Usage:** Collect regret over 100+ runs, optimize weights to minimize average regret.

### Grid Search Weight Optimization

**Pseudocode:**

```

best_weights = (0.7, 0.3)
min_regret = float('inf')

for w_tech in [0.5, 0.55, 0.6, ..., 0.9]:
 w_interact = 1.0 - w_tech

 total_regret = 0
 for run in historical_runs:
 consensus = run_weighted_consensus(run.artifacts, (w_tech,
w_interact))
 regret = calculate_regret(consensus, run.user_preferred)
 total_regret += regret

 avg_regret = total_regret / len(historical_runs)

 if avg_regret < min_regret:
 min_regret = avg_regret
 best_weights = (w_tech, w_interact)

print(f"Optimal weights: {best_weights} (regret: {min_regret})")

```

**Expected Result:** Optimal weights within  $\pm 0.05$  of 70/30 for most users.

---

## Success Metrics

### Phase 1 (Core)

- Weighted consensus selects best agent (validated on test cases)
- <50ms overhead (measured via benchmarks)
- 100% backward compatible (existing code works)
- Zero production incidents

### Phase 2 (Configuration)

- Users can configure weights via config.toml
- Stage-specific weights applied correctly
- Config validation prevents invalid weights
- Documentation guides users to optimal weights

### Phase 3 (Optimization)

- Regret tracking measures effectiveness
  - Grid search finds optimal weights per user
  - Average regret <0.05 (near-optimal selection)
- 

## Dependencies

**Phase 1:** - SPEC-PPP-004 (trajectory logging) - Must be implemented first - SPEC-PPP-002 (preference violations) - For R\_Pers calculation - consensus.rs (existing) - Refactor target

**Phase 2:** - config\_types.rs (existing) - Extend with PppWeights

**Phase 3:** - Historical consensus data (100+ runs) - User feedback mechanism (which agent was best?)

---

**Estimated Total Timeline:** - Phase 1: 1 day (solo engineer) - Phase 2: 0.75 days - Phase 3: 1.5 days (optional) - **Total: 2-3 days** (core + configuration)