

# ADR-004-002-async-batching-strategy

## ADR-004-002: Async Batching Strategy for Trajectory Logging

**Status:** Accepted **Date:** 2025-11-16 **Deciders:** Research Team  
**Related:** SPEC-PPP-004 (Trajectory Logging & MCP Integration)

### Context

Trajectory logging must record every user-agent turn during spec-kit execution. Requirements: - **Latency:** <1ms overhead per turn (agent execution is ~10-30 seconds, logging should be <0.01% overhead) - **Throughput:** Support 20-50 turns per agent execution - **Non-blocking:** Must not block tokio async runtime - **Data integrity:** No data loss on crashes

Four approaches were evaluated: 1. **Synchronous single write:** rusqlite, blocks on each turn 2. **Synchronous batched write:** rusqlite with transactions, blocks briefly 3. **Asynchronous single write:** tokio-rusqlite, non-blocking but overhead 4. **Asynchronous batched write:** tokio-rusqlite with buffering

### Decision

We will use **Asynchronous batched write** (Option 4) with the following parameters: - **Buffer size:** 5-10 turns - **Flush interval:** 500ms - **Library:** tokio-rusqlite (23,000 downloads/month)

### Rationale

#### Performance Comparison

Strategy	Latency/Turn	Throughput	Blocks Runtime	Implementation
Sync Single	1ms	1,000/sec	✗ Yes	Simple
Sync Batched	0.067ms	15,000/sec	✗ Yes	Medium
Async Single	0.5ms	2,000/sec	✓ No	Medium

<b>Async</b>	<b>0.1ms</b>	<b>10,000/sec</b>	<input checked="" type="checkbox"/> <b>No</b>	Complex
<b>Batched</b>				

**Source:** SPEC-PPP-004 findings.md, benchmark estimates from “15k inserts/s with Rust and SQLite”

**Winner:** Async batched achieves **0.1ms/turn** (10x better than sync single) without blocking runtime.

## Async Batched Implementation

```

use tokio::sync::mpsc;
use tokio_rusqlite::Connection;
use std::time::Duration;

pub struct TrajectoryLogger {
    tx: mpsc::Sender<LogEntry>,
}

impl TrajectoryLogger {
    pub fn new(db_path: String) -> Self {
        let (tx, rx) = mpsc::channel(100); // Bounded channel for
                                         // backpressure

        tokio::spawn(async move {
            writer_task(db_path, rx).await
        });
    }

    Self { tx }
}

pub async fn log_turn(&self, turn: Turn) -> Result<()> {
    self.tx.send(LogEntry::Turn(turn)).await?;
    Ok(())
}

async fn writer_task(
    db_path: String,
    mut rx: mpsc::Receiver<LogEntry>,
) -> Result<()> {
    let conn = Connection::open(&db_path).await?;
    let mut buffer = Vec::new();

    loop {
        tokio::select! {
            // Receive entries from channel
            Some(entry) = rx.recv() => {
                buffer.push(entry);

                // Flush if buffer full
                if buffer.len() >= 10 {
                    flush_buffer(&conn, &buffer).await?;
                    buffer.clear();
                }
            }
        }
    }
}

// Flush on timeout (every 500ms)

```

```

        _ = tokio::time::sleep(Duration::from_millis(500)) => {
            if !buffer.is_empty() {
                flush_buffer(&conn, &buffer).await?;
                buffer.clear();
            }
        }
    }
}

async fn flush_buffer(conn: &Connection, buffer: &[LogEntry]) ->
Result<()> {
    conn.call(move |conn| {
        let tx = conn.transaction()?;
        for entry in buffer {
            // Insert each entry
            tx.execute("INSERT INTO trajectory_turns (...) VALUES
(...)", ...)?;
        }
        tx.commit()
    }).await
}

```

**Key Features:** 1. **Bounded channel** (100 capacity) for backpressure  
- prevents memory overflow 2. **Dual flush triggers**: Buffer full (10 entries) OR timeout (500ms) 3. **Single writer task**: No lock contention (SQLite single-writer limitation) 4. **Transaction per batch**: Atomic writes, no partial data

---

## Buffer Size Selection

Buffer Size	Latency	Flush Frequency	Data Loss Risk
1 (no buffer)	0.5ms	Every turn	Minimal
5	0.2ms	Every 5 turns	Low
<b>10</b>	<b>0.1ms</b>	<b>Every 10 turns</b>	<b>Low</b>
50	0.02ms	Every 50 turns	Medium
100	0.01ms	Every 100 turns	High

**Decision:** Buffer size = **10**

**Rationale:** - Typical agent execution: 20-50 turns → 2-5 flushes per run  
- Data loss window: Max 10 turns (acceptable, <1 second of work)  
- Performance: 10x better than no buffering (0.1ms vs 1ms) - Memory: ~5KB per buffer (10 turns × 500 bytes/turn)

---

## Flush Interval Selection

Interval	Worst-Case Delay	Data Loss Window	Performance
100ms	100ms	100ms work	High overhead
250ms	250ms	250ms work	Medium overhead
<b>500ms</b>	<b>500ms</b>	<b>&lt;1sec work</b>	<b>Optimal</b>
1000ms	1s	1s work	Low overhead

5000ms	5s	5s work	Too high
--------	----	---------	----------

**Decision:** Flush interval = **500ms**

**Rationale:** - Agent turns are ~5-10 seconds apart (user typing time) - 500ms guarantees data persisted before next turn 99% of the time - Data loss window: <1 second of work (acceptable for logging) - Performance: Reduces write frequency by 5-10x vs 100ms

## Why Not Synchronous?

**Problem:** Synchronous writes block tokio runtime

```
// BAD: Blocks entire async runtime
pub fn log_turn_sync(conn: &Connection, turn: Turn) -> Result<()> {
    conn.execute(
        "INSERT INTO trajectory_turns (...) VALUES (...)",
        params![...],
    )?; // ← Blocks ALL async tasks for ~1ms
    Ok(())
}
```

**Impact:** - Blocks consensus synthesis (~1s per agent) - Blocks UI updates (TUI frozen during logging) - Cascading delays across all async tasks

**Evidence:** Rust async best practice - never block in async fn

**Source:** tokio documentation - “Blocking operations in async code”

## Why Not Async Single Write?

**Comparison:** Async batched vs async single

Metric	Async Single	Async Batched
Latency	0.5ms	<b>0.1ms (5x better)</b>
Lock contention	High (every turn)	<b>Low (every 10 turns)</b>
Transaction overhead	High	<b>Low (amortized)</b>
Complexity	Medium	Medium

**Decision:** 5x performance improvement justifies slightly more complex code.

## Consequences

### Positive

1. ✓ **10x performance:** 0.1ms/turn vs 1ms/turn (sync single)
2. ✓ **Non-blocking:** Tokio runtime stays responsive
3. ✓ **High throughput:** 10,000 turns/sec capacity
4. ✓ **Atomic batches:** Transaction guarantees no partial data
5. ✓ **Backpressure:** Bounded channel prevents memory overflow

## Negative

1. **△ Data loss window:** Up to 10 turns (500ms) if crash occurs
  - Mitigation: Acceptable for logging use case (not critical transactions)
  - Alternative: Set buffer size=1 for critical data (trade performance)
2. **△ Complexity:** More complex than synchronous writes
  - Mitigation: Well-tested pattern (used in production systems)
  - Code: ~50 lines vs ~10 lines (synchronous)
3. **△ Debugging:** Async stack traces harder to debug
  - Mitigation: Comprehensive logging in writer\_task
  - Error handling: Explicit Result propagation

## Neutral

1. **▀ Memory overhead:** ~5KB per buffer (10 turns)
    - Negligible: <0.01% of typical memory usage
- 

# Trade-offs Considered

## Flush on Agent Completion

**Alternative:** Don't use timeout, flush only when agent completes

```
// No timeout, manual flush
impl TrajectoryLogger {
    pub async fn flush(&self) -> Result<()> {
        // Signal writer to flush immediately
    }
}

// Usage
logger.log_turn(turn1).await?;
logger.log_turn(turn2).await?;
// ... agent completes
logger.flush().await?; // ← Manual flush
```

**Rejected because:** - Requires caller to remember to flush (error-prone)  
- If agent crashes, data is lost entirely (no timeout fallback)  
- More complex API (2 methods instead of 1)

**Accepted approach:** Automatic timeout ensures data persists even if flush() not called.

---

## WAL Mode

**Alternative:** Enable Write-Ahead Logging (WAL) for even better performance

```
conn.execute("PRAGMA journal_mode=WAL", []);  
conn.execute("PRAGMA synchronous=NORMAL", []);
```

**Performance:** - Latency: 0.02ms/turn (5x better than batched)  
- Throughput: 50,000/sec

**Deferred to Phase 3 because:** - Current performance (0.1ms) already meets requirements (<1ms target) - WAL adds complexity (separate WAL files to manage) - Marginal benefit (0.1ms → 0.02ms not critical)

**Decision:** Implement WAL in Phase 3 if benchmarks show bottleneck.

---

## Configuration

Add to config.toml:

```
[ppp.trajectory]
buffer_size = 10          # Turns to buffer before flush
flush_interval_ms = 500    # Flush interval in milliseconds
channel_capacity = 100     # Backpressure limit
```

**Tuning guidance:** - **Low latency priority:** Set buffer\_size=1, flush\_interval\_ms=100 - **High throughput priority:** Set buffer\_size=50, flush\_interval\_ms=1000 - **Balanced (default):** Use documented values

---

## Testing Strategy

### Unit Tests

```
#[tokio::test]
async fn test_async_batching() {
    let logger = TrajectoryLogger::new(":memory:".to_string());

    // Log 5 turns (below buffer size)
    for i in 0..5 {
        logger.log_turn(Turn { turn_number: i, ... });
    }

    // Wait for timeout flush (500ms)
    tokio::time::sleep(Duration::from_millis(600)).await;

    // Verify all 5 turns persisted
    let count = query_turn_count();
    assert_eq!(count, 5);
}

#[tokio::test]
async fn test_buffer_full_flush() {
    let logger = TrajectoryLogger::new(":memory:".to_string());

    // Log 10 turns (exactly buffer size)
    for i in 0..10 {
        logger.log_turn(Turn { turn_number: i, ... });
    }

    // Should flush immediately (no timeout wait needed)
    tokio::time::sleep(Duration::from_millis(10)).await;
```

```
// Verify all 10 turns persisted
let count = query_turn_count();
assert_eq!(count, 10);
}
```

## Benchmark Tests

```
#[tokio::test]
async fn bench_async_batched() {
    let logger = TrajectoryLogger::new("bench.db".to_string());

    let start = Instant::now();
    for i in 0..1000 {
        logger.log_turn(Turn { turn_number: i, ...
    }).await.unwrap();
    }
    let elapsed = start.elapsed();

    // Target: <0.1ms per turn
    assert!(elapsed < Duration::from_millis(100));
}
```

---

## References

1. Kerkour, S. (2024). "15k inserts/s with Rust and SQLite" - <https://kerkour.com/high-performance-rust-with-sqlite>
  2. tokio-rusqlite documentation - <https://lib.rs/crates/tokio-rusqlite>
  3. tokio async best practices - <https://tokio.rs/tokio/tutorial/spawning>
  4. SPEC-PPP-004 findings.md - Performance benchmarks
- 

## Decision Drivers

<b>Driver</b>	<b>Weight</b>	<b>Async Batched</b>	<b>Sync Single</b>	<b>Async Single</b>
Performance	40%	✓ 0.1ms	✗ 1ms	△ 0.5ms
Non-blocking	30%	✓ Yes	✗ No	✓ Yes
Data Integrity	20%	△ 500ms window	✓ Immediate	△ Immediate
Complexity	10%	△ Medium	✓ Simple	△ Medium

**Total Score:** Async Batched **85%**, Async Single **70%**, Sync Single **40%**

**Winner:** Async Batched by significant margin.