

ADR-004-003-turn-based-schema-design

ADR-004-003: Turn-Based Schema Design for Trajectory Storage

Status: Accepted **Date:** 2025-11-16 **Deciders:** Research Team
Related: SPEC-PPP-004 (Trajectory Logging & MCP Integration)

Context

PPP framework requires storing multi-turn conversations with metadata to calculate:
- **R_{Proact}**: Based on question effort (low/medium/high)
- **R_{Pers}**: Based on preference violations (minor/major/critical)

Four schema patterns were evaluated:
1. **Flat message history**: Single table with role-based messages
2. **Turn-based with metadata**: 4 tables (trajectories, turns, questions, violations)
3. **JSON blob storage**: Store turns as JSON in single column
4. **Hierarchical spans**: OpenTelemetry-style parent/child traces

Decision

We will use **Turn-based schema with metadata** (Option 2) with 4 normalized tables:
- trajectories: One per agent execution - trajectory_turns: One per user-agent exchange - trajectory_questions: Questions asked by agent - trajectory_violations: Preference violations detected

Rationale

Schema Design

```
-- Table 1: Trajectories (one per agent execution)
CREATE TABLE trajectories (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    spec_id TEXT NOT NULL, -- e.g., "SPEC-KIT-001"
    agent_name TEXT NOT NULL, -- e.g., "gemini-flash",
    run_id TEXT, -- Links to
    consensus_artifacts
        created_at TEXT DEFAULT (datetime('now')),
        INDEX idx_spec_agent (spec_id, agent_name),
        INDEX idx_run_id (run_id)
```

```

);

-- Table 2: Turns (one per user-agent exchange)
CREATE TABLE trajectory_turns (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    trajectory_id INTEGER NOT NULL,
    turn_number INTEGER NOT NULL,          -- 1, 2, 3, ...
    prompt TEXT NOT NULL,                 -- User input
    response TEXT NOT NULL,               -- Agent output
    token_count INTEGER,                  -- Optional: for cost
    tracking
        latency_ms INTEGER,              -- Optional: for performance
        timestamp TEXT DEFAULT (datetime('now')),
        FOREIGN KEY (trajectory_id) REFERENCES trajectories(id) ON
DELETE CASCADE,
    INDEX idx_trajectory (trajectory_id),
    UNIQUE (trajectory_id, turn_number)
);

-- Table 3: Questions (extracted from responses)
CREATE TABLE trajectory_questions (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    turn_id INTEGER NOT NULL,
    question_text TEXT NOT NULL,
    effort_level TEXT,                   -- 'low', 'medium', 'high'
    FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id) ON DELETE
CASCADE,
    INDEX idx_turn (turn_id)
);

-- Table 4: Violations (detected in responses)
CREATE TABLE trajectory_violations (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    turn_id INTEGER NOT NULL,
    preference_name TEXT NOT NULL,       -- e.g., 'require_json'
    expected TEXT NOT NULL,              -- e.g., 'Valid JSON'
    actual TEXT NOT NULL,                -- e.g., 'Plain text'
    severity TEXT NOT NULL,              -- 'minor', 'major',
    'critical'
    FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id) ON DELETE
CASCADE,
    INDEX idx_turn (turn_id),
    INDEX idx_preference (preference_name)
);

```

Comparison: 4 Schema Patterns

Feature	Flat	Turn-Based	JSON Blob	Hierarchical
Tables	1	4	1	2
Query Complexity	Low	Medium	High	Very High
Storage Efficiency	Good	Excellent	Poor	Medium
Extensibility	Poor	Excellent	Good	Excellent
PPP Support	20%	100%	50%	60%

Indexing	Limited	Full	Limited	Limited
-----------------	---------	-------------	---------	---------

Winner: Turn-based schema - Only option supporting all PPP calculations efficiently.

Why Turn-Based Wins

1. Supports PPP Calculations Natively

R_{Proact} Query (calculate from questions):

```

SELECT
CASE
    WHEN COUNT(q.id) = 0 THEN 0.05
    WHEN COUNT(CASE WHEN q.effort_level != 'low' THEN 1 END) = 0
THEN 0.05
    ELSE -0.1 * COUNT(CASE WHEN q.effort_level = 'medium' THEN 1
END)
    -0.5 * COUNT(CASE WHEN q.effort_level = 'high' THEN 1
END)
END AS r_proact
FROM trajectory_turns t
LEFT JOIN trajectory_questions q ON t.id = q.turn_id
WHERE t.trajectory_id = ?;
```

Performance: ~3-8ms for 100-turn trajectory

R_{Pers} Query (calculate from violations):

```

SELECT
CASE
    WHEN COUNT(v.id) = 0 THEN 0.05
    ELSE -0.01 * COUNT(CASE WHEN v.severity = 'minor' THEN 1
END)
    -0.03 * COUNT(CASE WHEN v.severity = 'major' THEN 1
END)
    -0.05 * COUNT(CASE WHEN v.severity = 'critical' THEN 1
END)
END AS r_pers
FROM trajectory_turns t
LEFT JOIN trajectoryViolations v ON t.id = v.turn_id
WHERE t.trajectory_id = ?;
```

Performance: ~3-8ms for 100-turn trajectory

Total Scoring Time: <20ms (both queries + merge)

2. Normalized Schema (No Duplication)

Flat Schema (redundant):

```

CREATE TABLE messages (
    id INTEGER PRIMARY KEY,
    session_id TEXT,          -- Repeated for every message
    spec_id TEXT,             -- Repeated for every message
    agent_name TEXT,          -- Repeated for every message
    role TEXT,
```

```
        content TEXT  
    );
```

Turn-Based Schema (normalized):

```
-- Metadata stored once per trajectory  
CREATE TABLE trajectories (  
    id INTEGER PRIMARY KEY,  
    spec_id TEXT,           -- ✓ Stored once  
    agent_name TEXT         -- ✓ Stored once  
);  
  
-- Only turn data stored per row  
CREATE TABLE trajectory_turns (  
    id INTEGER PRIMARY KEY,  
    trajectory_id INTEGER,   -- ✓ Reference, not duplicate  
    prompt TEXT,  
    response TEXT  
);
```

Storage Savings: ~30% less disk space (metadata not duplicated)

3. Extensible Without Migration

Add New Metadata (no schema change):

```
-- Want to track tool calls? Add new table  
CREATE TABLE trajectory_tool_calls (  
    id INTEGER PRIMARY KEY,  
    turn_id INTEGER,  
    tool_name TEXT,  
    arguments TEXT,  
    result TEXT,  
    FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id)  
);
```

Existing queries still work - no migration needed.

Flat/JSON Schema: Would require schema change or JSON parsing overhead.

Why NOT Flat Schema

Problem 1: Prompt + response = 2 rows

```
-- Flat schema forces 2 rows per turn  
INSERT INTO messages (role, content) VALUES ('user', 'Implement  
OAuth');  
      INSERT INTO messages (role, content) VALUES ('assistant', 'Which  
provider?');  
  
-- Querying requires grouping by turn (complex)  
SELECT  
    MAX(CASE WHEN role = 'user' THEN content END) AS prompt,  
    MAX(CASE WHEN role = 'assistant' THEN content END) AS response  
FROM messages  
GROUP BY turn_number;  -- ← turn_number not in schema!
```

Turn-Based: 1 row per turn (simpler)

```

    INSERT INTO trajectory_turns (prompt, response) VALUES ('Implement OAuth', 'Which provider?');

    SELECT prompt, response FROM trajectory_turns; -- ✓ Simple

```

Problem 2: No place to store metadata

- Where to store effort_level? (per question, not per message)
- Where to store violation details? (per violation, not per message)

Flat schema would require complex JSON columns or denormalized data.

Why NOT JSON Blob

Proposed Schema:

```

CREATE TABLE trajectories (
    id INTEGER PRIMARY KEY,
    spec_id TEXT,
    turns_json TEXT -- JSON: [{prompt, response, questions: [...],
    violations: [...]}, ...]
);

```

Example JSON:

```
{
  "turns": [
    {
      "prompt": "Implement OAuth",
      "response": "Which provider?",
      "questions": [{"text": "Which provider?", "effort": "low"}],
      "violations": []
    }
  ]
}
```

Problems:

1. **No indexing** (SQLite <3.38 limited JSON support)

```
-- Slow: Must scan entire JSON blob
SELECT * FROM trajectories WHERE json_extract(turns_json,
'$turns[*].questions[*].effort') = 'high';
```

2. **No foreign keys** (can't enforce referential integrity)

3. **Hard to query** (requires json_extract() everywhere)

```
-- Complex $R_{Proact}$ calculation
SELECT
  CASE
    WHEN json_extract(turns_json, '$.turns[*].questions') IS
      NULL THEN 0.05
    ELSE (
      SELECT -0.1 * COUNT(*) FROM json_each(...) -- ←
      Nested queries
    )
  END
FROM trajectories;
```

4. Storage overhead (JSON syntax: ~30% larger than normalized)

Decision: JSON blob rejected - too complex, no performance benefit.

Why NOT Hierarchical Spans (OpenTelemetry)

Proposed Schema:

```
CREATE TABLE spans (
    id INTEGER PRIMARY KEY,
    parent_id INTEGER,
    span_type TEXT, -- 'conversation', 'turn', 'llm_call',
    'tool_use'
    attributes_json TEXT,
    FOREIGN KEY (parent_id) REFERENCES spans(id)
);
```

Example Hierarchy:

```
Span (conversation, id=1)
  └─ Span (turn, id=2, parent=1)
    └─ Span (llm_call, id=3, parent=2)
      └─ Span (tool_use, id=4, parent=2)
    └─ Span (turn, id=5, parent=1)
```

Problems:

1. Overkill for CLI (no nested tool calls in codex-tui)

- Coding assistants don't have complex tool hierarchies
- Turns are flat (user → agent → user)

2. Complex queries (recursive CTEs required)

```
-- Get all turns for conversation (requires recursive CTE)
WITH RECURSIVE turn_hierarchy AS (
    SELECT * FROM spans WHERE id = ?
    UNION ALL
    SELECT s.* FROM spans s JOIN turn_hierarchy th ON s.parent_id
    = th.id
)
SELECT * FROM turn_hierarchy WHERE span_type = 'turn';
```

3. Storage overhead (~2x due to hierarchy metadata)

4. PPP calculations harder (questions/violations buried in attributes_json)

Decision: Hierarchical rejected - unnecessary complexity for flat turn structure.

Consequences

Positive

1. ✓ **100% PPP support:** All formulas map directly to SQL queries
2. ✓ **Normalized:** No data duplication, efficient storage

3. ✓ **Indexed**: Fast queries on spec_id, agent_name, trajectory_id
4. ✓ **Extensible**: Add tables without migrating existing data
5. ✓ **Foreign keys**: Cascading deletes maintain referential integrity
6. ✓ **Simple queries**: Standard SQL, no complex JOINs or CTEs

Negative

1. △ **4 tables**: More complex than flat schema
 - Mitigation: Standard normalization pattern, well-understood
 - Acceptable: Complexity justified by query performance
2. △ **JOINS required**: Queries must join tables
 - Mitigation: Indexed foreign keys make JOINs fast (<10ms)
 - Acceptable: Better than JSON parsing overhead
3. △ **More write operations**: 4 INSERTs per turn (worst case)
 - Mitigation: Async batching amortizes cost (0.1ms/turn total)
 - Acceptable: Write overhead << query benefit

Neutral

1. ■ **Storage**: ~500 bytes/turn (normalized)
 - Comparable to JSON blob (~600 bytes with overhead)
 - Acceptable: 20 turns × 500 bytes = 10 KB/trajectory
-

Migration Strategy

Phase 1: Schema Creation

```
pub fn init_trajectory_tables(conn: &Connection) -> Result<()> {
    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectories (...)",
        [],
    )?;

    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectory_turns (...)",
        [],
    )?;

    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectory_questions (...)",
        [],
    )?;

    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectory_violations (...)",
        [],
    )?;

    Ok(())
}
```

Phase 2: Versioned Migration

```
pub fn migrate_database(conn: &Connection) -> Result<()> {
    let version: i32 = conn.query_row("PRAGMA user_version", [],
```

```

| row| row.get(0))?;

    if version < 2 {
        init_trajectory_tables(conn)?;
        conn.execute("PRAGMA user_version = 2", []);
    }

    Ok(())
}

```

Phase 3: Rollback Support

```

-- If trajectory logging causes issues, can drop tables safely
DROP TABLE trajectory_violations;
DROP TABLE trajectory_questions;
DROP TABLE trajectory_turns;
DROP TABLE trajectories;

-- Consensus DB continues working (separate tables)

```

Alternative Considered: Hybrid (Turn-Based + JSON)

Idea: Store common fields in columns, metadata in JSON

```

CREATE TABLE trajectory_turns (
    id INTEGER PRIMARY KEY,
    trajectory_id INTEGER,
    prompt TEXT,
    response TEXT,
    metadata_json TEXT -- {questions: [...], violations: [...]})
);

```

Pros: - Simpler (3 tables instead of 4) - Flexible metadata

Cons: - Loses indexing on effort_level, severity - PPP queries become complex (JSON parsing) - Inconsistent: Some data normalized, some JSON

Rejected: Prefer full normalization for consistency.

Performance Validation

Benchmark Targets

Operation	Target	Measurement
Insert Turn	<0.1ms	Async batched
Query R_{Proact}	<10ms	100-turn trajectory
Query R_{Pers}	<10ms	100-turn trajectory
Full Scoring	<20ms	Both queries + merge

Test Queries

```

-- Test 1: R_Proact with 100 turns, 50 questions
SELECT ... FROM trajectory_turns t LEFT JOIN trajectory_questions q
...
WHERE t.trajectory_id = ?;
-- Expected: <10ms

-- Test 2: R_Pers with 100 turns, 10 violations
SELECT ... FROM trajectory_turns t LEFT JOIN trajectoryViolations v
...
WHERE t.trajectory_id = ?;
-- Expected: <10ms

```

References

1. SPEC-PPP-004 findings.md - Schema pattern analysis
 2. SPEC-PPP-004 comparison.md - Detailed schema comparison
 3. PPP Framework (arXiv:2511.02208) - Formula definitions
 4. SQLite normalization best practices
-

Decision Drivers

Driver	Weight	Turn-Based	Flat	JSON	Hierarchical
PPP Support	40%	✓ 100%	✗ 20%	△ 50%	△ 60%
Query Performance	30%	✓ Fast	△ Medium	✗ Slow	✗ Slow
Extensibility	20%	✓ High	✗ Low	△ Medium	✓ High
Complexity	10%	△ Medium	✓ Simple	△ Medium	✗ Complex

Total Score: Turn-Based **90%**, Hierarchical **55%**, JSON **45%**, Flat **35%**

Winner: Turn-Based by significant margin.