

comparison

SPEC-PPP-004: Comparative Analysis

Last Updated: 2025-11-16

Observability Platform Comparison

Feature	LangSmith	Helicone	Phoenix Arize	Prop (SQI)
Multi-turn Tracking	✓ Threads	✓ Sessions	✓ Sessions	✓ Traje
Trace Hierarchy	✓ Trace → Run	✓ Path syntax	✓ Spans	✓ Traje → Turn
Deployment	✗ Cloud only	✓ Self-hosted	✓ Self-hosted	✓ Local SQLite
License	Proprietary	Apache 2.0 (OSS)	Apache 2.0 (OSS)	N/A (pr code)
Pricing	\$39/mo+	Free (OSS)	Free (OSS)	FREE
Integration	LangChain required	HTTP proxy	OpenTelemetry	Native
Latency Overhead	50-100ms	10-20ms	30-50ms	<1ms (async)
Question Effort	✗ None	✗ None	✗ None	✓ Low/Med
Preference Violations	✗ None	✗ None	✗ None	✓ Full tracking
Interaction Scoring	✗ None	✗ None	△ Coherence	✓ $R_{Proa} + R_{Pers}$
Session-level Metrics	△ Multi-turn evals	△ Basic	✓ Comprehensive	✓ PPP formula
CLI-Specific Features	✗ None	✗ None	✗ None	✓ SPEC tracking
Storage Location	Cloud	Cloud/local	Local	Local & Cloud
Privacy	△ External	△ Depends	✓ Self-hosted	✓ 100%
Rust Support	△ Limited	△ HTTP API	△ Python-first	✓ Native
PPP Compliance	10%	15%	20%	100%

Best For	LangChain apps	Multi- provider apps	Research/analysis	PPP CI tools

Winner: Proposed (SQLite) - Only solution meeting all PPP requirements with zero external dependencies.

Async SQLite Crate Comparison

Feature	tokio- rusqlite	async- sqlite	sqlx	rusqlite (sync)
Maturity	✓ Mature	⚠ Young	✓ Very mature	✓ Very mature
Downloads/Month	23,000	<1,000	500,000+	1,000,000+
License	MIT	MIT	MIT/Apache 2.0	MIT
Async Runtime	tokio only	tokio + async_std	tokio + async_std	N/A (sync)
API Style	Callback-based	Direct async	SQL builder	Direct sync
Compile-time Checks	✗ None	✗ None	✓ Full	✗ None
Connection Pooling	⚠ Manual	✗ None	✓ Built-in	⚠ Manual
Migrations	✗ None	✗ None	✓ Built-in	✗ None
Binary Size	+50KB	+50KB	+2MB	+500KB
Latency (single op)	~0.5ms	~0.5ms	~0.3ms	~0.1ms
Latency (batched)	~0.1ms	~0.1ms	~0.05ms	~0.02ms
Learning Curve	⚠ Medium	✓ Low	⚠ Medium	✓ Low
Thread Safety	✓ Yes	✓ Yes	✓ Yes	⚠ Manual
Error Handling	✓ Good	⚠ Basic	✓ Excellent	✓ Good
Documentation	✓ Good	⚠ Limited	✓ Excellent	✓ Excellent
Community	⚠ Medium	⚠ Small	✓ Large	✓ Very large
Best For	Async logging	Simple cases	Full apps	Sync batch ops

Recommendation: **tokio-rusqlite** for Phase 1-2 (proven, async-friendly), consider **sqlx** for Phase 3 if migrations needed.

Database Schema Pattern Comparison

Query

Pattern	Tables	Complexity	Storage	Extensibility	S
Flat					
Message History	1	✓ Low	✓ Compact	✗ Poor	✗
Turn-Based (Proposed)	4	△ Medium (JOINs)	△ Normalized	✓ Excellent	✓
JSON Blob	1	✗ High (no index)	✗ Large	△ Good	△
Hierarchical Spans	2	✗ Very high (CTE)	△ Medium	✓ Excellent	△

Details:

1. Flat Message History

Schema:

```
CREATE TABLE messages (
    id INTEGER PRIMARY KEY,
    session_id TEXT,
    timestamp TEXT,
    role TEXT,
    content TEXT
);
```

Pros: - Simplest schema (1 table) - Easy chronological queries - Minimal storage overhead

Cons: - No turn grouping (prompt + response = 2 rows) - Can't store metadata (effort, violations) - Hard to calculate R_{Proact} (questions not tracked)

PPP Compliance: 20% (basic conversation history only)

2. Turn-Based with Metadata (Proposed)

Schema:

```
CREATE TABLE trajectories (
    id INTEGER PRIMARY KEY,
    spec_id TEXT NOT NULL,
    agent_name TEXT NOT NULL,
    run_id TEXT,
    created_at TEXT DEFAULT (datetime('now')),
    INDEX idx_spec_agent (spec_id, agent_name)
);

CREATE TABLE trajectory_turns (
    id INTEGER PRIMARY KEY,
    trajectory_id INTEGER NOT NULL,
    turn_number INTEGER NOT NULL,
    prompt TEXT NOT NULL,
    response TEXT NOT NULL,
    token_count INTEGER,
    latency_ms INTEGER,
```

```

    FOREIGN KEY (trajectory_id) REFERENCES trajectories(id) ON
DELETE CASCADE
);

CREATE TABLE trajectory_questions (
    id INTEGER PRIMARY KEY,
    turn_id INTEGER NOT NULL,
    question_text TEXT NOT NULL,
    effort_level TEXT, -- 'low', 'medium', 'high'
    FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id) ON DELETE
CASCADE
);

CREATE TABLE trajectory_violations (
    id INTEGER PRIMARY KEY,
    turn_id INTEGER NOT NULL,
    preference_name TEXT NOT NULL,
    expected TEXT NOT NULL,
    actual TEXT NOT NULL,
    severity TEXT NOT NULL,
    FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id) ON DELETE
CASCADE
);

```

Pros: - Turn-level granularity (prompt + response = 1 row)
- Normalized (no duplication)
- Supports all PPP calculations
- Extensible (add tables without migration)

Cons: - More complex queries (requires JOINS)
- 4 tables to manage
- Slightly higher write overhead

PPP Compliance: 100% (supports all 20 preferences + formulas)

Example Query (R_{Proact} calculation):

```

SELECT
CASE
    WHEN COUNT(q.id) = 0 THEN 0.05
    WHEN COUNT(CASE WHEN q.effort_level != 'low' THEN 1 END) = 0
THEN 0.05
    ELSE -0.1 * COUNT(CASE WHEN q.effort_level = 'medium' THEN 1
END)
        -0.5 * COUNT(CASE WHEN q.effort_level = 'high' THEN 1
END)
    END AS r_proact
FROM trajectory_turns t
LEFT JOIN trajectory_questions q ON t.id = q.turn_id
WHERE t.trajectory_id = ?;

```

3. JSON Blob Storage

Schema:

```

CREATE TABLE trajectories (
    id INTEGER PRIMARY KEY,
    spec_id TEXT,
    agent_name TEXT,
    turns_json TEXT -- JSON array: [{prompt, response, questions,
violations}, ...]

```

```
);
```

Example JSON:

```
{  
    "turns": [  
        {  
            "prompt": "Implement OAuth",  
            "response": "Which provider?",  
            "questions": [{"text": "Which provider?", "effort": "low"}],  
            "violations": []  
        }  
    ]  
}
```

Pros: - Flexible schema (no migrations) - Single INSERT per trajectory - Easy to serialize/deserialize

Cons: - SQLite JSON support limited (no indexing in <3.38) - Large storage footprint (JSON overhead ~30%) - Can't enforce foreign keys - Hard to query (requires json_extract())

PPP Compliance: 50% (can store data, but hard to calculate scores)

4. Hierarchical with Parent/Child Spans (OpenTelemetry)

Schema:

```
CREATE TABLE spans (  
    id INTEGER PRIMARY KEY,  
    parent_id INTEGER,  
    span_type TEXT, -- 'conversation', 'turn', 'llm_call',  
    'tool_use'  
    start_time TEXT,  
    end_time TEXT,  
    attributes_json TEXT,  
    FOREIGN KEY (parent_id) REFERENCES spans(id)  
);
```

Example Hierarchy:

```
Span (type: conversation, id: 1)  
  └── Span (type: turn, id: 2, parent: 1)  
    └── Span (type: llm_call, id: 3, parent: 2)  
      └── Span (type: tool_use, id: 4, parent: 2)  
    └── Span (type: turn, id: 5, parent: 1)
```

Pros: - Industry standard (OpenTelemetry) - Supports nested operations (tool calls, retries) - Migration path to observability platforms

Cons: - Overkill for CLI (no nested tools) - Complex queries (recursive CTEs) - Higher storage overhead (~2x)

PPP Compliance: 60% (can model turns, but not optimized for PPP formulas)

Example Query (recursive CTE to get all turns):

```

WITH RECURSIVE turn_hierarchy AS (
    SELECT * FROM spans WHERE span_type = 'conversation' AND id = ?
    UNION ALL
    SELECT s.* FROM spans s
    JOIN turn_hierarchy th ON s.parent_id = th.id
)
SELECT * FROM turn_hierarchy WHERE span_type = 'turn';

```

Performance Strategy Comparison

Strategy	Latency (per turn)	Throughput	Complexity	Concurrency
Sync (single)	~1ms	1,000/sec	✓ Low	✗ Blocks
Sync (batched)	~0.067ms	15,000/sec	△ Medium	✗ Blocks
Async (single)	~0.5ms	2,000/sec	△ Medium	✓ Good
Async (batched)	~0.1ms	10,000/sec	△ Medium	✓ Excellent
Async (WAL mode)	~0.02ms	50,000/sec	✗ High	✓ Excellent

Details:

1. Synchronous Single Insert

Implementation:

```

use rusqlite::Connection;

let conn = Connection::open("trajectory.db")?;
conn.execute(
    "INSERT INTO trajectory_turns (trajectory_id, prompt, response)
VALUES (?, ?, ?)",
    params![traj_id, prompt, response],
)?;

```

Performance: - Latency: ~1ms per turn - Throughput: ~1,000 turns/sec - Blocking: Yes (blocks tokio runtime)

Verdict: ✗ **Avoid** - Blocks async runtime, too slow for production.

2. Synchronous Batched Insert

Implementation:

```

let tx = conn.transaction()?;
for turn in turns {
    tx.execute(
        "INSERT INTO trajectory_turns (trajectory_id, prompt,

```

```

        response) VALUES (?1, ?2, ?3)",
        params![turn.traj_id, turn.prompt, turn.response],
    )?;
}
tx.commit()?;

```

Performance: - Latency: ~0.067ms per turn (15,000/sec) - Throughput: 15,000 turns/sec - Blocking: Yes (but brief)

Verdict: △ Acceptable - Good for background batch jobs, but still blocks.

3. Asynchronous Single Insert

Implementation:

```

use tokio_rusqlite::Connection;

let conn = Connection::open("trajectory.db").await?;
conn.call(|conn| {
    conn.execute(
        "INSERT INTO trajectory_turns (trajectory_id, prompt,
response) VALUES (?1, ?2, ?3)",
        params![traj_id, prompt, response],
    )
}).await?;

```

Performance: - Latency: ~0.5ms per turn (thread pool overhead) - Throughput: ~2,000 turns/sec - Blocking: No (async)

Verdict: △ Acceptable - Good for low-volume real-time logging.

4. Asynchronous Batched Insert (RECOMMENDED)

Implementation:

```

use tokio::sync::mpsc;
use tokio_rusqlite::Connection;

// Channel for buffered writes
let (tx, mut rx) = mpsc::channel::<Turn>(100);

// Background writer task
tokio::spawn(async move {
    let conn = Connection::open("trajectory.db").await?;
    let mut buffer = Vec::new();

    loop {
        tokio::select! {
            Some(turn) = rx.recv() => {
                buffer.push(turn);
                if buffer.len() >= 10 {
                    flush_batch(&conn, &buffer).await?;
                    buffer.clear();
                }
            }
            _ = tokio::time::sleep(Duration::from_millis(500)) => {
                if !buffer.is_empty() {

```

```

        flush_batch(&conn, &buffer).await?;
        buffer.clear();
    }
}
});
});

async fn flush_batch(conn: &Connection, turns: &[Turn]) ->
Result<()> {
    conn.call(move |conn| {
        let tx = conn.transaction()?;
        for turn in turns {
            tx.execute(
                "INSERT INTO trajectory_turns (trajectory_id,
prompt, response) VALUES (?1, ?2, ?3)",
                params![turn.traj_id, turn.prompt, turn.response],
            )?;
        }
        tx.commit()
    }).await
}

```

Performance: - Latency: ~0.1ms per turn (amortized) - Throughput: ~10,000 turns/sec - Blocking: No (async + buffered) - Lock contention: Minimal (batched writes)

Verdict: ✓ **RECOMMENDED** - Best balance of latency, throughput, and concurrency.

5. Asynchronous with WAL Mode

Implementation:

```

let conn = Connection::open("trajectory.db").await?;
conn.call(|conn| {
    conn.execute("PRAGMA journal_mode=WAL", [])?;
    conn.execute("PRAGMA synchronous=NORMAL", [])?;
    Ok(())
}).await?;

```

Performance: - Latency: ~0.02ms per turn - Throughput: 50,000+ turns/sec - Blocking: No (WAL allows concurrent reads)

Pros: - Fastest option (10x faster than batched) - Concurrent reads during writes - No lock contention

Cons: - More complex (WAL files to manage) - Slightly larger disk footprint (~30%)

Verdict: △ **Future optimization** - Use if Phase 1-2 benchmarks show bottleneck.

Integration Strategy Comparison

Approach	Complexity	Reuse	Migration	Latency	B F
----------	------------	-------	-----------	---------	--------

Extend consensus_db.rs	✓ Low	✓ High	✓ Easy	<1ms	Ph 1-2
Separate Database	△ Medium	△ Low	△ Manual	<1ms	Isol nee
MCP Server	✗ High	✗ None	✗ Hard	~5ms	Ext tool

Details:

1. Extend consensus_db.rs (RECOMMENDED)

Approach: Add trajectory tables to existing SQLite database managed by consensus_db.rs.

Implementation:

```
// codex-rs/tui/src/chatwidget/spec_kit/consensus_db.rs

pub fn init_trajectory_tables(conn: &Connection) -> Result<()> {
    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectories (
            id INTEGER PRIMARY KEY AUTOINCREMENT,
            spec_id TEXT NOT NULL,
            agent_name TEXT NOT NULL,
            run_id TEXT,
            created_at TEXT NOT NULL DEFAULT (datetime('now'))
        )",
        [],
    )?;

    conn.execute(
        "CREATE INDEX IF NOT EXISTS idx_trajectories_spec_agent
         ON trajectories(spec_id, agent_name)",
        [],
    )?;

    // ... other tables
}

Ok(())
}
```

Pros: - ✓ Reuses existing connection pool - ✓ Single database file (simpler deployment) - ✓ Can JOIN with consensus_artifacts (via run_id) - ✓ Minimal code changes

Cons: - △ Couples trajectory logging to consensus DB - △ Migrations must coordinate with existing schema

Verdict: ✓ RECOMMENDED - Simplest, fastest, most maintainable.

2. Separate Database

Approach: Create new trajectory.db separate from consensus.db.

Implementation:

```
// New module: codex-rs/tui/src/chatwidget/spec_kit/trajectory_db.rs
```

```

pub struct TrajectoryDb {
    conn: Arc<Mutex<Connection>>,
}

impl TrajectoryDb {
    pub fn open() -> Result<Self> {
        let path = dirs::data_dir()
            .unwrap()
            .join("codex-tui")
            .join("trajectory.db");
        let conn = Connection::open(path)?;
        init_tables(&conn)?;
        Ok(Self { conn: Arc::new(Mutex::new(conn)) })
    }
}

```

Pros: - ✓ Clean separation of concerns - ✓ Independent migrations - ✓ Can be disabled without affecting consensus

Cons: - ✗ Duplicate connection management - ✗ Can't JOIN across databases (require application-level merge) - ✗ More files to manage

Verdict: △ **Alternative** - Use only if trajectory logging needs to be independently deployable.

3. MCP Server (NOT RECOMMENDED)

Approach: Create new MCP server for trajectory logging (similar to consensus MCP).

Implementation:

```

// New crate: codex-rs/mcp-servers/mcp-trajectory-logger

#[mcp_tool]
async fn log_turn(
    trajectory_id: i64,
    prompt: String,
    response: String,
) -> Result<()> {
    // ... SQLite insert
}

```

Pros: - ✓ Reusable by other tools (not just codex-tui) - ✓ Standard MCP interface

Cons: - ✗ 5x slower (~5ms vs <1ms) - ✗ More complex (JSON-RPC overhead) - ✗ Requires MCP connection management - ✗ No benefit for CLI use case

Verdict: ✗ **NOT RECOMMENDED** - Unnecessary overhead for local logging.

Cost Analysis

Infrastructure Costs

Approach	Setup Cost	Monthly Cost	Storage Cost	Latency Cost
LangSmith	0 hours	\$39+	N/A (cloud)	High (50-100ms)
Helicone (cloud)	0 hours	\$0 (OSS)	N/A (cloud)	Medium (10-20ms)
Helicone (self-hosted)	4 hours	\$5 (VPS)	<1GB/year	Medium (10-20ms)
Phoenix (self-hosted)	4 hours	\$5 (VPS)	<1GB/year	Medium (30-50ms)
SQLite (local)	2 hours	\$0	<100MB/year	Low (<1ms)

Assumptions (1000 sessions/month, avg 20 turns/session): - Storage: 500 bytes/turn = 10MB/month = 120MB/year - Compute: Local SQLite uses existing machine (no incremental cost)

Winner: SQLite (local) - Zero ongoing cost, minimal storage.

Development Effort Costs

Task	Complexity	Estimated Effort	Dependencies
1. Schema Design	△ MEDIUM	4 hours	None
2. Extend consensus_db.rs	✓ LOW	2 hours	consensus_db.rs
3. Async Logging API	△ MEDIUM	6 hours	tokio-rusqlite
4. Question Effort Classifier	✗ HIGH	8 hours	SPEC-PPP-001
5. Preference Violation Detector	△ MEDIUM	6 hours	SPEC-PPP-002
6. R_{Proact} Calculator	△ MEDIUM	4 hours	SPEC-PPP-003
7. R_{Pers} Calculator	△ MEDIUM	4 hours	SPEC-PPP-002, 003
8. Integration Tests	△ MEDIUM	4 hours	Test harness
9. Documentation	✓ LOW	2 hours	None

Total Effort: ~40 hours (~1 week for 1 engineer)

Phase 1 Target (Schema + basic logging): ~12 hours (1.5 days)

Performance Benchmarks (Estimated)

Component	Operation	Latency	Throughput	Bottleneck
Schema Creation	Init tables	<1ms	N/A	Disk I/O
Insert Turn (sync)	Single INSERT	1ms	1,000/sec	Lock contention
Insert Turn (async)	Single INSERT	0.5ms	2,000/sec	Thread pool
Insert Turn (batched)	Batched INSERT	0.1ms	10,000/sec	Minimal
Query Trajectory	SELECT with JOINs	2-5ms	200/sec	JOIN complexity
Calculate R_{Proact}	Aggregate query	3-8ms	100/sec	COUNT + CASE
Calculate R_{Pers}	Aggregate query	3-8ms	100/sec	COUNT + violations
Full Weighted Score	Multi-query	10-20ms	50/sec	Multiple JOINs

Expected Overhead (per agent execution): - Logging: ~0.1ms/turn × 20 turns = 2ms - Scoring: ~20ms (once at end) - **Total:** ~22ms (0.1% of typical agent execution time)

Recommendations Summary

Decision	Recommended Option	Alternative	Rationale
Observability Platform	Local SQLite	Phoenix Arize (self-hosted)	Zero cost, <1ms overhead, 100% privacy
Async SQLite Crate	tokio-rusqlite	sqlx	Proven (23k/mo), async-friendly, lightweight
Schema Pattern	Turn-based (4 tables)	JSON blob	Normalized, supports PPP formulas, indexable
Performance Strategy	Async batched	WAL mode	0.1ms/turn, 10K/sec throughput, simple
Integration	Extend consensus_db.rs	Separate DB	Reuses connection, single file,

Buffering	5-10 turns, 500ms flush	Immediate write	simplest Reduces lock contention, <1ms amortized
Indexes	spec_id, agent_name, trajectory_id	None	Fast lookups for scoring queries
WAL Mode	Phase 3 optimization	Default	Defer until benchmarks show bottleneck

Next Steps

1. **Validate schema design** with trajectory examples from real agent runs
2. **Prototype** async batching with tokio-rusqlite (benchmark actual latency)
3. **Benchmark** scoring queries with 1K+ trajectories (validate <20ms target)
4. **Phase 1 implementation:** Schema + basic logging (defer scoring to Phase 2)
5. **Phase 2 integration:** Connect to SPEC-PPP-003 weighted consensus