

recommendations

SPEC-PPP-004: Implementation Recommendations

Last Updated: 2025-11-16 **Target:** Trajectory logging infrastructure for PPP framework **Total Effort:** ~40 hours (1 week, 1 engineer)

Executive Summary

Implement PPP trajectory logging using **local SQLite extension** (NOT MCP server) with **async batching** for <1ms overhead. Extend existing consensus_db.rs to add 4 new tables tracking multi-turn conversations, question effort, and preference violations.

Key Decisions: 1. ✓ SQLite (NOT MCP server) - 5x faster, zero cost, simpler deployment 2. ✓ tokio-rusqlite - Proven async SQLite library (23k downloads/month) 3. ✓ Async batching - Buffer 5-10 turns, flush every 500ms (0.1ms/turn overhead) 4. ✓ Extend consensus_db.rs - Reuse existing connection, single database file 5. ✓ Turn-based schema - Normalized 4-table design supporting PPP formulas

Phased Rollout

Phase 1: Foundation (12 hours, 1.5 days)

Goal: Basic trajectory logging without scoring

Deliverables: - ✓ SQLite schema (4 tables: trajectories, turns, questions, violations) - ✓ Async logging API with batching - ✓ Integration with existing consensus_db.rs - ✓ Unit tests (CRUD operations)

Capabilities: - Log multi-turn conversations - Track question effort (manual classification) - Record preference violations - Query trajectories by spec_id, agent_name

NOT Included (defer to Phase 2): - Automatic question effort classification - R_{Proact} and R_{Pers} calculation - Weighted consensus integration

Phase 2: Scoring Integration (16 hours, 2 days)

Goal: Enable PPP interaction scoring

Deliverables: - ✓ R_{Proact} calculation queries - ✓ R_{Pers} calculation queries - ✓ Integration with SPEC-PPP-003 weighted consensus - ✓ Question effort classifier (heuristic-based, 75-85% accuracy) - ✓ Integration tests (full scoring pipeline)

Capabilities: - Calculate proactivity scores from trajectories - Calculate personalization scores - Use scores in weighted consensus - Automatic question effort detection (basic)

NOT Included (defer to Phase 3): - LLM-based question effort classification (90-95% accuracy) - Observability platform export (OpenTelemetry) - WAL mode optimization

Phase 3: Production Hardening (12 hours, 1.5 days)

Goal: Optimize for production workloads

Deliverables: - ✓ WAL mode for concurrent reads - ✓ LLM-based question effort classifier (optional upgrade) - ✓ Archival/cleanup strategy (auto-delete old trajectories) - ✓ Observability export (OpenTelemetry spans, if needed) - ✓ Load testing (10K+ trajectories) - ✓ Performance benchmarks

Capabilities: - <0.02ms/turn latency (WAL mode) - 50,000+ turns/sec throughput - Automatic cleanup (keep last 30 days) - Export to Phoenix Arize (if desired)

Detailed Task Breakdown

Task 1: Database Schema Design (4 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/consensus_db.rs

Implementation:

```
// Add to existing init_database() function

pub fn init_trajectory_tables(conn: &Connection) -> Result<()> {
    // Table 1: Trajectories (one per agent execution)
    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectories (
            id INTEGER PRIMARY KEY AUTOINCREMENT,
            spec_id TEXT NOT NULL,
            agent_name TEXT NOT NULL,
            run_id TEXT, -- Links to consensus_artifacts.run_id
            created_at TEXT NOT NULL DEFAULT (datetime('now')),
            INDEX idx_trajectories_spec_agent (spec_id, agent_name),
            INDEX idx_trajectories_run_id (run_id)
        )",
        [],
    )?;

    // Table 2: Trajectory Turns (one per user-agent exchange)
    conn.execute(
        "CREATE TABLE IF NOT EXISTS trajectory_turns (
            id INTEGER PRIMARY KEY AUTOINCREMENT,

```

```

        trajectory_id INTEGER NOT NULL,
        turn_number INTEGER NOT NULL,
        prompt TEXT NOT NULL,
        response TEXT NOT NULL,
        token_count INTEGER,
        latency_ms INTEGER,
        timestamp TEXT NOT NULL DEFAULT (datetime('now')),
        FOREIGN KEY (trajectory_id) REFERENCES trajectories(id)
    ON DELETE CASCADE,
        INDEX idx_turns_trajectory (trajectory_id),
        UNIQUE (trajectory_id, turn_number)
    )",
    [],
);
)?;

// Table 3: Questions Asked (extracted from responses)
conn.execute(
    "CREATE TABLE IF NOT EXISTS trajectory_questions (
        id INTEGER PRIMARY KEY AUTOINCREMENT,
        turn_id INTEGER NOT NULL,
        question_text TEXT NOT NULL,
        effort_level TEXT, -- 'low', 'medium', 'high'
        FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id) ON
DELETE CASCADE,
        INDEX idx_questions_turn (turn_id)
    )",
    [],
);
)?;

// Table 4: Preference Violations (detected in responses)
conn.execute(
    "CREATE TABLE IF NOT EXISTS trajectoryViolations (
        id INTEGER PRIMARY KEY AUTOINCREMENT,
        turn_id INTEGER NOT NULL,
        preference_name TEXT NOT NULL, -- e.g., 'no_ask',
'require_json'
        expected TEXT NOT NULL, -- e.g., 'no questions',
'valid JSON'
        actual TEXT NOT NULL, -- e.g., 'asked
question', 'malformed JSON'
        severity TEXT NOT NULL, -- 'minor', 'major',
'critical'
        FOREIGN KEY (turn_id) REFERENCES trajectory_turns(id) ON
DELETE CASCADE,
        INDEX idx_violations_turn (turn_id),
        INDEX idx_violations_preference (preference_name)
    )",
    [],
);
)?;

Ok(())
}

```

Migration:

```

pub fn migrate_database(conn: &Connection) -> Result<()> {
    let version: i32 = conn.query_row(
        "PRAGMA user_version",
        [],
        |row| row.get(0)
    ).unwrap_or(0);

```

```

    if version < 2 {
        init_trajectory_tables(conn)?;
        conn.execute("PRAGMA user_version = 2", [])?;
    }

    Ok(())
}

```

Testing:

```

#[cfg(test)]
mod tests {
    #[test]
    fn test_trajectory_schema() {
        let conn = Connection::open_in_memory().unwrap();
        init_trajectory_tables(&conn).unwrap();

        // Verify tables exist
        let count: i32 = conn.query_row(
            "SELECT COUNT(*) FROM sqlite_master WHERE type='table'
AND name LIKE 'trajectory%'",

            [],
            |row| row.get(0)
        ).unwrap();
        assert_eq!(count, 4);
    }
}

```

Acceptance Criteria: - [] Schema creates 4 tables with foreign keys -
[] Indexes created on spec_id, agent_name, run_id, trajectory_id - []
Migration from version 1 to 2 succeeds - [] All tests pass

Task 2: Async Logging API with Batching (6 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/trajectory_logger.rs (new)

Implementation:

```

use tokio::sync::mpsc;
use tokio_rusqlite::Connection;
use std::time::Duration;

#[derive(Debug, Clone)]
pub struct Turn {
    pub trajectory_id: i64,
    pub turn_number: i32,
    pub prompt: String,
    pub response: String,
    pub token_count: Option<i32>,
    pub latency_ms: Option<i32>,
}

#[derive(Debug, Clone)]
pub struct Question {
    pub turn_id: i64,
    pub text: String,
    pub effort_level: EffortLevel,
}

```

```

#[derive(Debug, Clone, Copy, PartialEq)]
pub enum EffortLevel {
    Low,      // Selection questions, accessible context
    Medium,   // Some research needed
    High,     // Blocking, deep investigation
}

pub struct TrajectoryLogger {
    tx: mpsc::Sender<LogEntry>,
}

enum LogEntry {
    Turn(Turn),
    Question(Question),
    Violation(Violation),
}

impl TrajectoryLogger {
    pub fn new(db_path: String) -> Self {
        let (tx, rx) = mpsc::channel(100);

        // Spawn background writer task
        tokio::spawn(async move {
            if let Err(e) = writer_task(db_path, rx).await {
                eprintln!("Trajectory logger error: {}", e);
            }
        });
    }

    Self { tx }
}

pub async fn log_turn(&self, turn: Turn) -> Result<()> {
    self.tx.send(LogEntry::Turn(turn)).await
        .map_err(|e| anyhow!("Failed to send turn: {}", e))
}

pub async fn log_question(&self, question: Question) ->
Result<()> {
    self.tx.send(LogEntry::Question(question)).await
        .map_err(|e| anyhow!("Failed to send question: {}", e))
}

async fn writer_task(
    db_path: String,
    mut rx: mpsc::Receiver<LogEntry>,
) -> Result<()> {
    let conn = Connection::open(&db_path).await?;
    let mut buffer = Vec::new();

    loop {
        tokio::select! {
            Some(entry) = rx.recv() => {
                buffer.push(entry);

                // Flush if buffer full
                if buffer.len() >= 10 {
                    flush_buffer(&conn, &buffer).await?;
                    buffer.clear();
                }
            }
        }
    }
}

```

```

        }
    }

    // Flush every 500ms
    _ = tokio::time::sleep(Duration::from_millis(500)) => {
        if !buffer.is_empty() {
            flush_buffer(&conn, &buffer).await?;
            buffer.clear();
        }
    }
}

async fn flush_buffer(
    conn: &Connection,
    buffer: &[LogEntry],
) -> Result<()> {
    conn.call(move |conn| {
        let tx = conn.transaction()?;
        for entry in buffer {
            match entry {
                LogEntry::Turn(turn) => {
                    tx.execute(
                        "INSERT INTO trajectory_turns
                            (trajectory_id, turn_number, prompt,
                            response, token_count, latency_ms)
                            VALUES (?1, ?2, ?3, ?4, ?5, ?6)",
                        params![
                            turn.trajectory_id,
                            turn.turn_number,
                            turn.prompt,
                            turn.response,
                            turn.token_count,
                            turn.latency_ms,
                        ],
                    )?;
                }
                LogEntry::Question(q) => {
                    tx.execute(
                        "INSERT INTO trajectory_questions
                            (turn_id, question_text, effort_level)
                            VALUES (?1, ?2, ?3)",
                        params![
                            q.turn_id,
                            q.text,
                            q.effort_level.to_string(),
                        ],
                    )?;
                }
            }
            // ... handle violations
        }
        tx.commit()
    }).await
}

```

Usage Example:

```

// In consensus.rs:220-249

let logger = TrajectoryLogger::new(db_path.clone());

// Start trajectory
let trajectory_id = create_trajectory(spec_id, agent_name, run_id)?;

// Log each turn
for (i, (prompt, response)) in conversation.iter().enumerate() {
    logger.log_turn(Turn {
        trajectory_id,
        turn_number: i as i32,
        prompt: prompt.clone(),
        response: response.clone(),
        token_count: Some(count_tokens(response)),
        latency_ms: Some(latency),
    }).await?;
}

// Extract and log questions
let questions = extract_questions(response);
for q in questions {
    logger.log_question(Question {
        turn_id: get_last_turn_id(trajectory_id, i)?,
        text: q,
        effort_level: classify_effort(&q), // Heuristic
    }).await?;
}
}

```

Acceptance Criteria: - [] Async API with <1ms latency (buffered) - [] Batches 5-10 entries, flushes every 500ms - [] No blocking of tokio runtime - [] Handles backpressure (bounded channel) - [] Unit tests for buffering logic

Task 3: Question Extraction & Effort Classification (8 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/question_classifier.rs
(new)

Heuristic Implementation (Phase 1):

```

use regex::Regex;

pub fn extract_questions(text: &str) -> Vec<String> {
    let question_re = Regex::new(r"(?m)^.*\?$").unwrap();
    question_re.find_iter(text)
        .map(|m| m.as_str().trim().to_string())
        .collect()
}

pub fn classify_effort(question: &str) -> EffortLevel {
    let question_lower = question.to_lowercase();

    // Low effort indicators
    let low_indicators = [
        "which",           // Selection: "Which option?"
        "do you prefer", // Preference: "Do you prefer A or B?"
        "would you like",
    ];
}

```

```

        "choose",
        "select",
        "pick",
    ];
}

// High effort indicators
let high_indicators = [
    "investigate", // Deep research: "Should I investigate X?"
    "research",
    "analyze",
    "determine", // Blocking: "How should I determine...?"
    "what is the best way to",
    "how should i",
    "before proceeding",
];
}

// Check high effort first (more specific)
for indicator in &high_indicators {
    if question_lower.contains(indicator) {
        return EffortLevel::High;
    }
}

// Check low effort
for indicator in &low_indicators {
    if question_lower.contains(indicator) {
        return EffortLevel::Low;
    }
}

// Default to medium
EffortLevel::Medium
}

```

LLM-Based Implementation (Phase 3 upgrade):

```

pub async fn classify_effort_llm(question: &str) ->
Result<EffortLevel> {
    let prompt = format!(
        "Classify the following question by effort level required
from the user:

Question: {}

Effort levels:
- Low: Selection from options, accessible context (e.g.,
'Which provider?')
- Medium: Some research needed, not blocking (e.g., 'What
format do you prefer?')
- High: Deep investigation or blocking decision (e.g.,
'Should I investigate caching strategies?')

Respond with ONLY one word: Low, Medium, or High",
question
);

    let response = call_llmCheap(prompt).await?; // Use
haiku/flash for cost

    match response.trim().to_lowercase().as_str() {
        "low" => Ok(EffortLevel::Low),

```

```

        "medium" => Ok(EffortLevel::Medium),
        "high" => Ok(EffortLevel::High),
        _ => Ok(EffortLevel::Medium), // Fallback
    }
}

```

Benchmark (Phase 1 vs Phase 3):

Method	Accuracy	Latency	Cost/1000
Heuristic	75-85%	<0.1ms	\$0
LLM (Haiku)	90-95%	~200ms	~\$0.05

Acceptance Criteria: - [] Heuristic classifier: 75%+ accuracy on test set (SPEC-PPP-001) - [] Extracts questions from agent responses - [] Handles edge cases (rhetorical questions, multi-sentence) - [] Unit tests with labeled examples

Task 4: *R_{Proact}* Calculation Queries (4 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/PPP_scoring.rs (new)

Implementation:

```

use rusqlite::{params, Connection};

pub struct ProactivityScore {
    pub r_proact: f32,
    pub total_questions: usize,
    pub low_effort: usize,
    pub medium_effort: usize,
    pub high_effort: usize,
}

pub fn calculate_r_proact(
    conn: &Connection,
    trajectory_id: i64,
) -> Result<ProactivityScore> {
    let query = "
        SELECT
            COUNT(q.id) as total_questions,
            SUM(CASE WHEN q.effort_level = 'low' THEN 1 ELSE 0 END)
        as low_effort,
            SUM(CASE WHEN q.effort_level = 'medium' THEN 1 ELSE 0
        END) as medium_effort,
            SUM(CASE WHEN q.effort_level = 'high' THEN 1 ELSE 0 END)
        as high_effort
        FROM trajectory_turns t
        LEFT JOIN trajectory_questions q ON t.id = q.turn_id
        WHERE t.trajectory_id = ?
    ";
    let mut stmt = conn.prepare(query)?;
    let row = stmt.query_row(params![trajectory_id], |row| {
        Ok((
            row.get::<_, i32>(0)?,
            row.get::<_, i32>(1)?,
            row.get::<_, i32>(2)?,
            row.get::<_, i32>(3)?,
        )
    })
}
```

```

        ))
})?;

let (total, low, medium, high) = row;

// Apply PPP formula
let r_proact = if total == 0 {
    0.05 // No questions asked
} else if low == total {
    0.05 // All questions low-effort
} else {
    -0.1 * (medium as f32) - 0.5 * (high as f32)
};

Ok(ProactivityScore {
    r_proact,
    total_questions: total as usize,
    low_effort: low as usize,
    medium_effort: medium as usize,
    high_effort: high as usize,
})
}
}

```

Testing:

```

#[test]
fn test_r_proact_calculation() {
    let conn = Connection::open_in_memory().unwrap();
    init_trajectory_tables(&conn).unwrap();

    // Scenario 1: No questions → +0.05
    let traj_id = insert_test_trajectory(&conn, "SPEC-001",
                                         "agent1");
    let score = calculate_r_proact(&conn, traj_id).unwrap();
    assert_eq!(score.r_proact, 0.05);

    // Scenario 2: 2 low effort → +0.05
    insert_test_question(&conn, traj_id, "Which?", EffortLevel::Low);
    insert_test_question(&conn, traj_id, "Pick one?", EffortLevel::Low);
    let score = calculate_r_proact(&conn, traj_id).unwrap();
    assert_eq!(score.r_proact, 0.05);

    // Scenario 3: 1 high effort → -0.5
    insert_test_question(&conn, traj_id, "Investigate?", EffortLevel::High);
    let score = calculate_r_proact(&conn, traj_id).unwrap();
    assert_eq!(score.r_proact, -0.5);
}

```

Acceptance Criteria: - [] Correct calculation for all 3 scenarios (none, all-low, mixed) - [] Query executes in <10ms for 100-turn trajectory - [] Unit tests cover edge cases

Task 5: R_{Pers} Calculation Queries (4 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/ppp_scoring.rs

Implementation:

```
pub struct PersonalizationScore {
    pub r_pers: f32,
    pub total_violations: usize,
    pub minor_violations: usize,
    pub major_violations: usize,
    pub critical_violations: usize,
}

pub fn calculate_r_pers(
    conn: &Connection,
    trajectory_id: i64,
) -> Result<PersonalizationScore> {
    let query = "
        SELECT
            COUNT(v.id) as total_violations,
            SUM(CASE WHEN v.severity = 'minor' THEN 1 ELSE 0 END) as
        minor,
            SUM(CASE WHEN v.severity = 'major' THEN 1 ELSE 0 END) as
        major,
            SUM(CASE WHEN v.severity = 'critical' THEN 1 ELSE 0 END)
        as critical
        FROM trajectory_turns t
        LEFT JOIN trajectoryViolations v ON t.id = v.turn_id
        WHERE t.trajectory_id = ?
    ";
}

let mut stmt = conn.prepare(query)?;
let row = stmt.query_row(params![trajectory_id], |row| {
    Ok((
        row.get::<_, i32>(0)?,
        row.get::<_, i32>(1)?,
        row.get::<_, i32>(2)?,
        row.get::<_, i32>(3)?,
    ))
})?;

let (total, minor, major, critical) = row;

// Apply PPP formula
let r_pers = if total == 0 {
    0.05 // No violations → full compliance
} else {
    -0.01 * (minor as f32) - 0.03 * (major as f32) - 0.05 *
(critical as f32)
};

Ok(PersonalizationScore {
    r_pers,
    total_violations: total as usize,
    minor_violations: minor as usize,
    major_violations: major as usize,
    critical_violations: critical as usize,
})
}
```

Acceptance Criteria: - [] Correct calculation for violation scenarios -
[] Query executes in <10ms - [] Unit tests for edge cases

Task 6: Weighted Consensus Integration (6 hours)

File: codex-rs/tui/src/chatwidget/spec_kit/consensus.rs:681-958

Refactor run_spec_consensus():

```
pub async fn run_spec_consensus_weighted(
    spec_id: &str,
    stage: &str,
    artifacts: Vec<ConsensusArtifactData>,
    weights: (f32, f32), // (technical, interaction) - default
    (0.7, 0.3)
) -> Result<WeightedConsensus> {
    let db = open_consensus_db()?;
    let mut scores = Vec::new();

    for artifact in artifacts {
        // Technical score (existing logic)
        let technical = calculate_technical_score(&artifact)?;

        // Interaction score (new: from trajectory)
        let trajectory_id = get_trajectory_id(&db, spec_id,
        &artifact.agent)?;
        let proact = calculate_r_proact(&db, trajectory_id)?;
        let pers = calculate_r_pers(&db, trajectory_id)?;
        let interaction = proact.r_proact + pers.r_pers;

        // Weighted combination
        let (w_tech, w_interact) = weights;
        let final_score = (w_tech * technical) + (w_interact * interaction);

        scores.push(AgentScore {
            agent_name: artifact.agent.clone(),
            technical_score: technical,
            interaction_score: interaction,
            final_score,
            details: ScoreDetails {
                proactivity: proact,
                personalization: pers,
            },
        });
    }

    // Sort by final_score descending
    scores.sort_by(|a, b|
        b.final_score.partial_cmp(&a.final_score).unwrap());
}

Ok(WeightedConsensus {
    best_agent: scores[0].agent_name.clone(),
    confidence: scores[0].final_score,
    scores,
})
}
```

Configuration (add to config.toml):

```
[ppp.weights]
technical = 0.7 # Technical quality weight
interaction = 0.3 # Interaction quality weight (proactivity +
```

personalization)

Acceptance Criteria: - [] Weighted consensus uses both technical + interaction scores - [] Configurable weights via config.toml - [] Integration tests with mock trajectories - [] Backward compatible (defaults to technical-only if no trajectory)

Configuration Changes

config.toml Additions

```
# PPP Framework Settings
[ppp]
enabled = true # Enable PPP trajectory logging and scoring

[ppp.trajectory]
# Logging settings
buffer_size = 10          # Turns to buffer before flush
flush_interval_ms = 500    # Flush interval in milliseconds
enable_wal = false         # Enable WAL mode (Phase 3)

[ppp.question_classifier]
method = "heuristic"      # "heuristic" or "llm" (Phase 3)
llm_model = "claude-haiku" # If method=llm

[ppp.weights]
# Consensus scoring weights
technical = 0.7           # Technical quality weight
interaction = 0.3          # Interaction quality (R_Proact +
R_Pers)

[ppp.archival]
# Cleanup settings (Phase 3)
enabled = false
retention_days = 30         # Keep last N days
auto_cleanup = true         # Auto-delete old trajectories
```

Testing Strategy

Unit Tests

```
// codex-rs/tui/src/chatwidget/spec_kit/trajectory_logger_tests.rs

#[tokio::test]
async fn test_trajectory_logging() {
    let logger = TrajectoryLogger::new(":memory:".to_string());

    let turn = Turn {
        trajectory_id: 1,
        turn_number: 1,
        prompt: "Implement OAuth".to_string(),
        response: "Which provider?".to_string(),
        token_count: Some(50),
        latency_ms: Some(200),
    };
}
```

```

        logger.log_turn(turn).await.unwrap();

        // Verify logged
        tokio::time::sleep(Duration::from_millis(600)).await; // Wait
    for flush
        // ... query and assert
    }

#[test]
fn test_question_classifier() {
    assert_eq!(
        classify_effort("Which provider?"),
        EffortLevel::Low
    );

    assert_eq!(
        classify_effort("Should I investigate caching?"),
        EffortLevel::High
    );
}

#[test]
fn test_r_proact_formula() {
    // Test all scenarios from PPP paper
    // ...
}

```

Integration Tests

```

// codex-rs/tui/tests/ppp_integration_test.rs

#[tokio::test]
async fn test_full_ppp_pipeline() {
    // 1. Create test trajectory
    let db = open_test_db();
    let traj_id = create_trajectory(&db, "SPEC-TEST", "agent1",
"run-1")?;

    // 2. Log turns with questions
    log_turn(traj_id, 1, "Implement X", "Which approach?").await?;
    log_question(traj_id, 1, "Which approach?", EffortLevel::Low).await?;

    // 3. Calculate scores
    let proact = calculate_r_proact(&db, traj_id)?;
    assert_eq!(proact.r_proact, 0.05);

    let pers = calculate_r_pers(&db, traj_id)?;
    assert_eq!(pers.r_pers, 0.05);

    // 4. Weighted consensus
    let consensus = run_spec_consensus_weighted("SPEC-TEST", "plan",
artifacts, (0.7, 0.3)).await?;
    assert!(consensus.scores[0].interaction_score > 0.0);
}

```

Performance Benchmarks

Benchmark Suite

```
// codex-rs/tui/benches/trajectory_bench.rs

use criterion::{black_box, criterion_group, criterion_main,
Criterion};

fn bench_log_turn(c: &mut Criterion) {
    c.bench_function("log_turn_async_batched", |b| {
        let logger = TrajectoryLogger::new(":memory:".to_string());
        b.iter(|| {
            logger.log_turn(black_box(Turn { /* ... */ }));
        });
    });
}

fn bench_r_proact_query(c: &mut Criterion) {
    let db = create_test_db_with_1000_turns();
    c.bench_function("calculate_r_proact_1000_turns", |b| {
        b.iter(|| {
            calculate_r_proact(&db, black_box(1));
        });
    });
}

criterion_group!(benches, bench_log_turn, bench_r_proact_query);
criterion_main!(benches);
```

Expected Results: - log_turn: <0.1ms (buffered) - calculate_r_proact (100 turns): <5ms - calculate_r_pers (100 turns): <5ms - full_weighted_consensus: <20ms

Migration Guide

Existing Projects

Step 1: Update database schema

```
# Automatic migration on first run
cargo run -- /speckit.status SPEC-KIT-001
# Detects old schema, runs migration
```

Step 2: Enable PPP in config

```
# config.toml
[ppp]
enabled = true
```

Step 3: Verify migration

```
sqlite3 ~/.local/share/codex-tui/consensus.db
> SELECT name FROM sqlite_master WHERE type='table' AND name LIKE
'trajectory%';
# Should show 4 tables
```

Rollback Plan

If PPP trajectory logging causes issues:

Step 1: Disable in config

```
[ppp]
enabled = false
```

Step 2: System continues with technical-only scoring (backward compatible)

Step 3: Optional cleanup

```
DROP TABLE trajectory_violations;
DROP TABLE trajectory_questions;
DROP TABLE trajectory_turns;
DROP TABLE trajectories;
```

Success Metrics

Phase 1 (Foundation)

- Database schema created successfully
- Async logging achieves <1ms latency (buffered)
- 100% test coverage for CRUD operations
- No crashes or data loss in stress tests

Phase 2 (Scoring)

- R_{Proact} and R_{Pers} calculations correct (match PPP paper formulas)
- Weighted consensus selects best agent (validated on test cases)
- Question classifier: 75%+ accuracy
- Integration tests pass

Phase 3 (Production)

- WAL mode reduces latency to <0.02ms
 - Load test: 10K+ trajectories, queries <20ms
 - Archival cleanup works (auto-delete old data)
 - Zero production incidents
-

Next Steps

1. **Implement Phase 1** (Foundation) - 12 hours
 - Start with Task 1 (schema)
 - Then Task 2 (async logging)
 - Unit tests for both
2. **Validate with real agents** - Use /speckit.plan to generate test trajectories
3. **Measure baseline performance** - Benchmark before optimization
4. **Proceed to Phase 2** - Only after Phase 1 validated

5. Integration with SPEC-PPP-003 - Connect weighted consensus

Dependencies

Phase 1: - `tokio-rusqlite = "0.5"` (add to `Cargo.toml`) - regex (already in workspace)

Phase 2: - SPEC-PPP-001 (question effort classifier research) - SPEC-PPP-002 (preference violation detection) - SPEC-PPP-003 (weighted consensus formulas)

Phase 3: - Optional: opentelemetry crates for Phoenix export -
Optional: LLM API access for question classification

Estimated Total Timeline: - Phase 1: 1.5 days (solo engineer) -
Phase 2: 2 days - Phase 3: 1.5 days - **Total: 5 days** (1 week with
buffer)